

Lecture 10: Machine Translation I

Alan Ritter

(many slides from Greg Durrett)

Final Project

- Groups of 3-4
- Sign up here:
 - <https://forms.gle/WTAp1FQFjSRuss6i6>
- **Highly Recommended:** stop by the office hours to chat and get feedback on your project.
- **Scope:** on the order of one of the programming assignments
 - But, need to define the problem, come up with right feature representation, write up results in a formal report.

Selecting a Topic

- Part of your thesis? Great!
- Find a problem you are interested in where you think NLP can help.
- Experiment with one of the algorithms we discussed in class.
- **First question:** what is the dataset?

Datasets

- Various Semeval Tasks:
 - <http://alt.qcri.org/semeval2018/index.php?id=tasks>
- Fake News Challenge:
 - <http://www.fakenewschallenge.org/>
- Machine Translation:
 - <http://www.statmt.org/wmt19/robustness.html>
- Dialogue:
 - <https://github.com/mgalley/DSTC7-End-to-End-Conversation-Modeling>
- Many more...

Requirements

- 4 Page Report
- Include empirical analysis of your approach
 - Report performance on dev / test set
 - Compare against some reasonable baseline method.

Advice

- **First question:** is data available?
- Try to get a simple baseline working as early as possible to determine whether your project idea is feasible.
- Start with a manageable-sized dataset
 - Then scale up...

Group Formation Time (5 minutes)

This Lecture

- ▶ MT and evaluation
- ▶ Word alignment
- ▶ Language models
- ▶ Phrase-based decoders
- ▶ Syntax-based decoders (probably next time)

MT Basics

MT Basics



< 2/8

特朗普偕家人在白宫阳台观看百年一遇日全食

>

People's Daily, August 30, 2017

MT Basics



The image shows a woman with blonde hair wearing 3D glasses, looking upwards. A man's head is partially visible behind her. Overlaid on the image is a white rectangular box containing a translation interface. The word "Translate" is written in red at the top left. Below it is a row of language buttons: English, French, Spanish, Chinese - detected, and a dropdown arrow. To the right of these buttons is a double-headed arrow icon. At the bottom of the box, there is a large orange-bordered text area containing the Chinese translation: "特朗普偕家人在白宫阳台观看百年一遇日全食". In the bottom left corner of the main image, there is a navigation bar with a left arrow, the number "2/8", and the text "特朗普偕家人在白宫阳台观看百年一遇日全食".

Translate

English French Spanish Chinese - detected ▾

特朗普偕家人在白宫阳台观看百年一遇日全食

2/8

特朗普偕家人在白宫阳台观看百年一遇日全食

People's Daily, August 30, 2017

MT Basics



People's Daily, August 30, 2017

Trump Pope family watch a hundred years a year in the White House balcony

MT Basics



People's Daily, August 30, 2017

Trump Pope family watch a hundred years a year in the White House balcony

MT Ideally

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami
J'ai une amie

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation
- ▶ Hard for semantic representations to cover everything

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation
- ▶ Hard for semantic representations to cover everything
- ▶ Everyone has a friend =>

MT Ideally

- ▶ I have a friend $\Rightarrow \exists x \text{ friend}(x, \text{self}) \Rightarrow$ J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation
- ▶ Hard for semantic representations to cover everything
- ▶ Everyone has a friend $\Rightarrow \exists x \forall y \text{ friend}(x, y)$
 $\forall x \exists y \text{ friend}(x, y)$

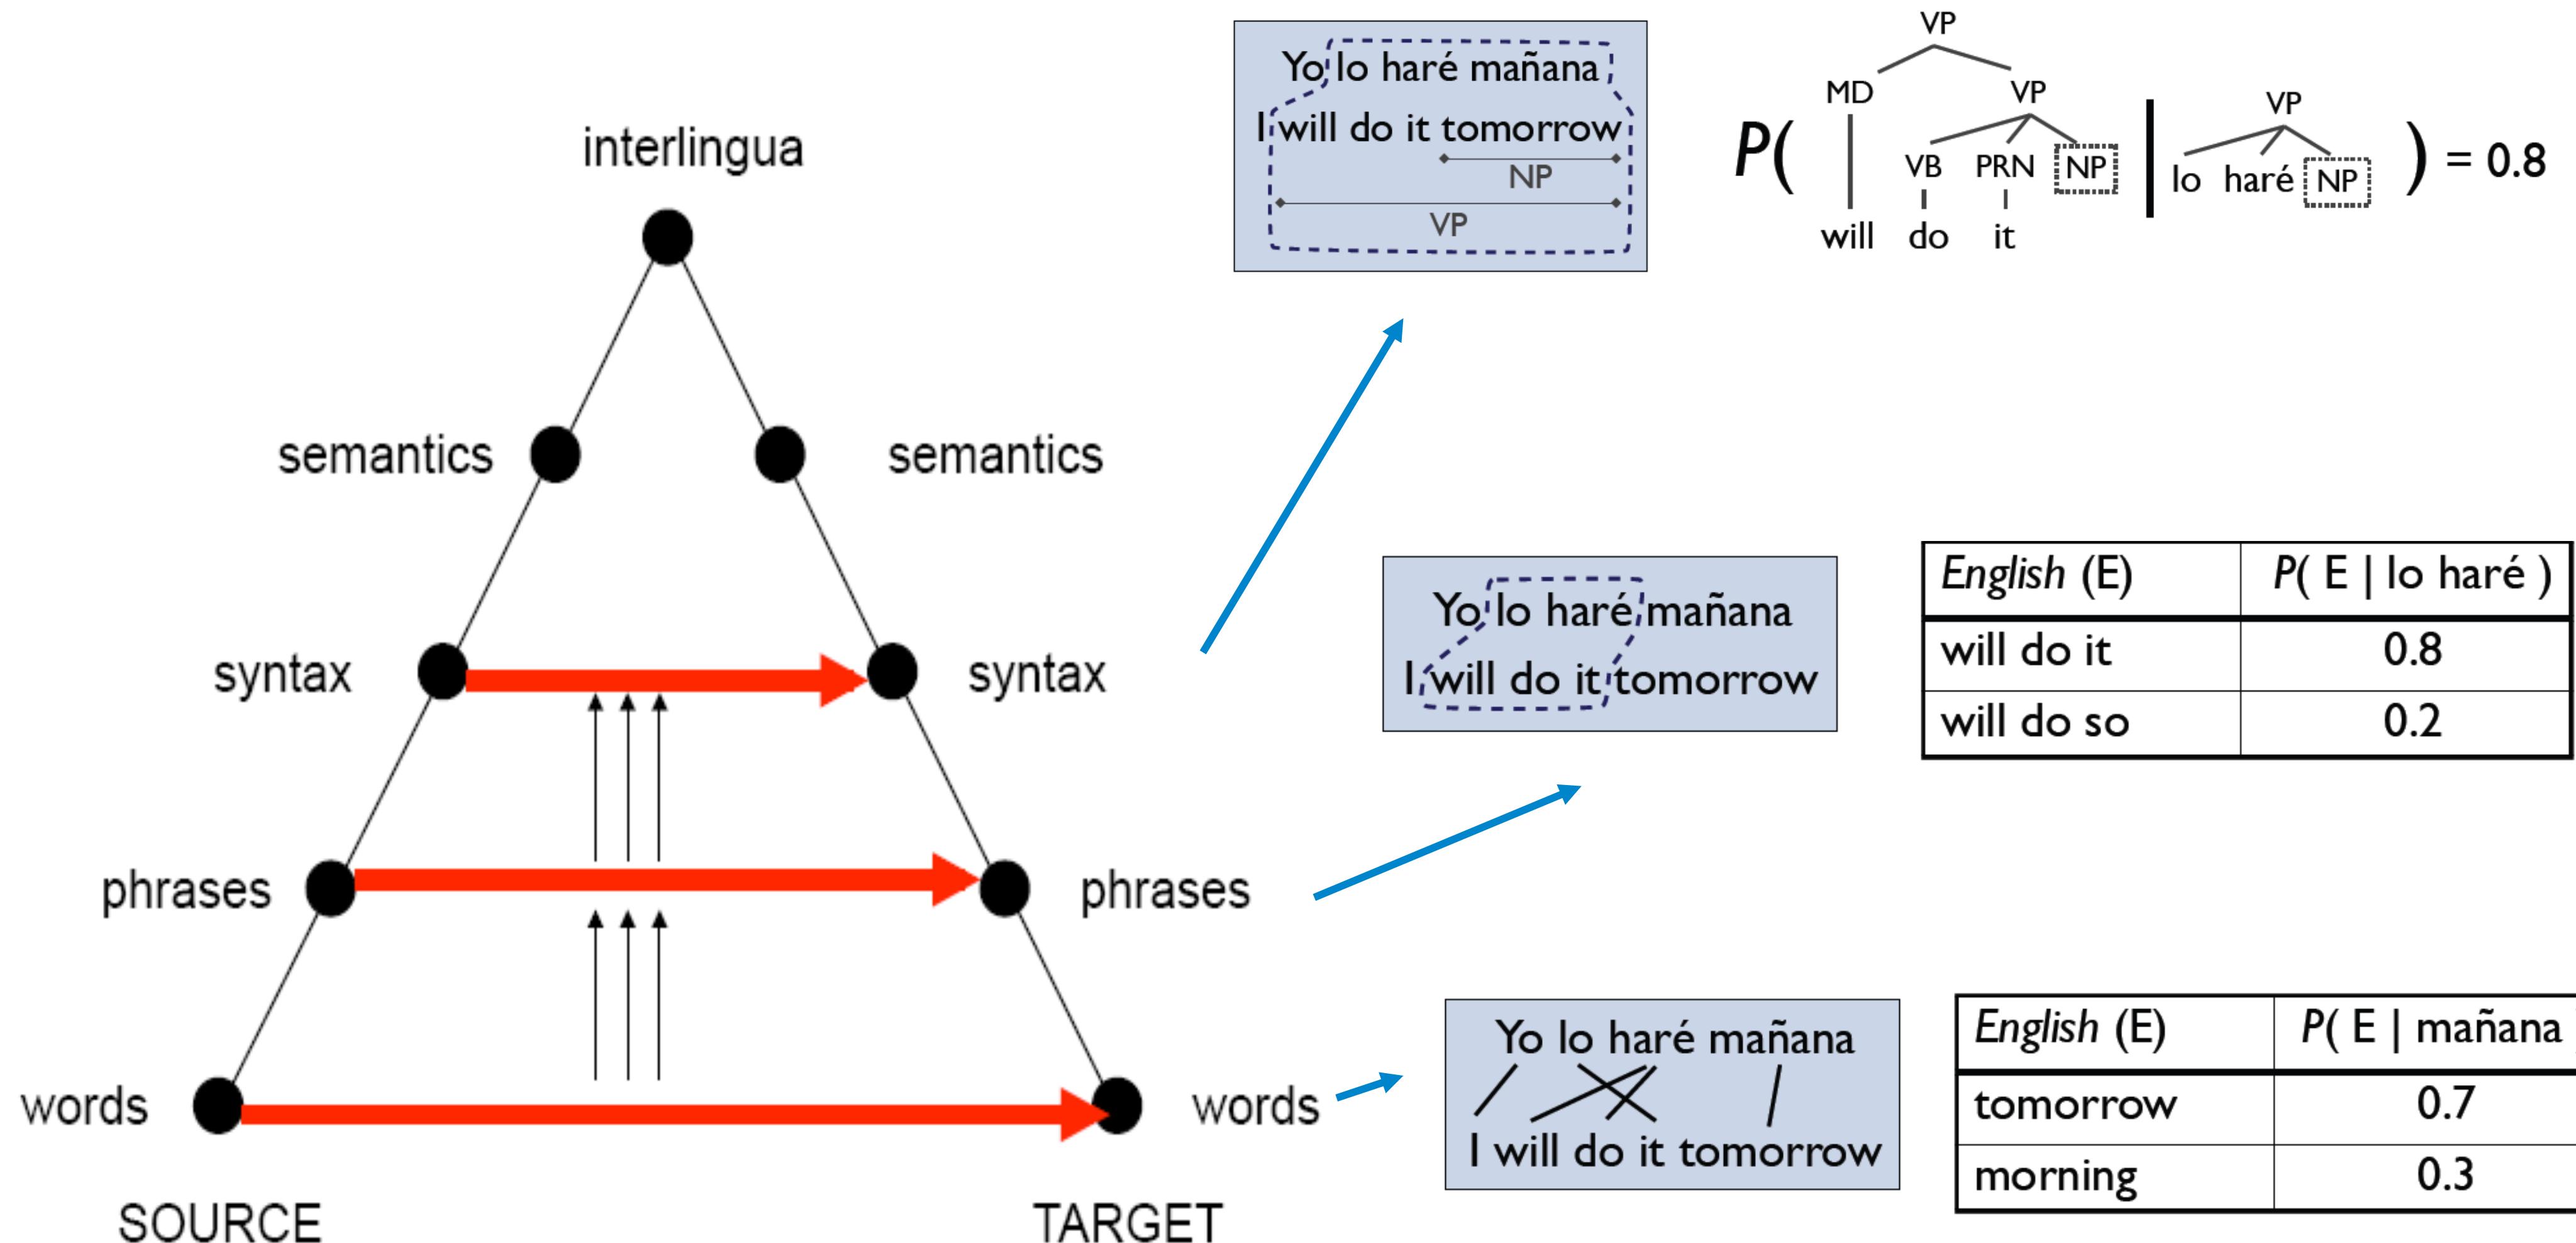
MT Ideally

- ▶ I have a friend $\Rightarrow \exists x \text{ friend}(x, \text{self}) \Rightarrow$ J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation
- ▶ Hard for semantic representations to cover everything
- ▶ Everyone has a friend $\Rightarrow \exists x \forall y \text{ friend}(x, y) \Rightarrow$ Tous a un ami
 $\forall x \exists y \text{ friend}(x, y)$

MT Ideally

- ▶ I have a friend => $\exists x \text{ friend}(x, \text{self})$ => J'ai un ami
J'ai une amie
- ▶ May need information you didn't think about in your representation
- ▶ Hard for semantic representations to cover everything
- ▶ Everyone has a friend => $\exists x \forall y \text{ friend}(x, y)$ => Tous a un ami
 $\forall x \exists y \text{ friend}(x, y)$
 - ▶ Can often get away without doing all disambiguation — same ambiguities may exist in both languages

Levels of Transfer: Vauquois Triangle



- Today: mostly phrase-based, some syntax

Slide credit: Dan Klein

Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use

Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use
- ▶ Remember phrases from training data, translate piece-by-piece and stitch those pieces together to translate

Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use
- ▶ Remember phrases from training data, translate piece-by-piece and stitch those pieces together to translate
 - ▶ How to identify phrases? Word alignment over source-target bitext

Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use
- ▶ Remember phrases from training data, translate piece-by-piece and stitch those pieces together to translate
 - ▶ How to identify phrases? Word alignment over source-target bitext
 - ▶ How to stitch together? Language model over target language

Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use
- ▶ Remember phrases from training data, translate piece-by-piece and stitch those pieces together to translate
 - ▶ How to identify phrases? Word alignment over source-target bitext
 - ▶ How to stitch together? Language model over target language
 - ▶ Decoder takes phrases and a language model and searches over possible translations

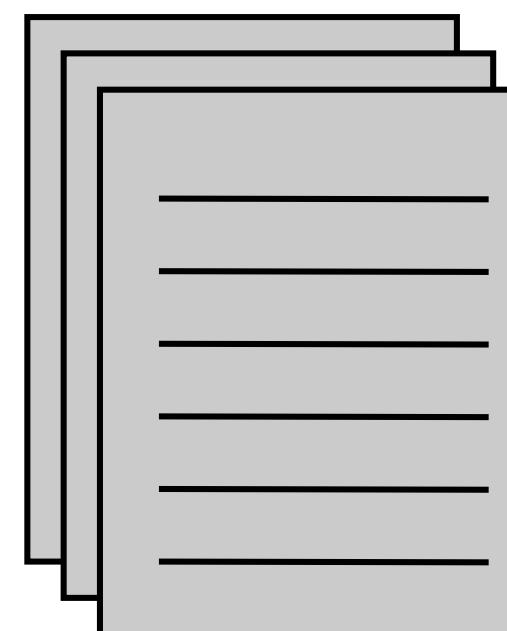
Phrase-Based MT

- ▶ Key idea: translation works better the bigger chunks you use
- ▶ Remember phrases from training data, translate piece-by-piece and stitch those pieces together to translate
 - ▶ How to identify phrases? Word alignment over source-target bitext
 - ▶ How to stitch together? Language model over target language
 - ▶ Decoder takes phrases and a language model and searches over possible translations
- ▶ NOT like standard discriminative models (take a bunch of translation pairs, learn a ton of parameters in an end-to-end way)

Phrase-Based MT

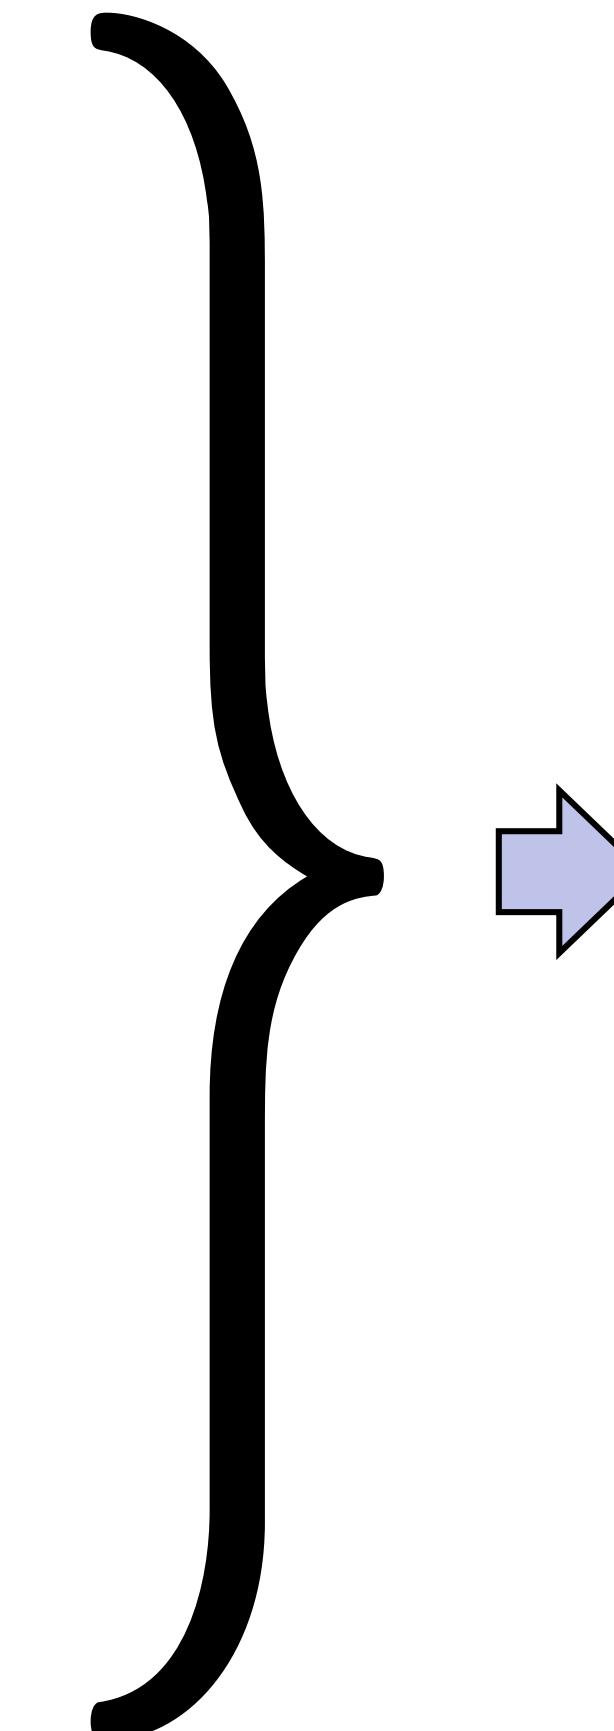
```
cat ||| chat ||| 0.9  
the cat ||| le chat ||| 0.8  
dog ||| chien ||| 0.8  
house ||| maison ||| 0.6  
my house ||| ma maison ||| 0.9  
language ||| langue ||| 0.9  
...
```

Phrase table $P(f|e)$



Unlabeled English data

Language model $P(e)$



$$P(e|f) \propto P(f|e)P(e)$$

Noisy channel model:
combine scores from
translation model +
language model to
translate foreign to
English

“Translate faithfully but make fluent English”

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

		1-gram	2-gram	3-gram
	hypothesis 1	I am exhausted ==		
	hypothesis 2	Tired is I -		
	hypothesis 3	I I I -		
	reference 1	I am tired ==		
	reference 2	I am ready to sleep now and so exhausted ==		

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right)$$

hypothesis 1

I am exhausted
==

hypothesis 2

Tired is I
-

hypothesis 3

I I I
-

reference 1

I am tired
==

reference 2

I am ready to sleep now and so exhausted
==

	1-gram	2-gram	3-gram
3/3	1/2	0/1	
1/3	0/2	0/1	
1/3	0/2	0/1	

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right).$$

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right). \quad \blacktriangleright \text{Typically } n = 4, w_i = 1/4$$

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right). \quad \blacktriangleright \text{ Typically } n = 4, w_i = 1/4$$

$$\text{BP} = \begin{cases} 1 & \text{if } c > r \\ e^{(1-r/c)} & \text{if } c \leq r \end{cases}.$$

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right).$$

- ▶ Typically $n = 4$, $w_i = 1/4$

$$\text{BP} = \begin{cases} 1 & \text{if } c > r \\ e^{(1-r/c)} & \text{if } c \leq r \end{cases}.$$

- ▶ r = length of reference
- ▶ c = length of prediction

Evaluating MT

- ▶ Fluency: does it sound good in the target language?
- ▶ Fidelity/adequacy: does it capture the meaning of the original?
- ▶ BLEU score: geometric mean of 1-, 2-, 3-, and 4-gram precision vs. a reference, multiplied by brevity penalty

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right).$$

- ▶ Typically $n = 4$, $w_i = 1/4$

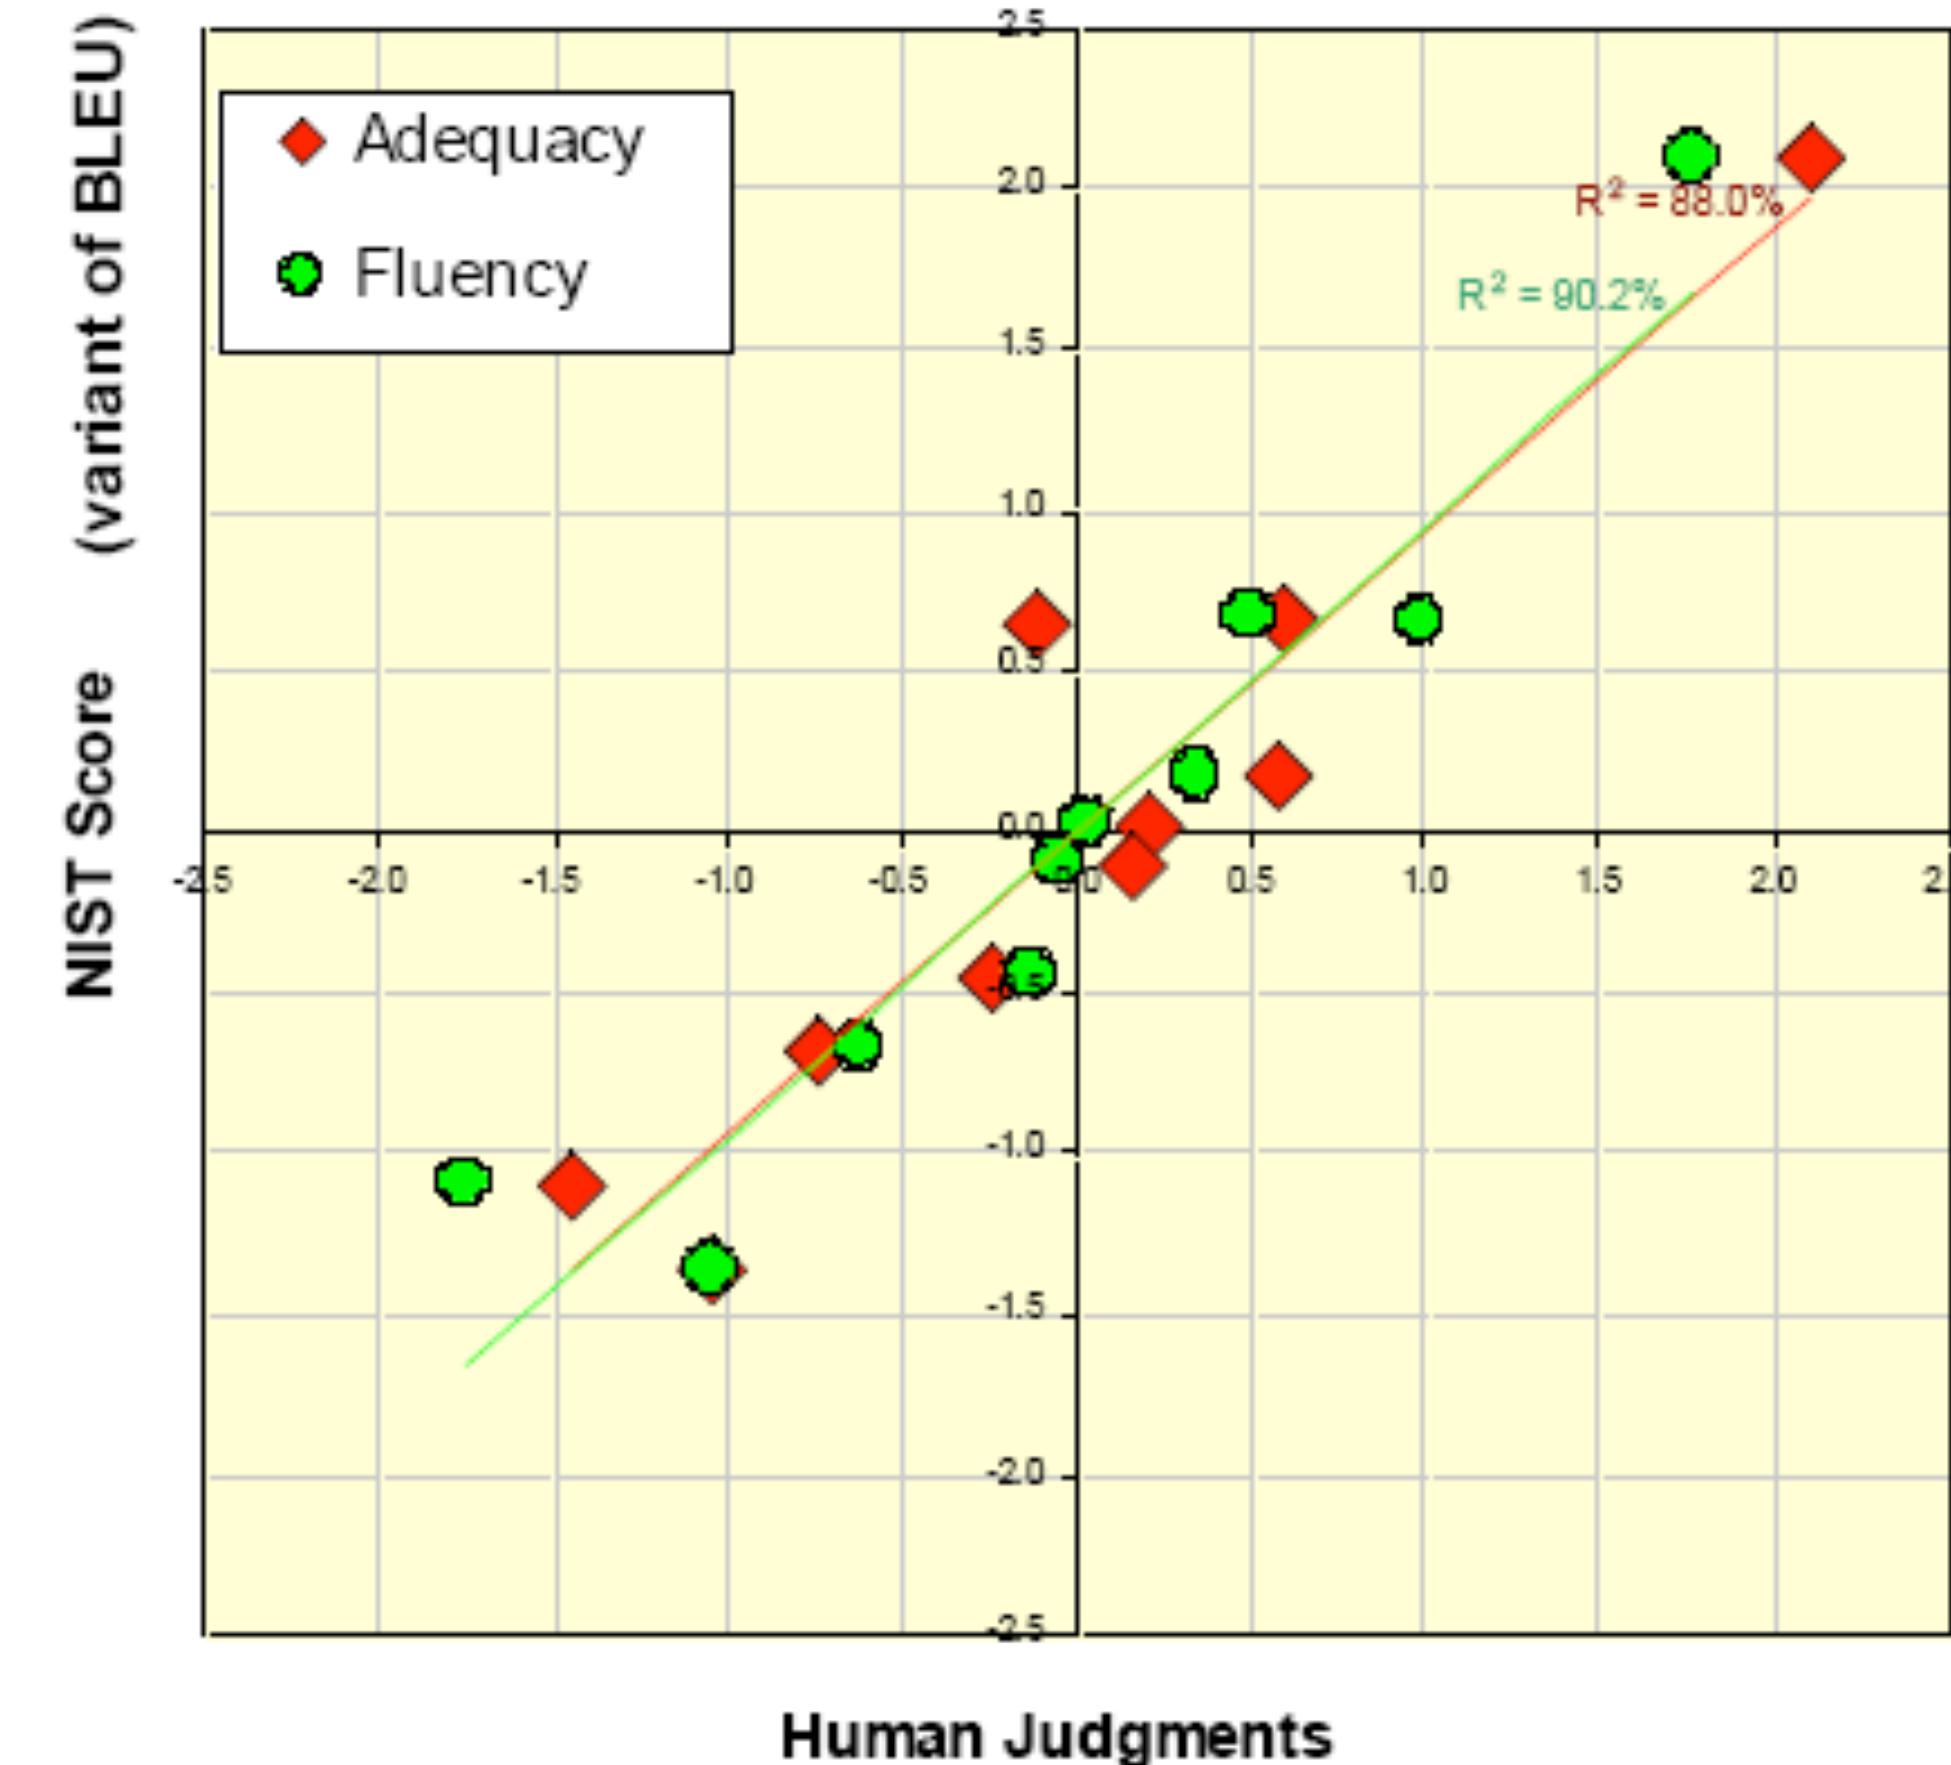
$$\text{BP} = \begin{cases} 1 & \text{if } c > r \\ e^{(1-r/c)} & \text{if } c \leq r \end{cases}.$$

- ▶ r = length of reference
- ▶ c = length of prediction

- ▶ Does this capture fluency and adequacy?

BLEU Score

- ▶ Better methods with human-in-the-loop
- ▶ HTER: human-assisted translation error rate
- ▶ If you're building real MT systems, you do user studies.
In academia, you mostly use BLEU



Word Alignment

Word Alignment

- ▶ Input: a bitext, pairs of translated sentences

nous acceptons votre opinion . ||| we accept your view

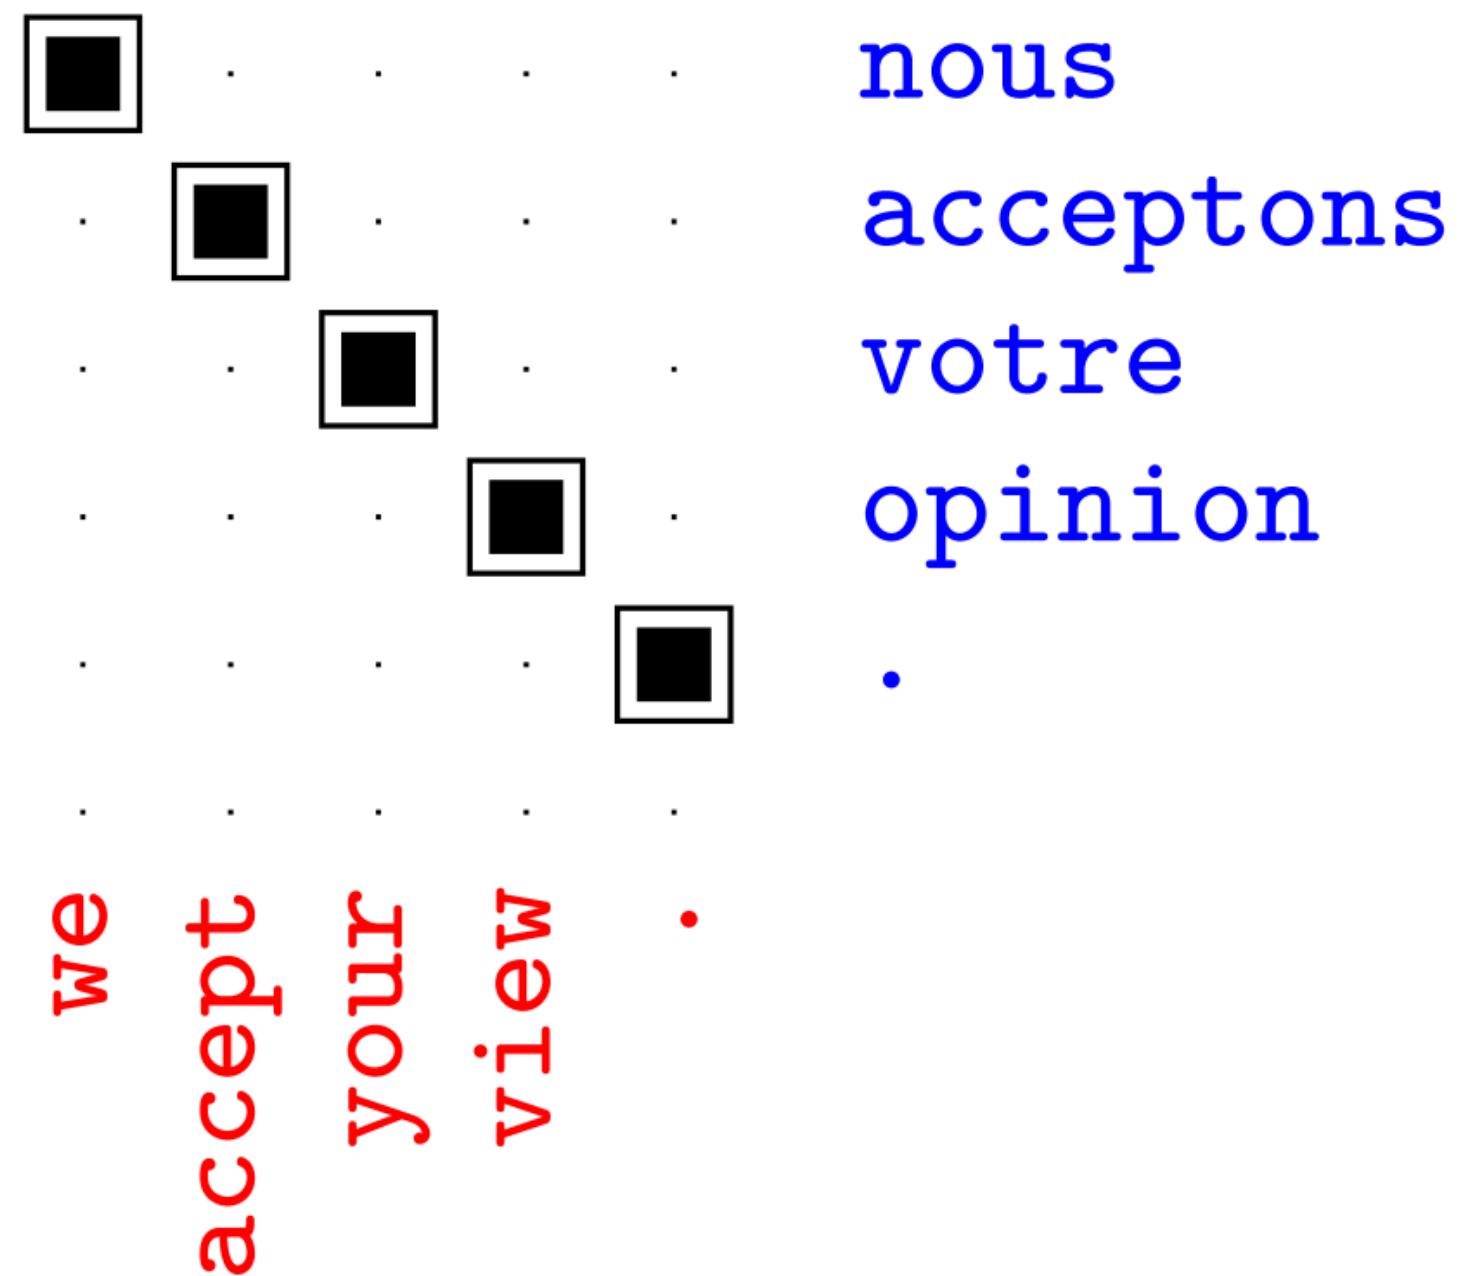
nous allons changer d'avis ||| we are going to change our minds

Word Alignment

- ▶ Input: a bitext, pairs of translated sentences

nous acceptons votre opinion . ||| we accept your view

nous allons changer d'avis ||| we are going to change our minds



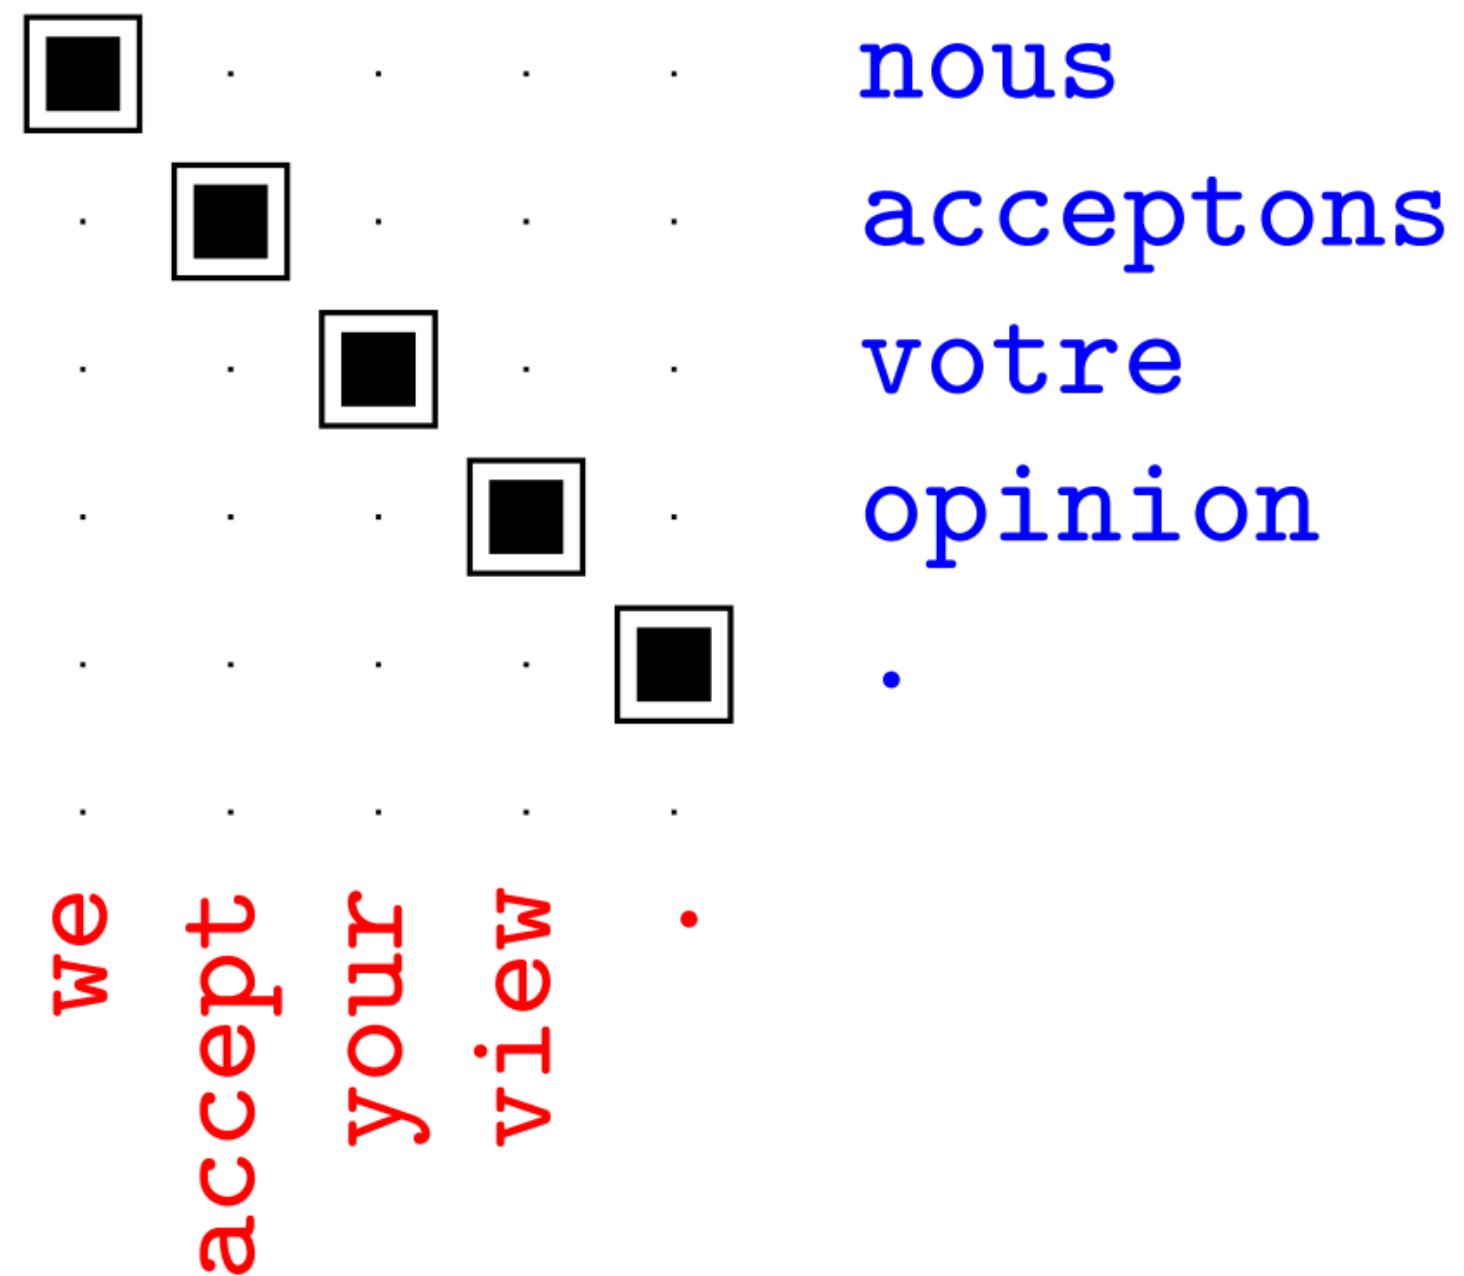
Word Alignment

- ▶ Input: a bitext, pairs of translated sentences

nous acceptons votre opinion . ||| we accept your view

nous allons changer d'avis ||| we are going to change our minds

- ▶ Output: alignments between words in each sentence



Word Alignment

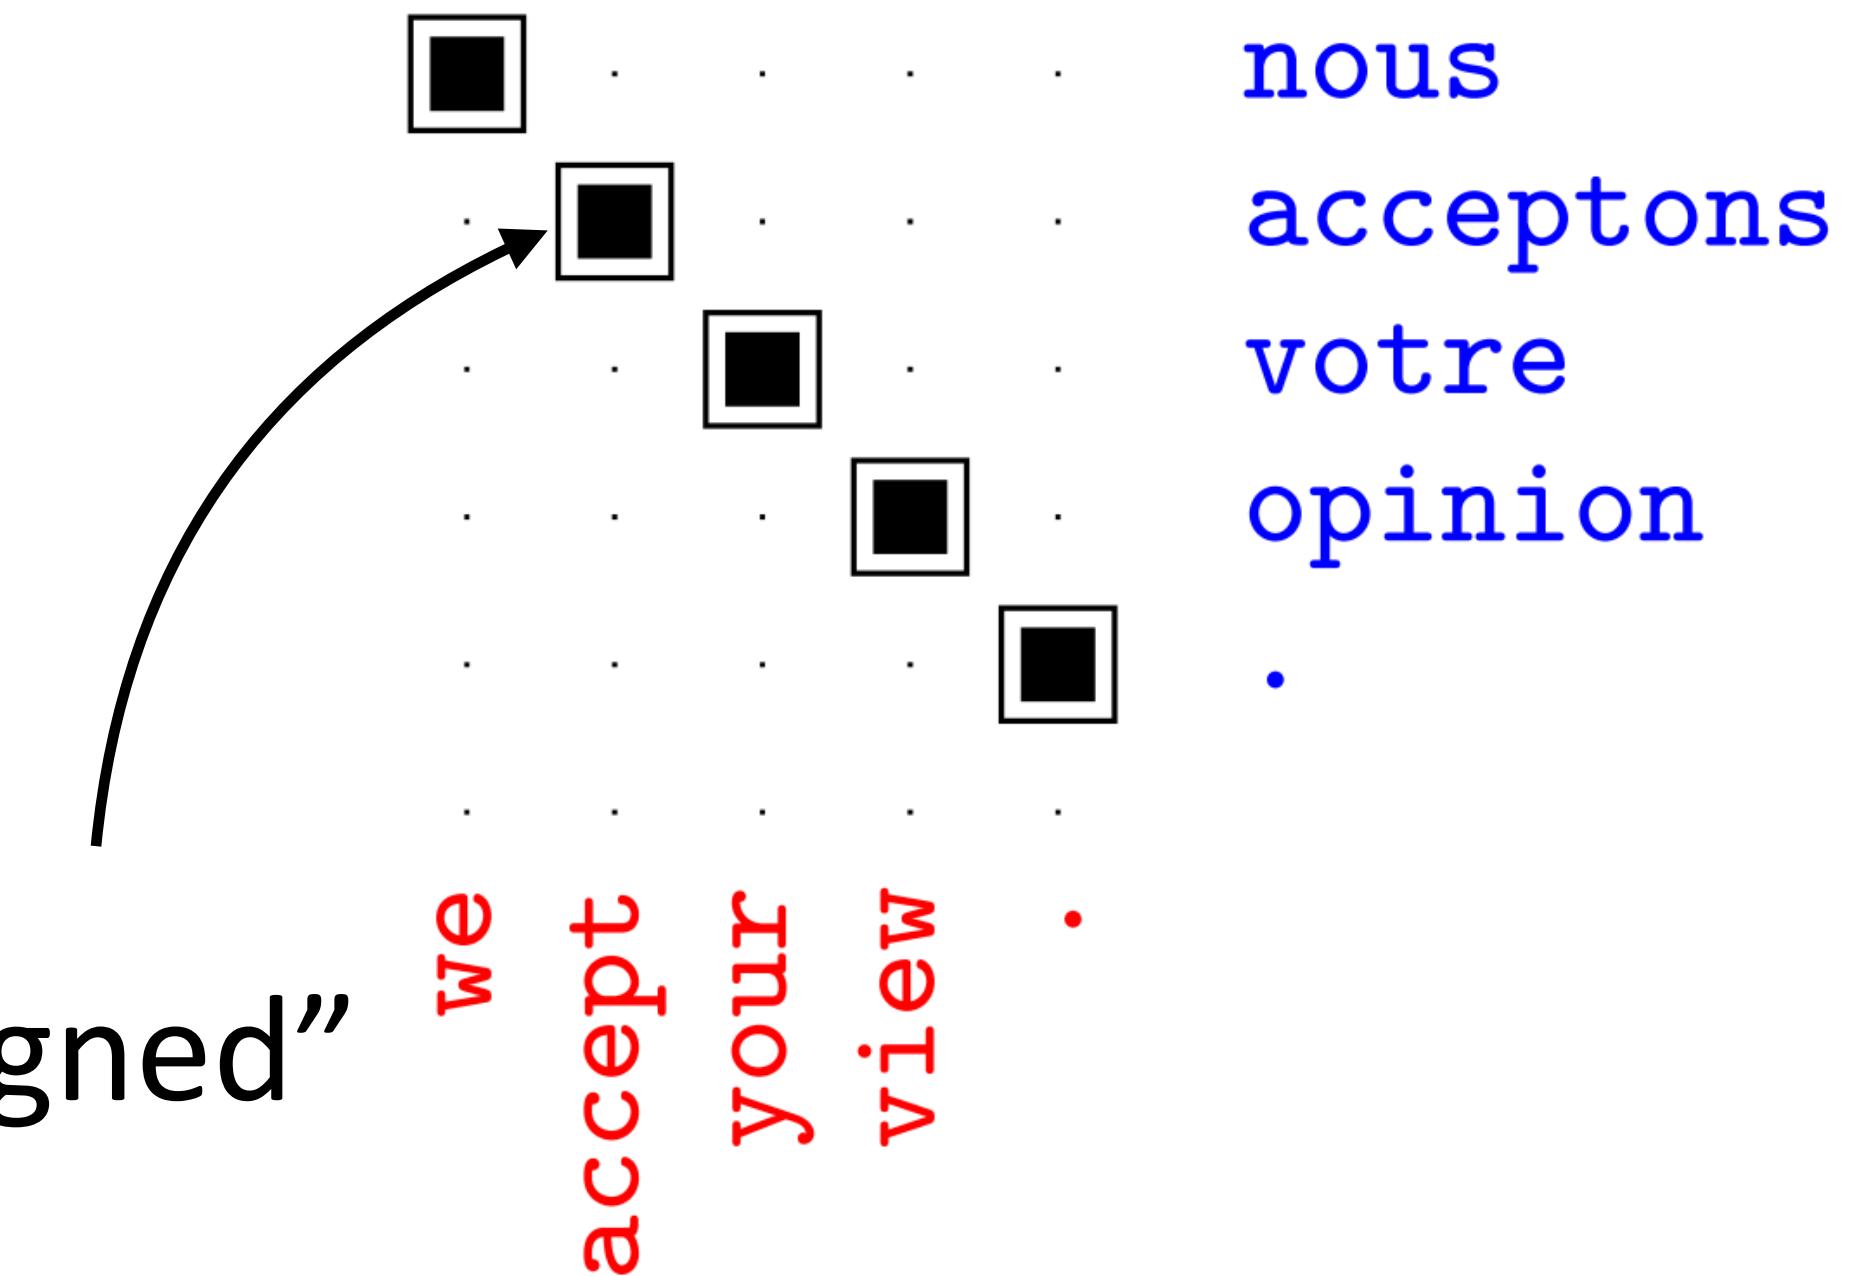
- ▶ Input: a bitext, pairs of translated sentences

nous acceptons votre opinion . ||| we accept your view

nous allons changer d'avis ||| we are going to change our minds

- ▶ Output: alignments between words in each sentence

“accept and acceptons are aligned”



Word Alignment

- ▶ Input: a bitext, pairs of translated sentences

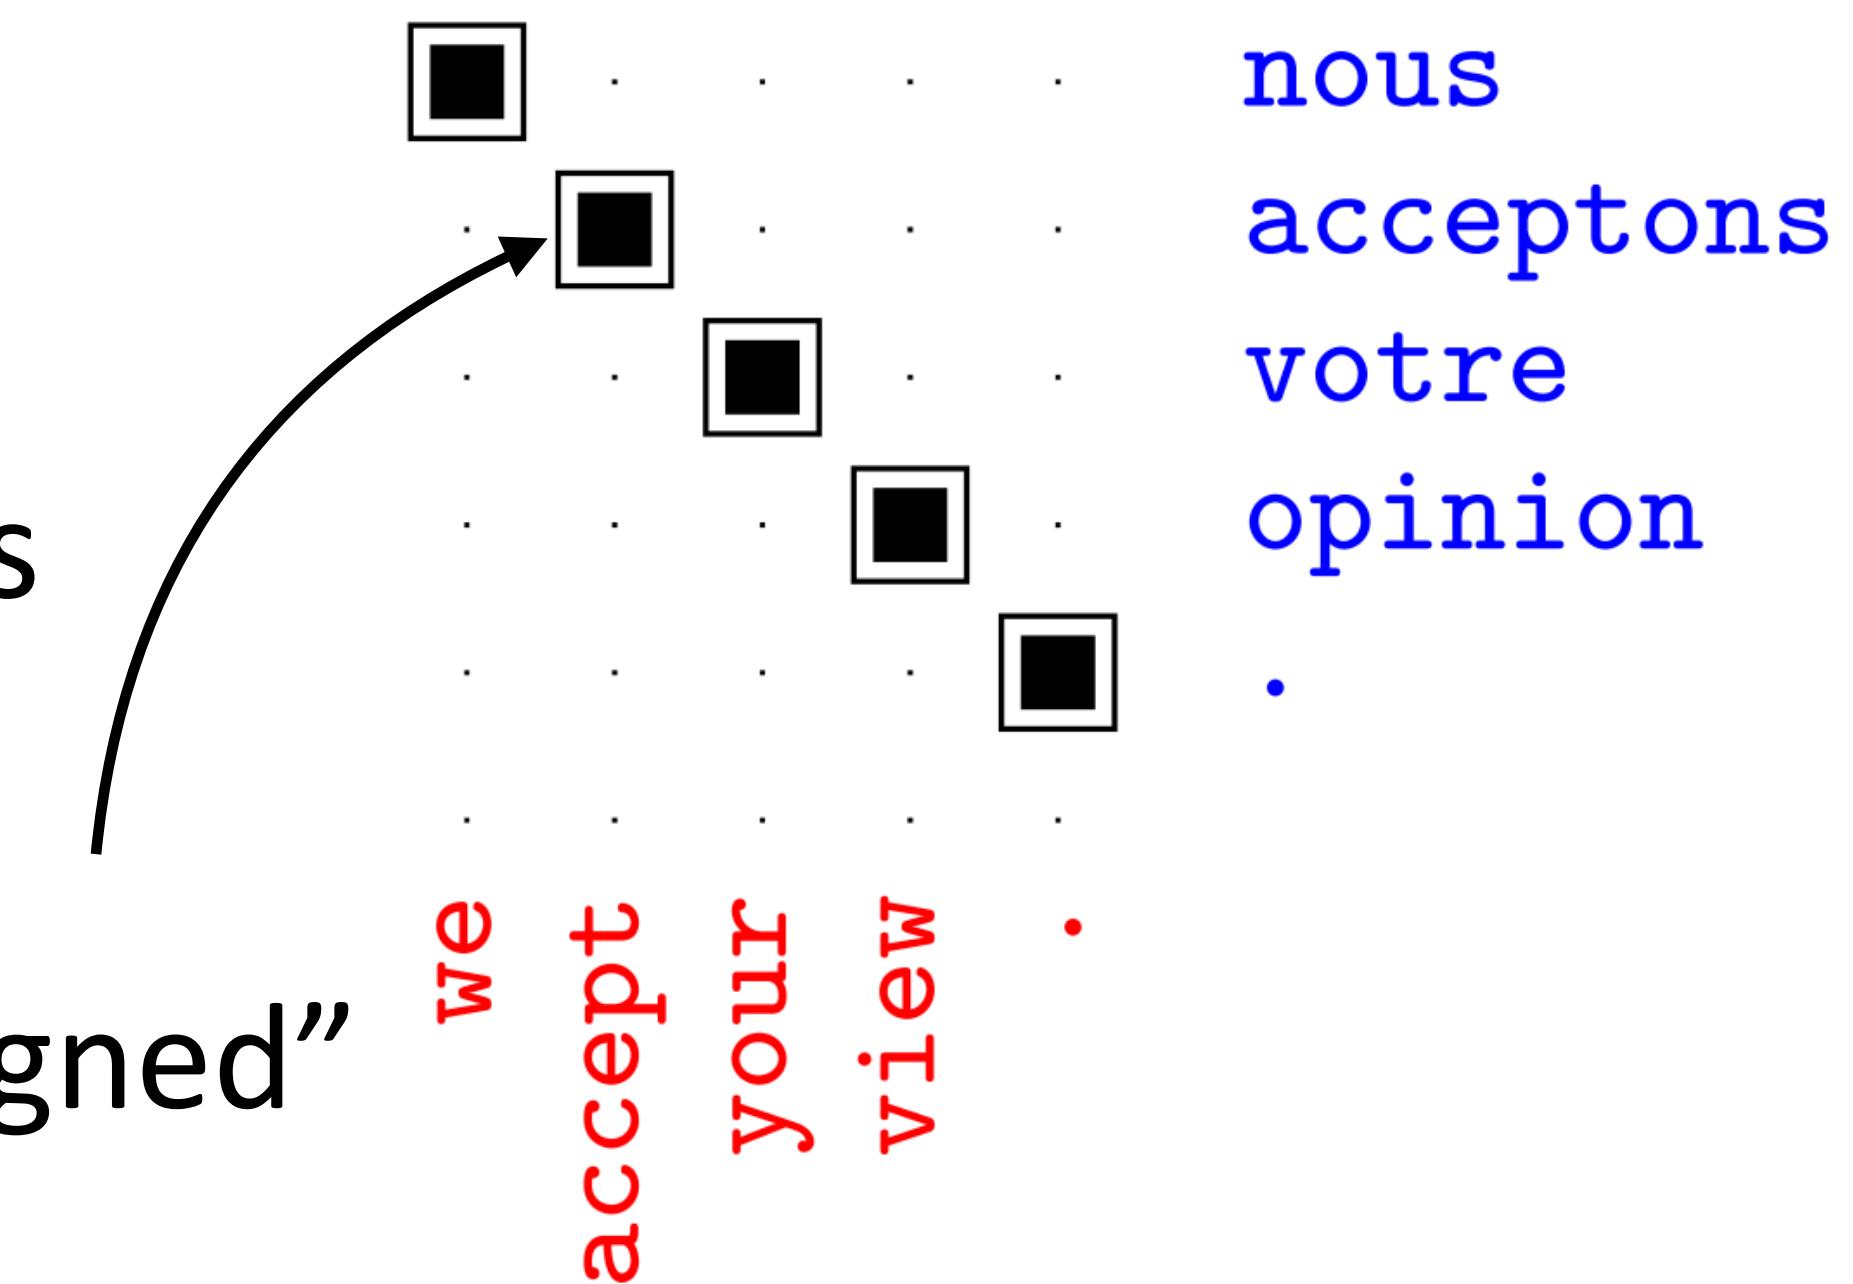
nous acceptons votre opinion . ||| we accept your view

nous allons changer d'avis ||| we are going to change our minds

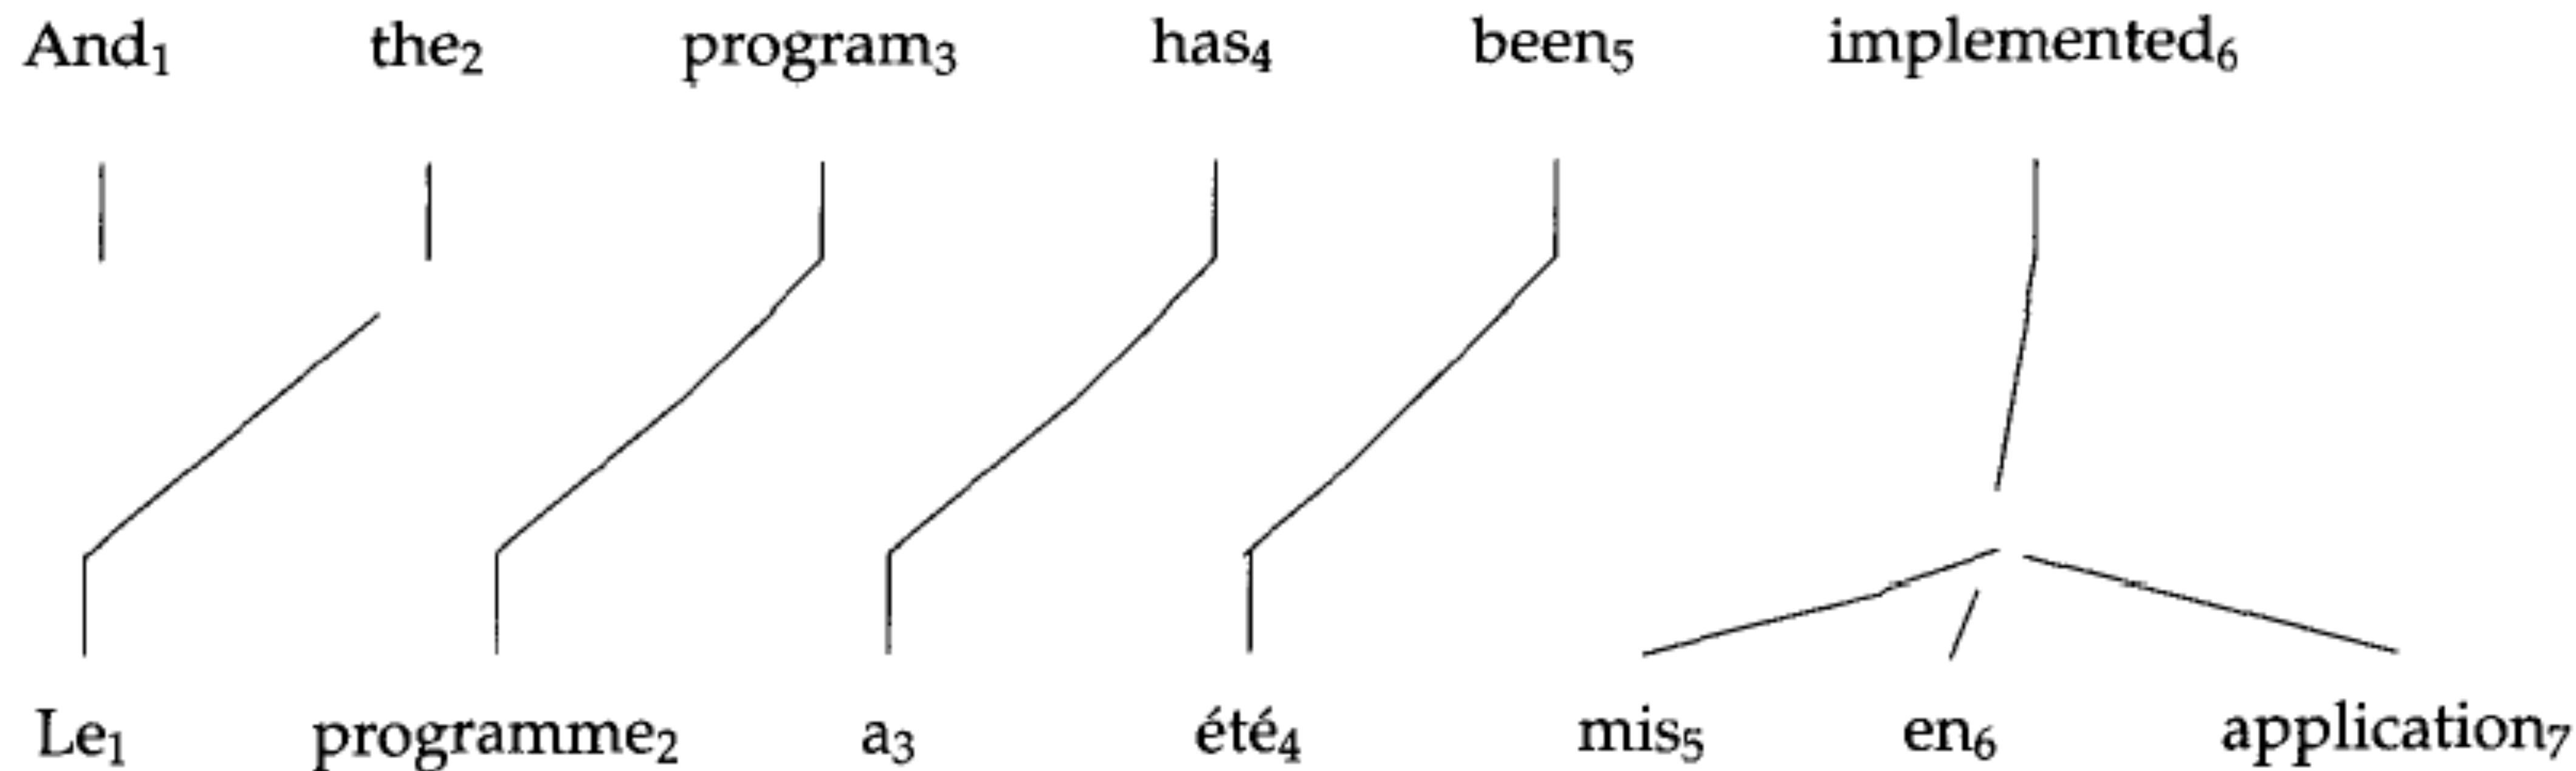
- ▶ Output: alignments between words in each sentence

- ▶ We will see how to turn these into phrases

“accept and acceptons are aligned”



1-to-Many Alignments



Word Alignment

- Models $P(f|e)$: probability of “French” sentence being generated from “English” sentence according to a model

Word Alignment

- ▶ Models $P(f|e)$: probability of “French” sentence being generated from “English” sentence according to a model
- ▶ Latent variable model:
$$P(f|e) = \sum_a P(f, a|e) = \sum_a P(f|a, e)P(a)$$

Word Alignment

- ▶ Models $P(f|e)$: probability of “French” sentence being generated from “English” sentence according to a model
- ▶ Latent variable model:
$$P(f|e) = \sum_a P(f, a|e) = \sum_a P(f|a, e)P(a)$$
- ▶ Correct alignments should lead to higher-likelihood generations, so by optimizing this objective we will learn correct alignments

IBM Model 1

- ▶ Each French word is aligned to *at most* one English word

$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i)$$

IBM Model 1

- ▶ Each French word is aligned to *at most* one English word

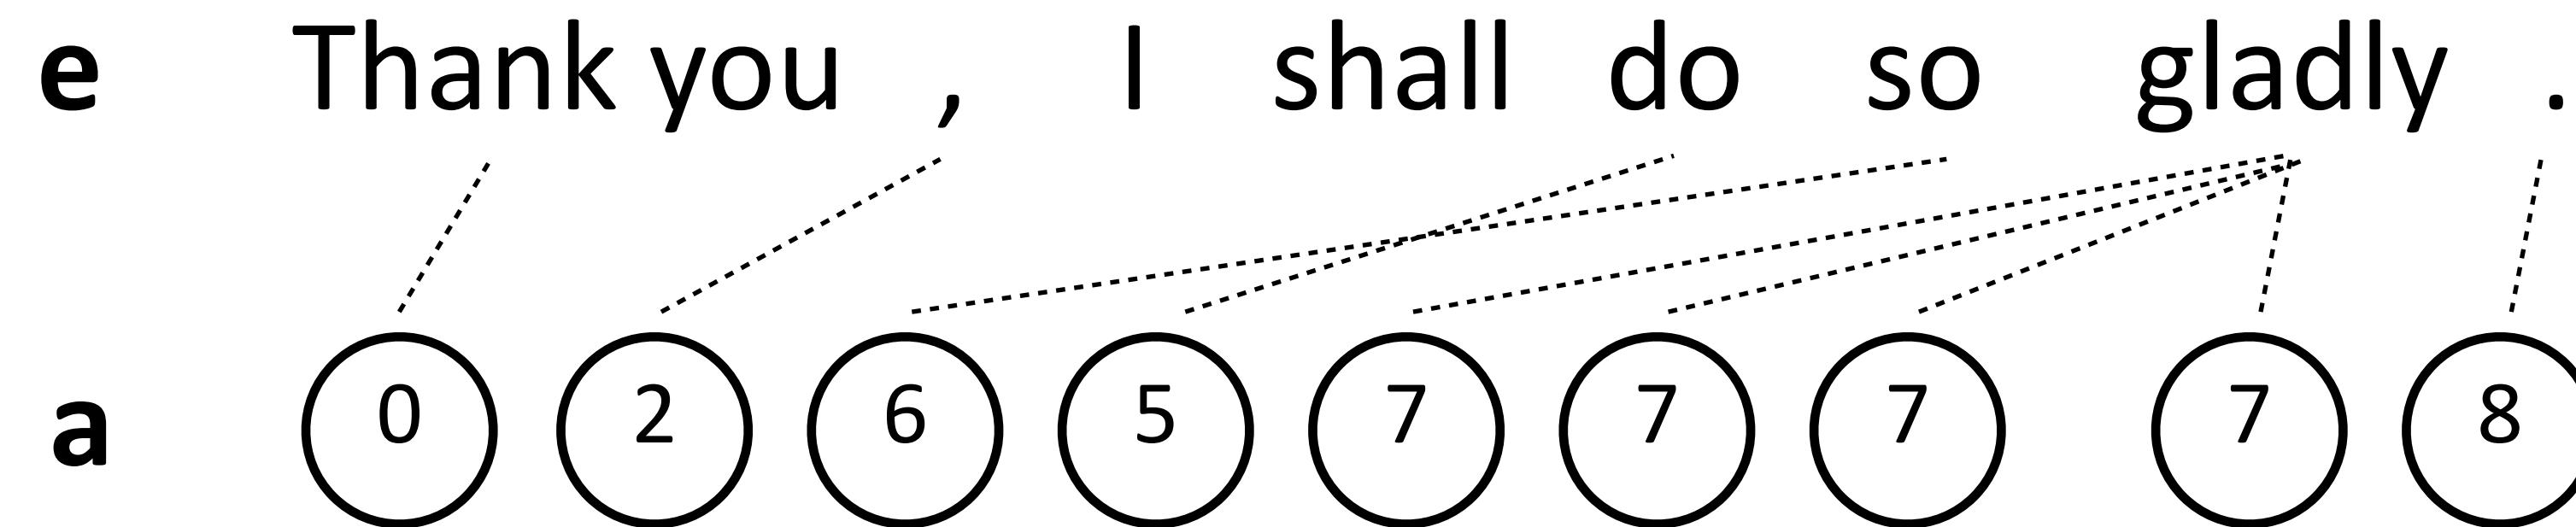
$$P(\mathbf{f}, \mathbf{a}|\mathbf{e}) = \prod_{i=1}^n P(f_i|e_{a_i})P(a_i)$$

e Thank you , I shall do so gladly .

IBM Model 1

- ▶ Each French word is aligned to *at most* one English word

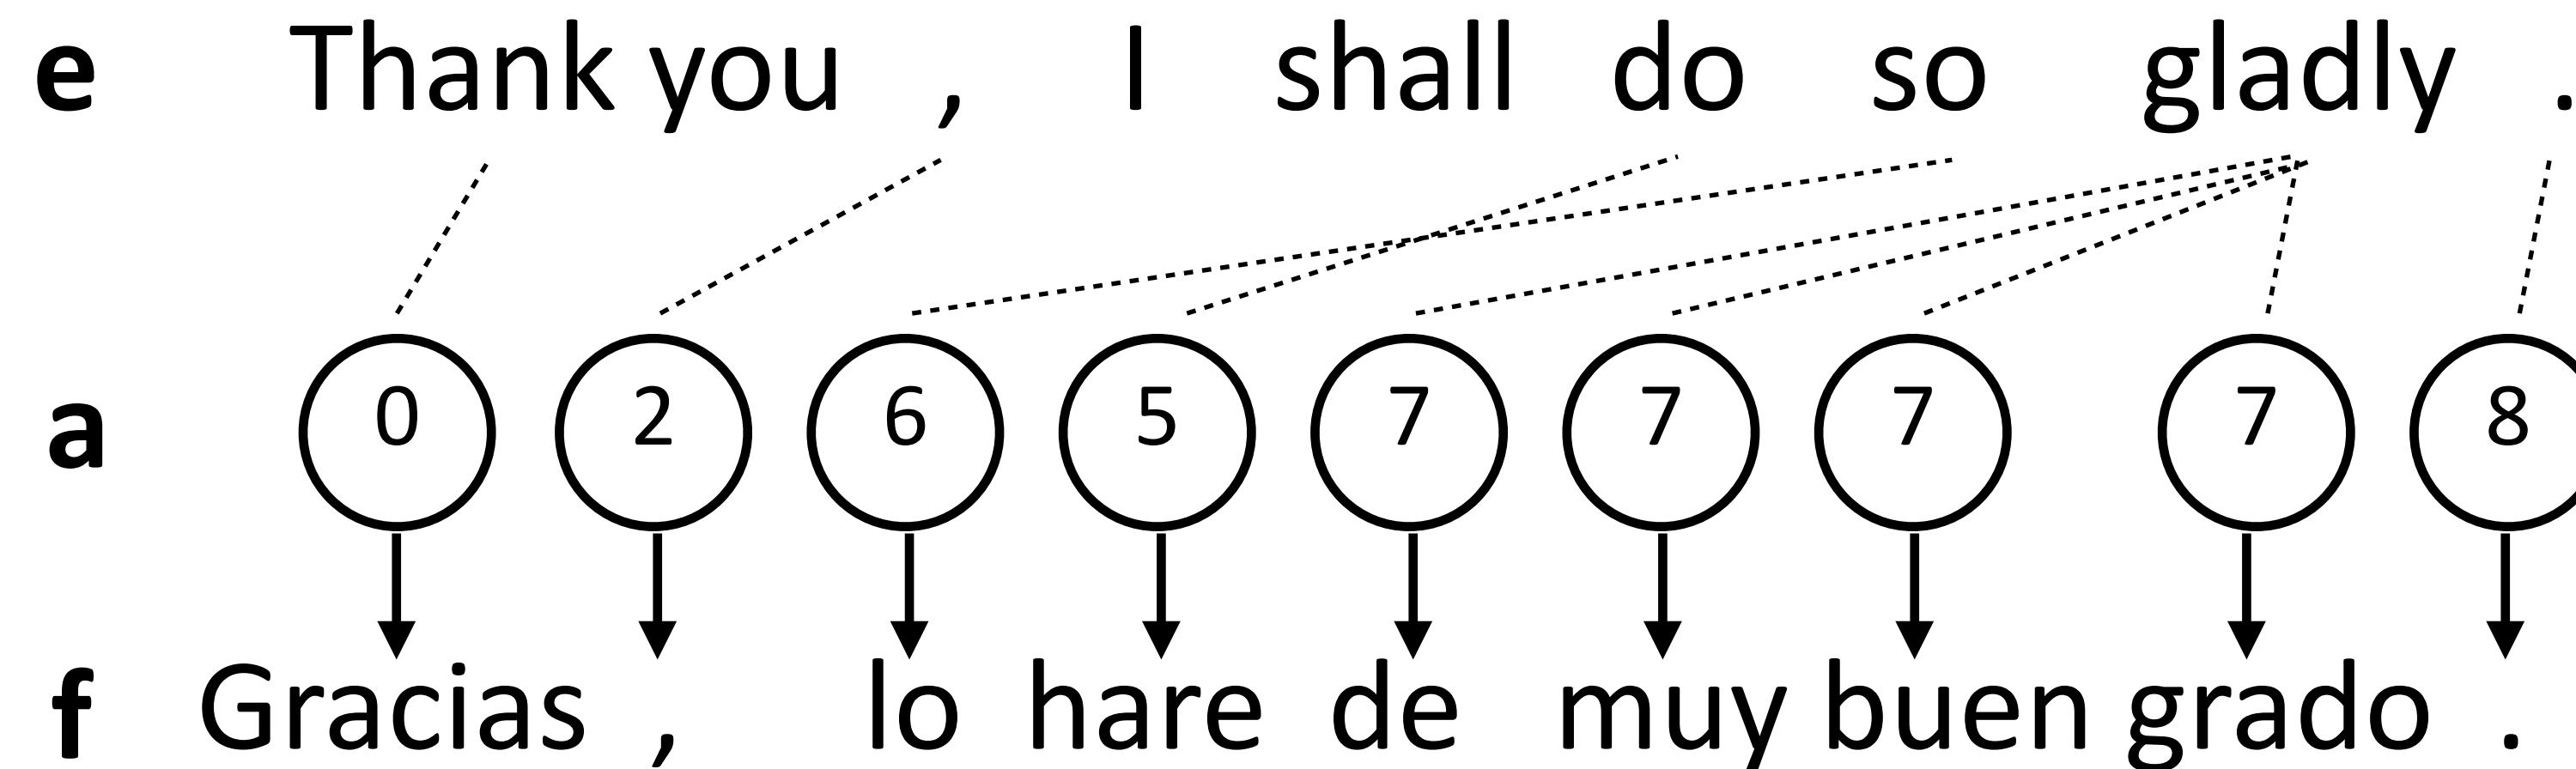
$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i)$$



IBM Model 1

- Each French word is aligned to *at most* one English word

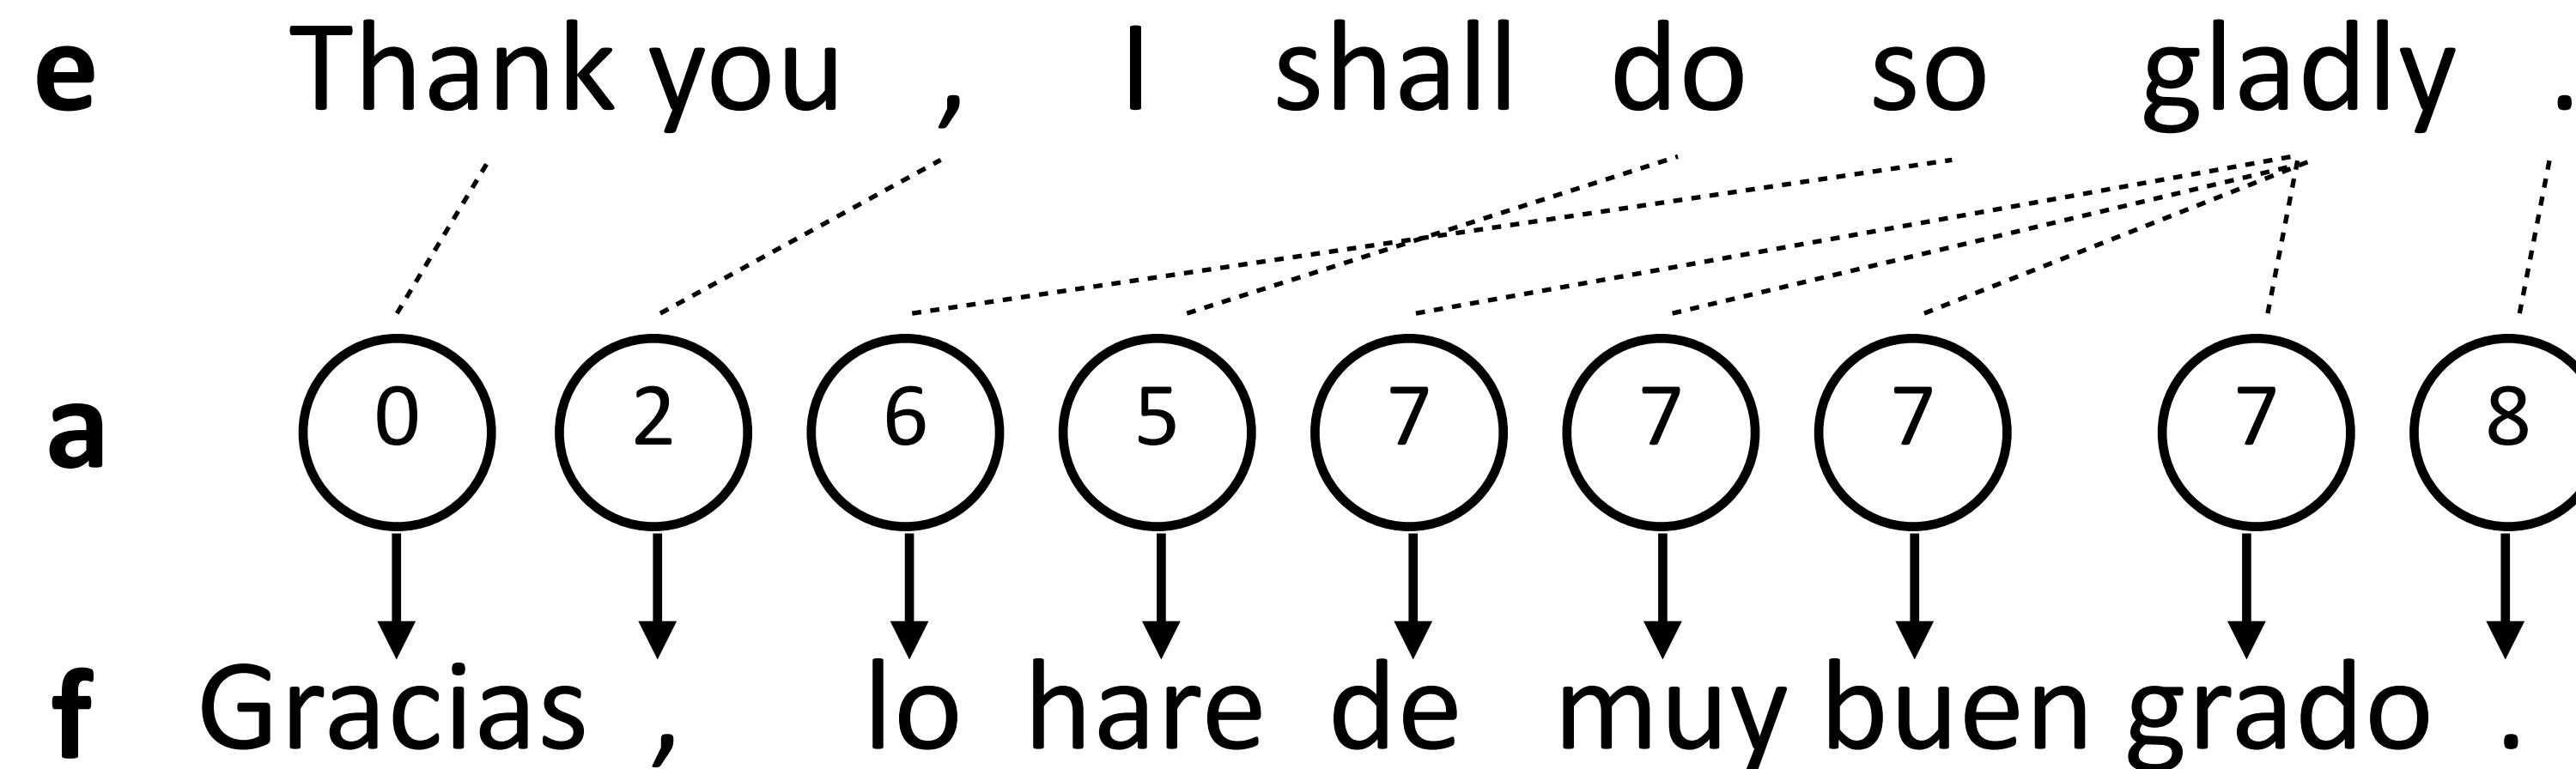
$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i)$$



IBM Model 1

- Each French word is aligned to *at most* one English word

$$P(\mathbf{f}, \mathbf{a}|\mathbf{e}) = \prod_{i=1}^n P(f_i|e_{a_i})P(a_i)$$

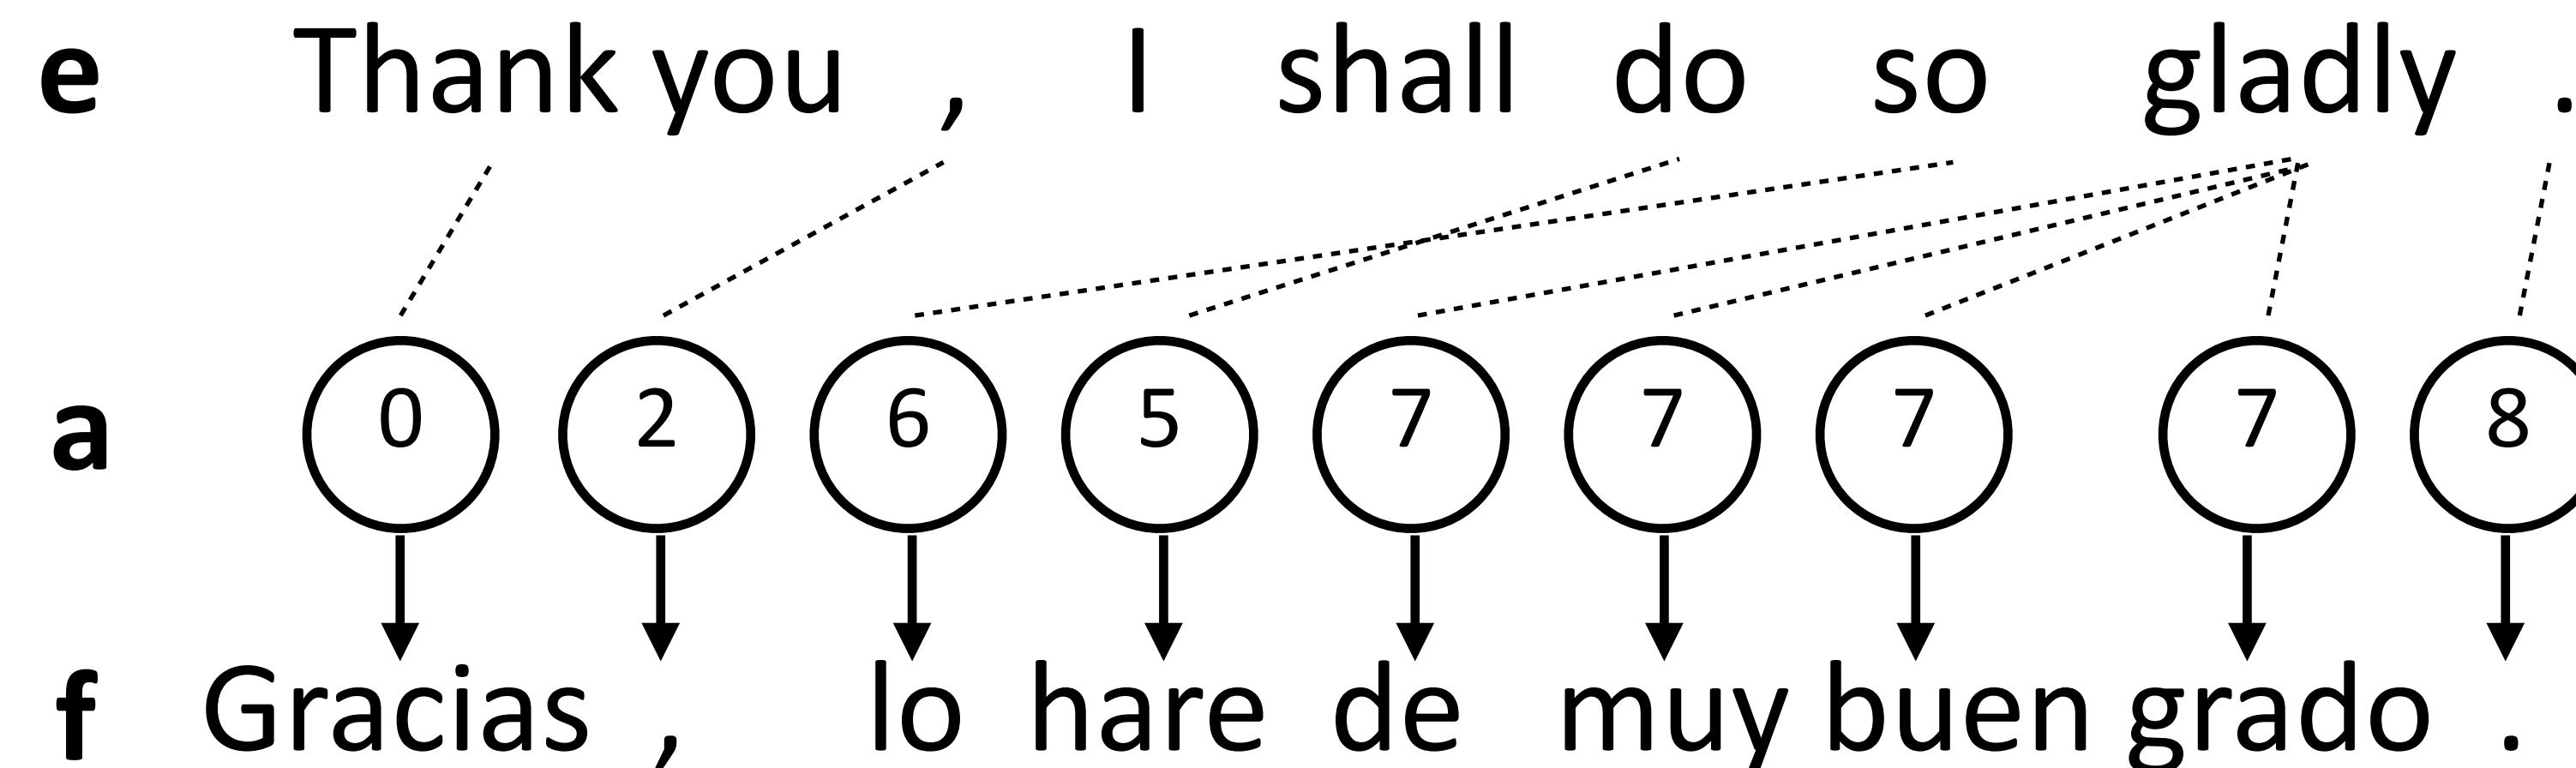


- Set $P(a)$ uniformly (no prior over good alignments)

IBM Model 1

- Each French word is aligned to *at most* one English word

$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i)$$



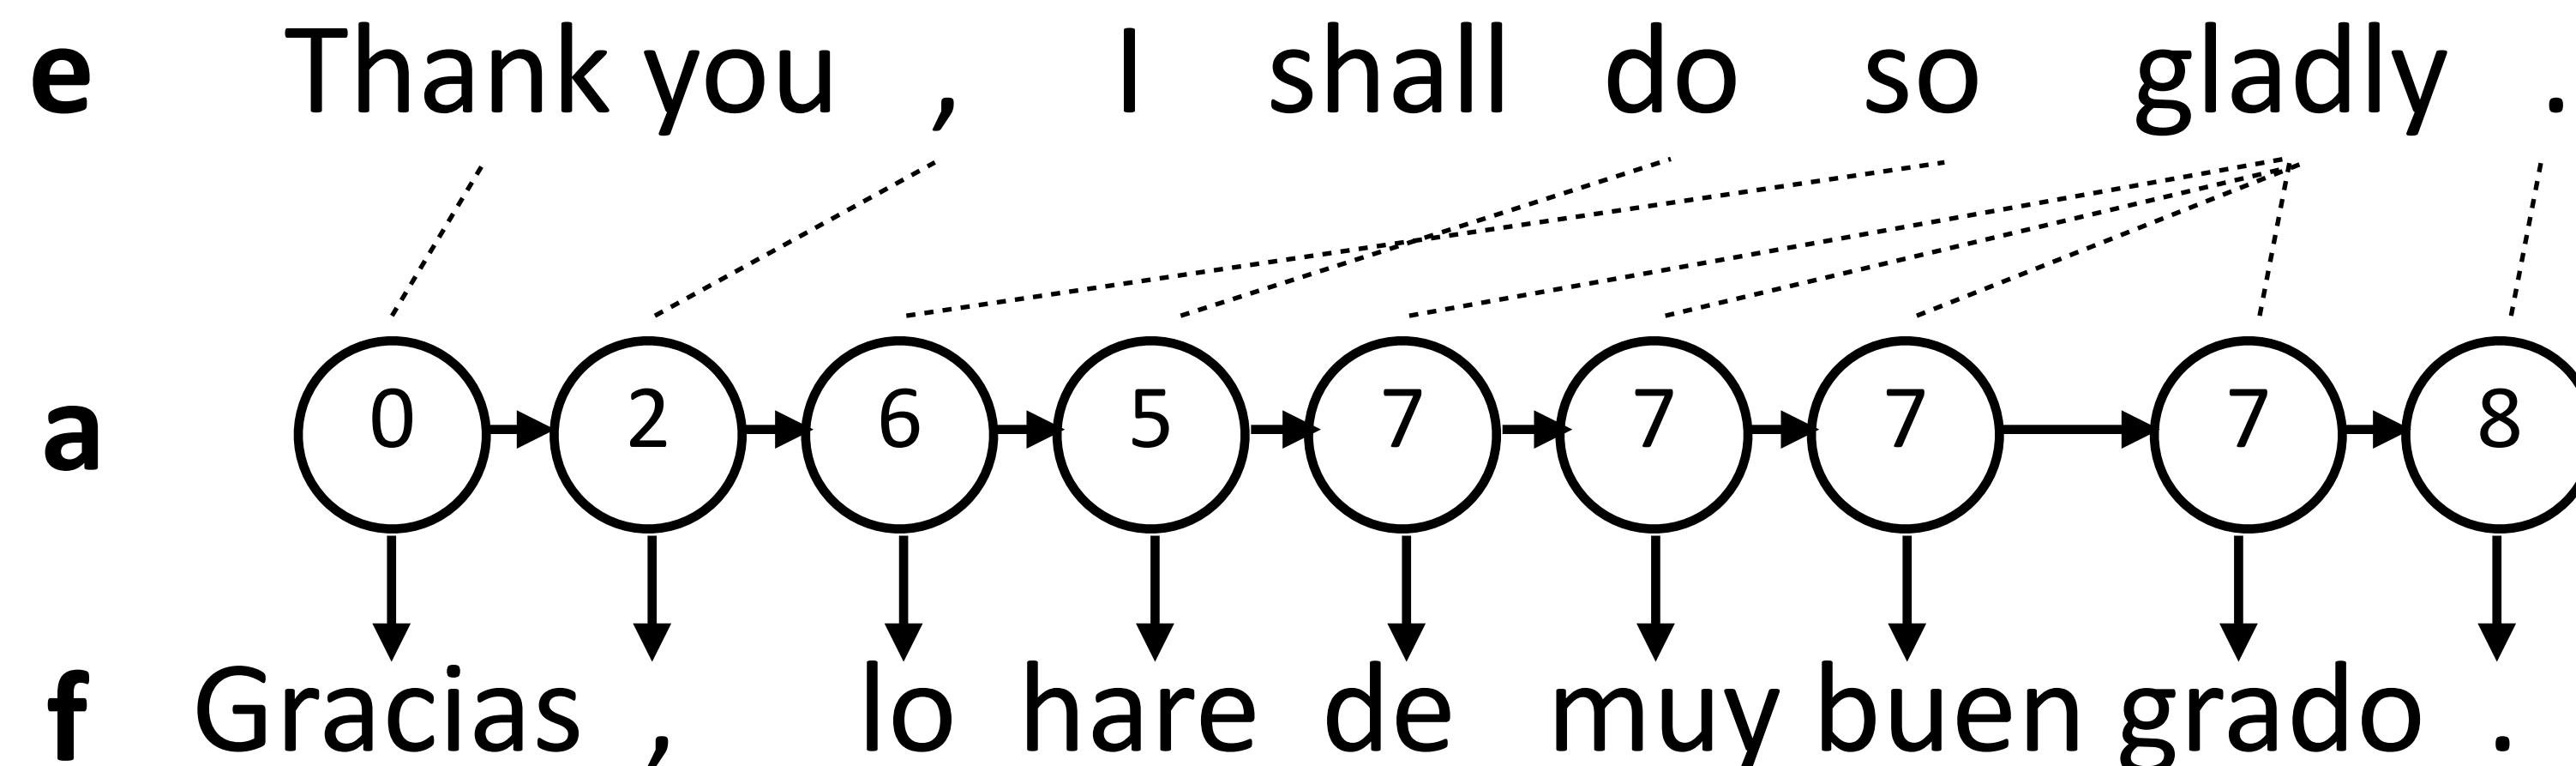
- Set $P(a)$ uniformly (no prior over good alignments)
- $P(f_i | e_{a_i})$: word translation probability table

Brown et al. (1993)

HMM for Alignment

- ▶ Sequential dependence between a's to capture monotonicity

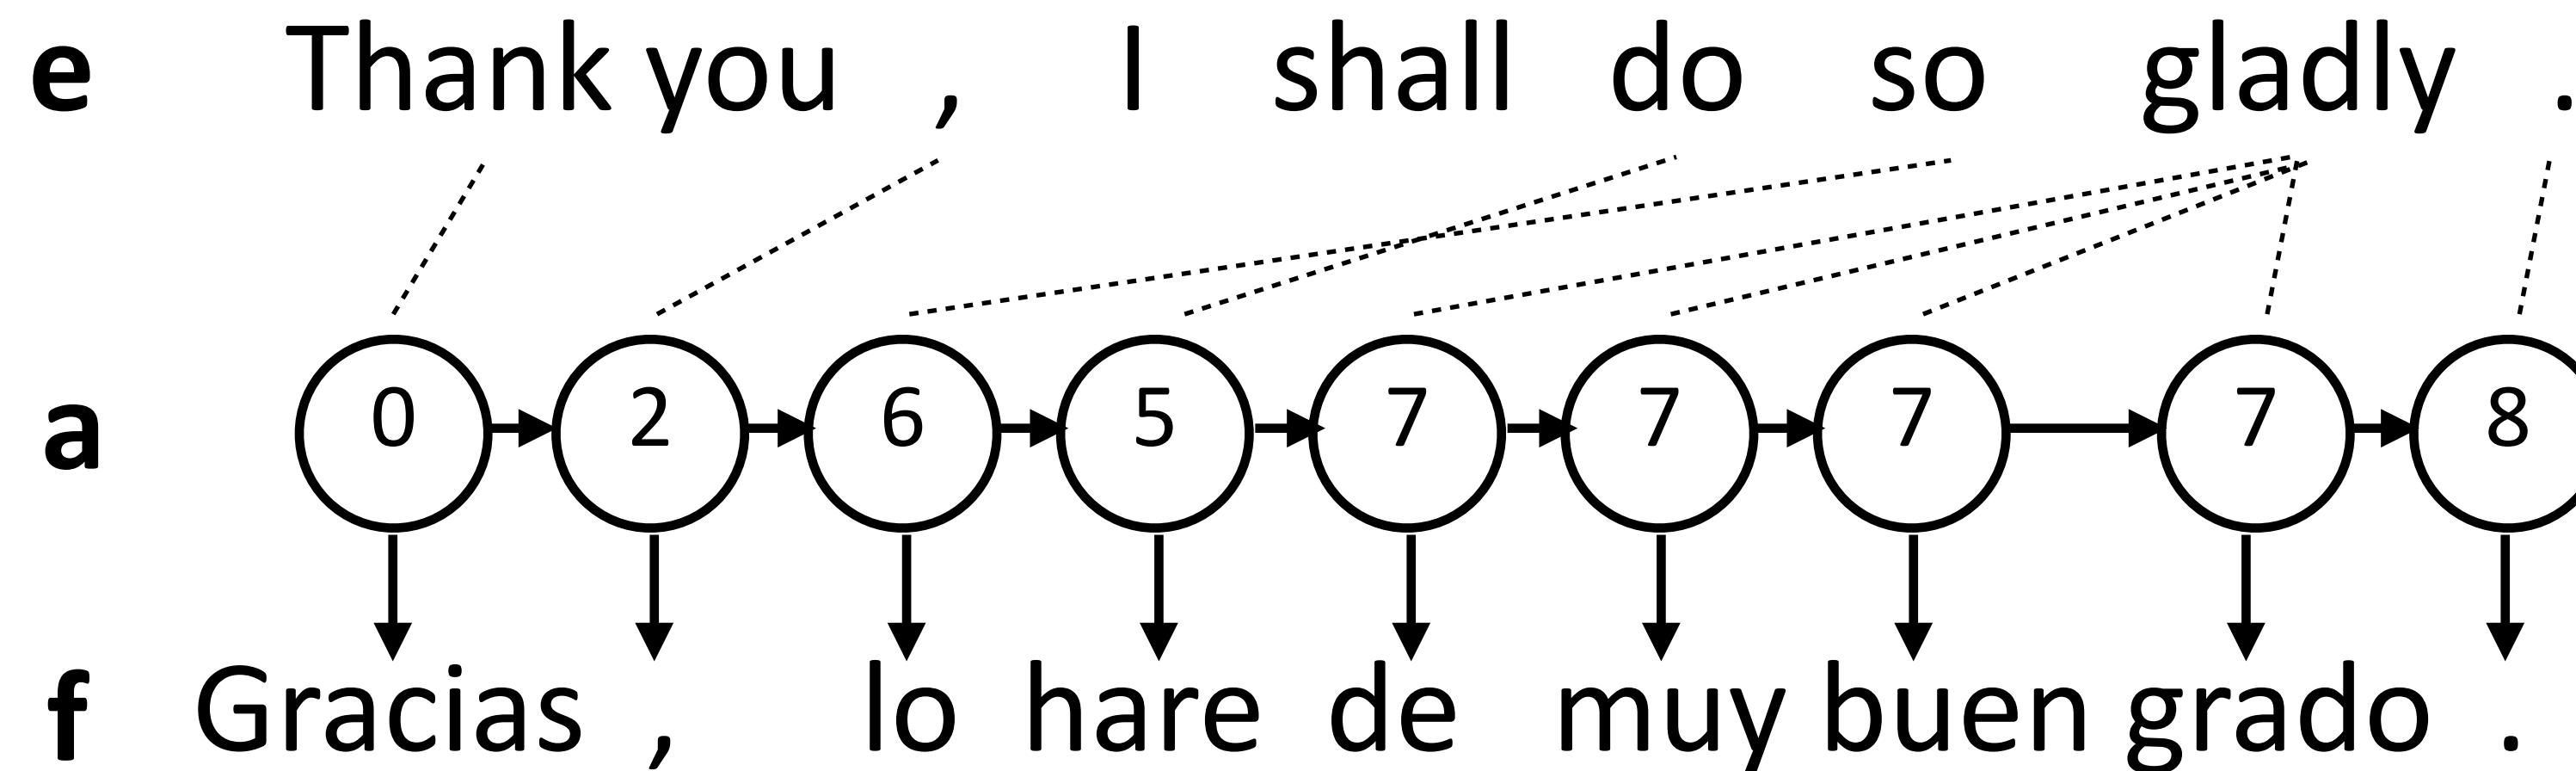
$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i | a_{i-1})$$



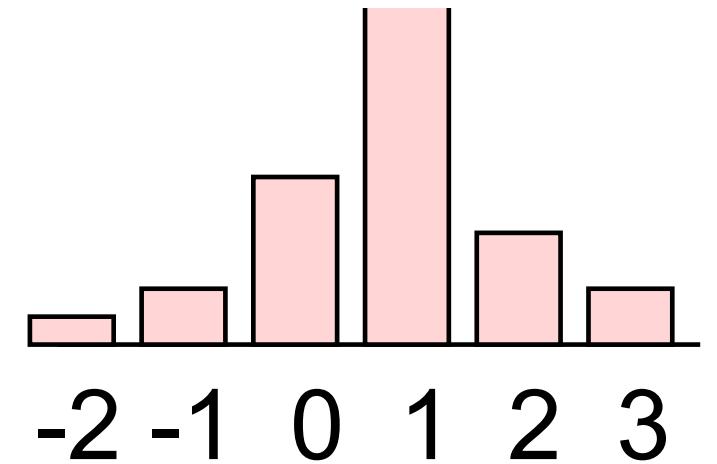
HMM for Alignment

- ▶ Sequential dependence between a's to capture monotonicity

$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i | a_{i-1})$$



- ▶ Alignment dist parameterized by jump size: $P(a_j - a_{j-1})$

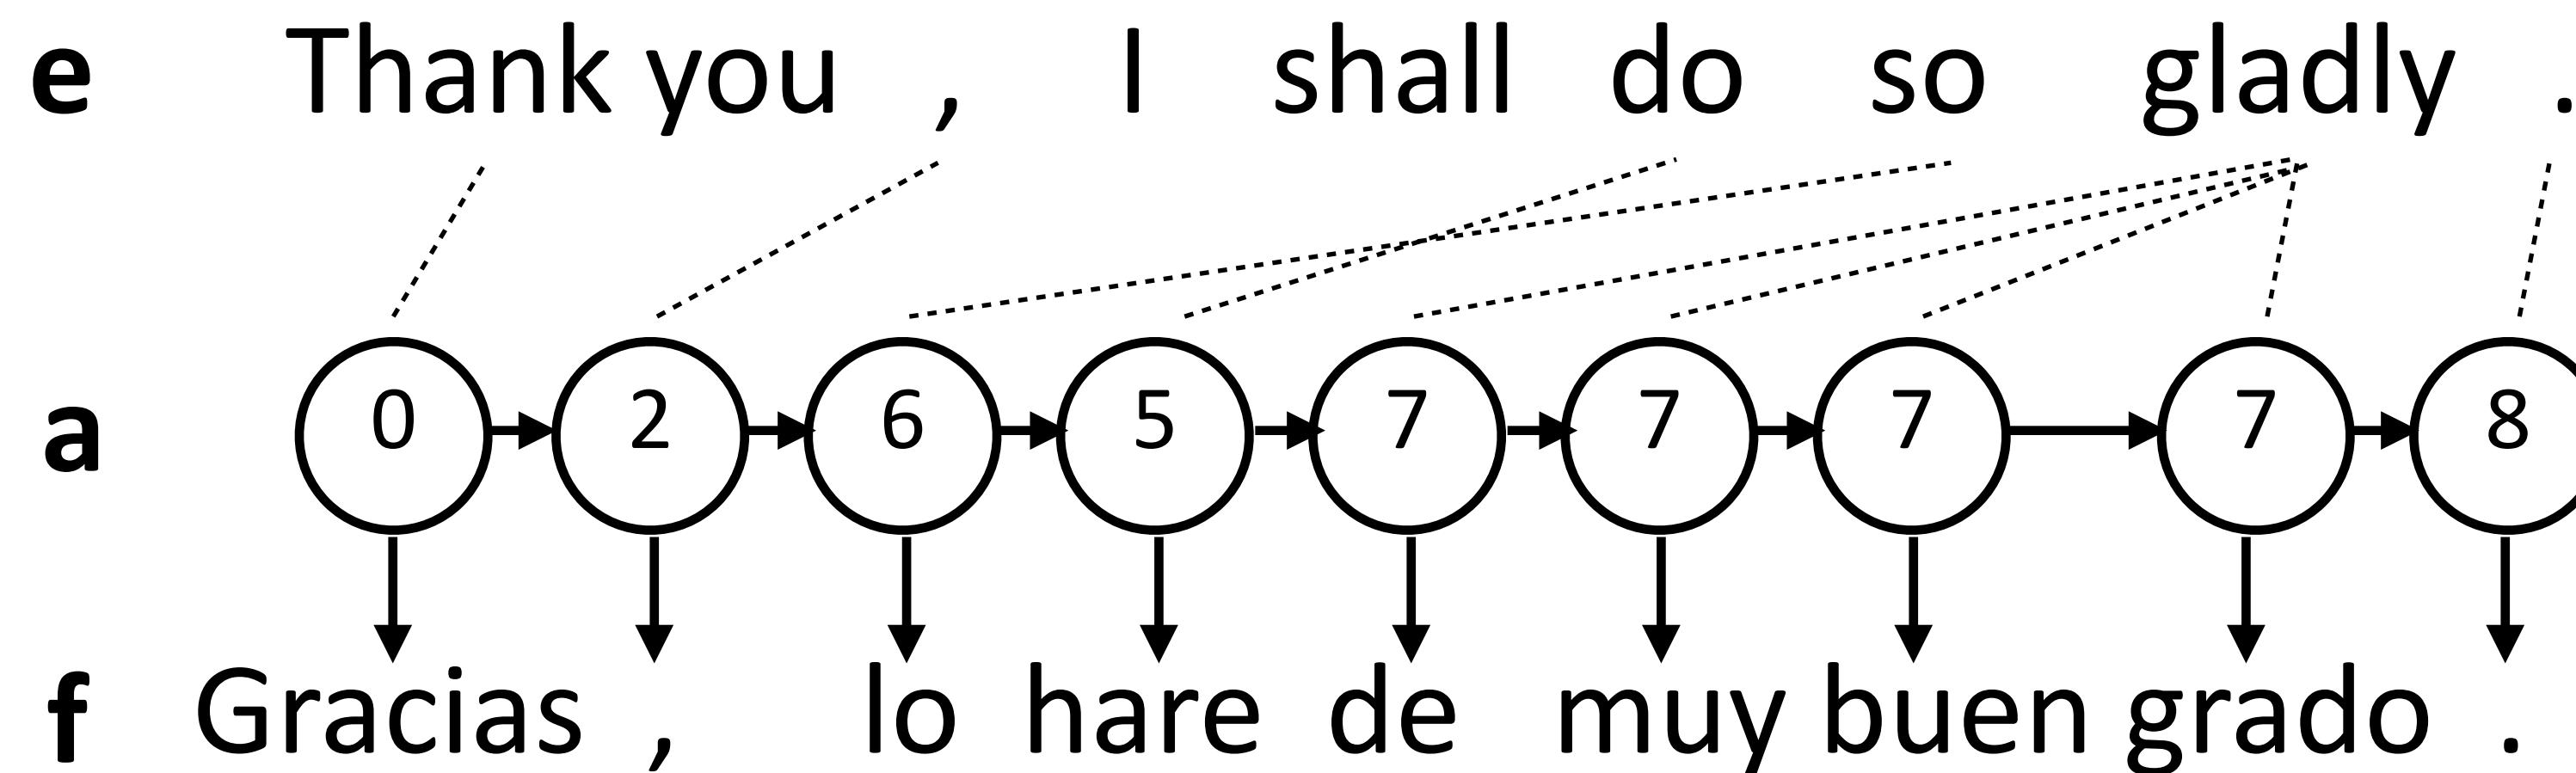


Brown et al. (1993)

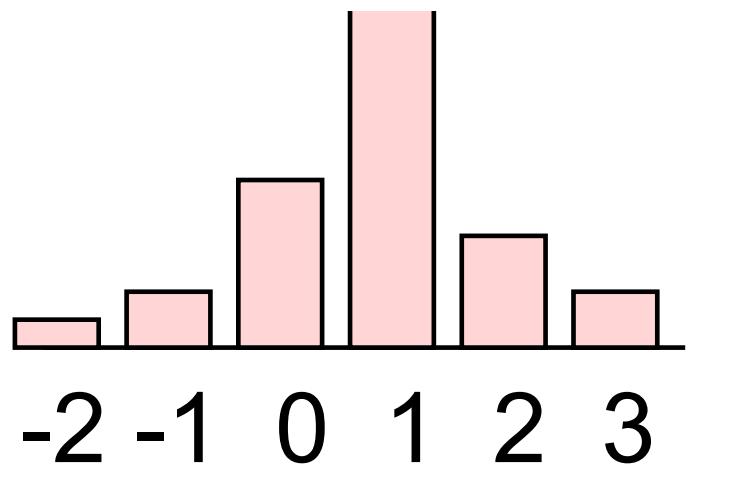
HMM for Alignment

- ▶ Sequential dependence between a's to capture monotonicity

$$P(\mathbf{f}, \mathbf{a} | \mathbf{e}) = \prod_{i=1}^n P(f_i | e_{a_i}) P(a_i | a_{i-1})$$



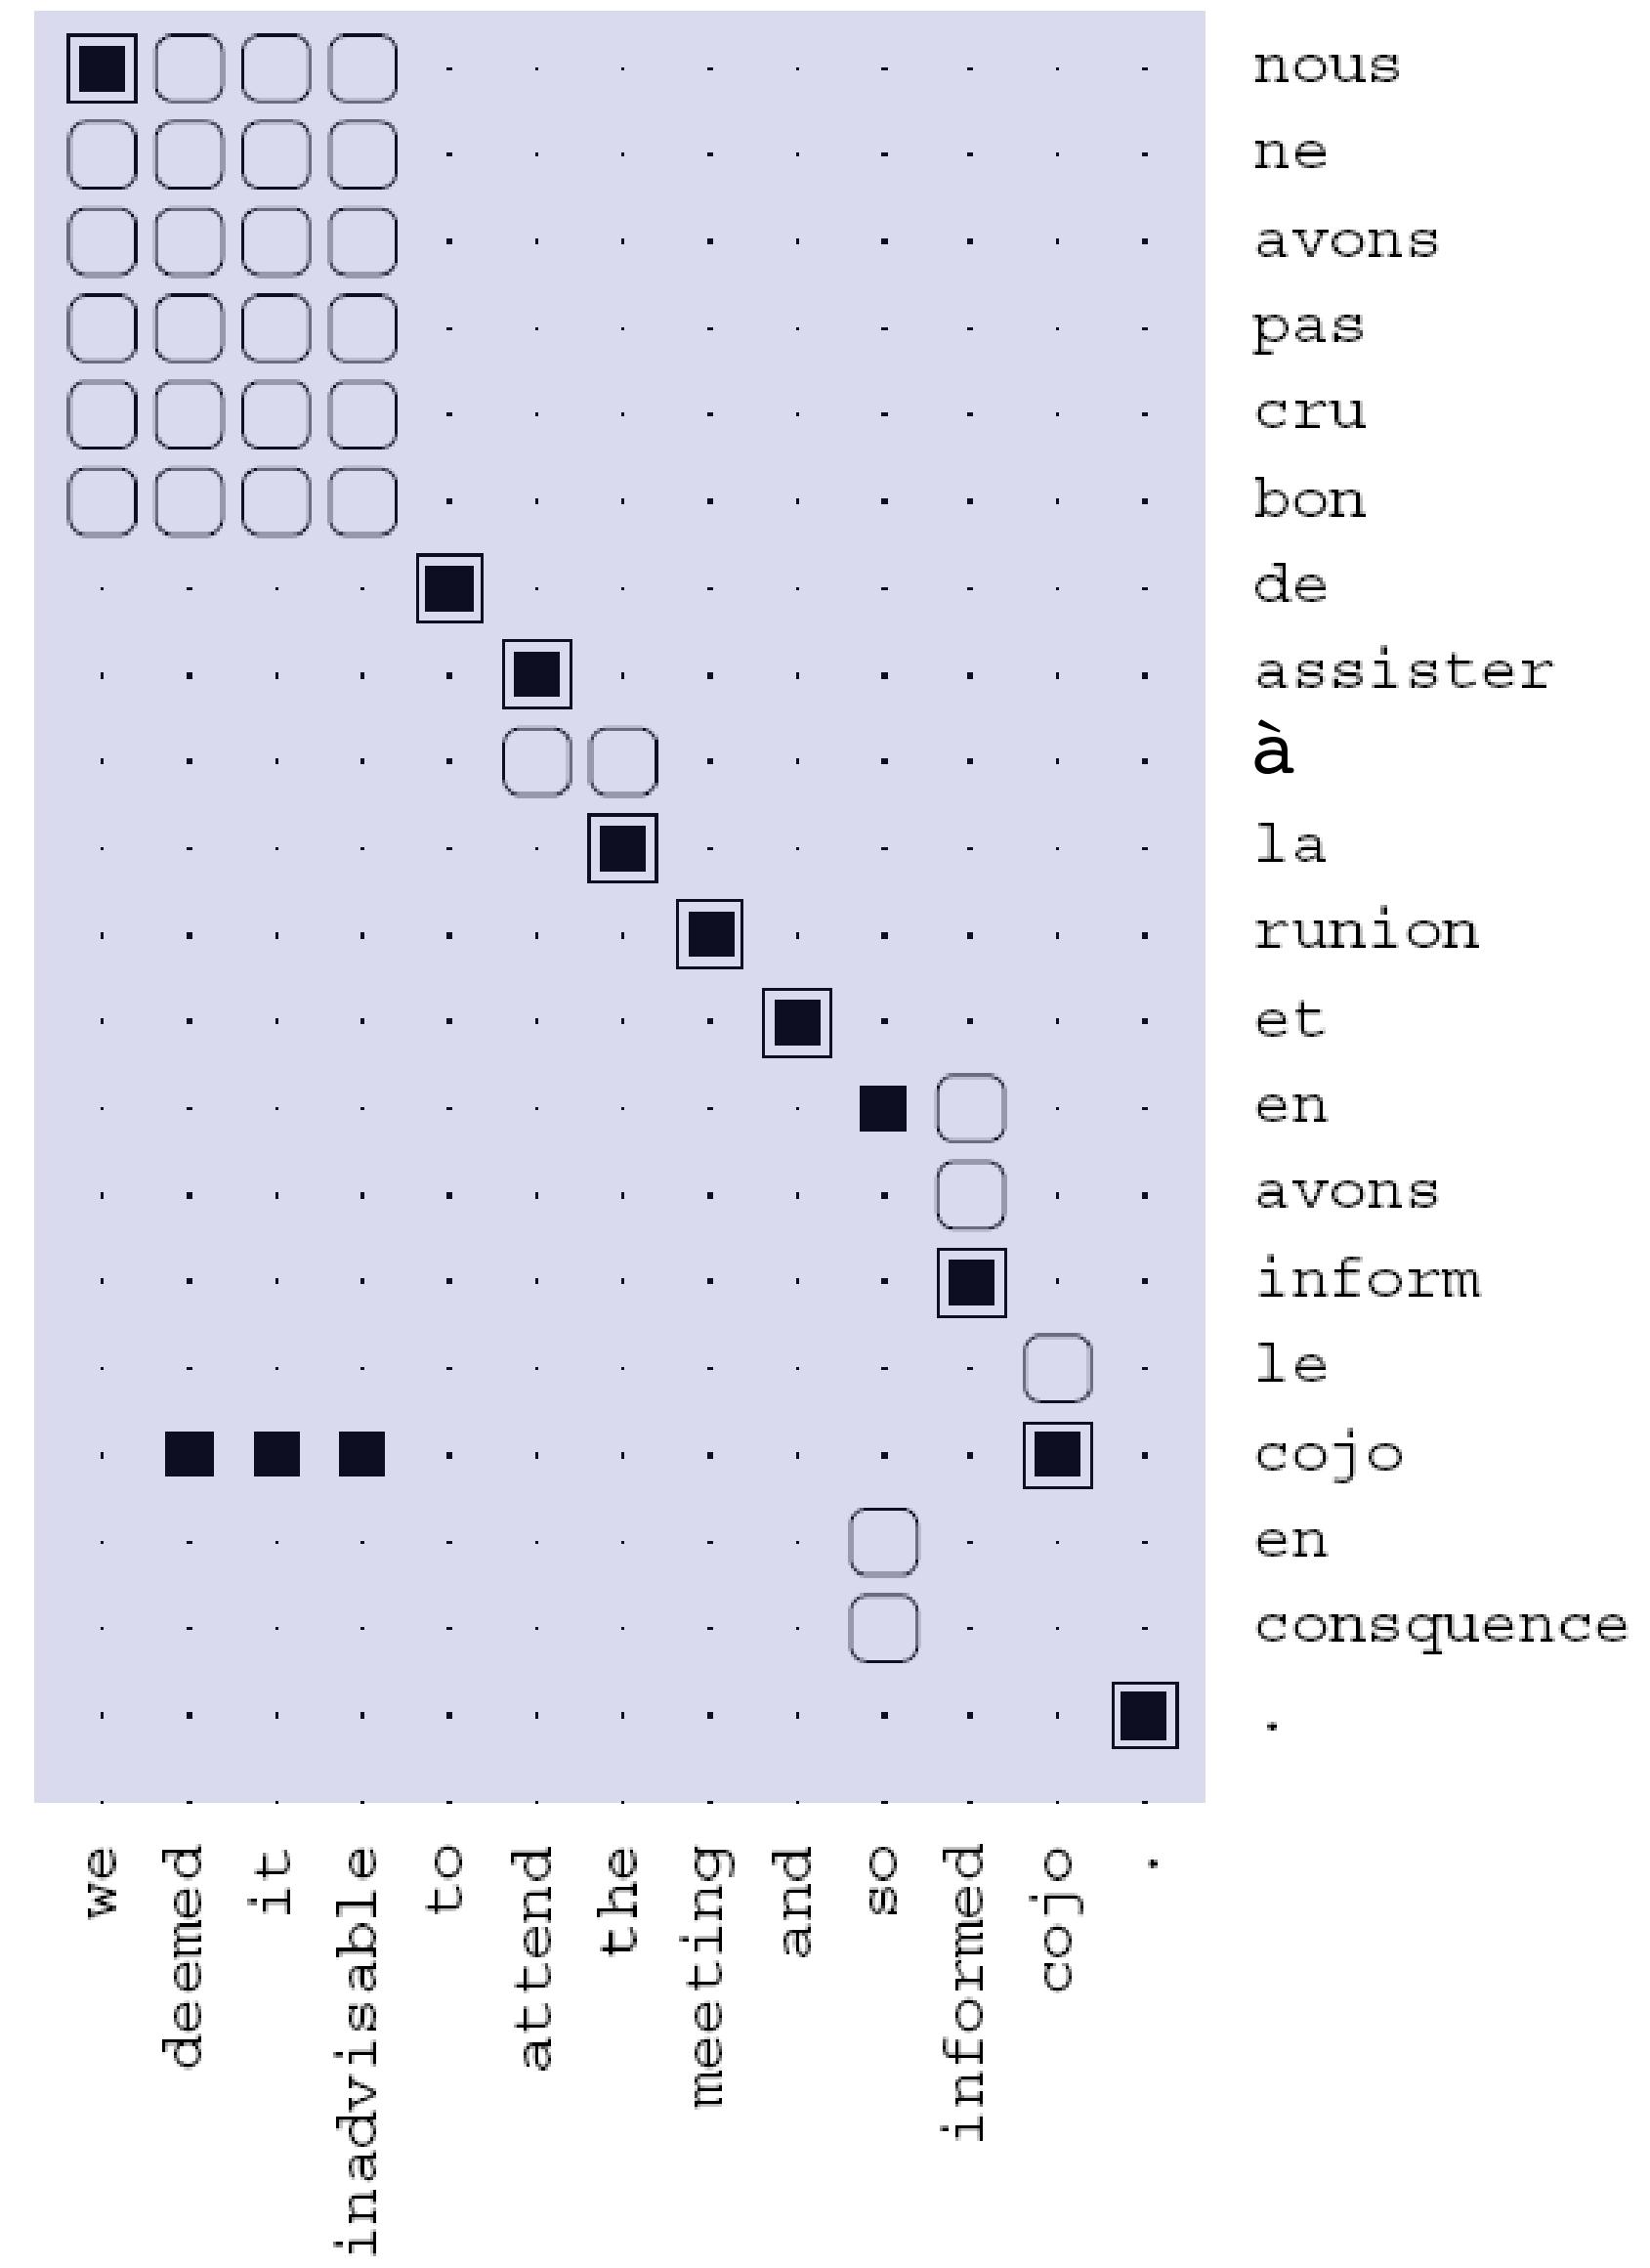
- ▶ Alignment dist parameterized by jump size: $P(a_j - a_{j-1})$
- ▶ $P(f_i | e_{a_i})$: same as before



Brown et al. (1993)

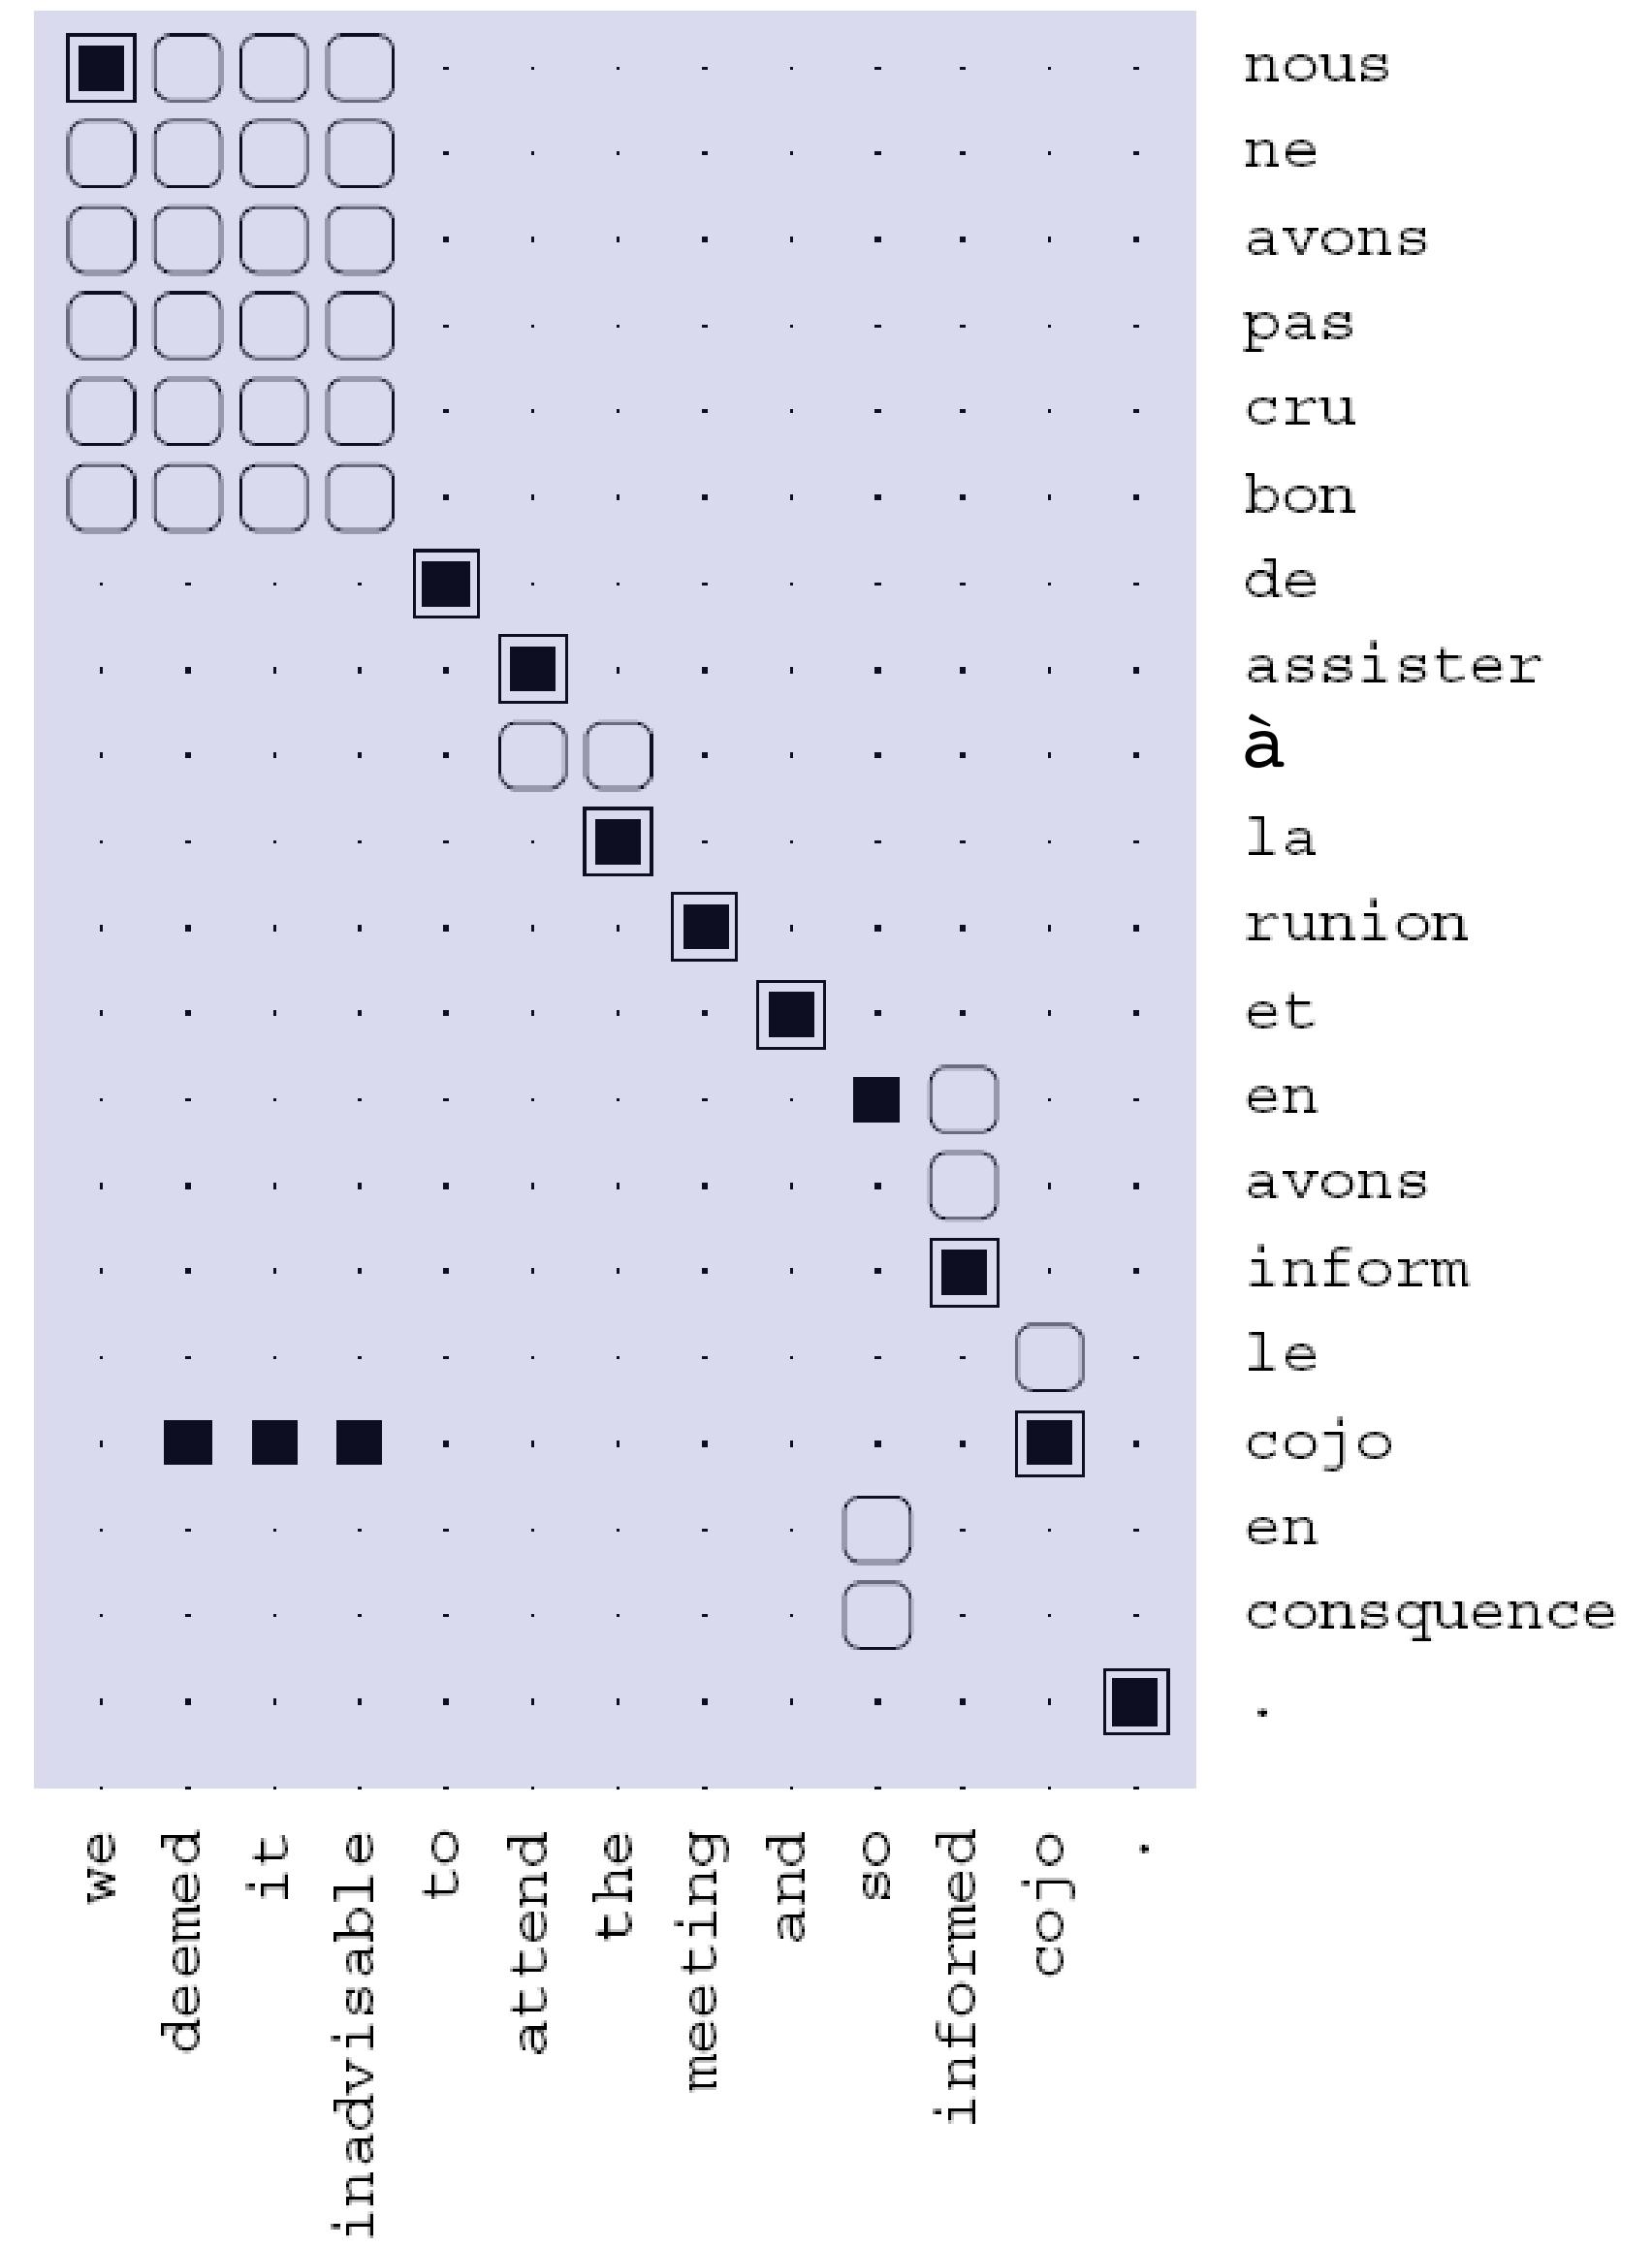
HMM Model

- # ► Which direction is this?



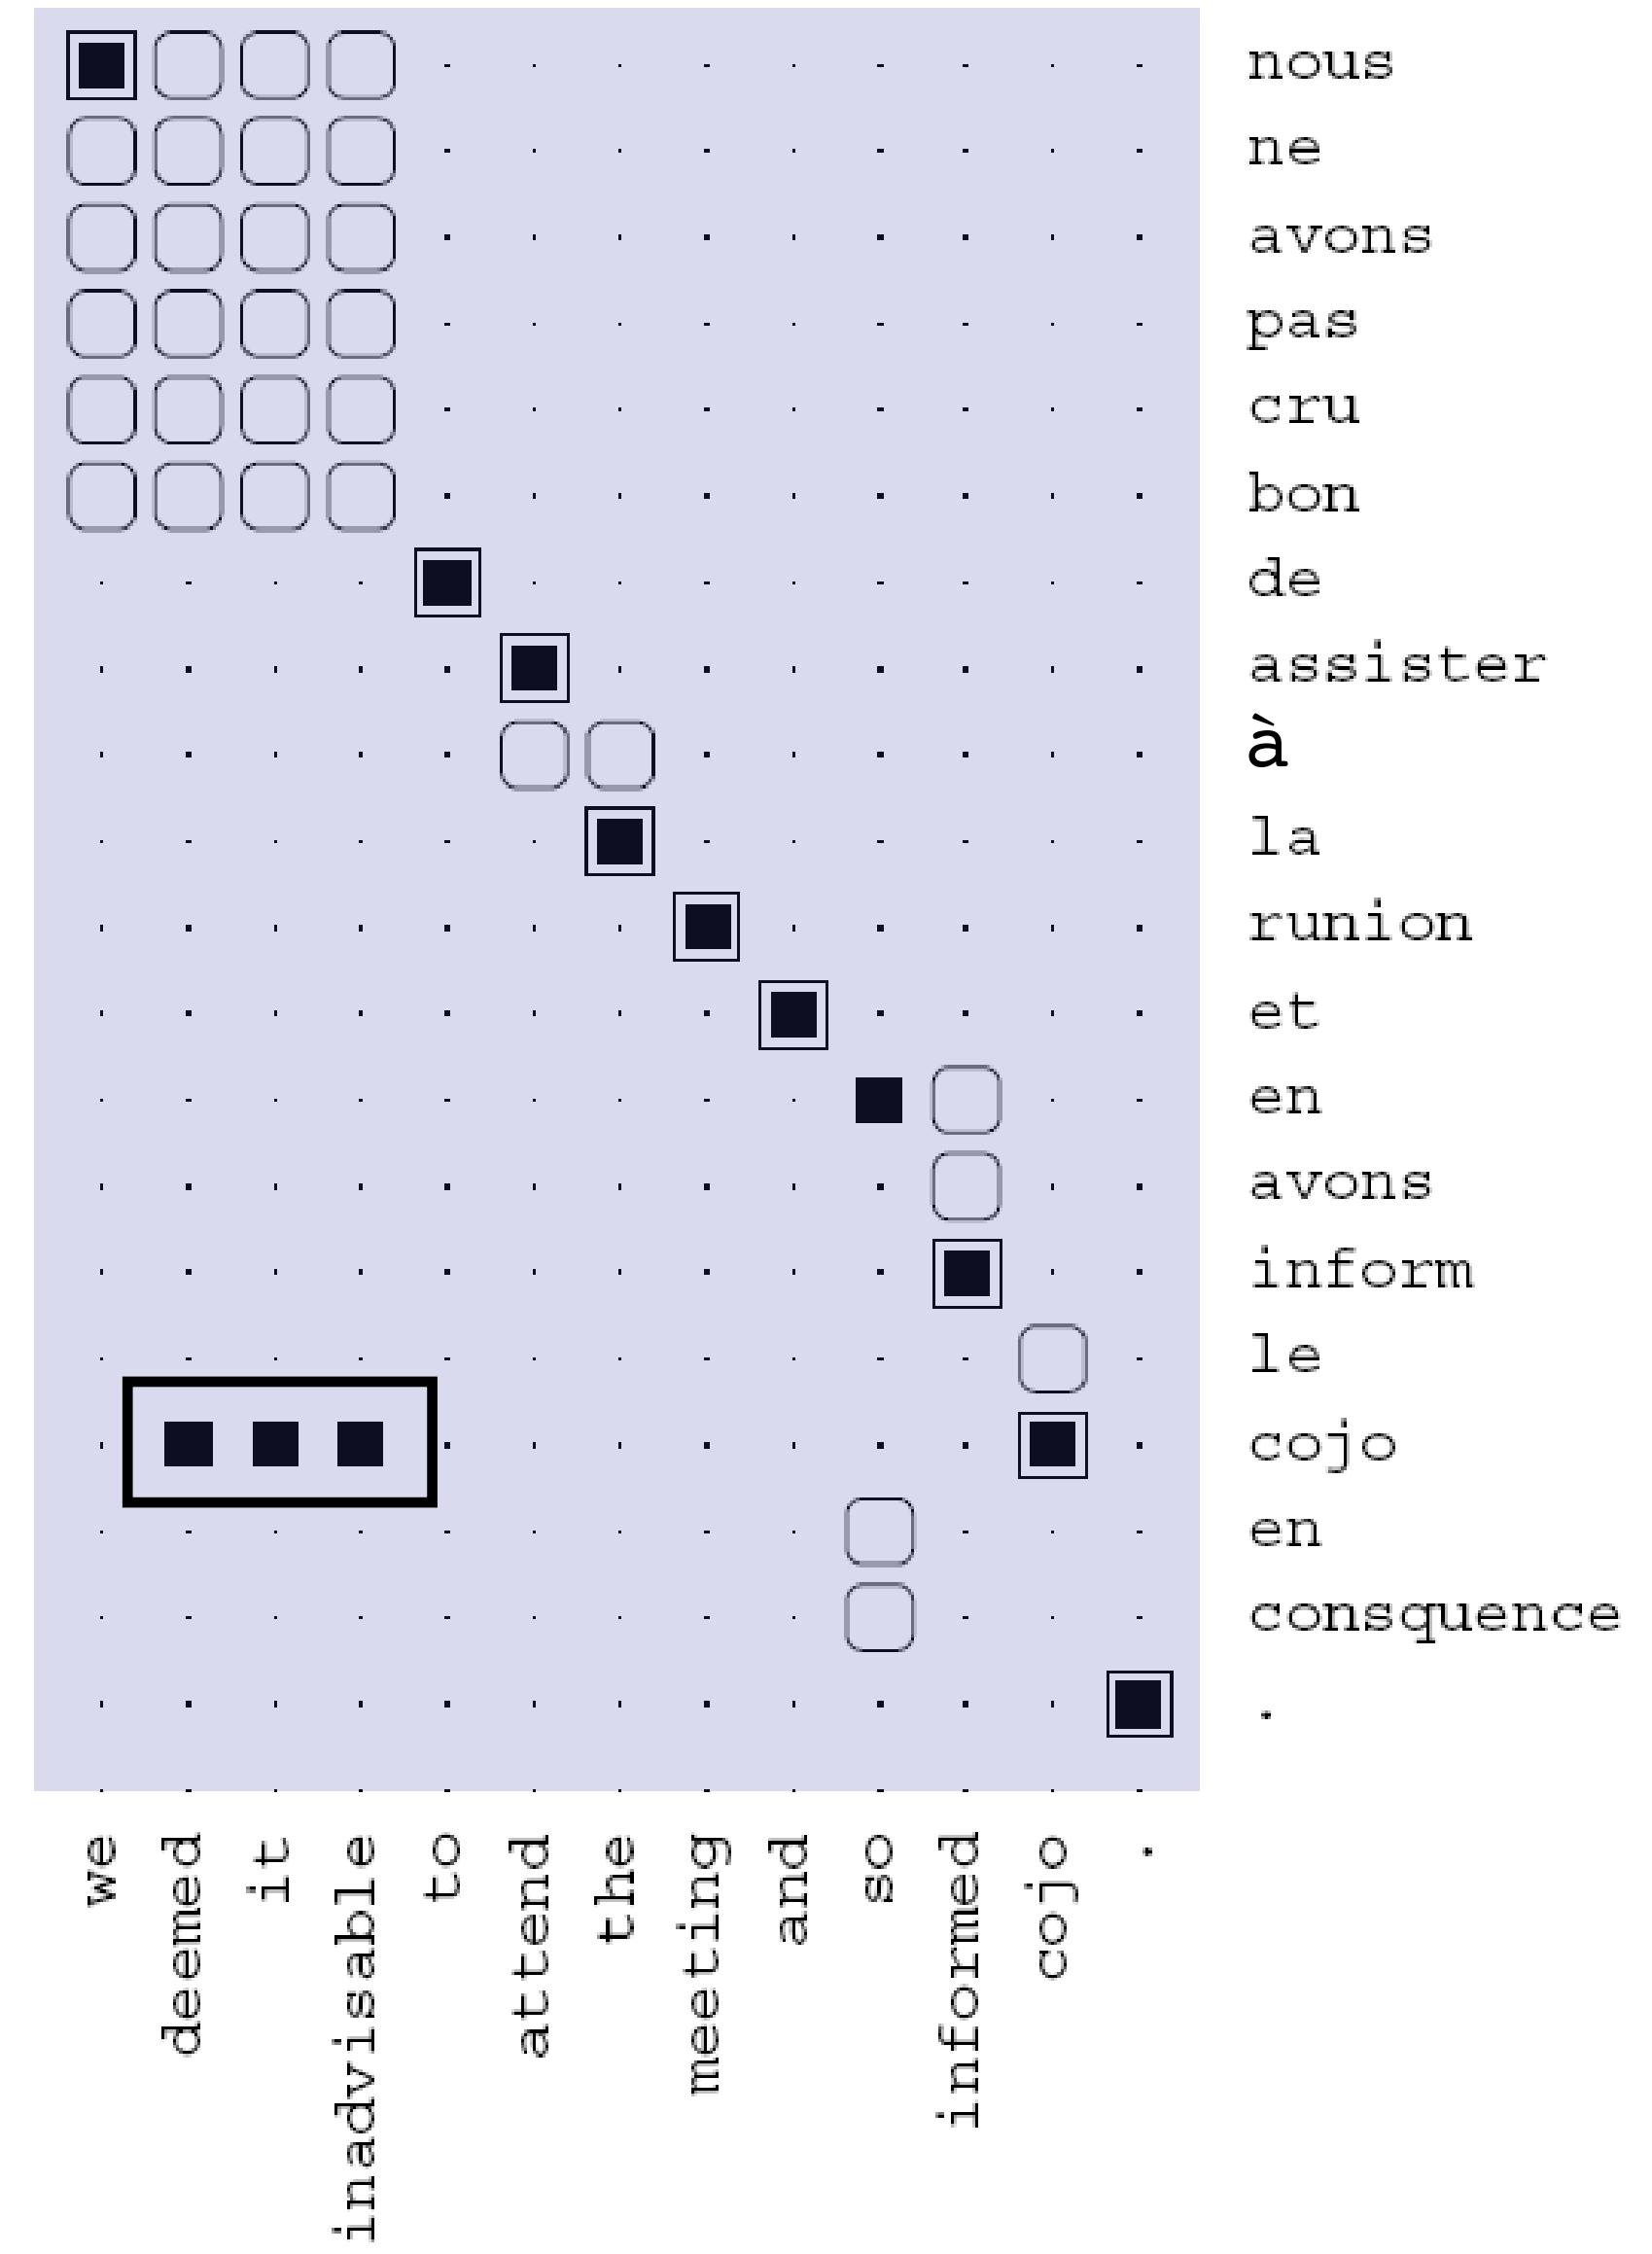
HMM Model

- ▶ Which direction is this?
- ▶ Alignments are generally monotonic (along diagonal)



HMM Model

- ▶ Which direction is this?
- ▶ Alignments are generally monotonic (along diagonal)
- ▶ Some mistakes, especially when you have rare words (*garbage collection*)



Evaluating Word Alignment

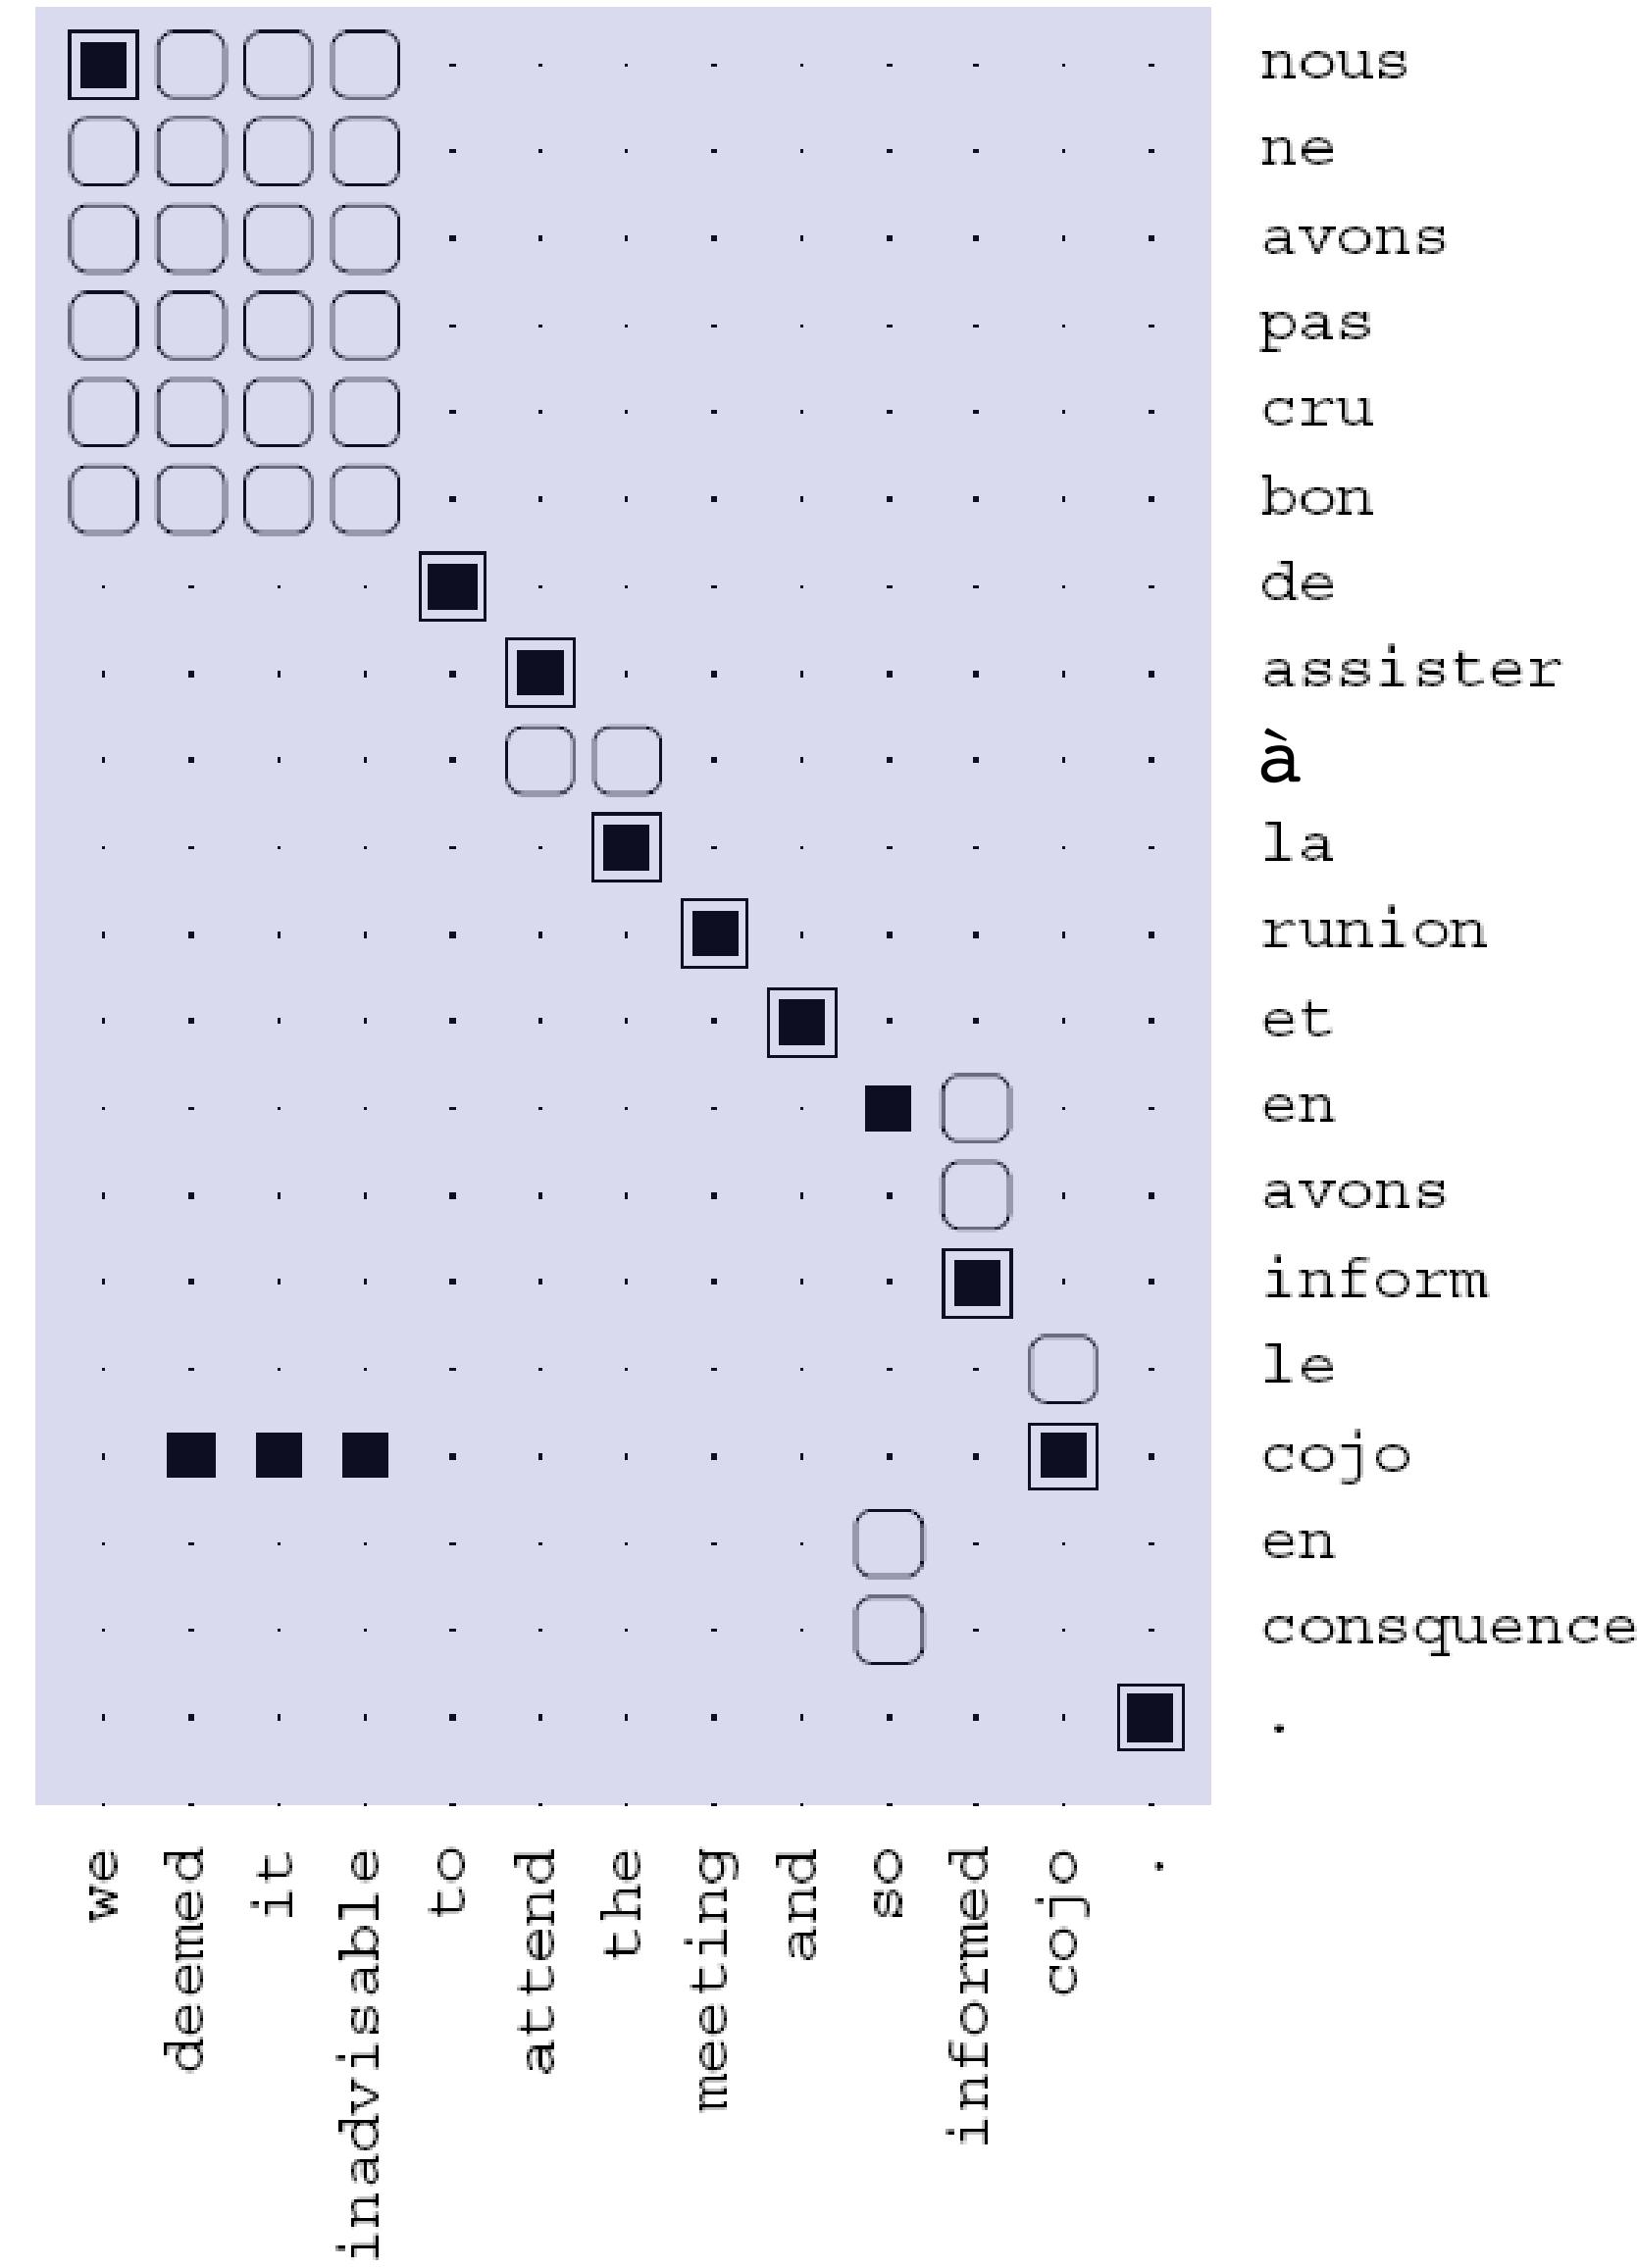
- ▶ “Alignment error rate”: use labeled alignments on small corpus

Model	AER
Model 1 INT	19.5
HMM E→F	11.4
HMM F→E	10.8
HMM AND	7.1
HMM INT	4.7
GIZA M4 AND	6.9

- ▶ Run Model 1 in both directions and intersect “intelligently”
- ▶ Run HMM model in both directions and intersect “intelligently”

Phrase Extraction

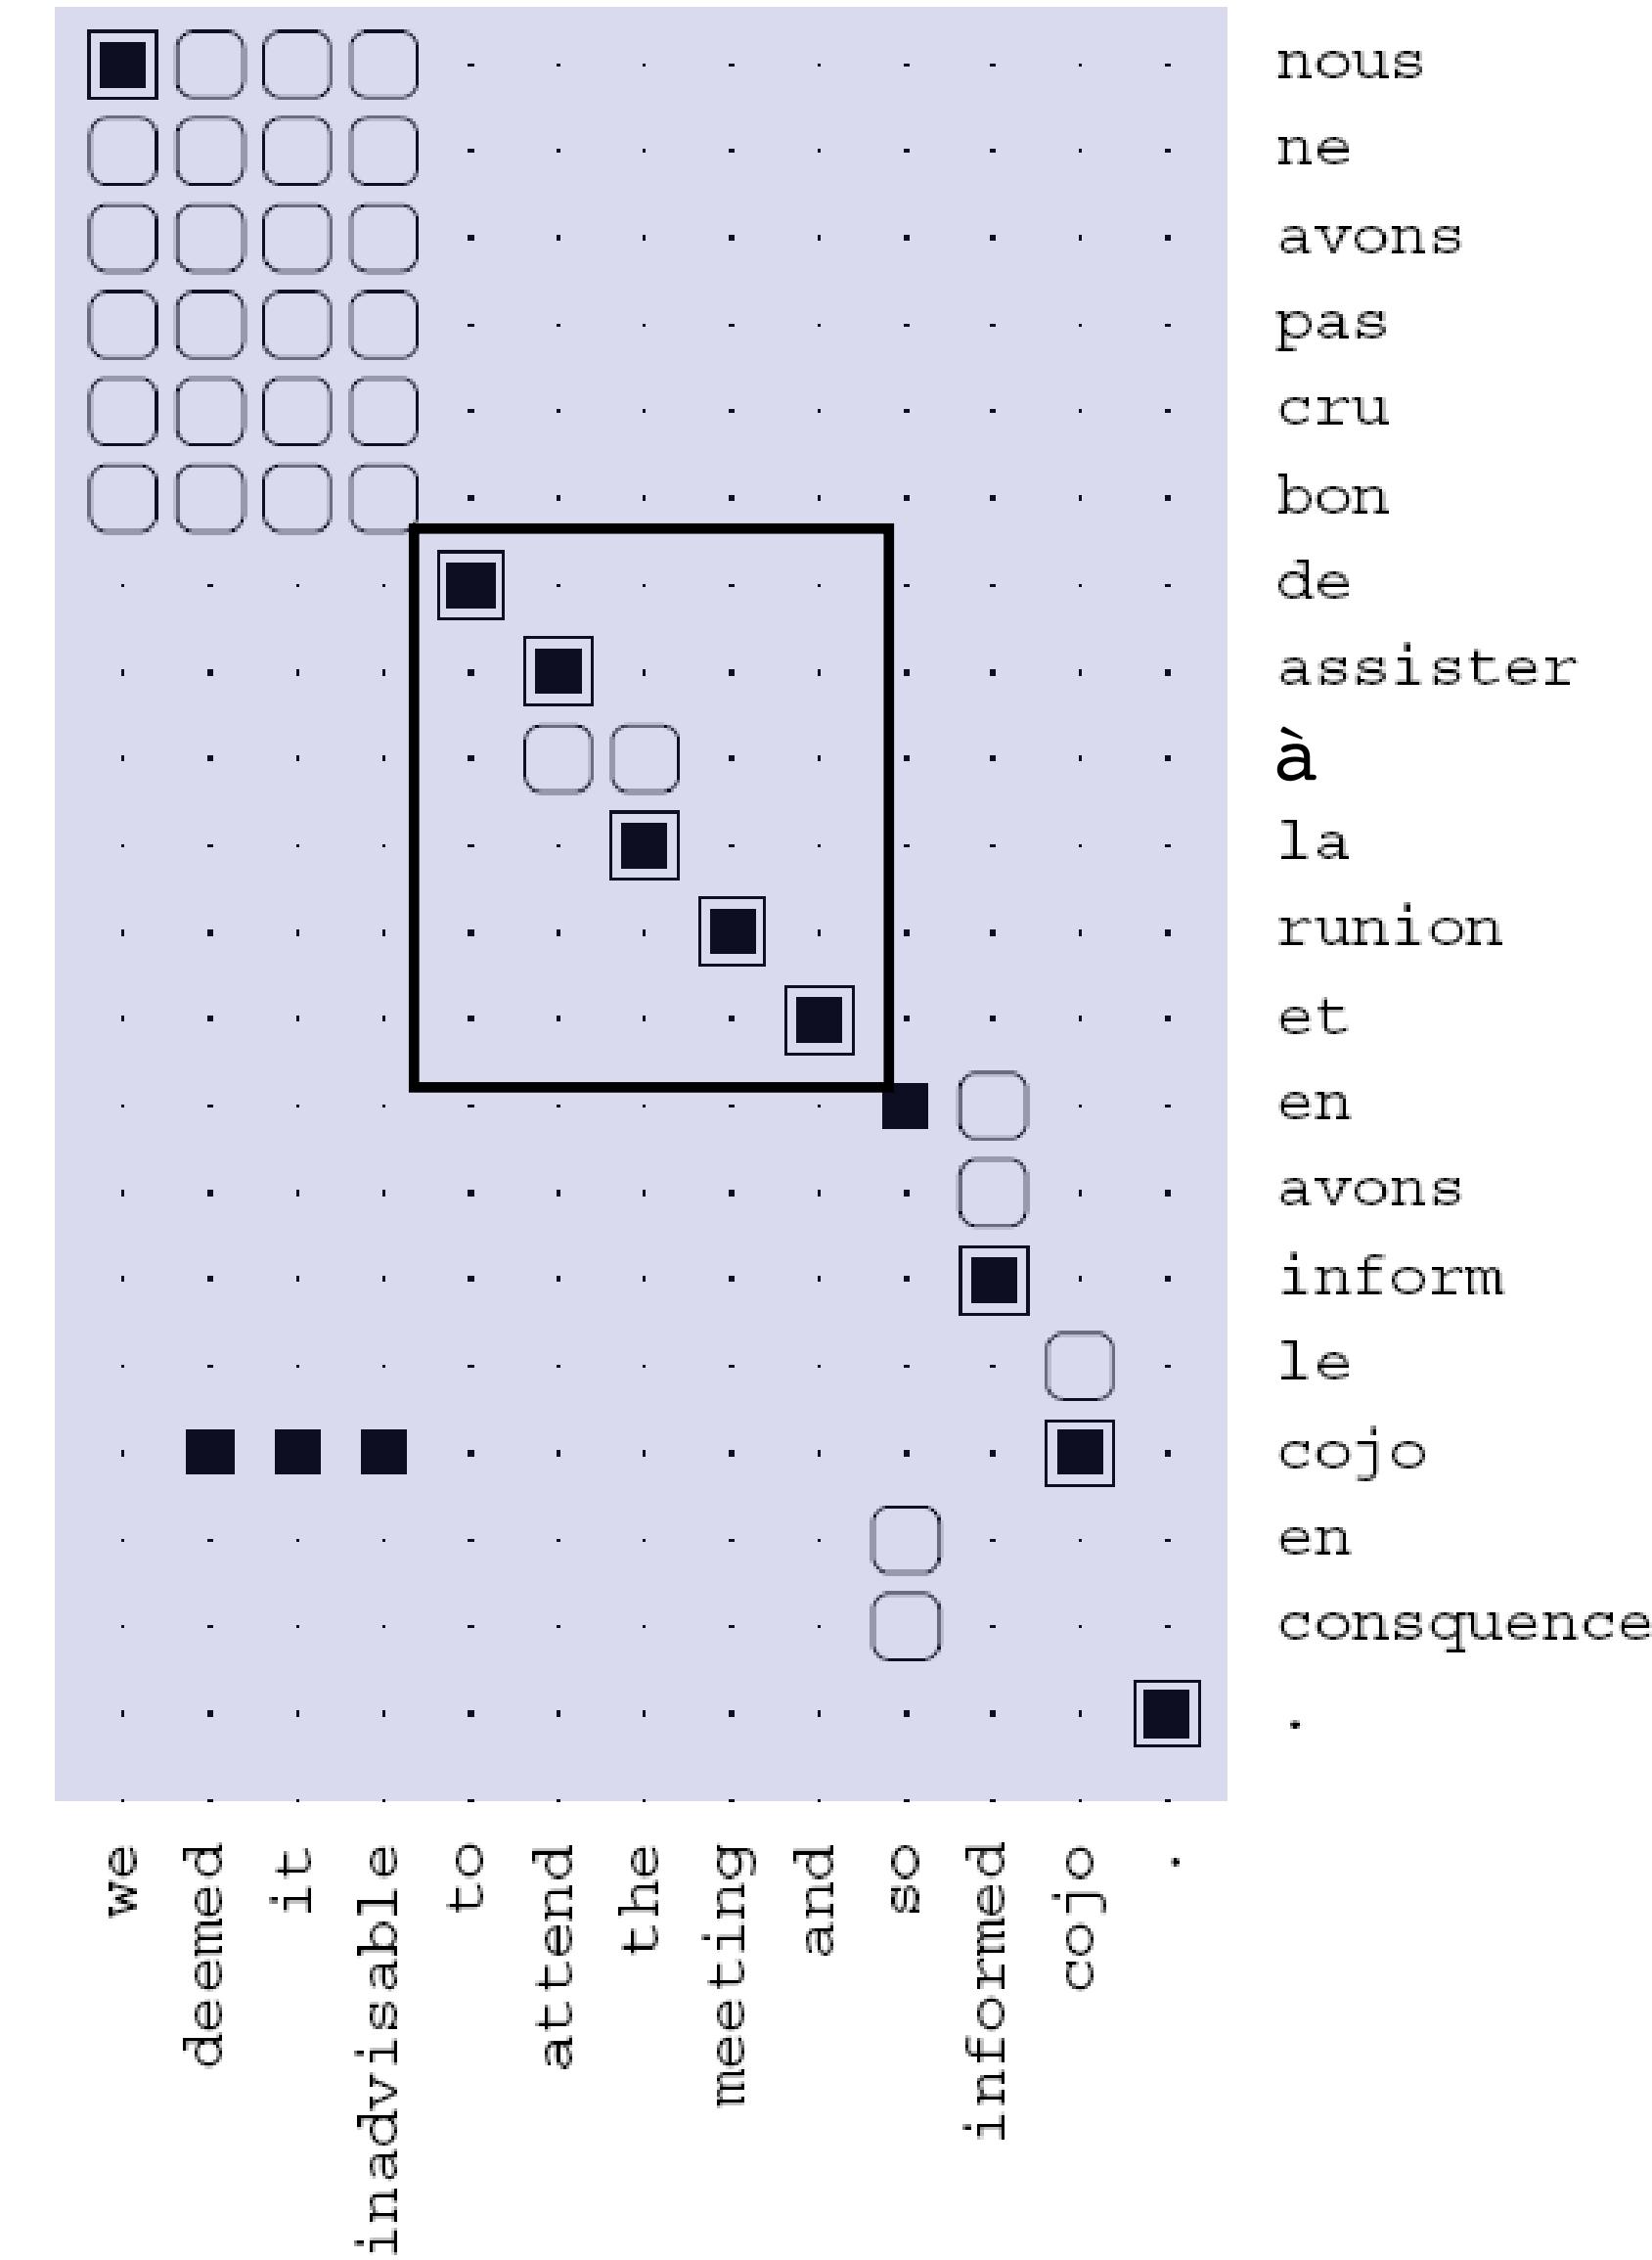
- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words



Phrase Extraction

- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

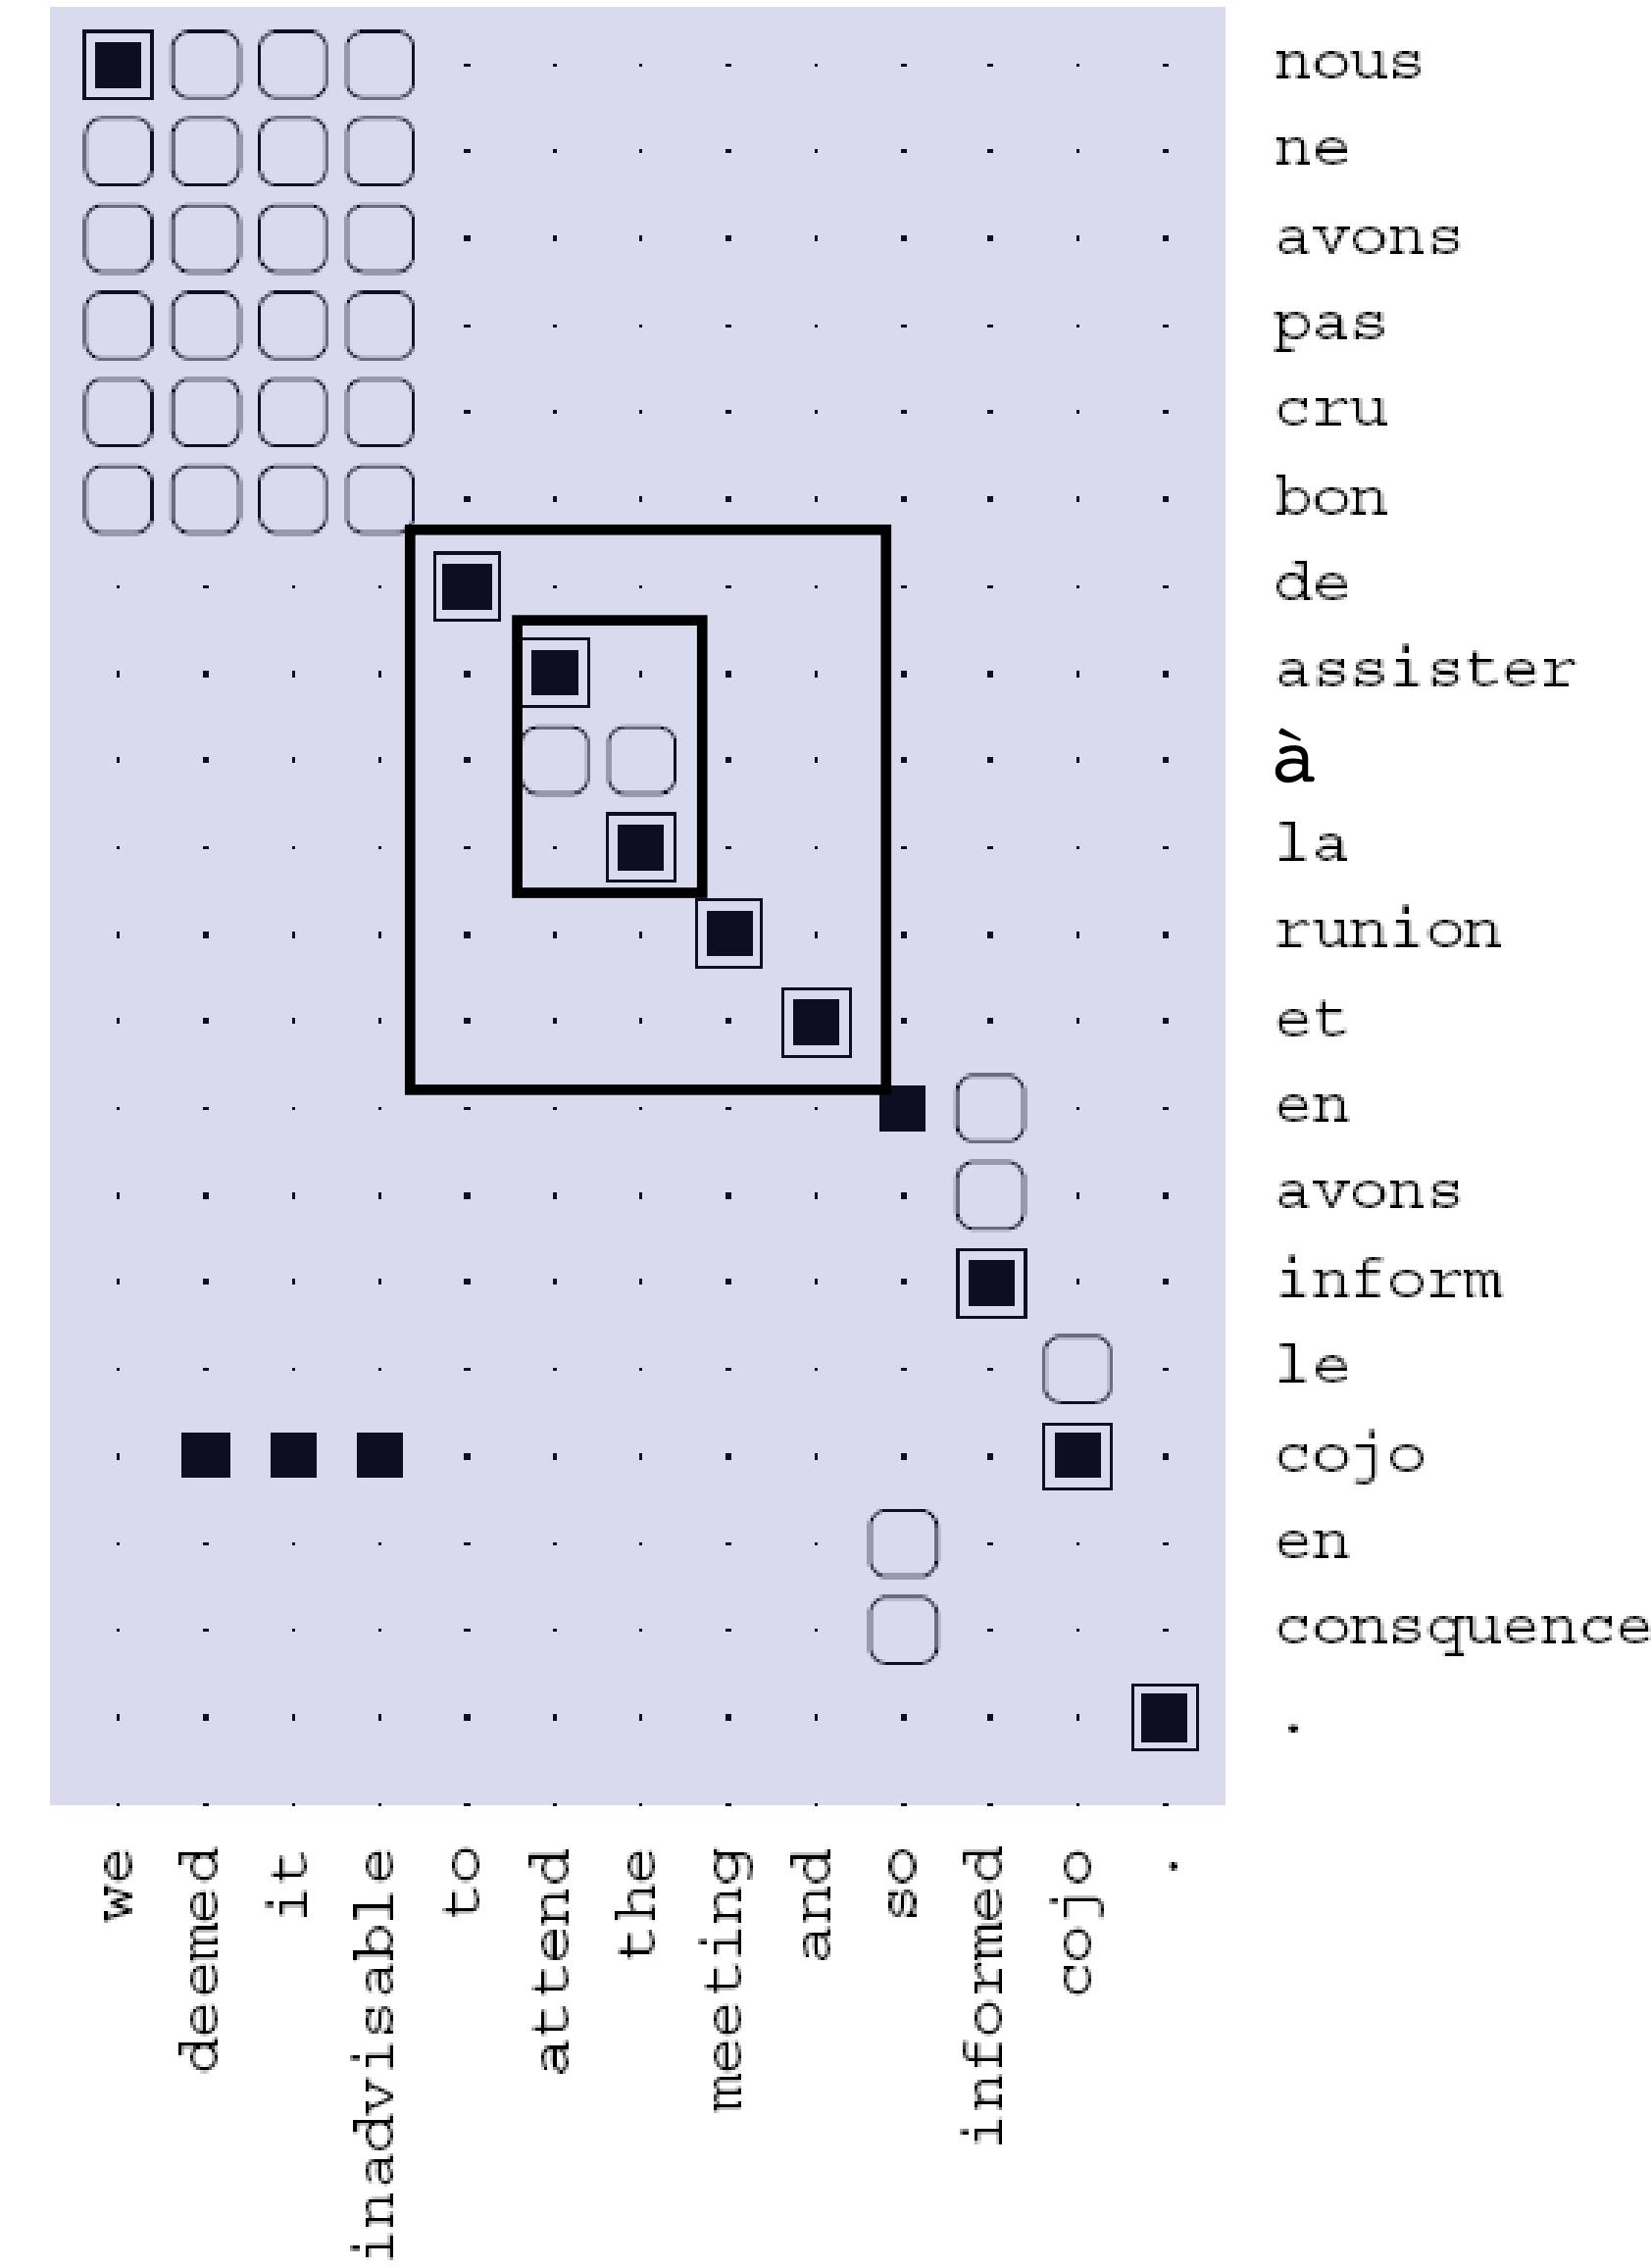
d'assister à la reunion et ||| to attend the meeting and



Phrase Extraction

- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

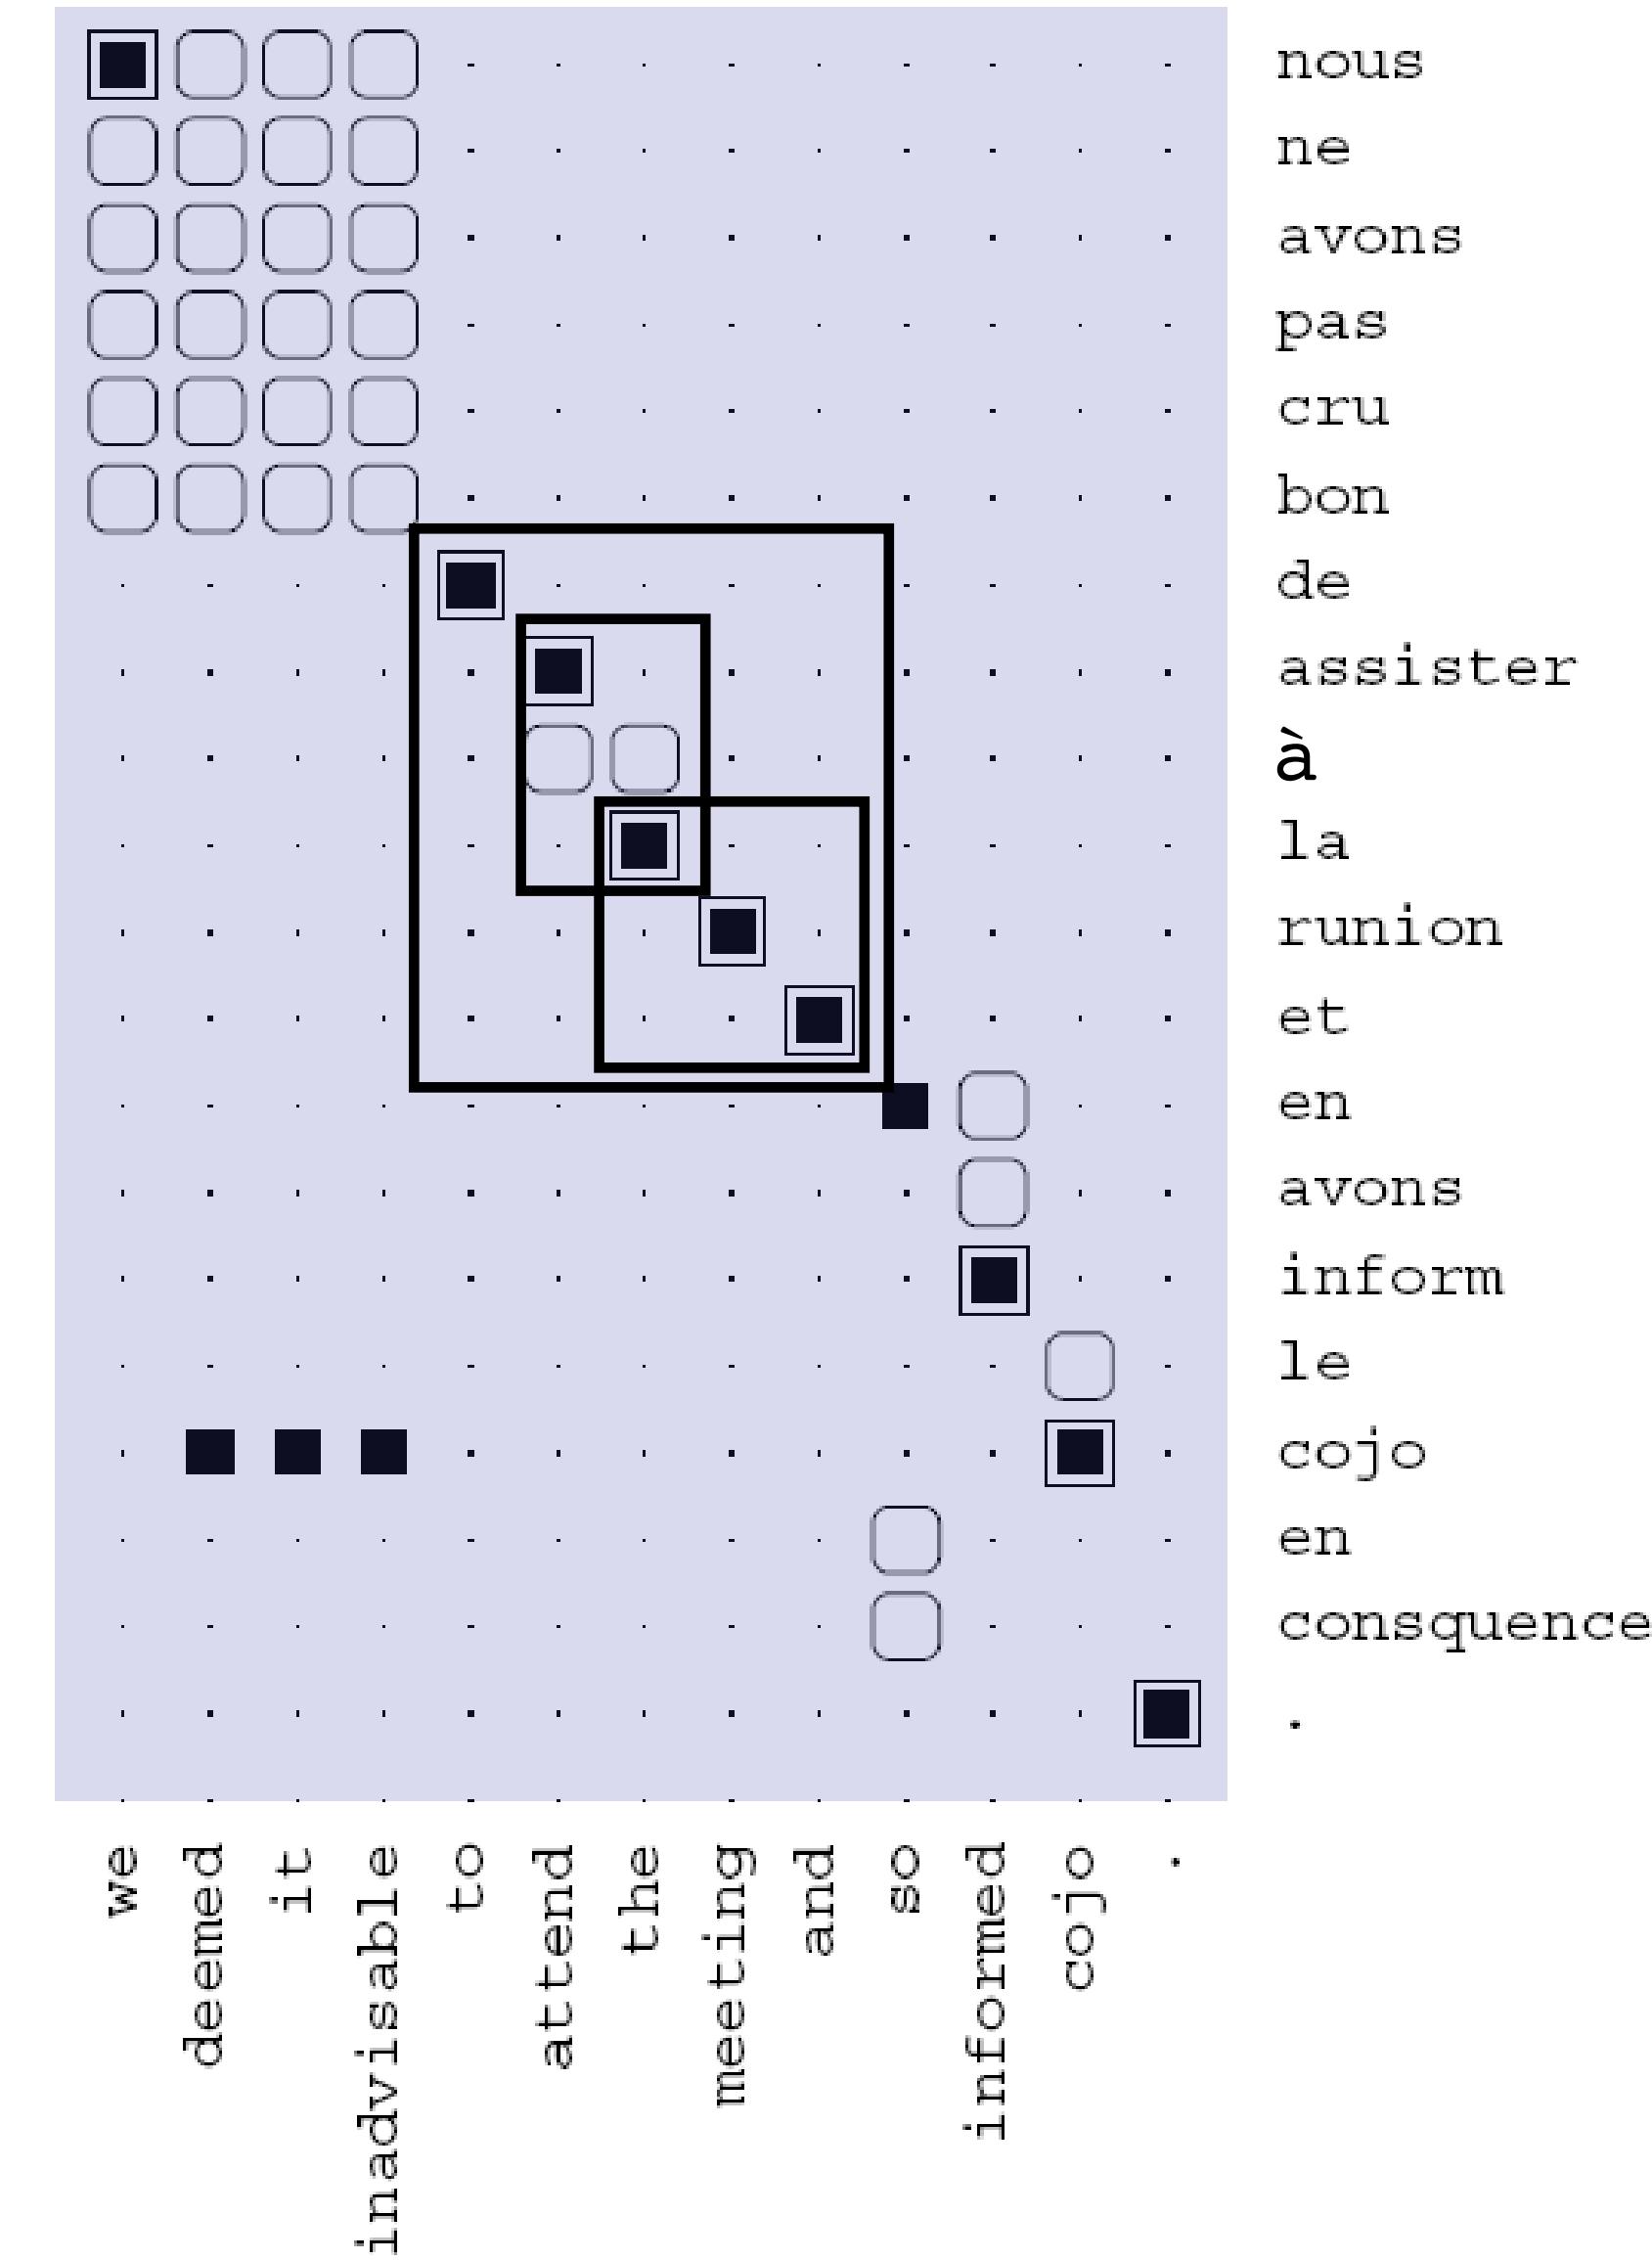
d'assister à la reunion et ||| to attend the meeting and
assister à la reunion ||| attend the meeting



Phrase Extraction

- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

d'assister à la reunion et ||| to attend the meeting and
assister à la reunion ||| attend the meeting
la reunion and ||| the meeting and



Phrase Extraction

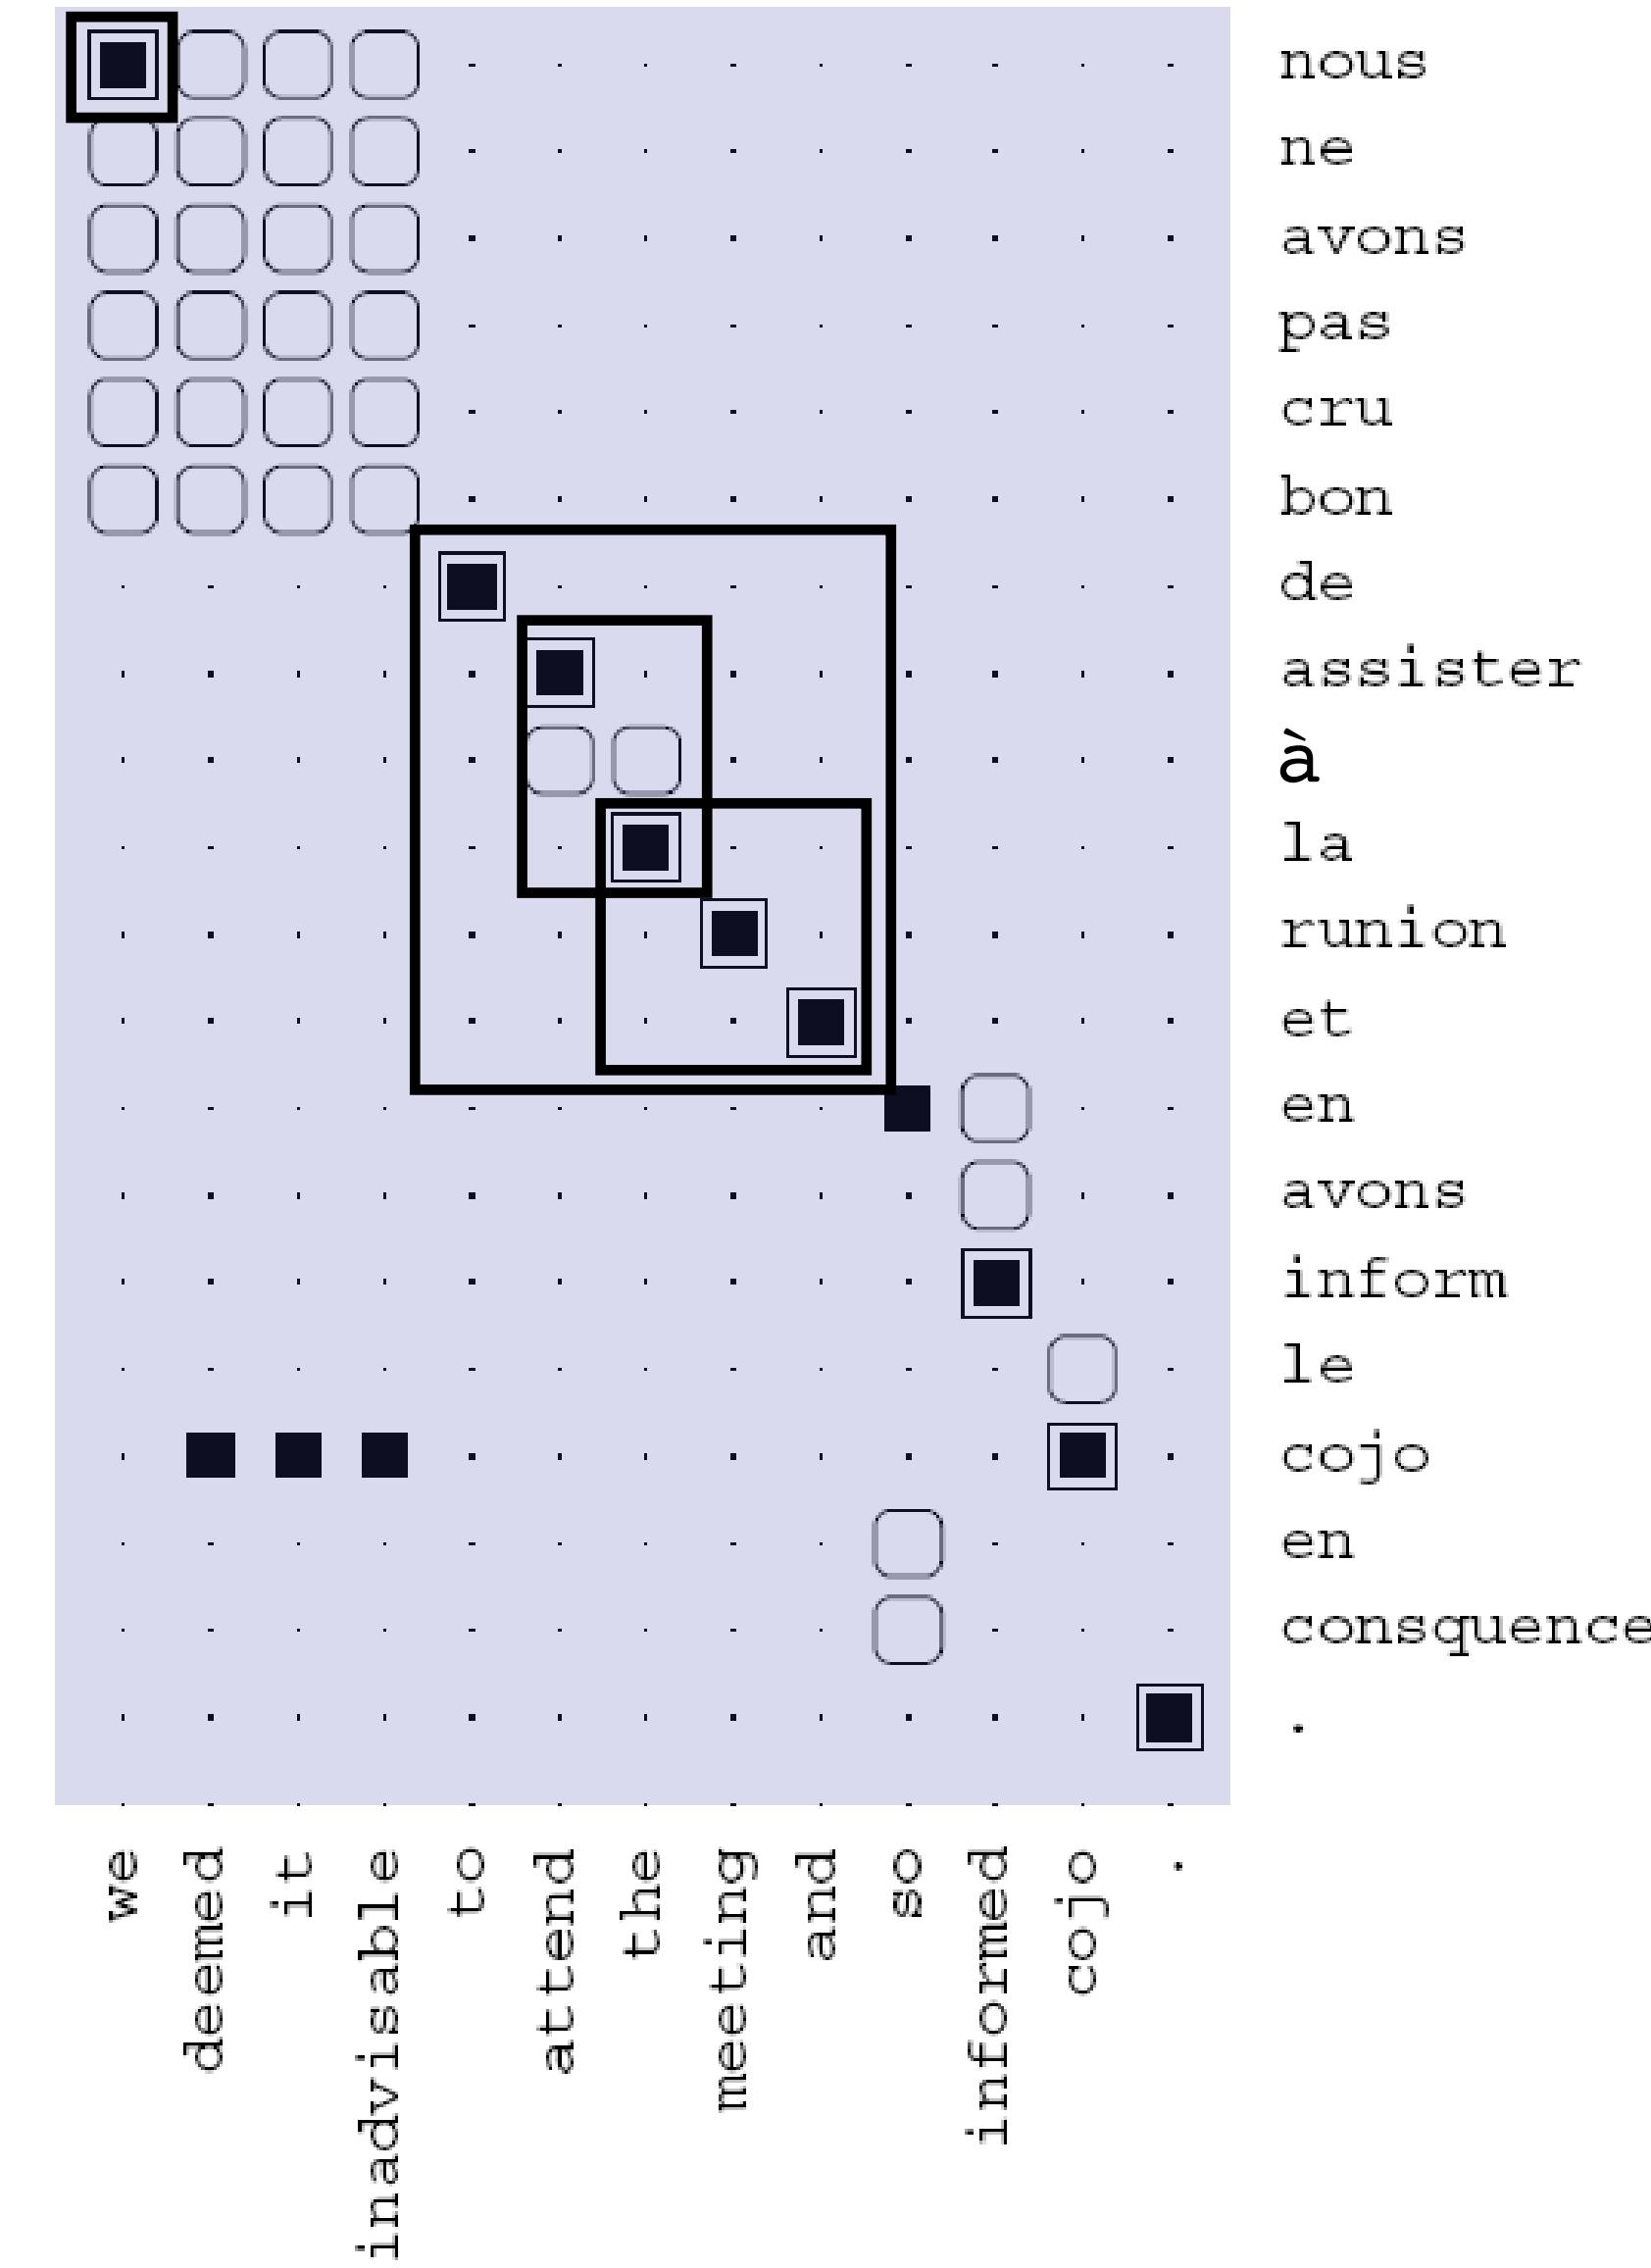
- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

d'assister à la reunion et ||| to attend the meeting and

assister à la reunion ||| attend the meeting

la reunion and ||| the meeting and

nous ||| we



Phrase Extraction

- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

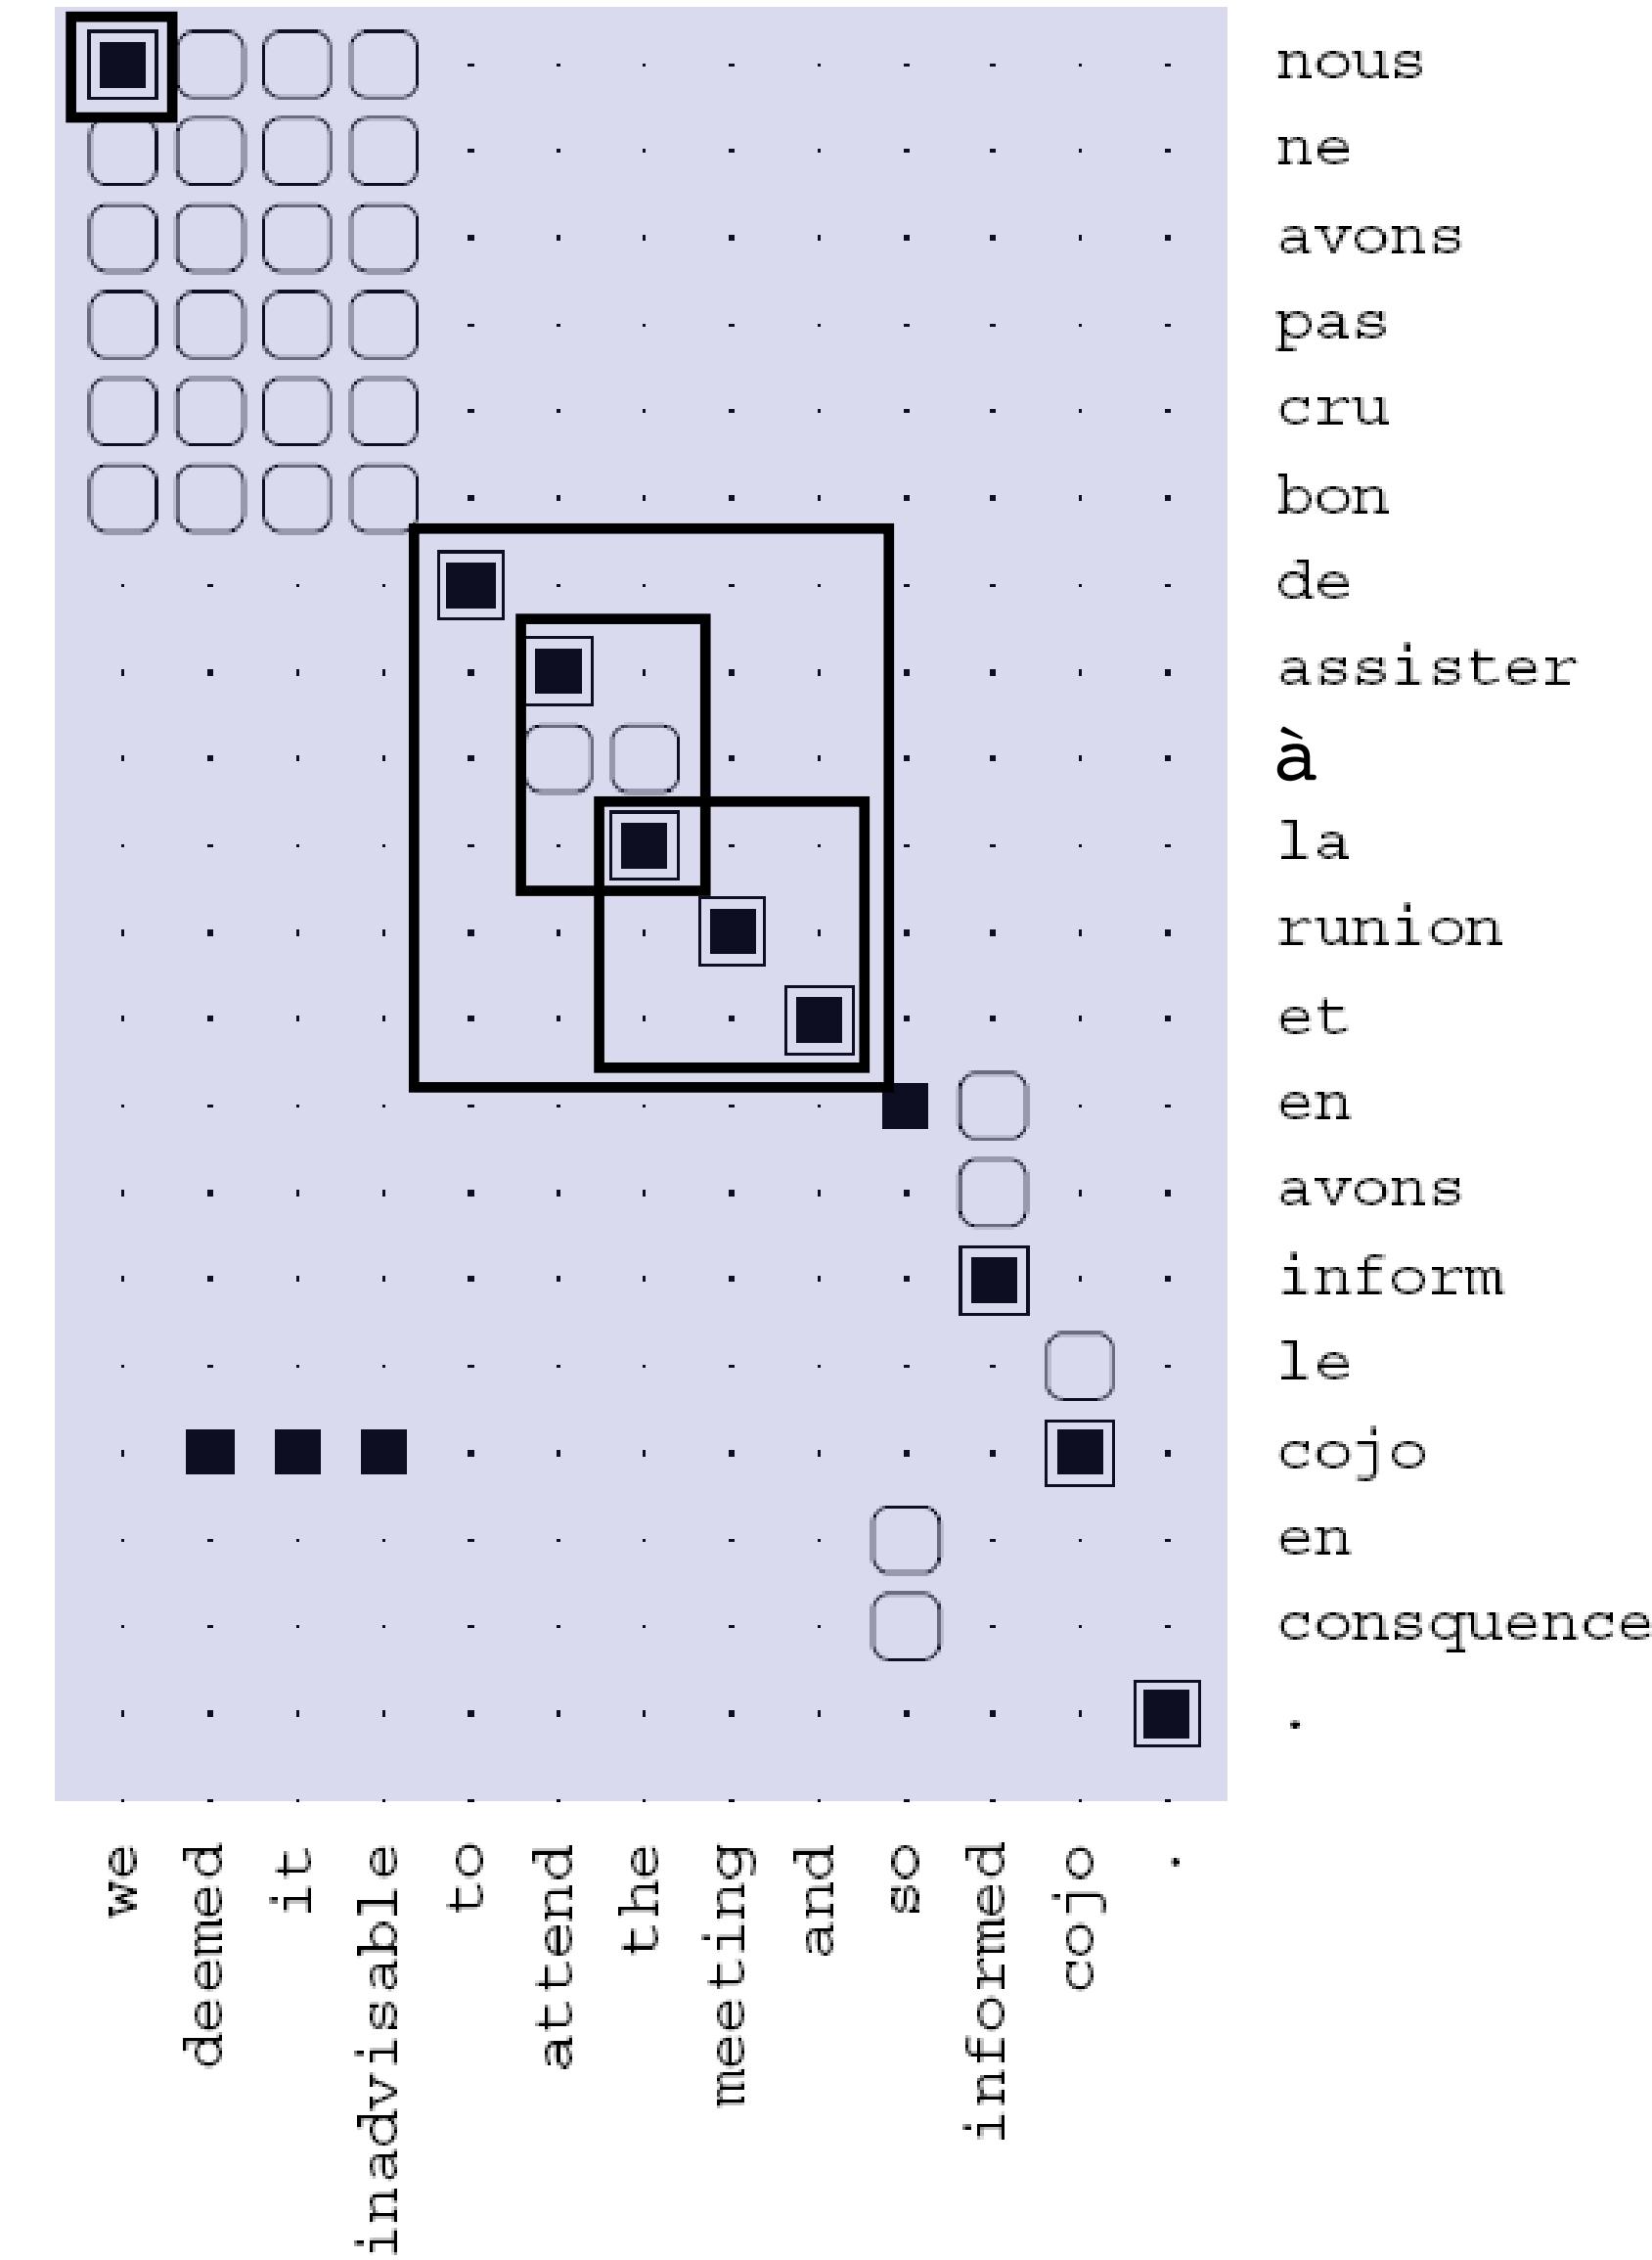
d'assister à la reunion et ||| to attend the meeting and

assister à la reunion ||| attend the meeting

la reunion and ||| the meeting and

nous ||| we

...



Phrase Extraction

- ▶ Find contiguous sets of aligned words in the two languages that don't have alignments to other words

d'assister à la reunion et ||| to attend the meeting and

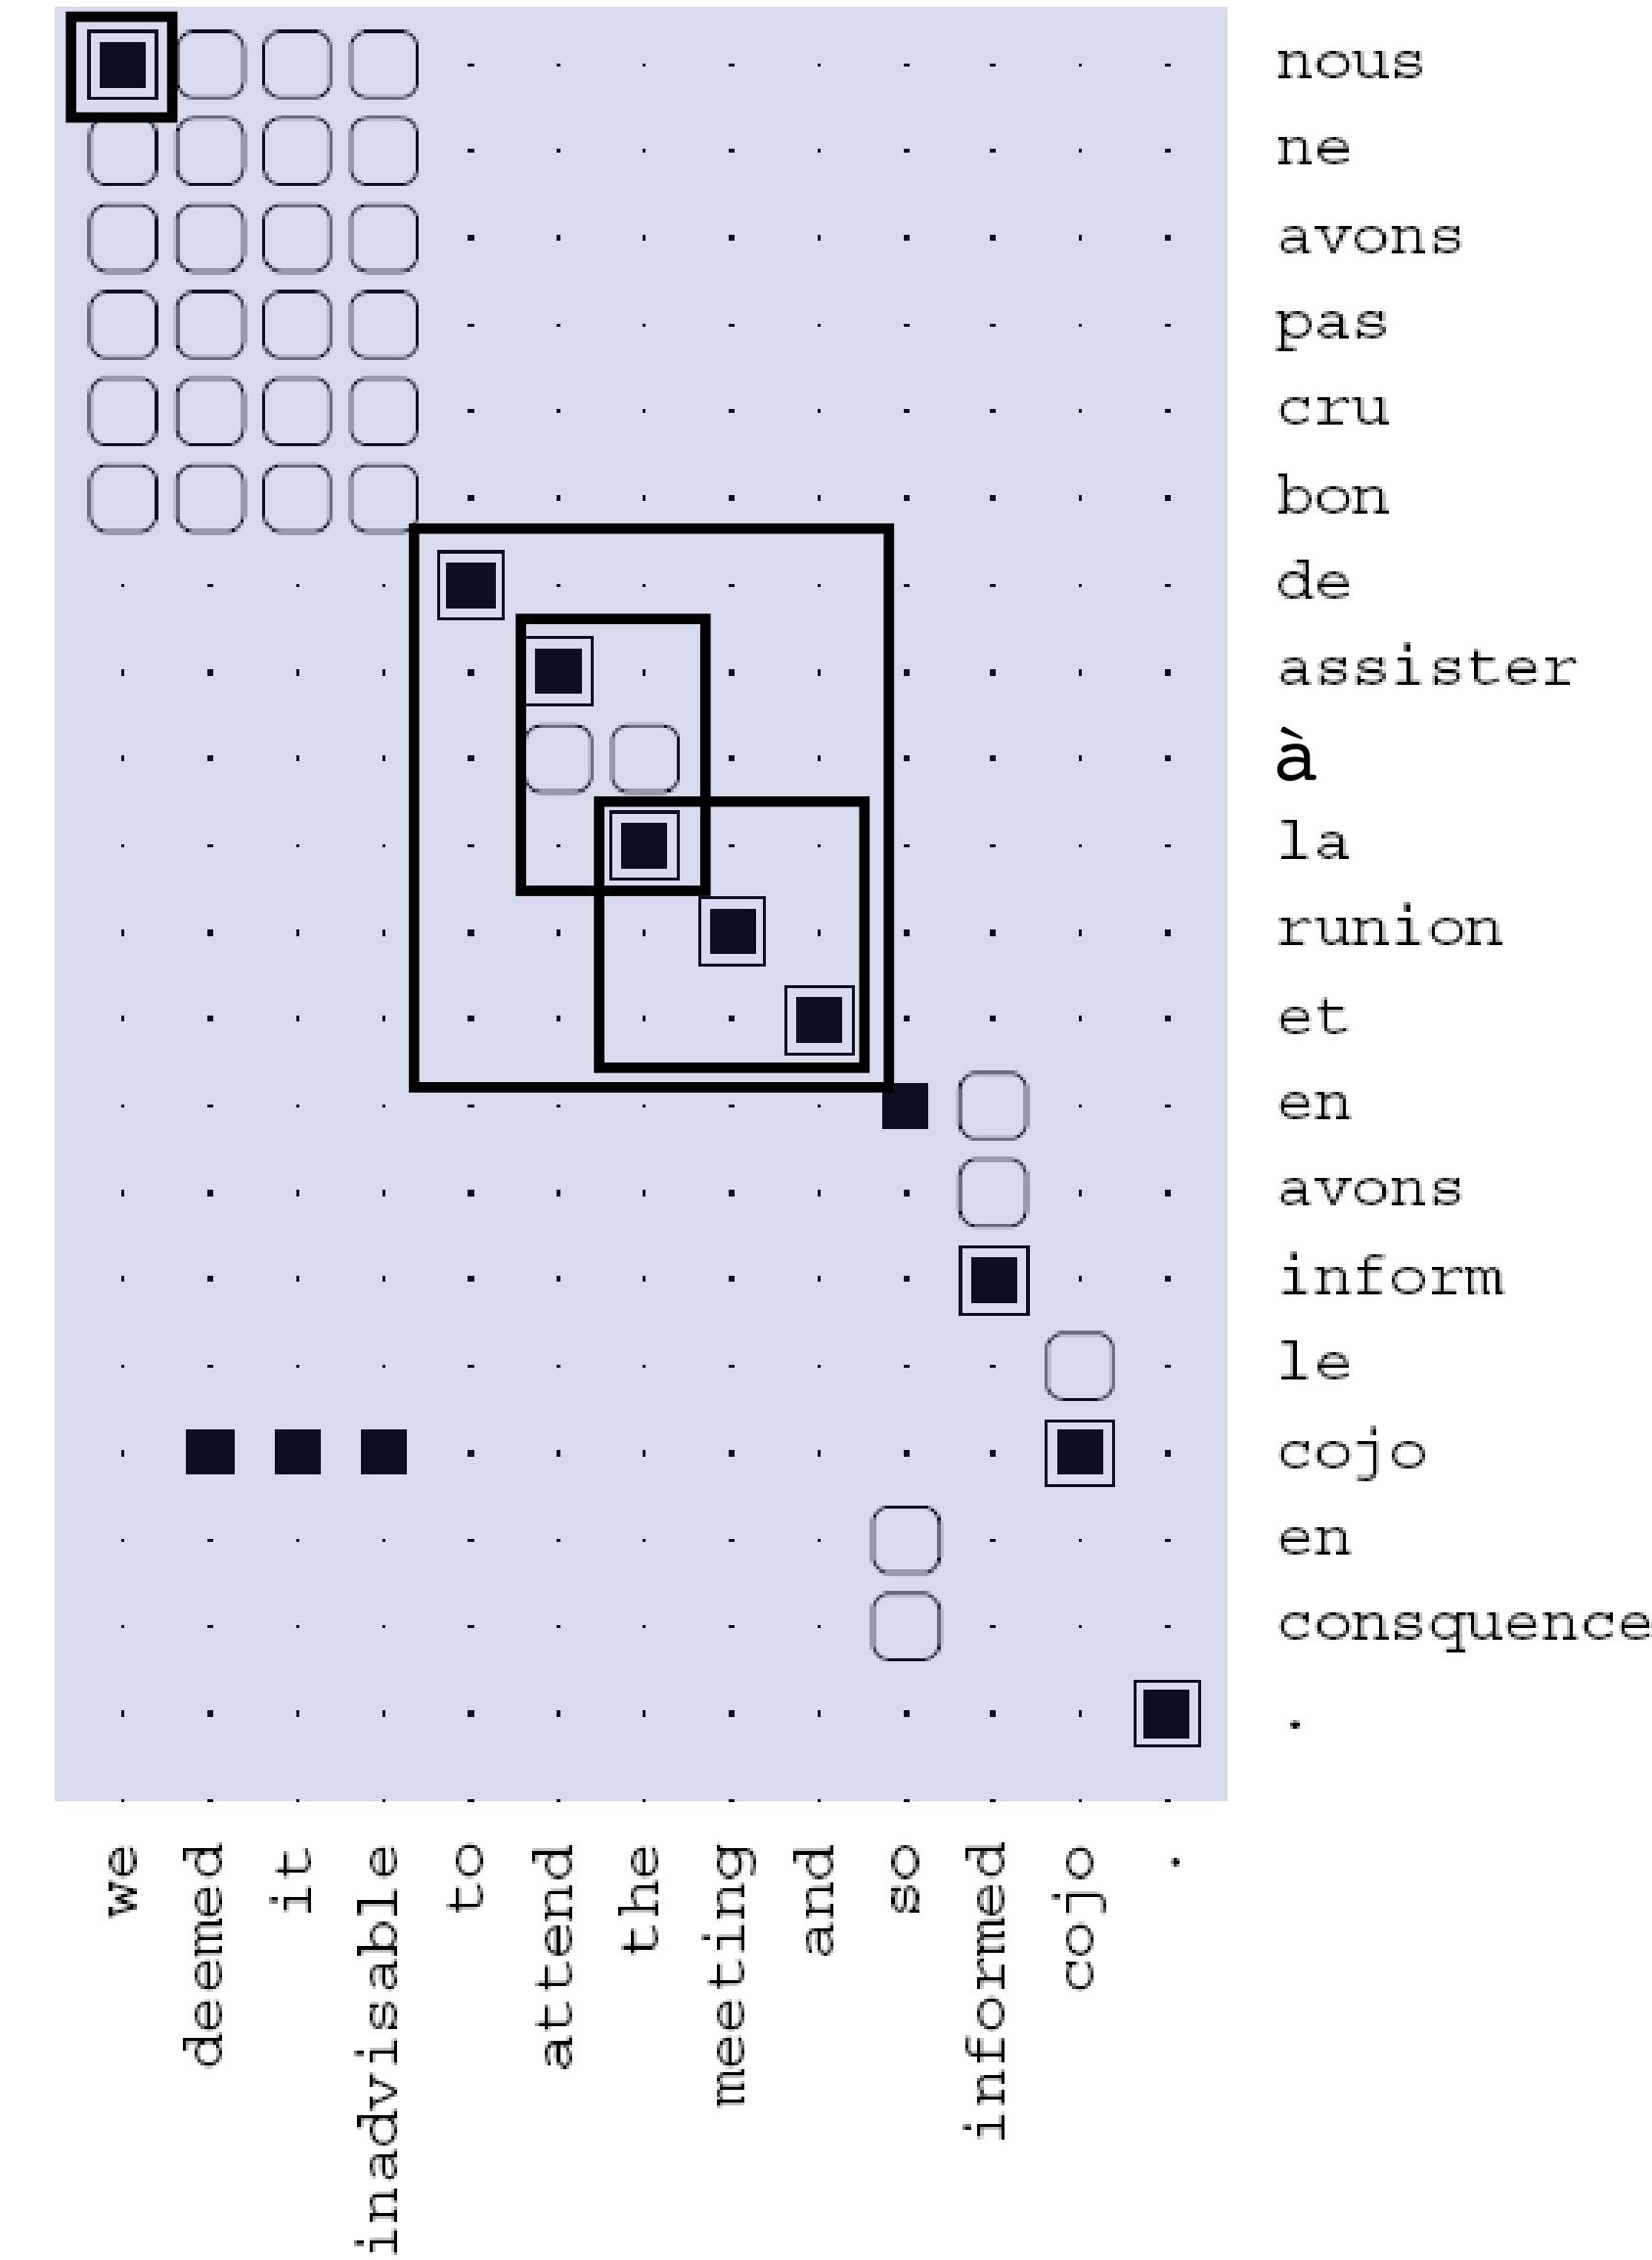
assister à la reunion ||| attend the meeting

la reunion and ||| the meeting and

nous ||| we

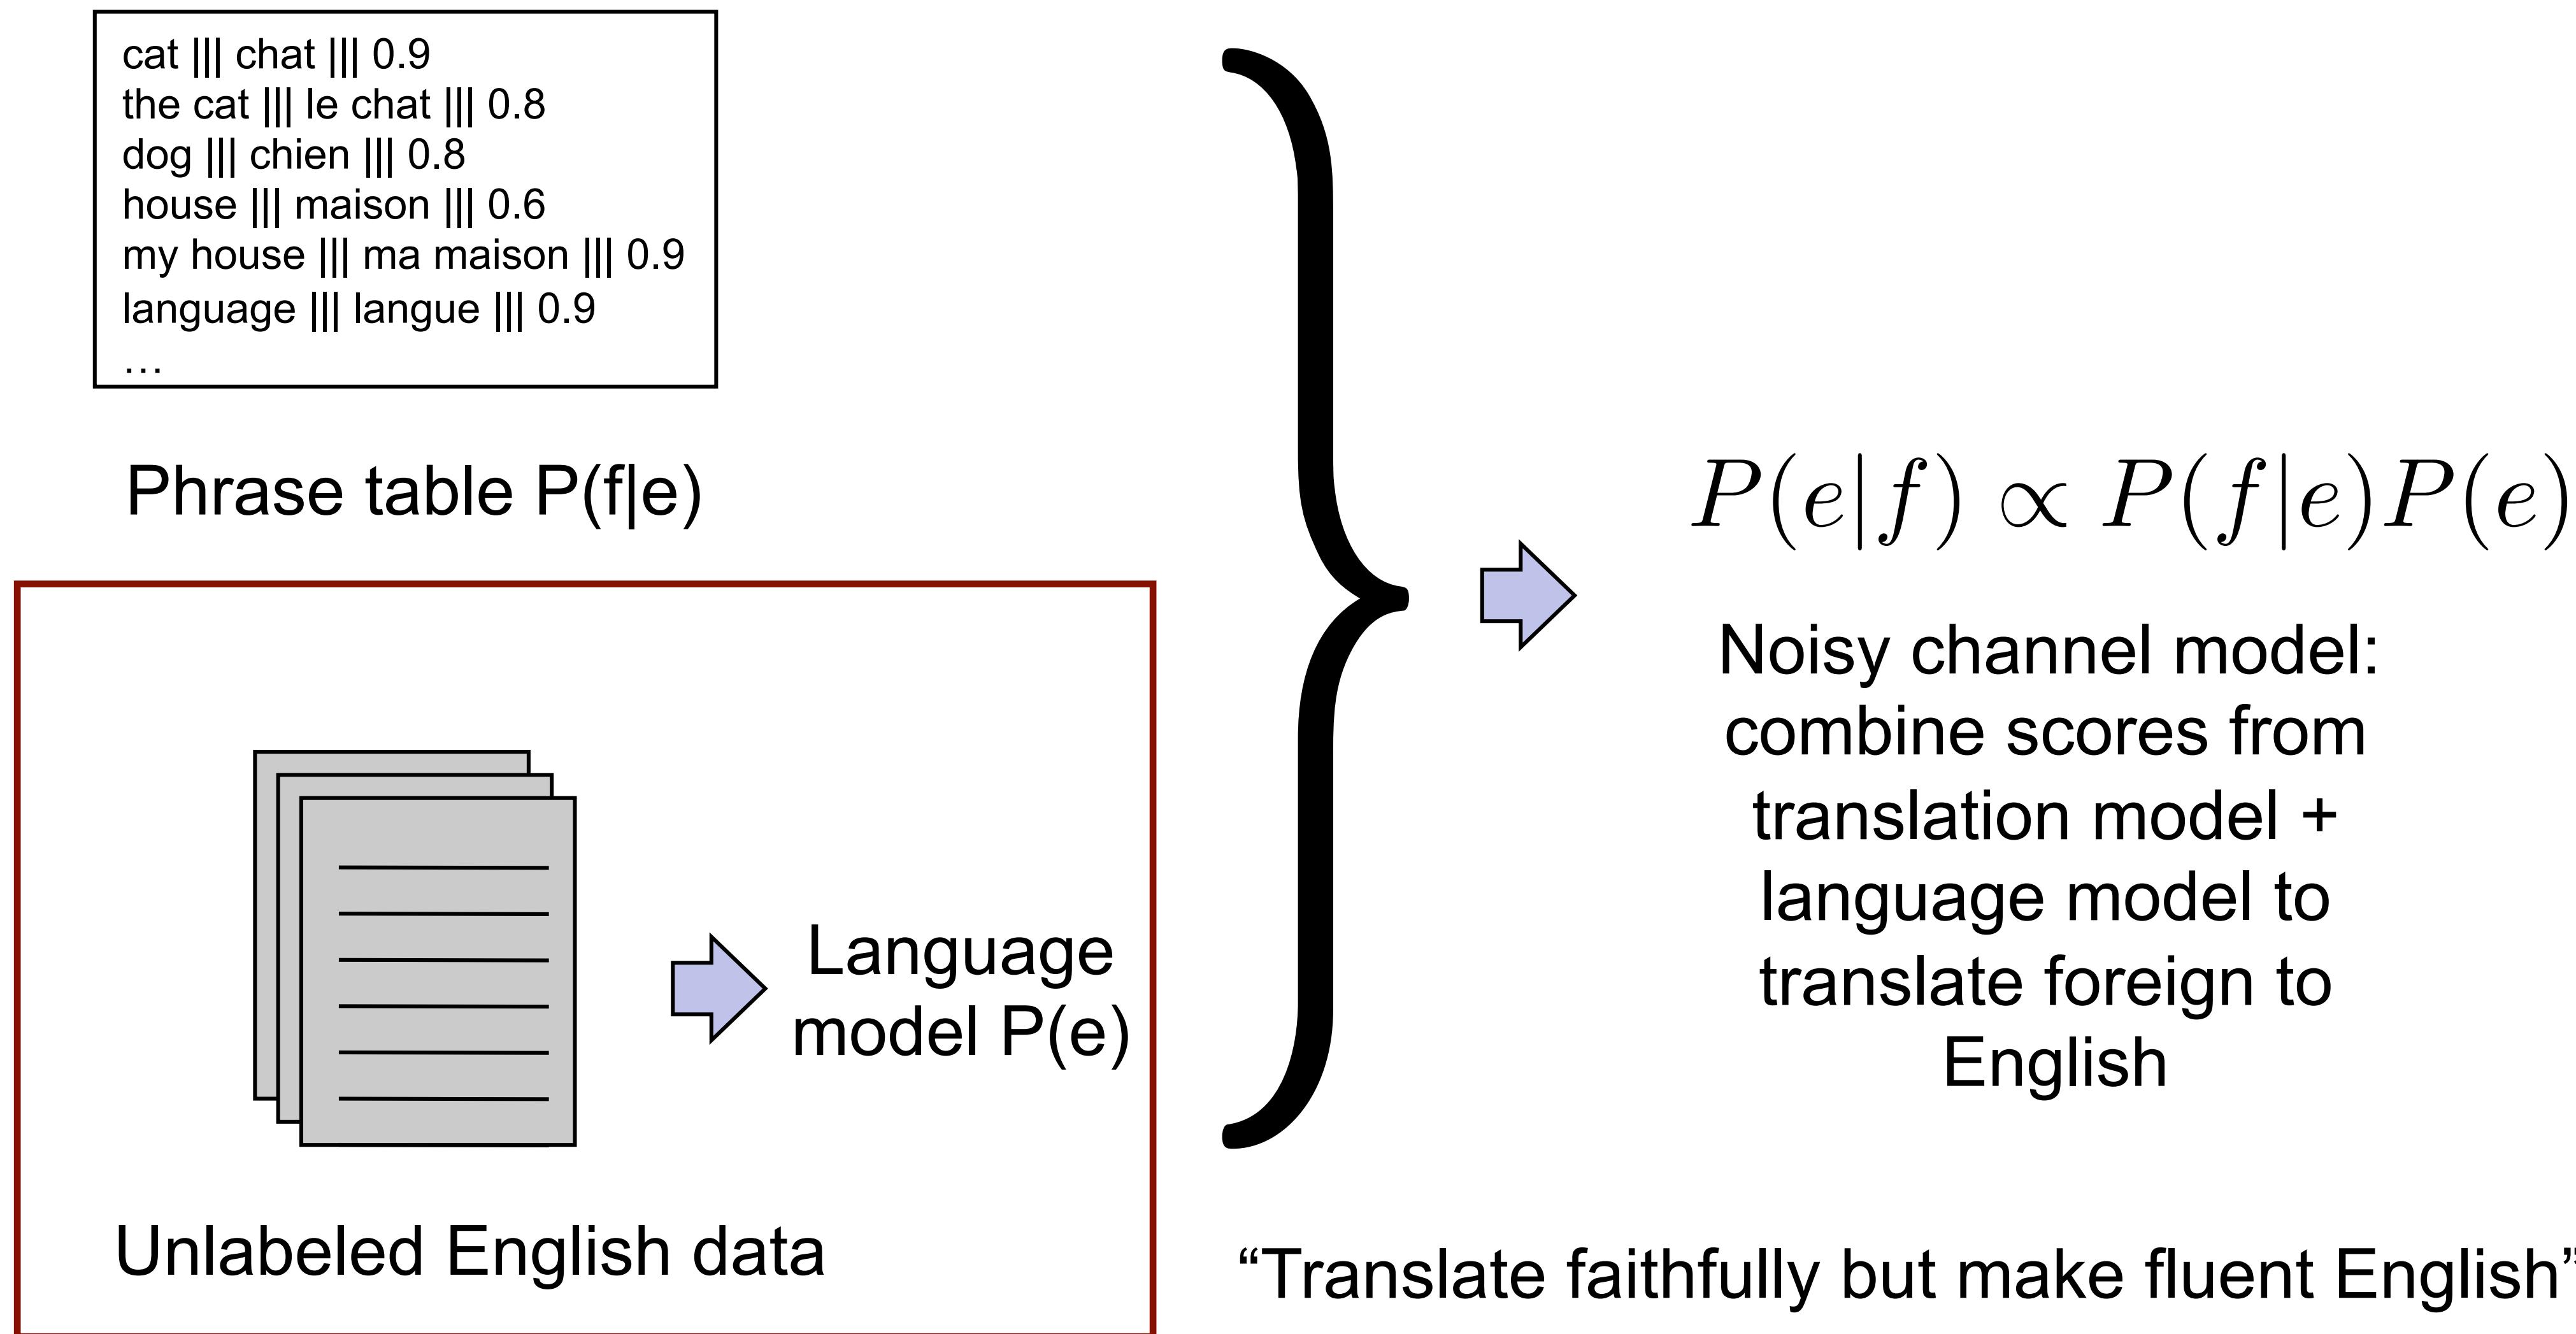
...

- ▶ Lots of phrases possible, count across all sentences and score by frequency



Language Modeling

Phrase-Based MT



N-gram Language Models

I visited San _____ put a distribution over the next word

N-gram Language Models

I visited San _____ put a distribution over the next word

- ▶ Simple generative model: distribution of next word is a multinomial distribution conditioned on previous n-1 words

N-gram Language Models

I visited San _____ put a distribution over the next word

- ▶ Simple generative model: distribution of next word is a multinomial distribution conditioned on previous n-1 words

$$P(x|\text{visited San}) = \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})}$$

N-gram Language Models

I visited San _____ put a distribution over the next word

- ▶ Simple generative model: distribution of next word is a multinomial distribution conditioned on previous n-1 words

$$P(x|\text{visited San}) = \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})}$$

Maximum likelihood estimate of this probability from a corpus

N-gram Language Models

I visited San _____ put a distribution over the next word

- ▶ Simple generative model: distribution of next word is a multinomial distribution conditioned on previous n-1 words

$$P(x|\text{visited San}) = \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})}$$

Maximum likelihood estimate of this probability from a corpus

- ▶ Just relies on counts, even in 2008 could scale up to 1.3M word types, 4B n-grams (all 5-grams occurring >40 times on the Web)

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

$$P(x|\text{visited San}) = (1 - \lambda) \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

$$P(x|\text{visited San}) = (1 - \lambda) \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

smooth
this
too!

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

$$P(x|\text{visited San}) = (1 - \lambda) \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

smooth
this
too!

- ▶ One technique is “absolute discounting:” subtract off constant k from numerator, set lambda to make this normalize ($k=1$ is like leave-one-out)

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

$$P(x|\text{visited San}) = (1 - \lambda) \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

smooth
this
too!

- ▶ One technique is “absolute discounting:” subtract off constant k from numerator, set lambda to make this normalize ($k=1$ is like leave-one-out)

$$P(x|\text{visited San}) = \frac{\text{count}(\text{visited San}, x) - k}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

Smoothing N-gram Language Models

I visited San _____ put a distribution over the next word!

- ▶ Smoothing is very important, particularly when using 4+ gram models

$$P(x|\text{visited San}) = (1 - \lambda) \frac{\text{count}(\text{visited San}, x)}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

smooth
this
too!

- ▶ One technique is “absolute discounting:” subtract off constant k from numerator, set lambda to make this normalize ($k=1$ is like leave-one-out)

$$P(x|\text{visited San}) = \frac{\text{count}(\text{visited San}, x) - k}{\text{count}(\text{visited San})} + \lambda \frac{\text{count}(\text{San}, x)}{\text{count}(\text{San})}$$

- ▶ Kneser-Ney smoothing: this trick, plus low-order distributions modified to capture fertilities (how many distinct words appear in a context)

Engineering N-gram Models

- ▶ For 5+-gram models, need to store between 100M and 10B context-word-count triples

(a) Context-Encoding

w	c	val
1933	15176585	3
1933	15176587	2
1933	15176593	1
1933	15176613	8
1933	15179801	1
1935	15176585	298
1935	15176589	1

(b) Context Deltas

Δw	Δc	val
1933	15176585	3
+0	+2	1
+0	+5	1
+0	+40	8
+0	+188	1
+2	15176585	298
+0	+4	1

(c) Bits Required

$ \Delta w $	$ \Delta c $	$ val $
24	40	3
2	3	3
2	3	3
2	9	6
2	12	3
4	36	15
2	6	3

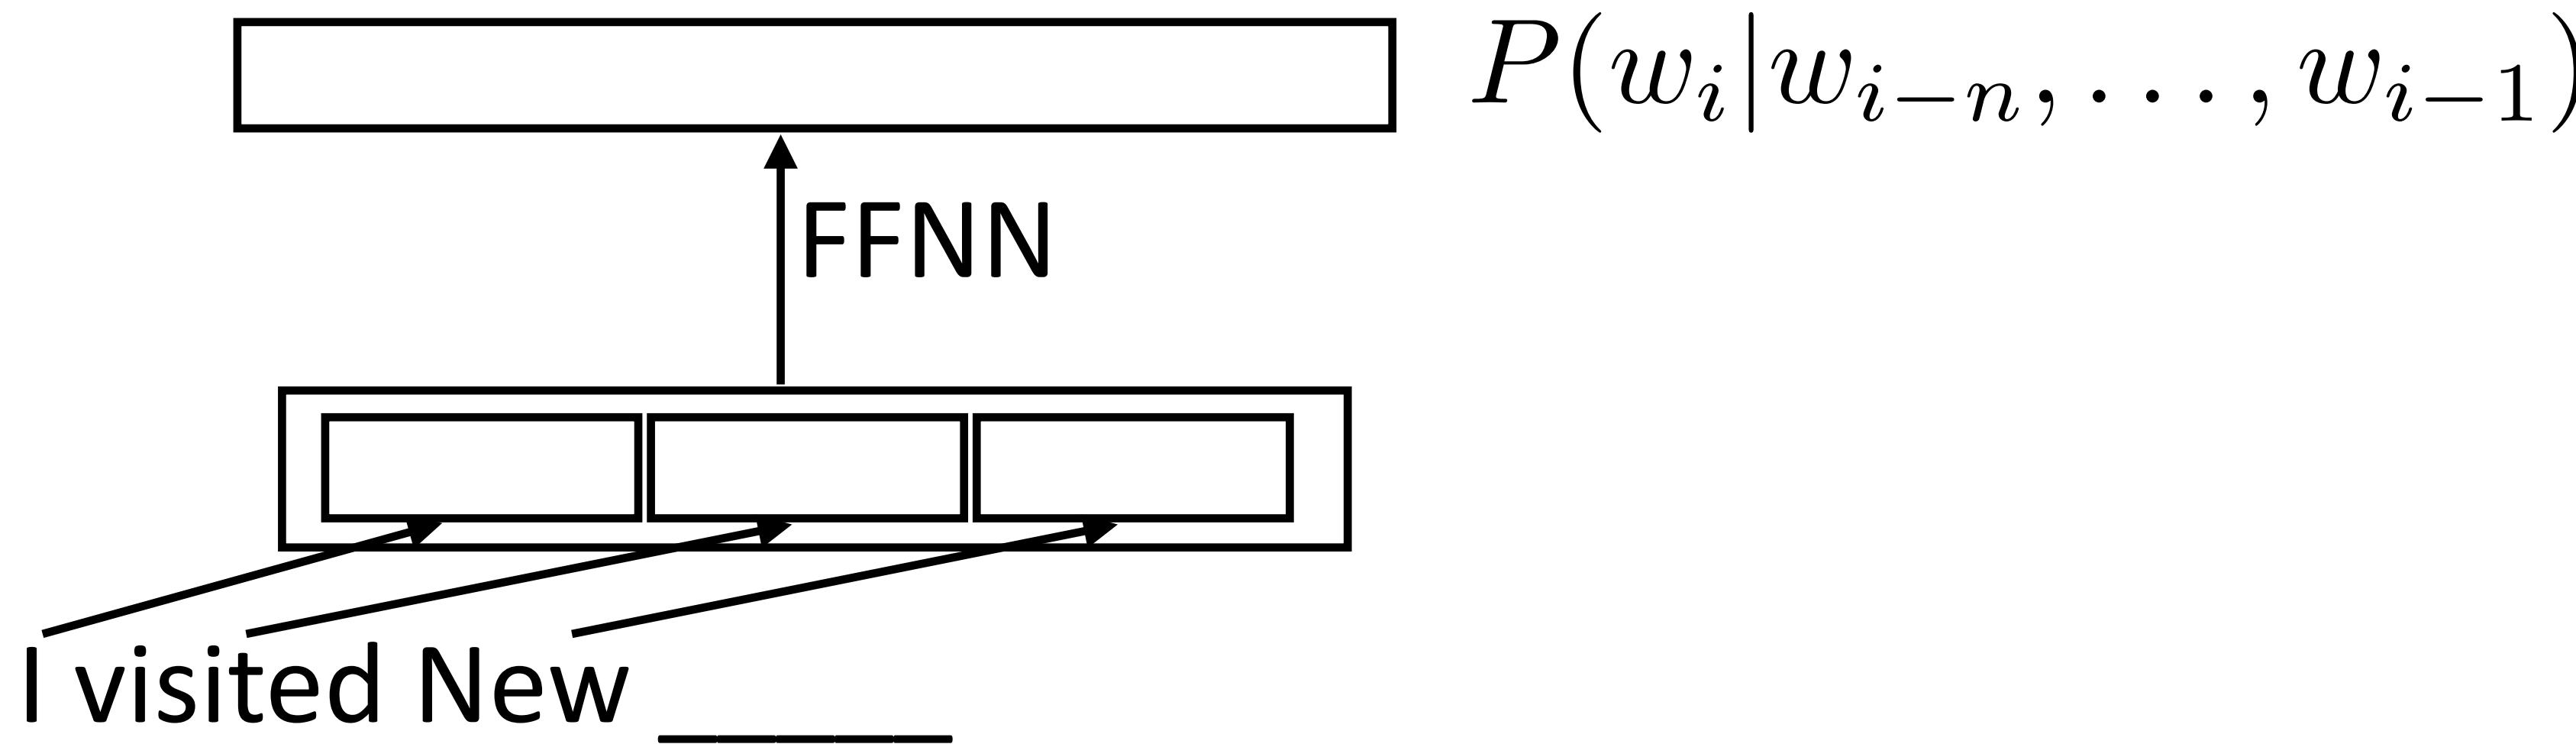
- ▶ Make it fit in memory by *delta encoding* scheme: store deltas instead of values and use variable-length encoding

Neural Language Models

- ▶ Early work: feedforward neural networks looking at context

Neural Language Models

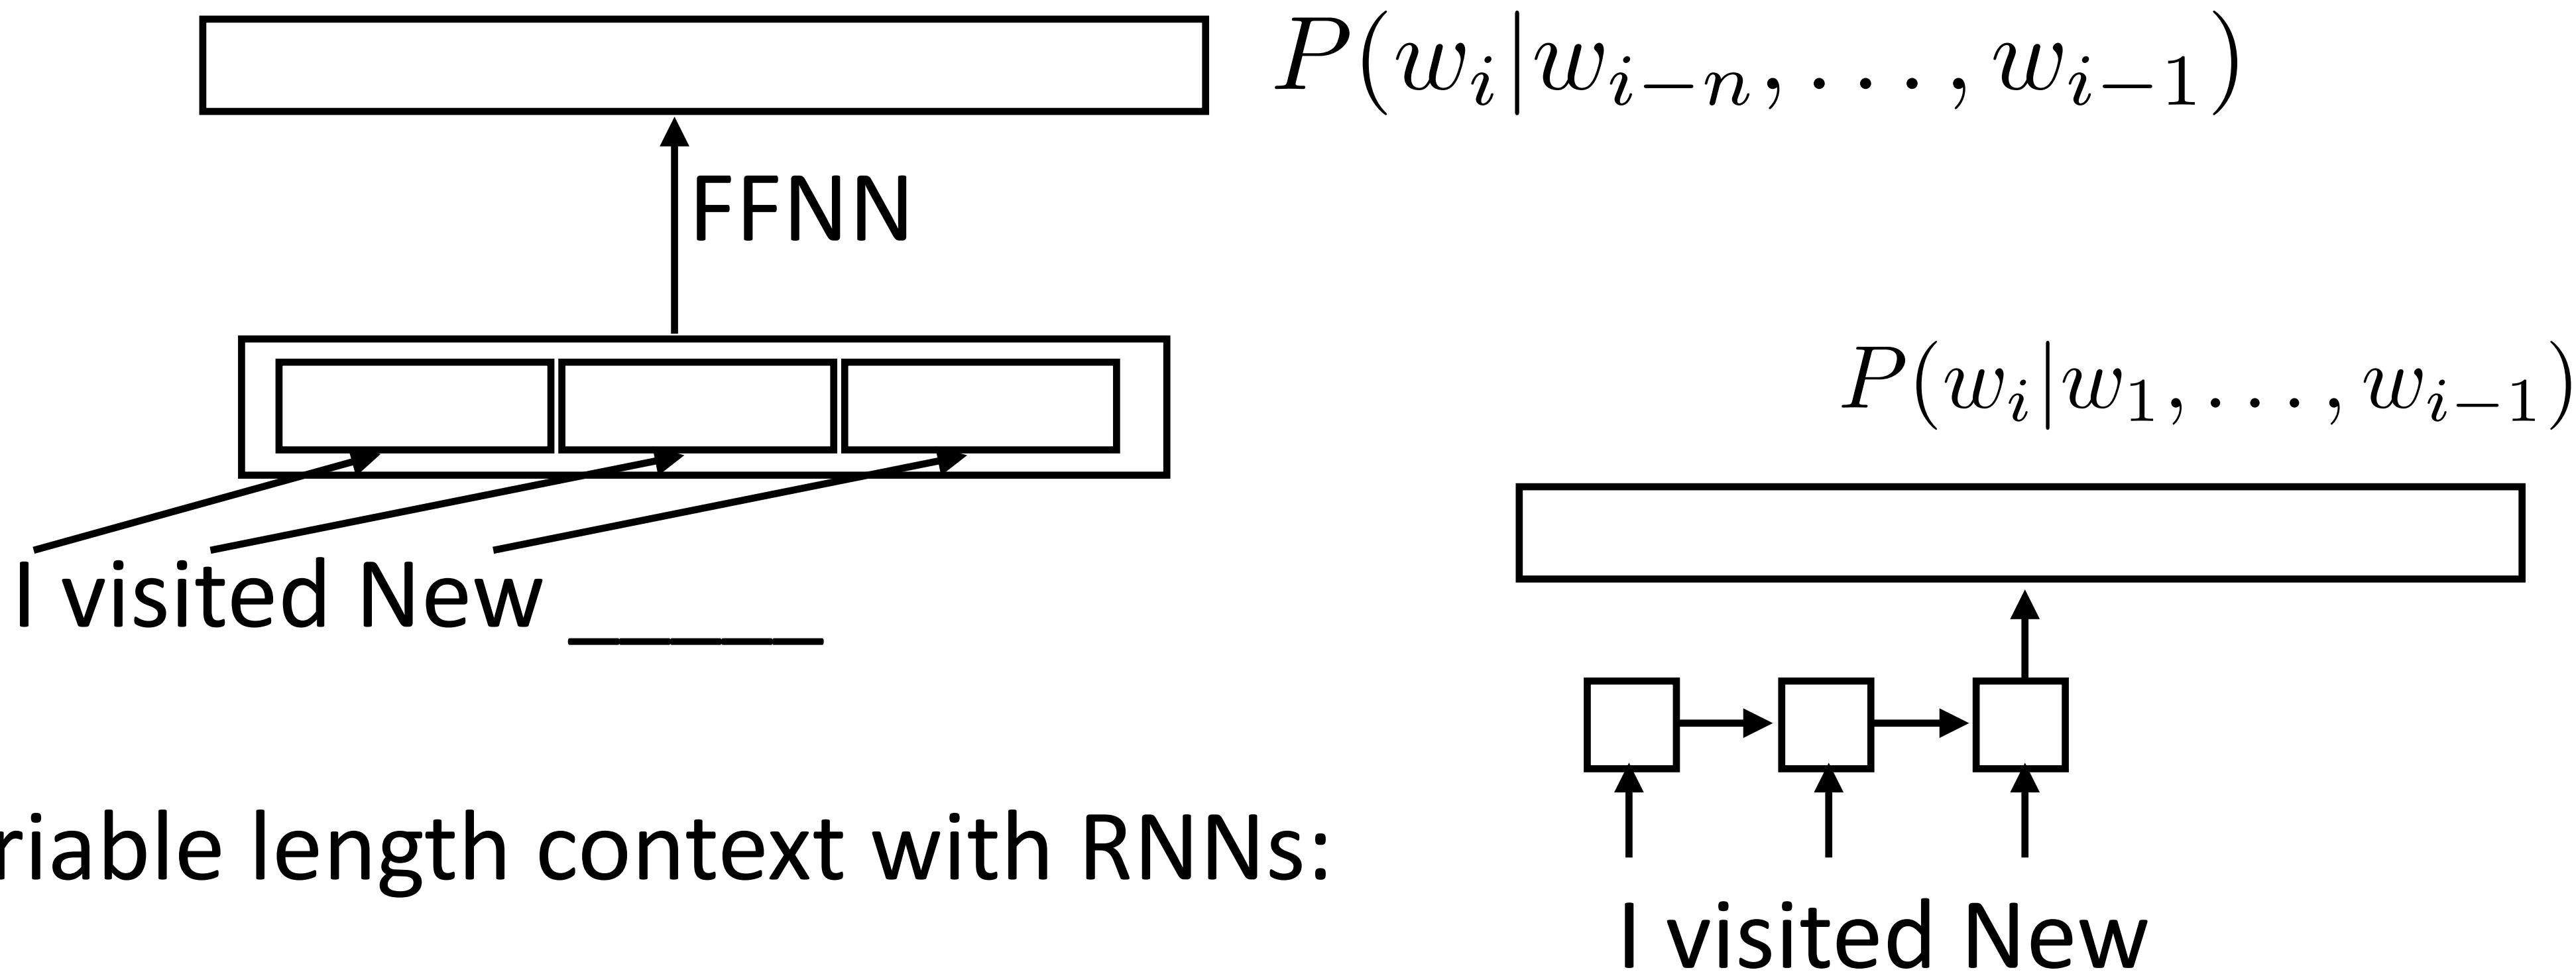
- ▶ Early work: feedforward neural networks looking at context



Mnih and Hinton (2003)

Neural Language Models

- ▶ Early work: feedforward neural networks looking at context

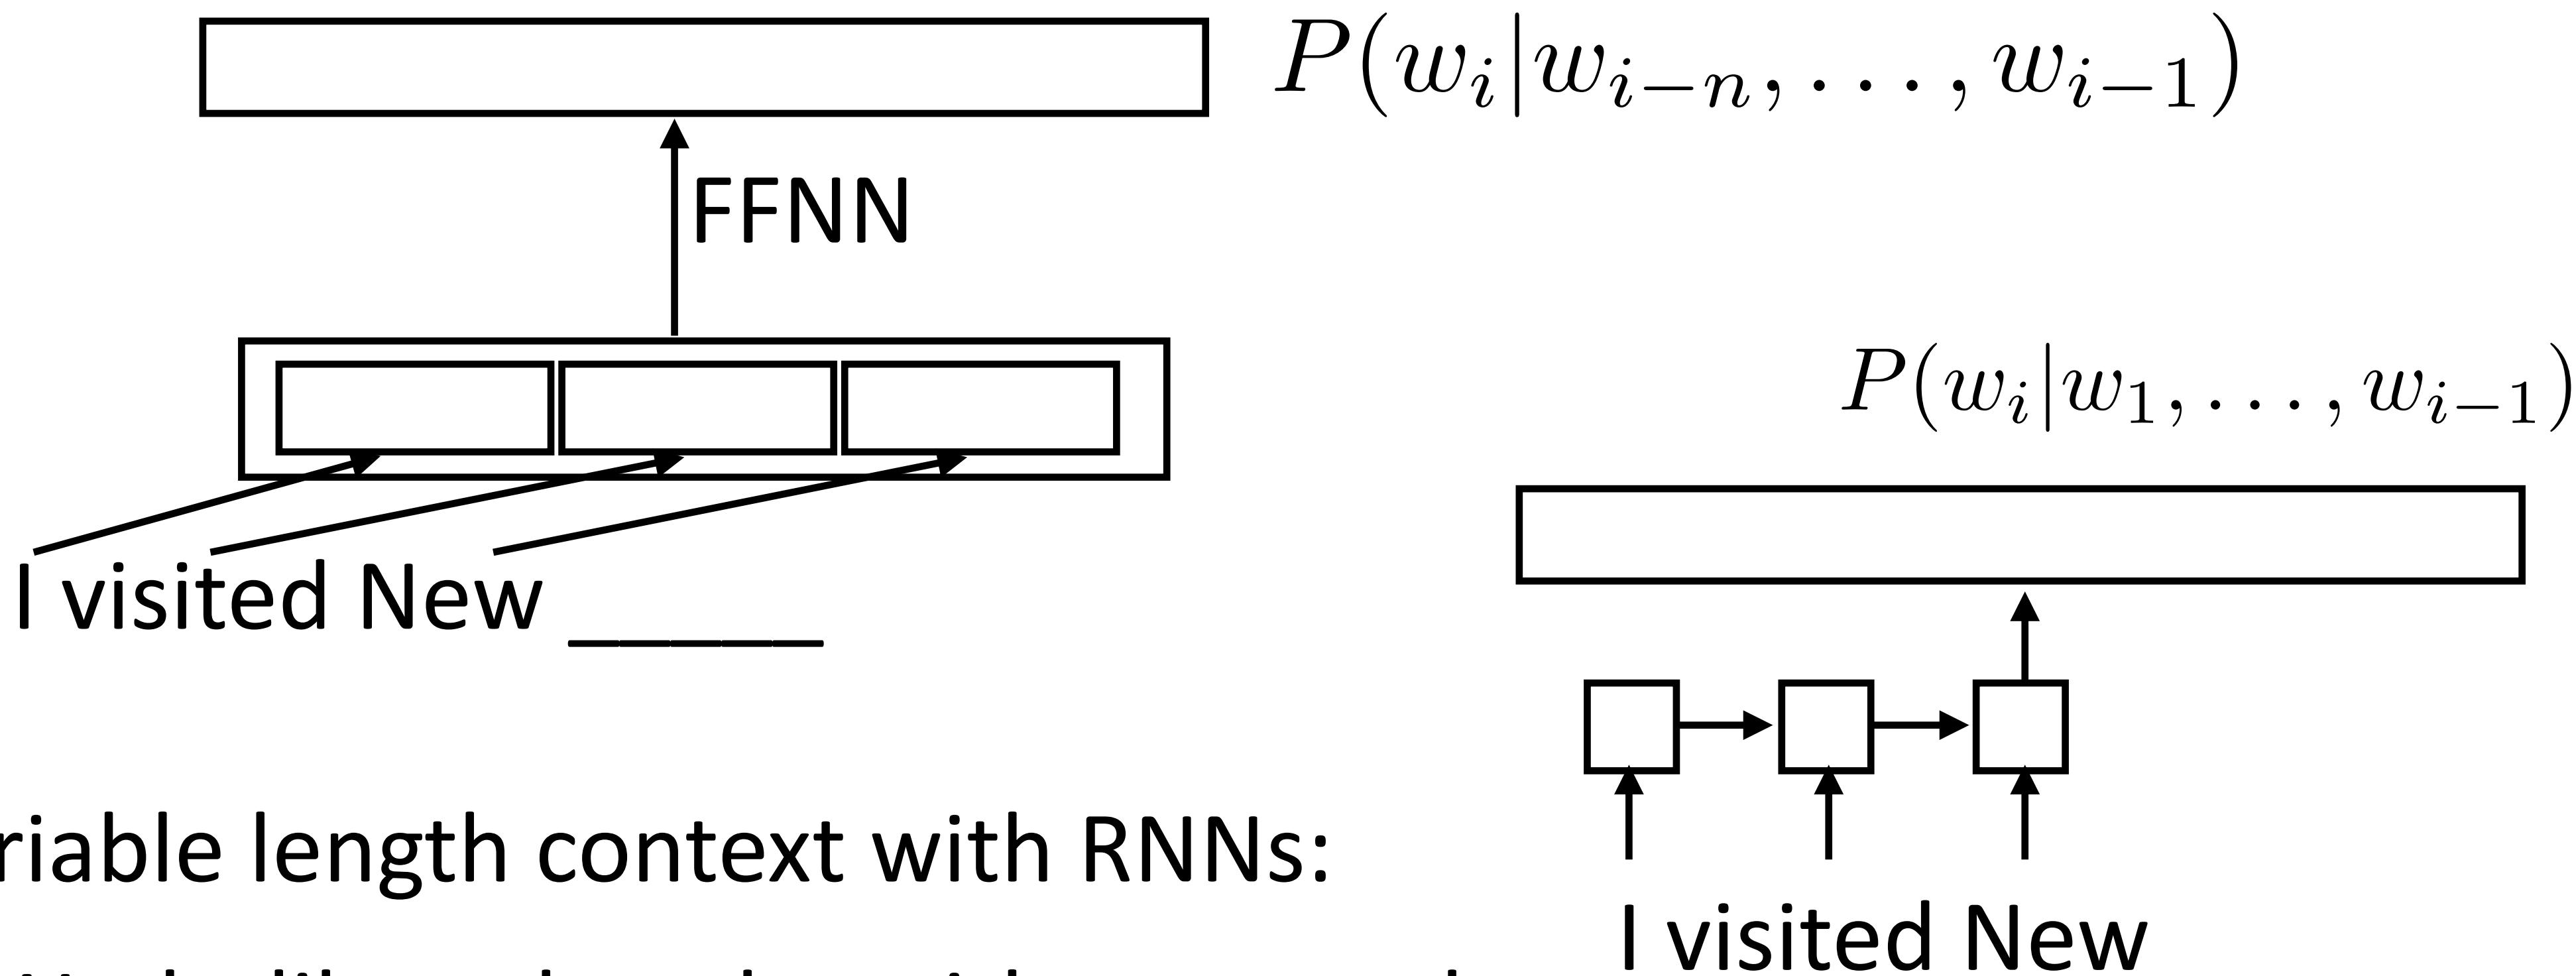


- ▶ Variable length context with RNNs:

Mnih and Hinton (2003)

Neural Language Models

- ▶ Early work: feedforward neural networks looking at context

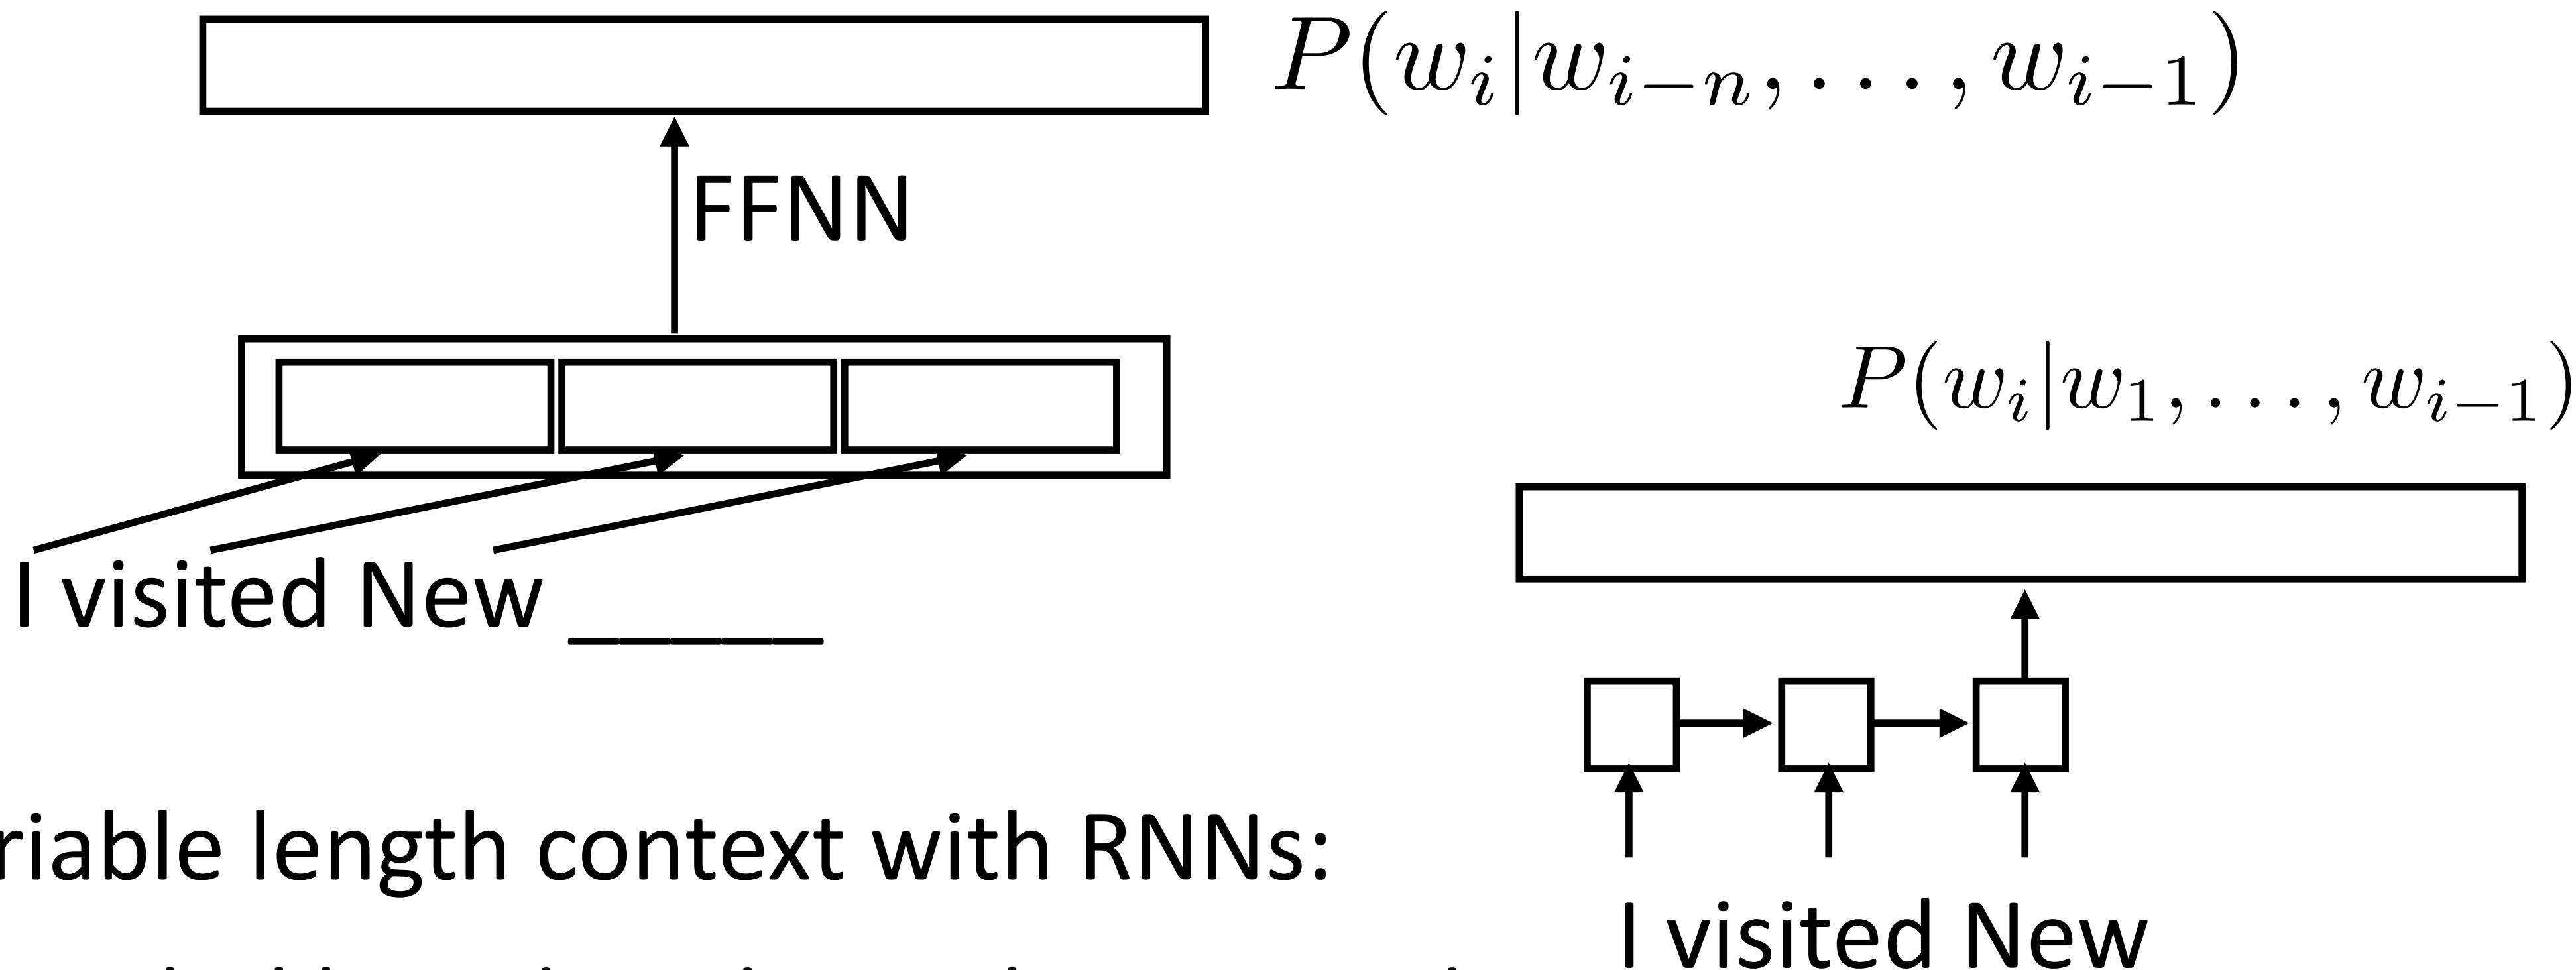


- ▶ Variable length context with RNNs:
 - ▶ Works like a decoder with no encoder

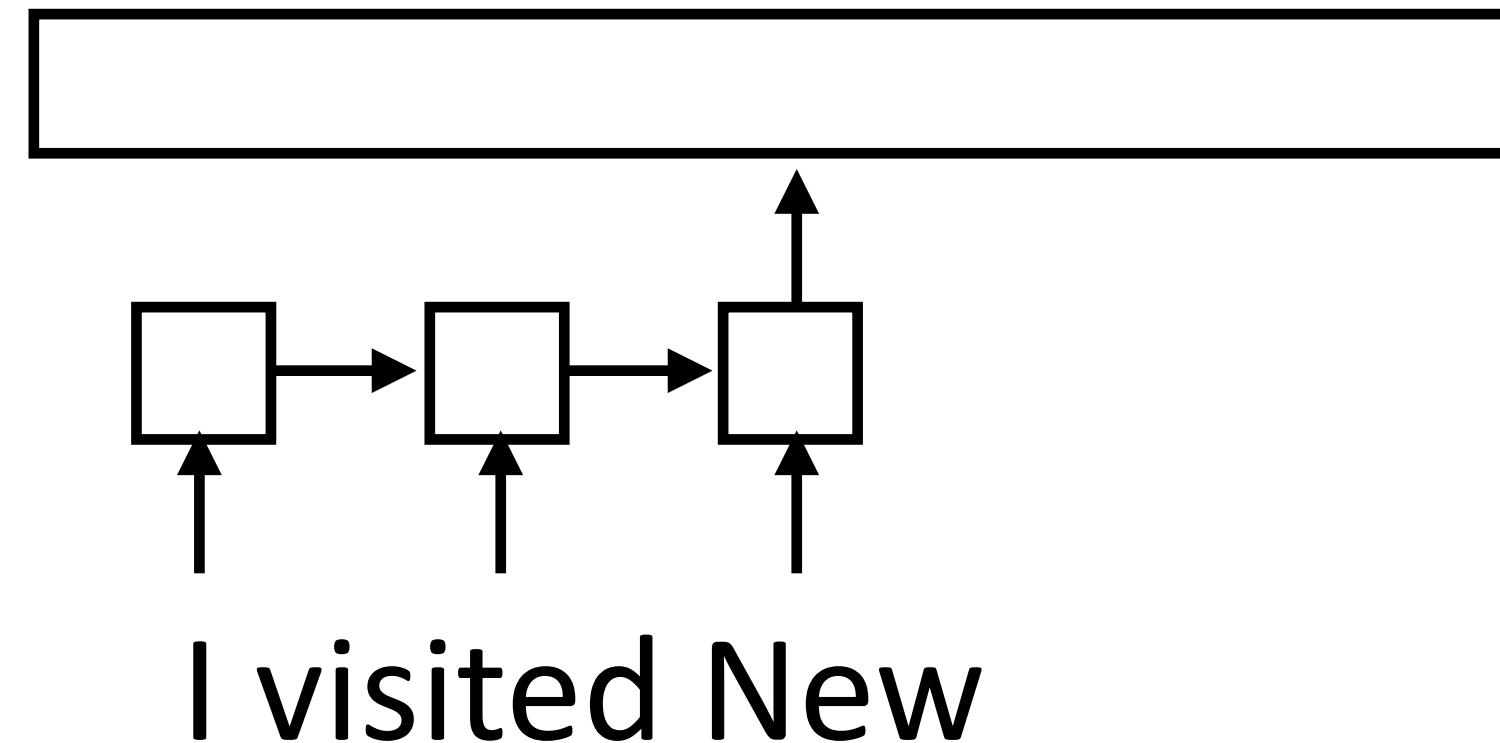
Mnih and Hinton (2003)

Neural Language Models

- ▶ Early work: feedforward neural networks looking at context



- ▶ Variable length context with RNNs:
 - ▶ Works like a decoder with no encoder
 - ▶ Slow to train over lots of data!



Mnih and Hinton (2003)

Evaluation

Evaluation

- (One sentence) negative log likelihood: $\sum_{i=1}^n \log p(x_i|x_1, \dots, x_{i-1})$

Evaluation

- ▶ (One sentence) negative log likelihood: $\sum_{i=1}^n \log p(x_i|x_1, \dots, x_{i-1})$
- ▶ Perplexity: $2^{-\frac{1}{n} \sum_{i=1}^n \log_2 p(x_i|x_1, \dots, x_{i-1})}$

Evaluation

- ▶ (One sentence) negative log likelihood: $\sum_{i=1}^n \log p(x_i|x_1, \dots, x_{i-1})$
- ▶ Perplexity: $2^{-\frac{1}{n} \sum_{i=1}^n \log_2 p(x_i|x_1, \dots, x_{i-1})}$
 - ▶ NLL (base 2) averaged over the sentence, exponentiated

Evaluation

- ▶ (One sentence) negative log likelihood: $\sum_{i=1}^n \log p(x_i|x_1, \dots, x_{i-1})$
- ▶ Perplexity: $2^{-\frac{1}{n} \sum_{i=1}^n \log_2 p(x_i|x_1, \dots, x_{i-1})}$
 - ▶ NLL (base 2) averaged over the sentence, exponentiated
 - ▶ NLL = -2 -> on average, correct thing has prob 1/4 -> PPL = 4. PPL is sort of like branching factor

Results

Merity et al. (2017), Melis et al. (2017)

Results

- ▶ Evaluate on Penn Treebank: small dataset (1M words) compared to what's used in MT, but common benchmark

Results

- ▶ Evaluate on Penn Treebank: small dataset (1M words) compared to what's used in MT, but common benchmark
- ▶ Kneser-Ney 5-gram model with cache: PPL = 125.7

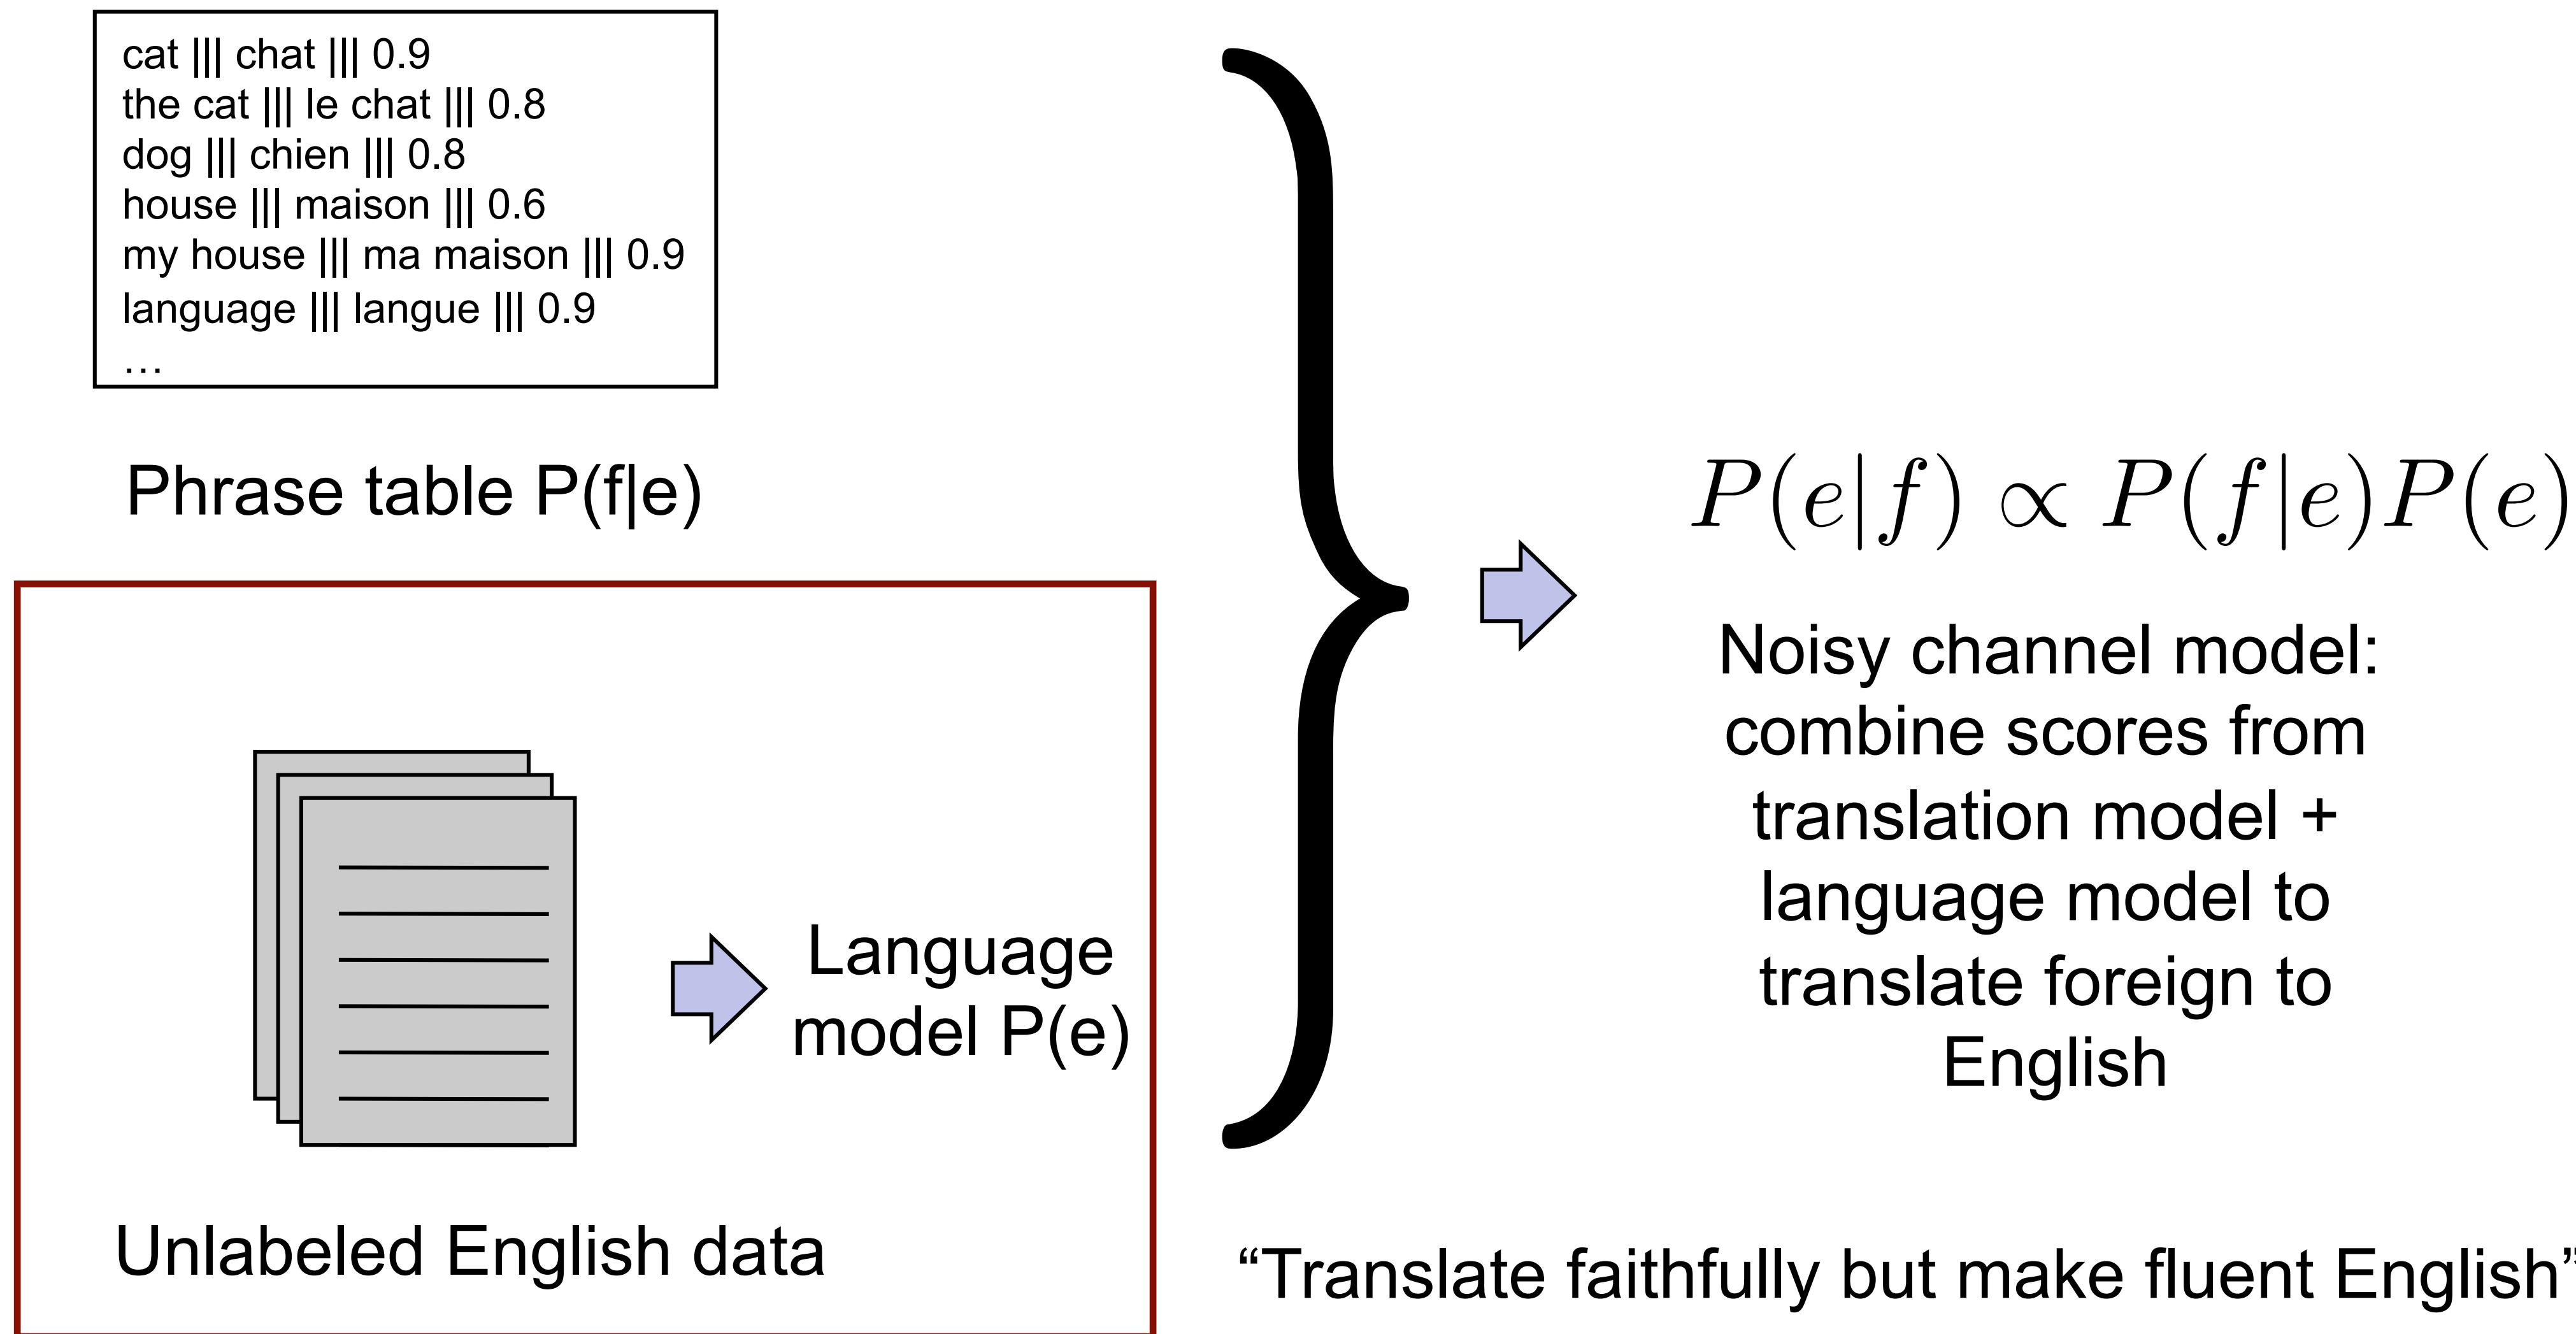
Results

- ▶ Evaluate on Penn Treebank: small dataset (1M words) compared to what's used in MT, but common benchmark
- ▶ Kneser-Ney 5-gram model with cache: PPL = 125.7
- ▶ LSTM: PPL ~ 60-80 (depending on how much you optimize it)

Results

- ▶ Evaluate on Penn Treebank: small dataset (1M words) compared to what's used in MT, but common benchmark
- ▶ Kneser-Ney 5-gram model with cache: PPL = 125.7
- ▶ LSTM: PPL ~ 60-80 (depending on how much you optimize it)
- ▶ Melis et al.: many neural LM improvements from 2014-2017 are subsumed by just using the right regularization (right dropout settings). So LSTMs are pretty good

Phrase-Based MT



Decoding

Phrase-Based Decoding

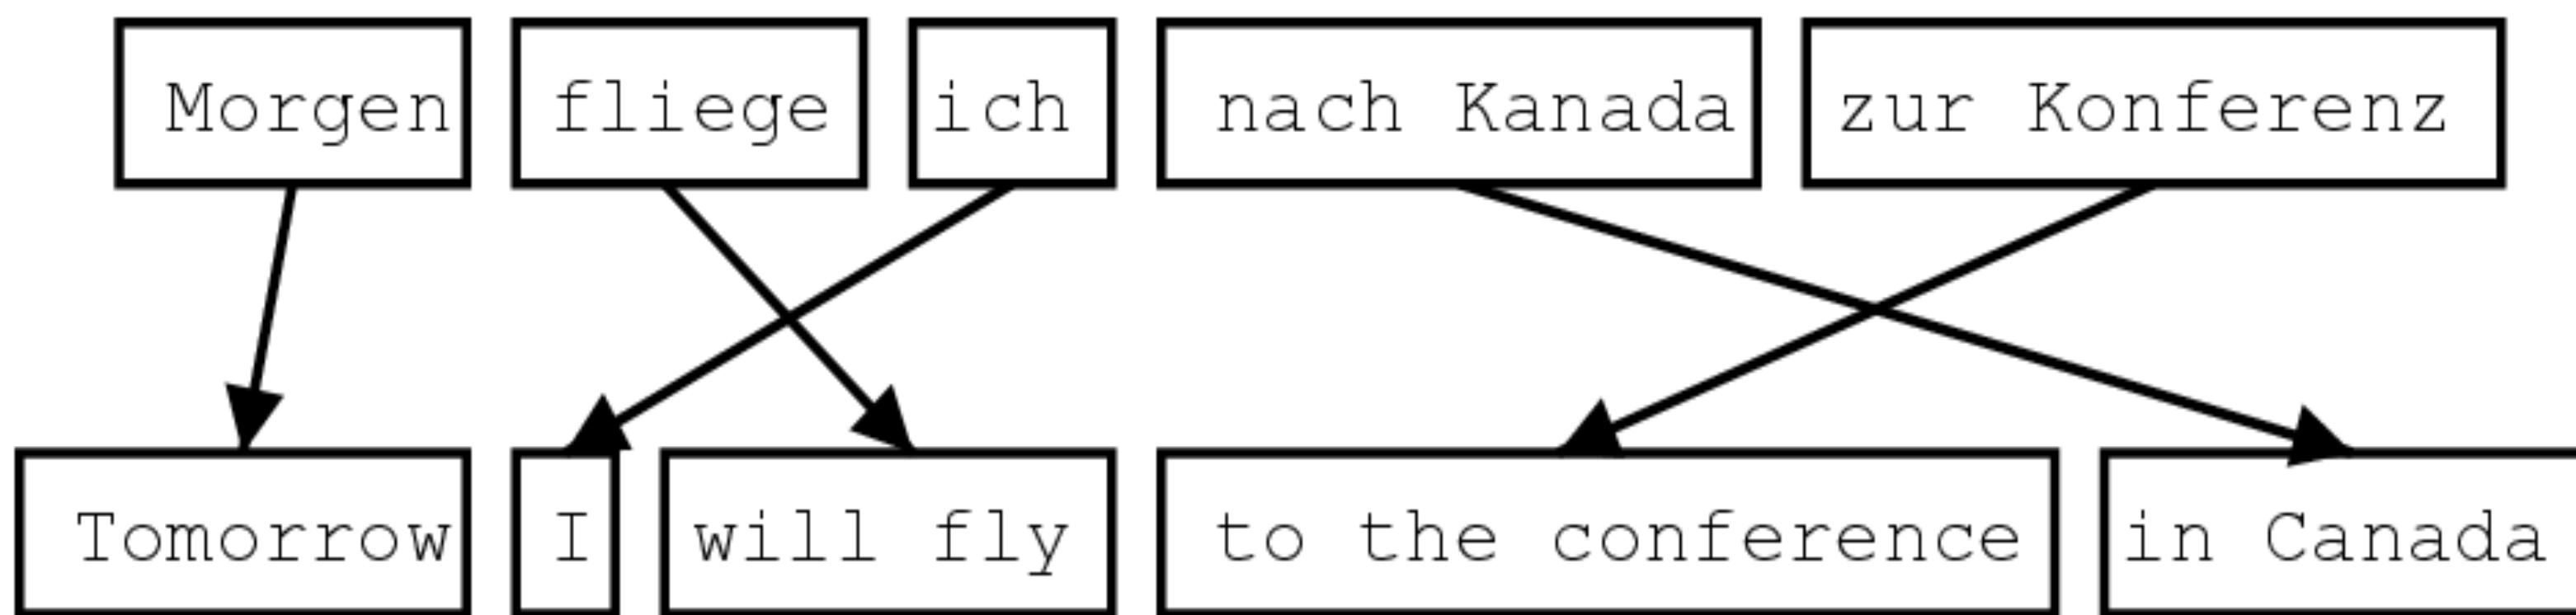
- ▶ Inputs:
 - ▶ Language model that scores $P(e_i|e_1, \dots, e_{i-1}) \approx P(e_i|e_{i-n-1}, \dots, e_{i-1})$
 - ▶ Phrase table: set of phrase pairs **(e, f)** with probabilities $P(f|e)$

Phrase-Based Decoding

- ▶ Inputs:
 - ▶ Language model that scores $P(e_i|e_1, \dots, e_{i-1}) \approx P(e_i|e_{i-n-1}, \dots, e_{i-1})$
 - ▶ Phrase table: set of phrase pairs **(e, f)** with probabilities $P(f|e)$
- ▶ What we want to find: **e** produced by a series of phrase-by-phrase translations from an input **f**, possibly with reordering:

Phrase-Based Decoding

- ▶ Inputs:
 - ▶ Language model that scores $P(e_i|e_1, \dots, e_{i-1}) \approx P(e_i|e_{i-n-1}, \dots, e_{i-1})$
 - ▶ Phrase table: set of phrase pairs **(e, f)** with probabilities $P(f|e)$
- ▶ What we want to find: **e** produced by a series of phrase-by-phrase translations from an input **f**, possibly with reordering:



Phrase lattices are big!

这 7人 中包括 来自 法国 和 俄罗斯 的 宇航 员 .

the	7 people	including	by some	and	the russian	the	the astronauts	,
it	7 people included	by france		and the	the russian		international astronautical	of rapporteur .
this	7 out	including the	from	the french	and the russian	the fifth		.
these	7 among	including from		the french and	of the russian	of	space	members .
that	7 persons	including from the		of france	and to	russian	of the	aerospace members .
	7 include	from the	of france and		russian	astronauts		. the
	7 numbers include	from france		and russian	of astronauts who			."
	7 populations include	those from france		and russian		astronauts .		
	7 deportees included	come from	france	and russia	in	astronautical	personnel	;
7 philtrum	including those from		france and	russia	a space		member	
	including representatives from		france and the	russia	astronaut			
	include	came from	france and russia		by cosmonauts			
	include representatives from	french	and russia		cosmonauts			
	include	came from france	and russia 's		cosmonauts .			
	includes	coming from	french and	russia 's	cosmonaut			
			french and russian	's	astronavigation	member .		
			french	and russia	astronauts			
				and russia 's			special rapporteur	
				, and russia			rapporteur	
				, and russia			rapporteur .	
				, and russia				
				or	russia 's			

Phrase-Based Decoding

- ▶ Input

lo haré | rápidamente |.

- ▶ Translations

I'll do it | quickly |.

quickly | I'll do it |.

- ▶ Decoding objective (for 3-gram LM)

$$\arg \max_{\mathbf{e}} [P(\mathbf{f}|\mathbf{e}) \cdot P(\mathbf{e})]$$

$$\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, f \rangle} P(\bar{f}|\bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i|e_{i-1}, e_{i-2}) \right]$$

The decoder...

tries different segmentations,

translates phrase by phrase,

and considers reorderings.

Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
				<u>slap</u>		<u>the</u>		
						<u>the</u>		
							<u>the witch</u>	

$$\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, \bar{f} \rangle} P(\bar{f} | \bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i | e_{i-1}, e_{i-2}) \right]$$

Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

- If we translate with beam search, what state do we need to keep in the beam?

$$\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, \bar{f} \rangle} P(\bar{f} | \bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i | e_{i-1}, e_{i-2}) \right]$$

Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
				<u>slap</u>		<u>the</u>		
						<u>the</u>		
							<u>the witch</u>	

- ▶ If we translate with beam search, what state do we need to keep in the beam?

- ▶ What have we translated so far?

$$\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, \bar{f} \rangle} P(\bar{f} | \bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i | e_{i-1}, e_{i-2}) \right]$$

Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

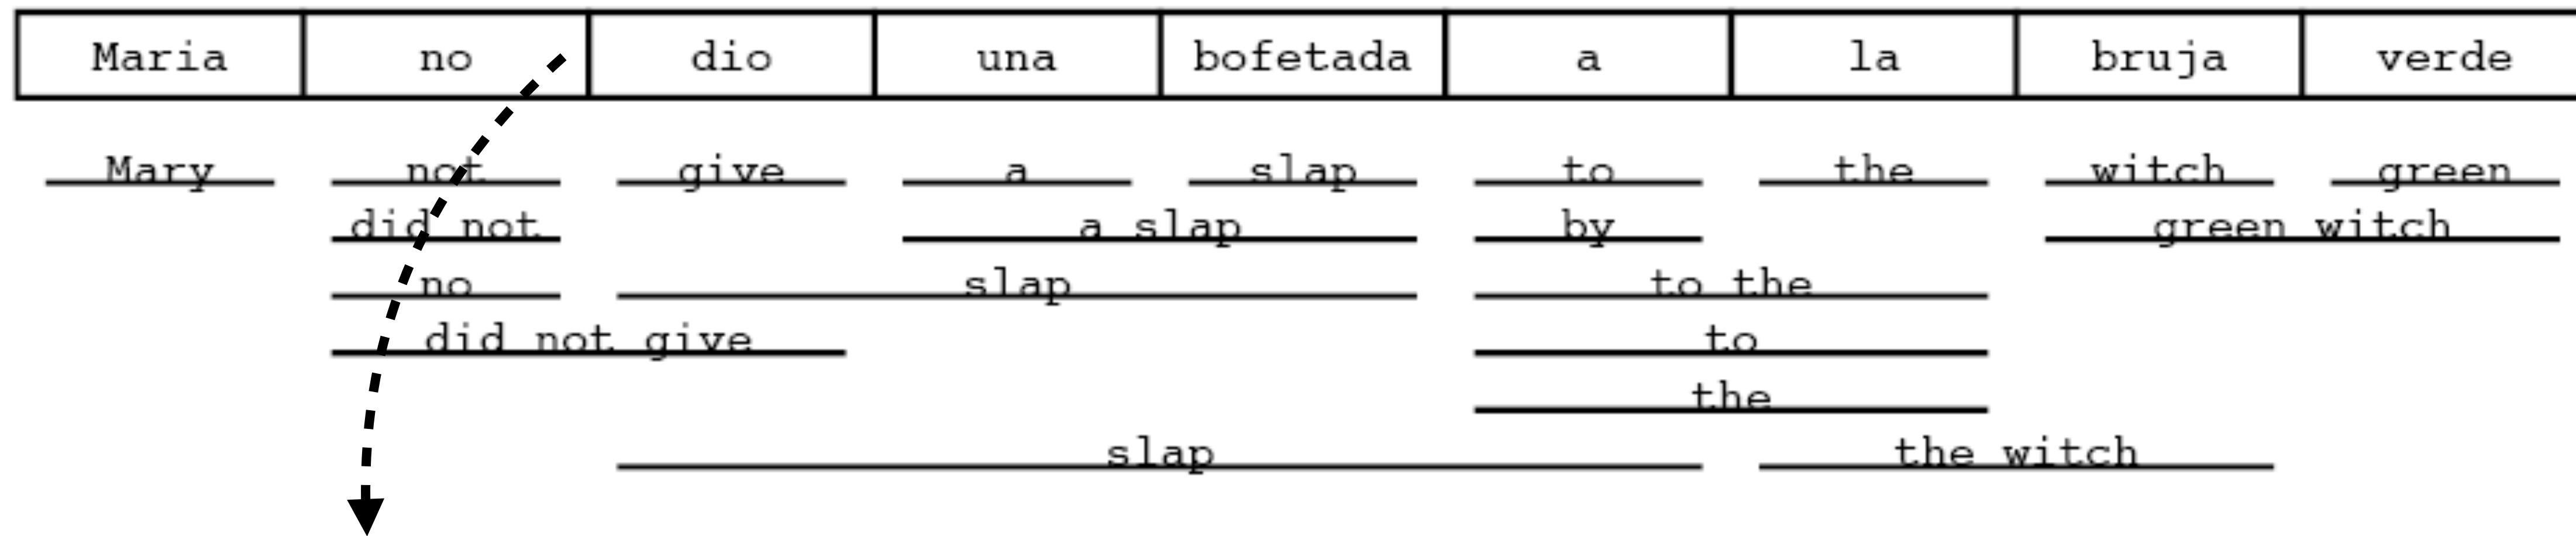
- ▶ If we translate with beam search, what state do we need to keep in the beam?
- ▶ What have we translated so far? $\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, \bar{f} \rangle} P(\bar{f} | \bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i | e_{i-1}, e_{i-2}) \right]$
- ▶ What words have we produced so far?

Monotonic Translation

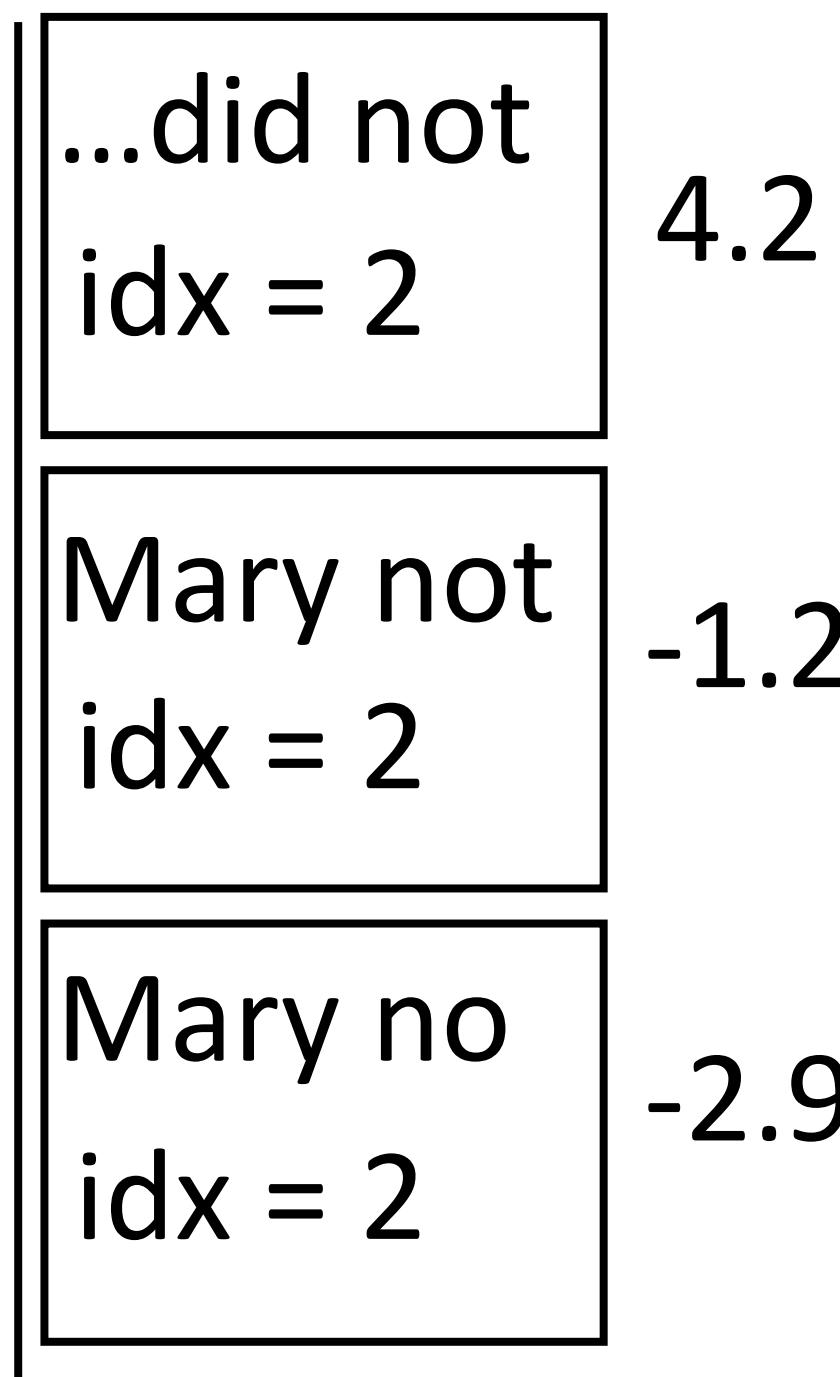
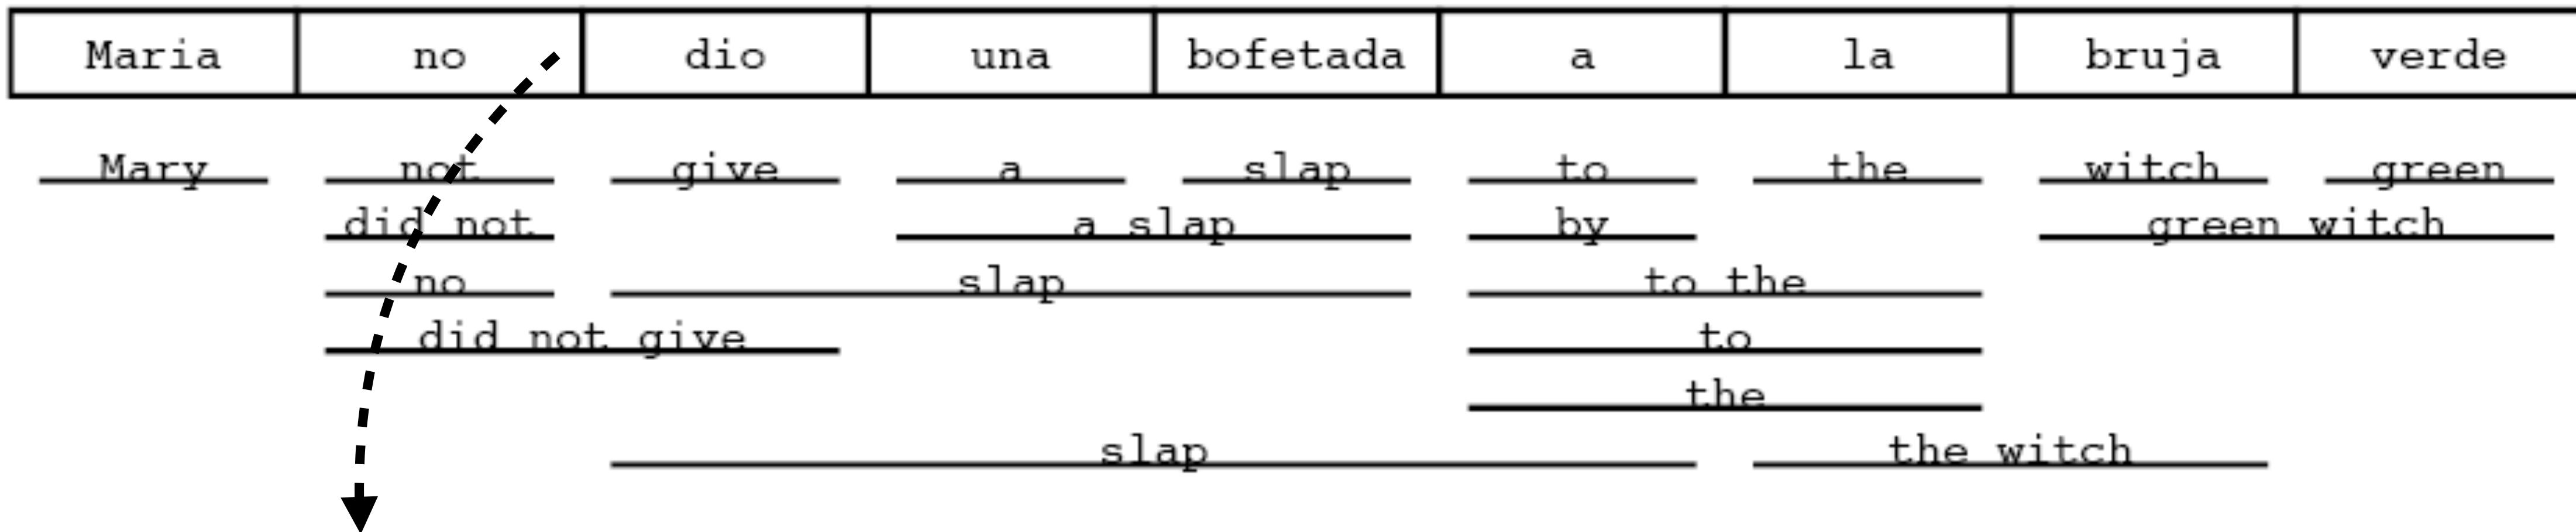
Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

- ▶ If we translate with beam search, what state do we need to keep in the beam?
- ▶ What have we translated so far? $\arg \max_{\mathbf{e}} \left[\prod_{\langle \bar{e}, \bar{f} \rangle} P(\bar{f} | \bar{e}) \cdot \prod_{i=1}^{|\mathbf{e}|} P(e_i | e_{i-1}, e_{i-2}) \right]$
- ▶ What words have we produced so far?
- ▶ When using a 3-gram LM, only need to remember the last 2 words!

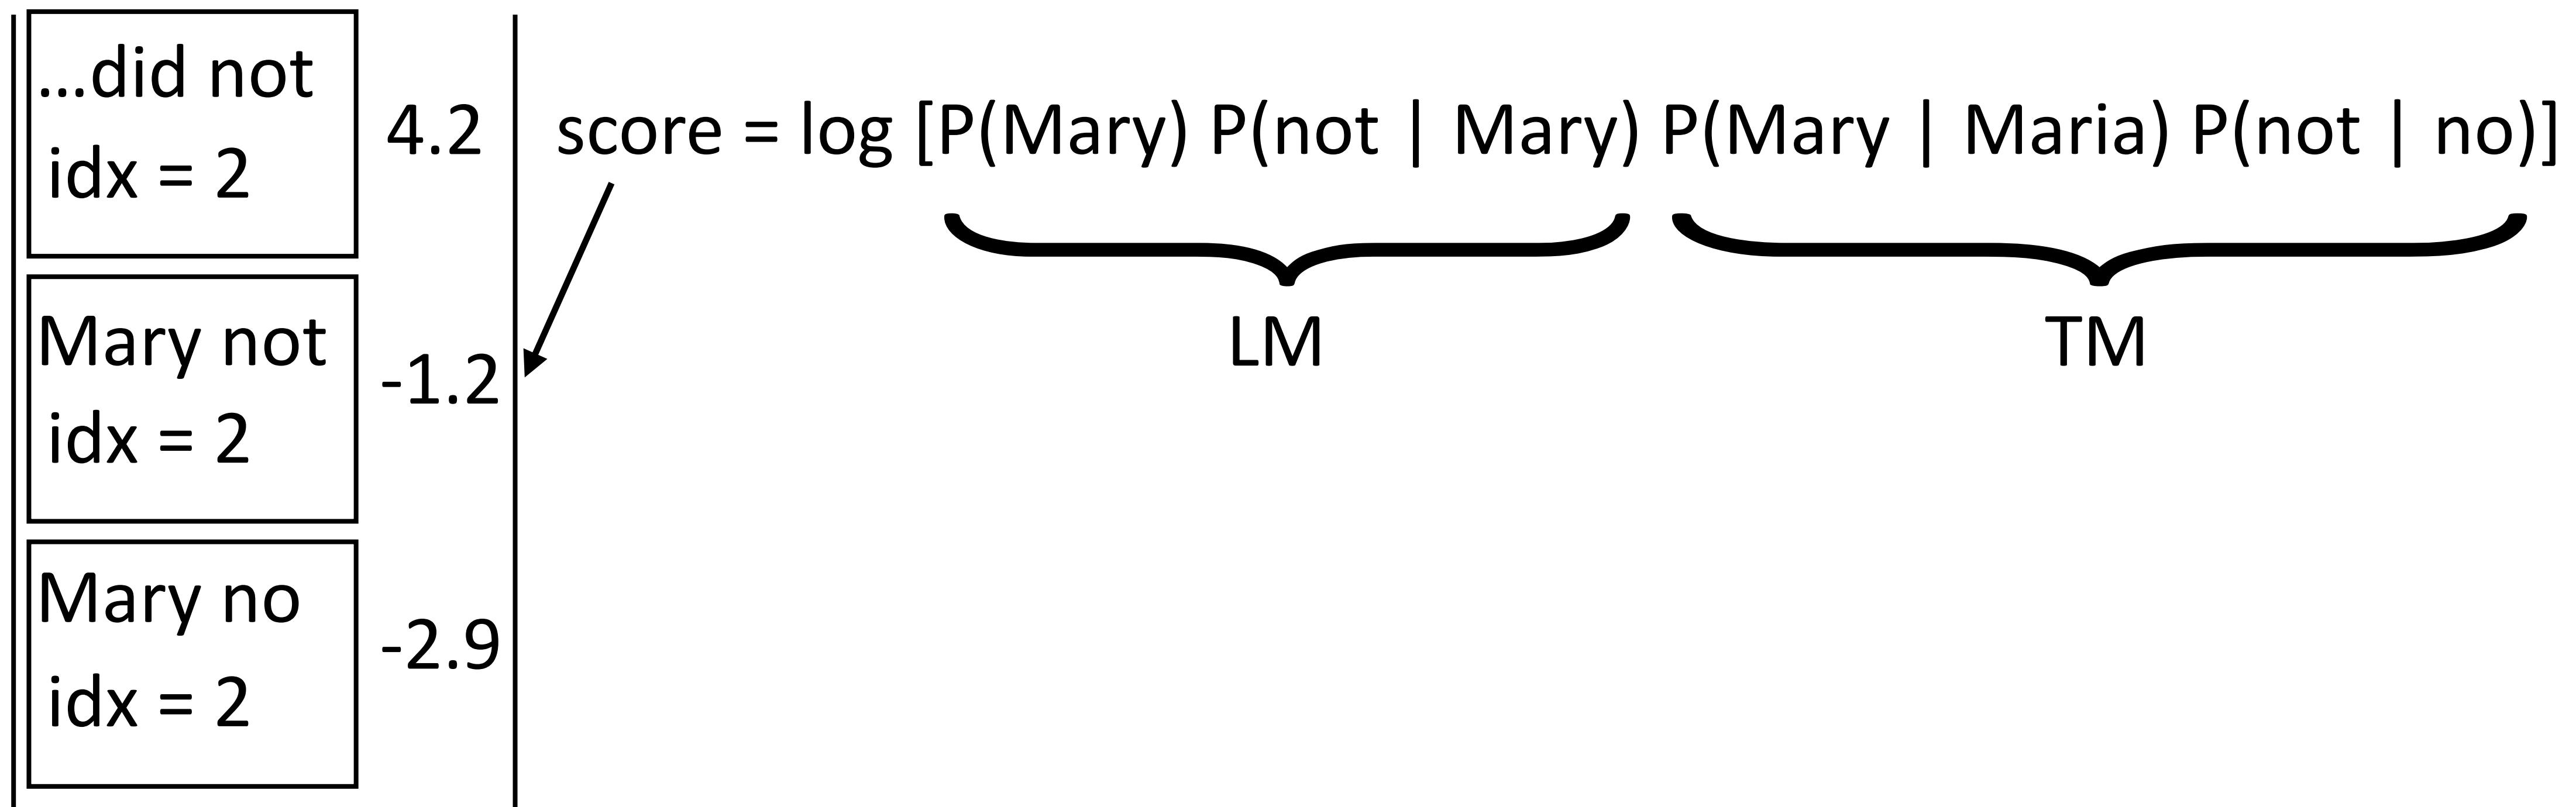
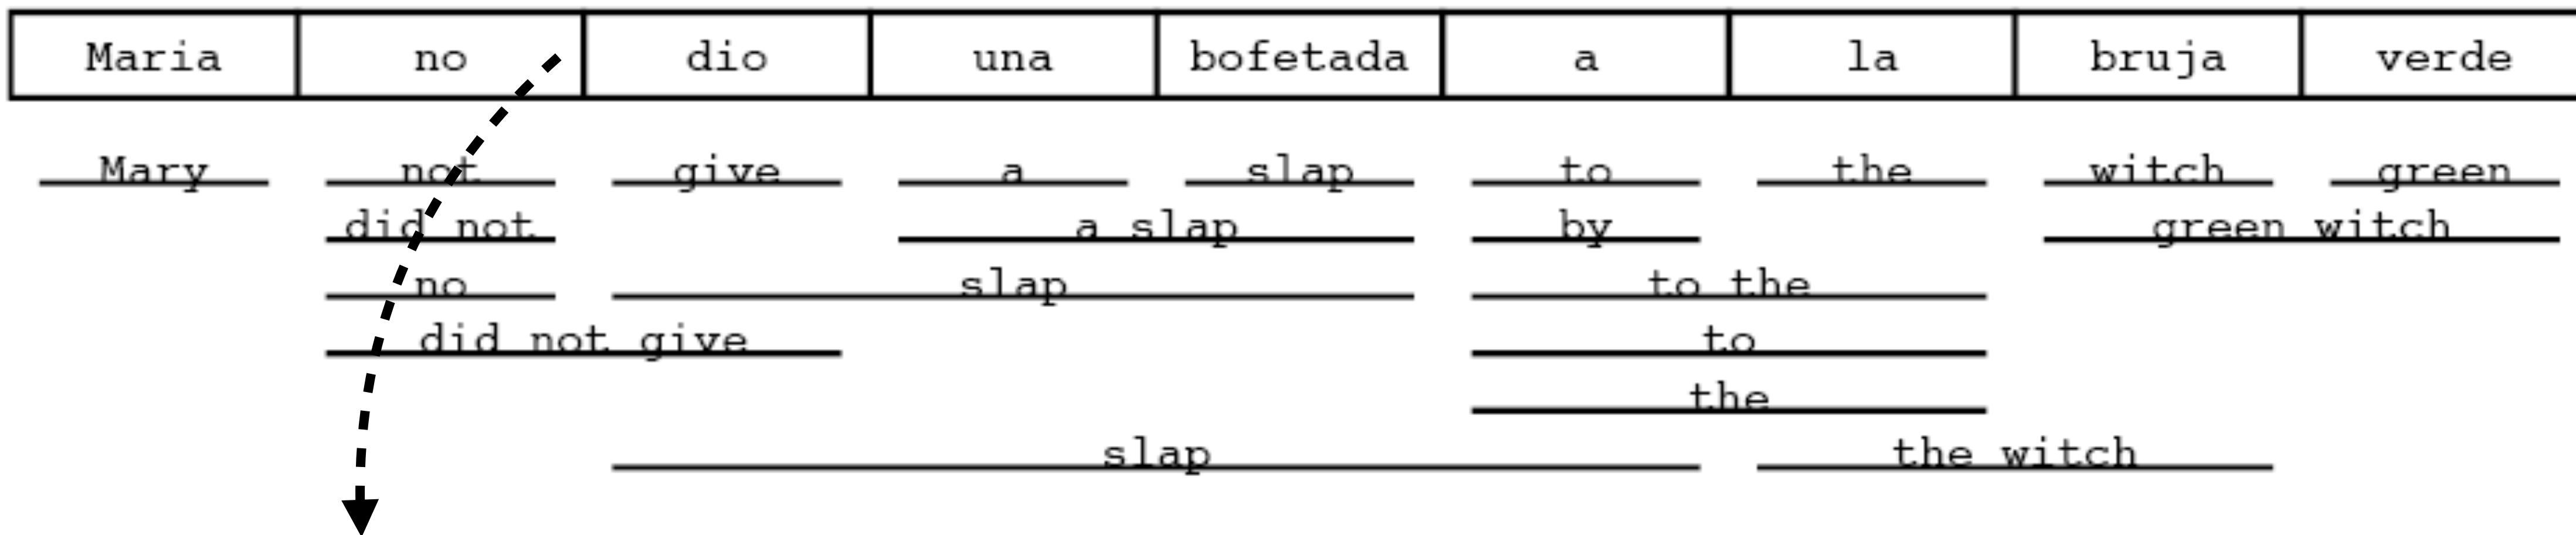
Monotonic Translation



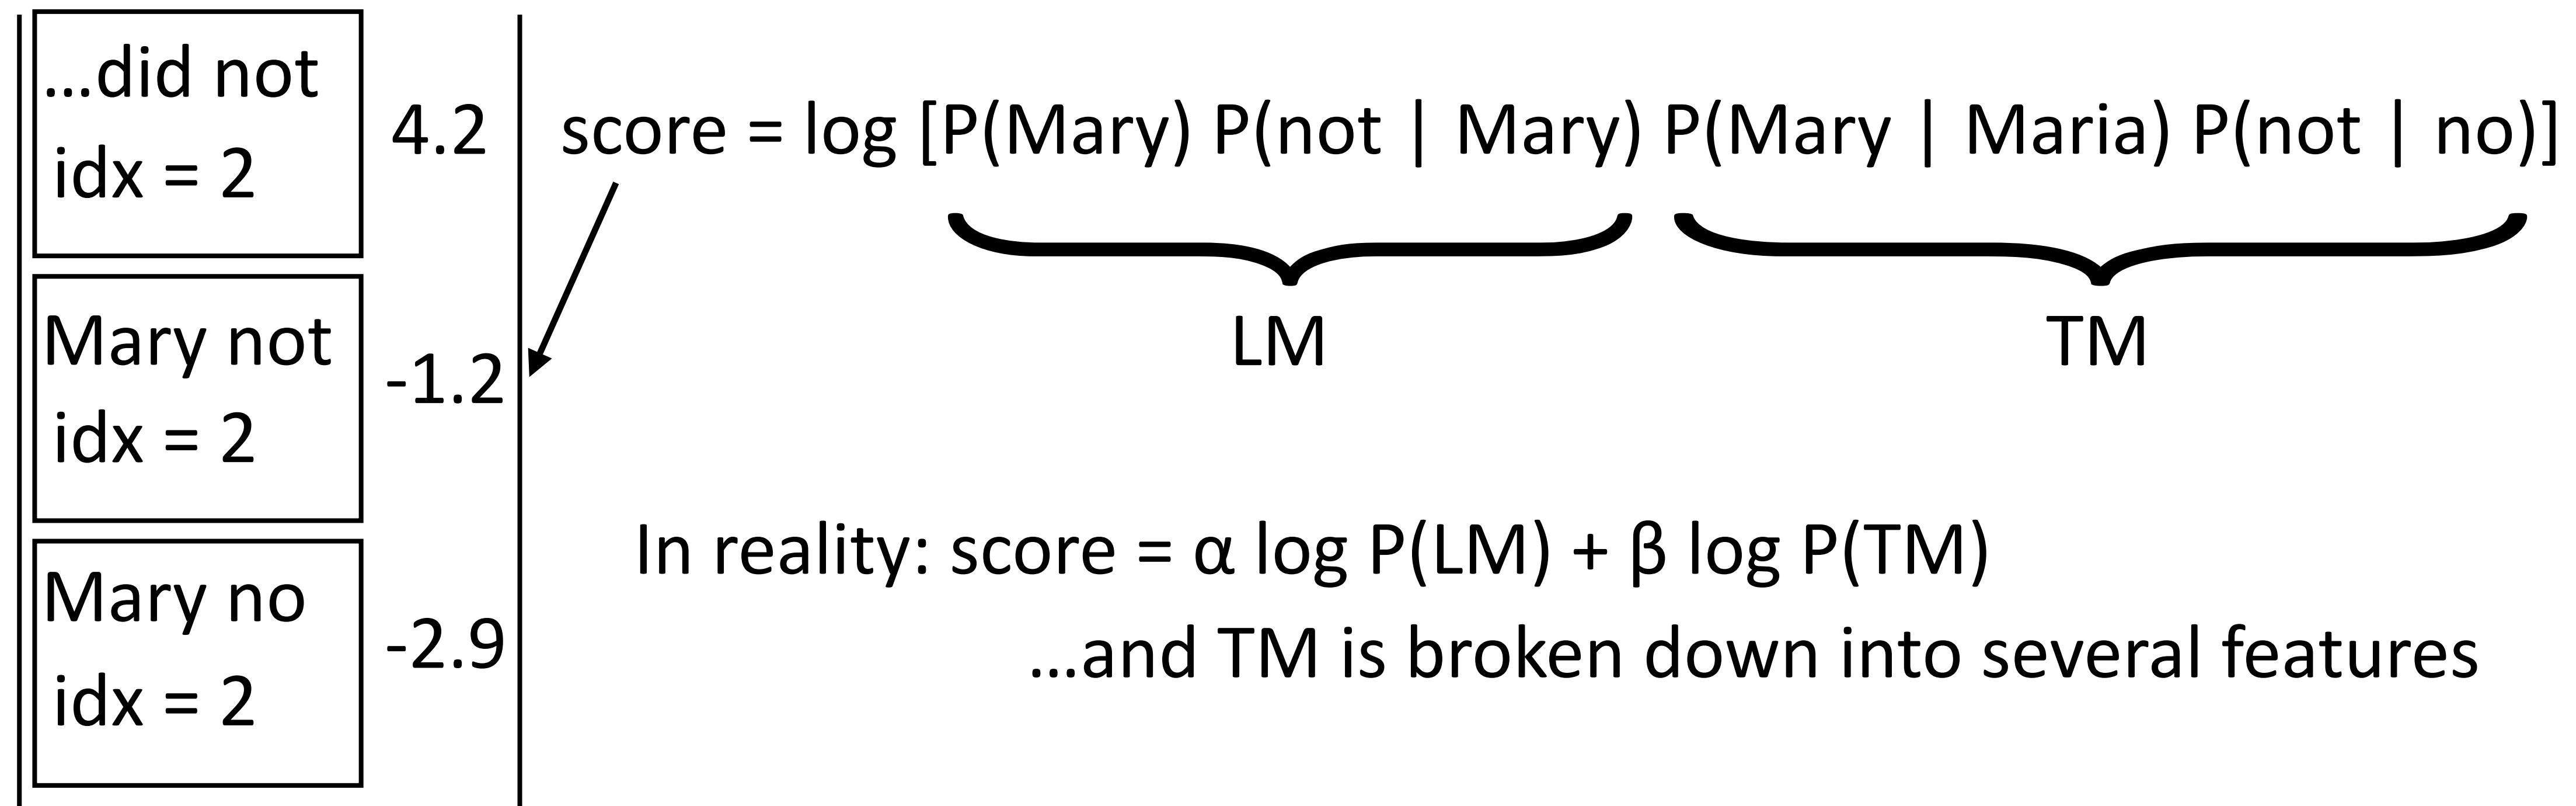
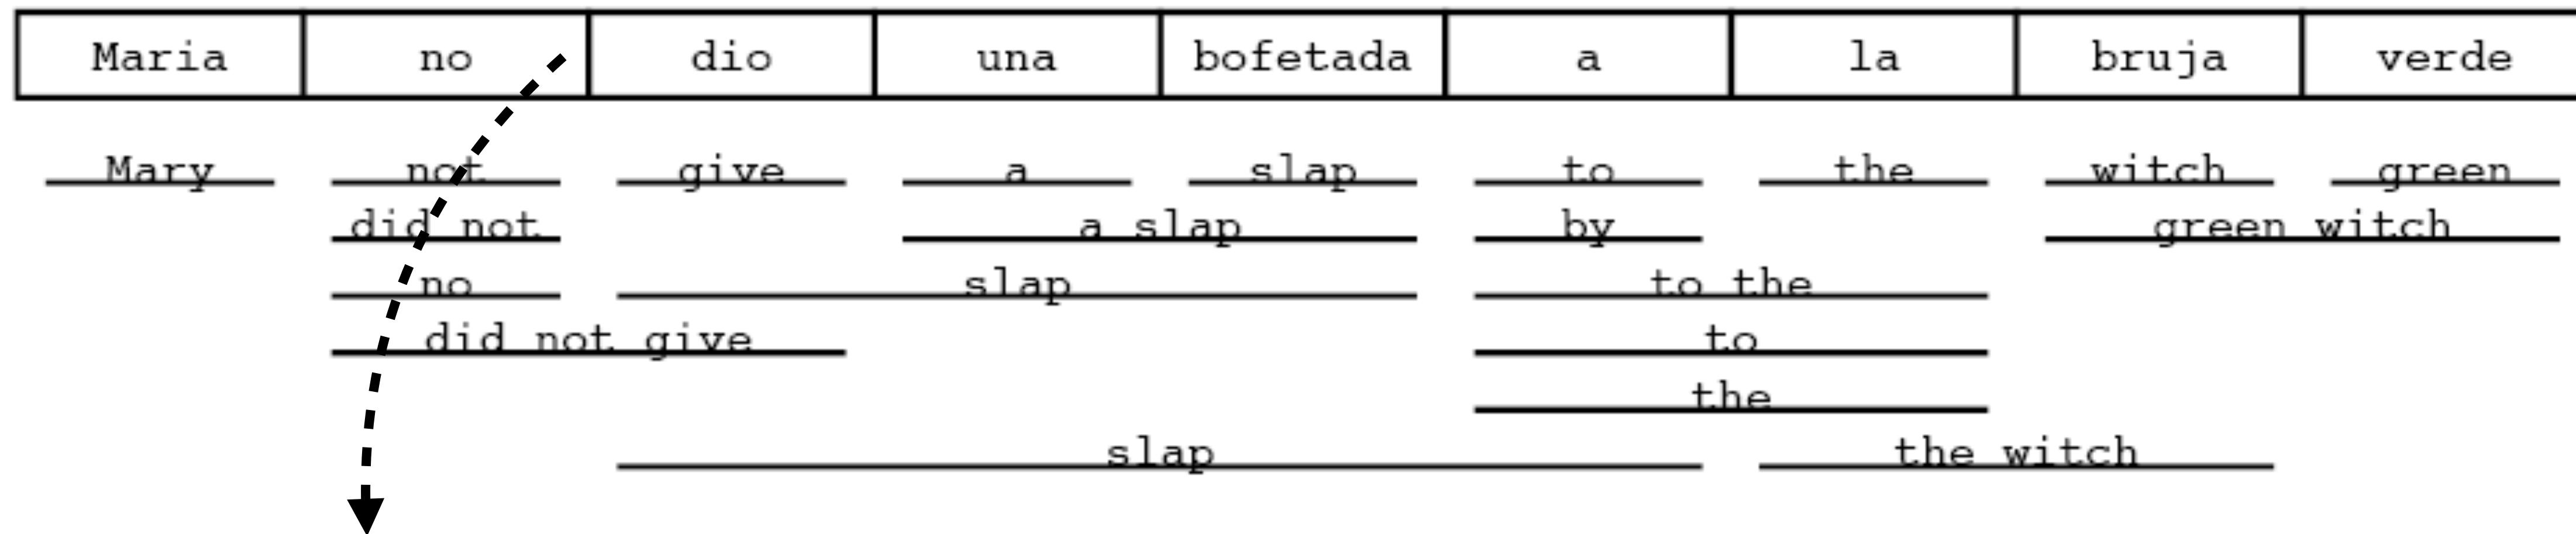
Monotonic Translation



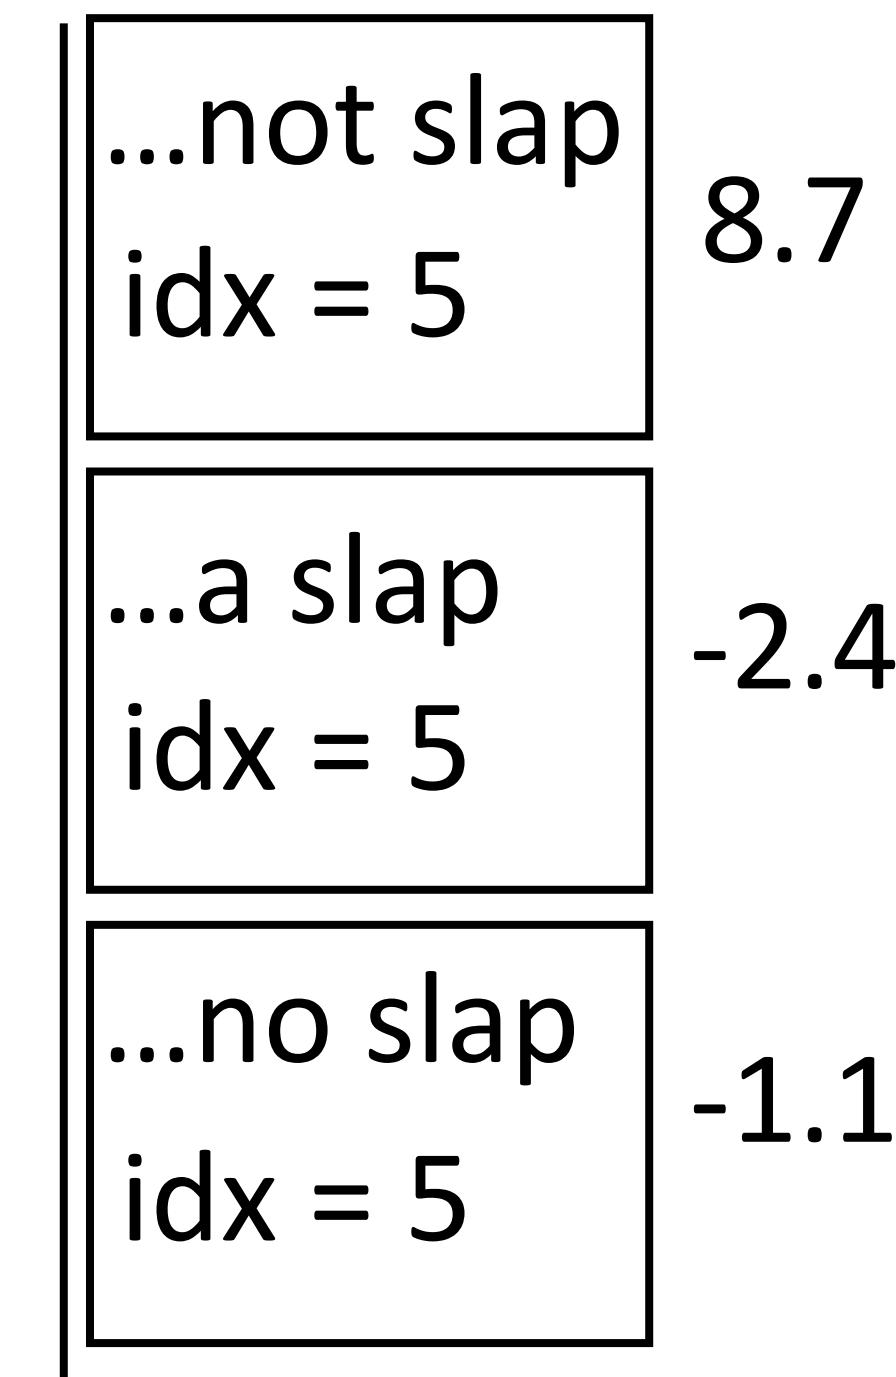
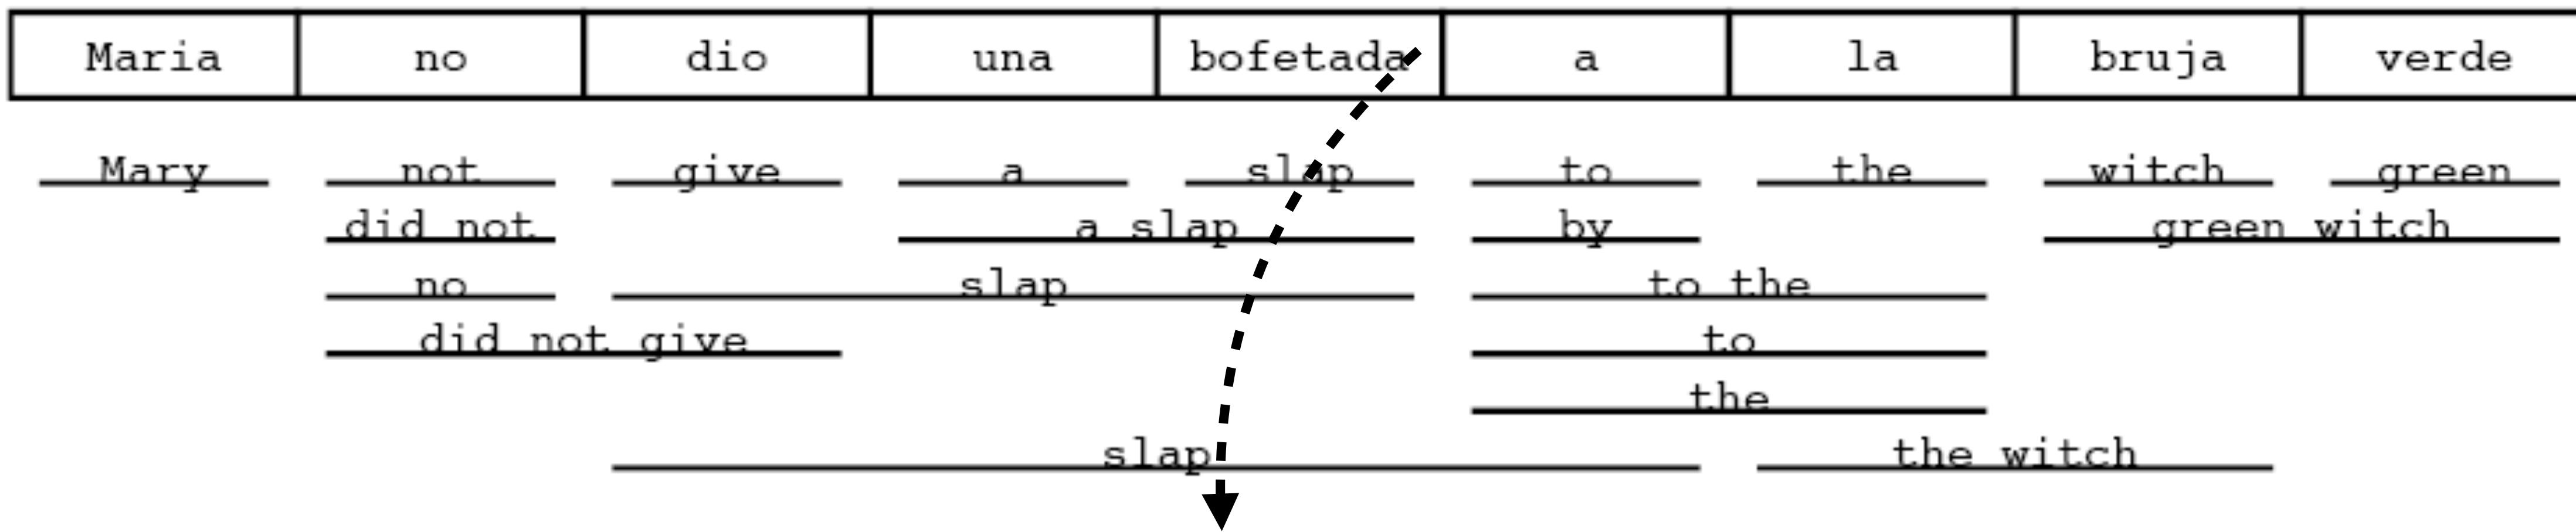
Monotonic Translation



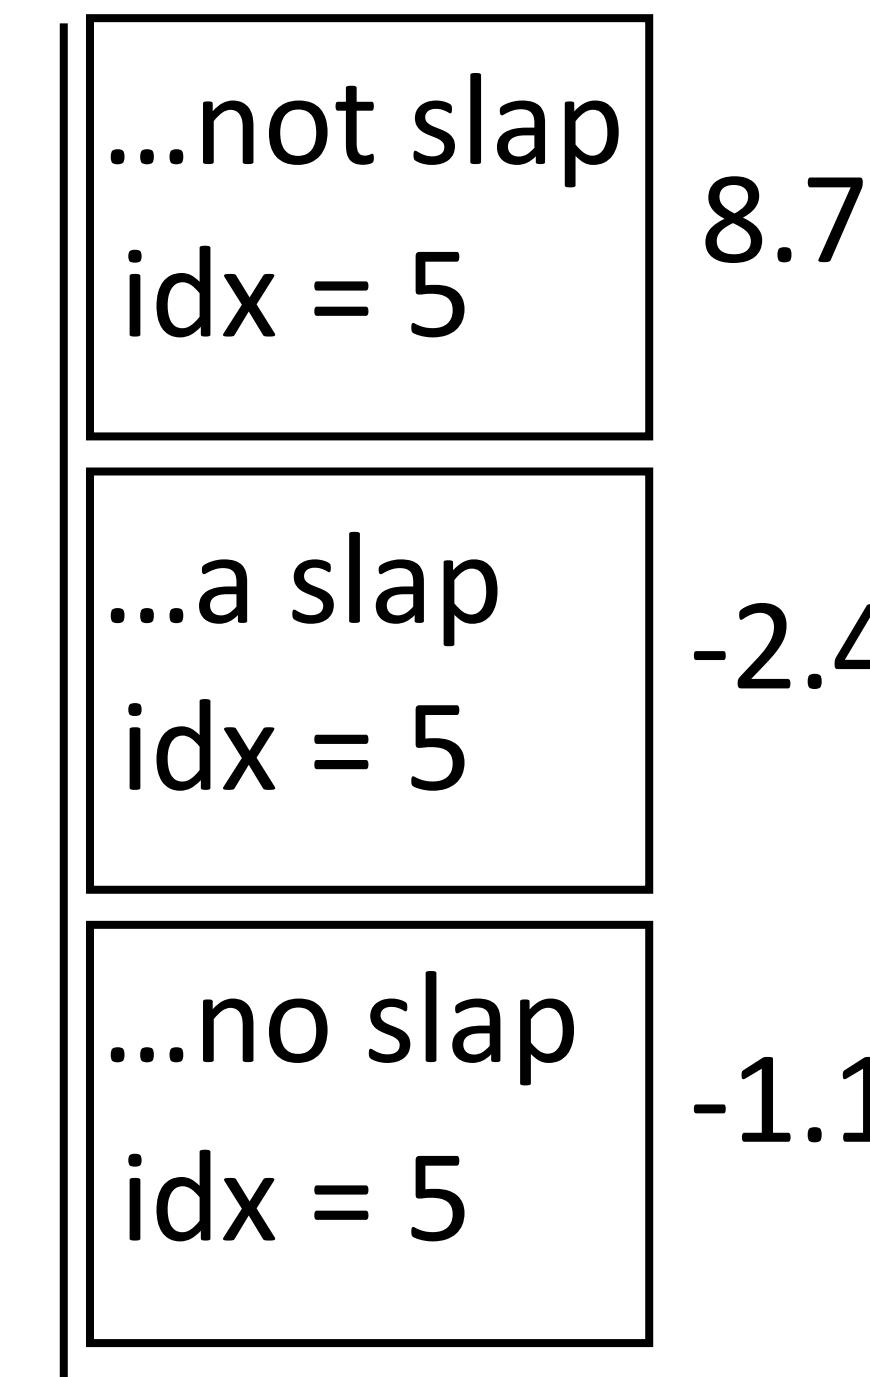
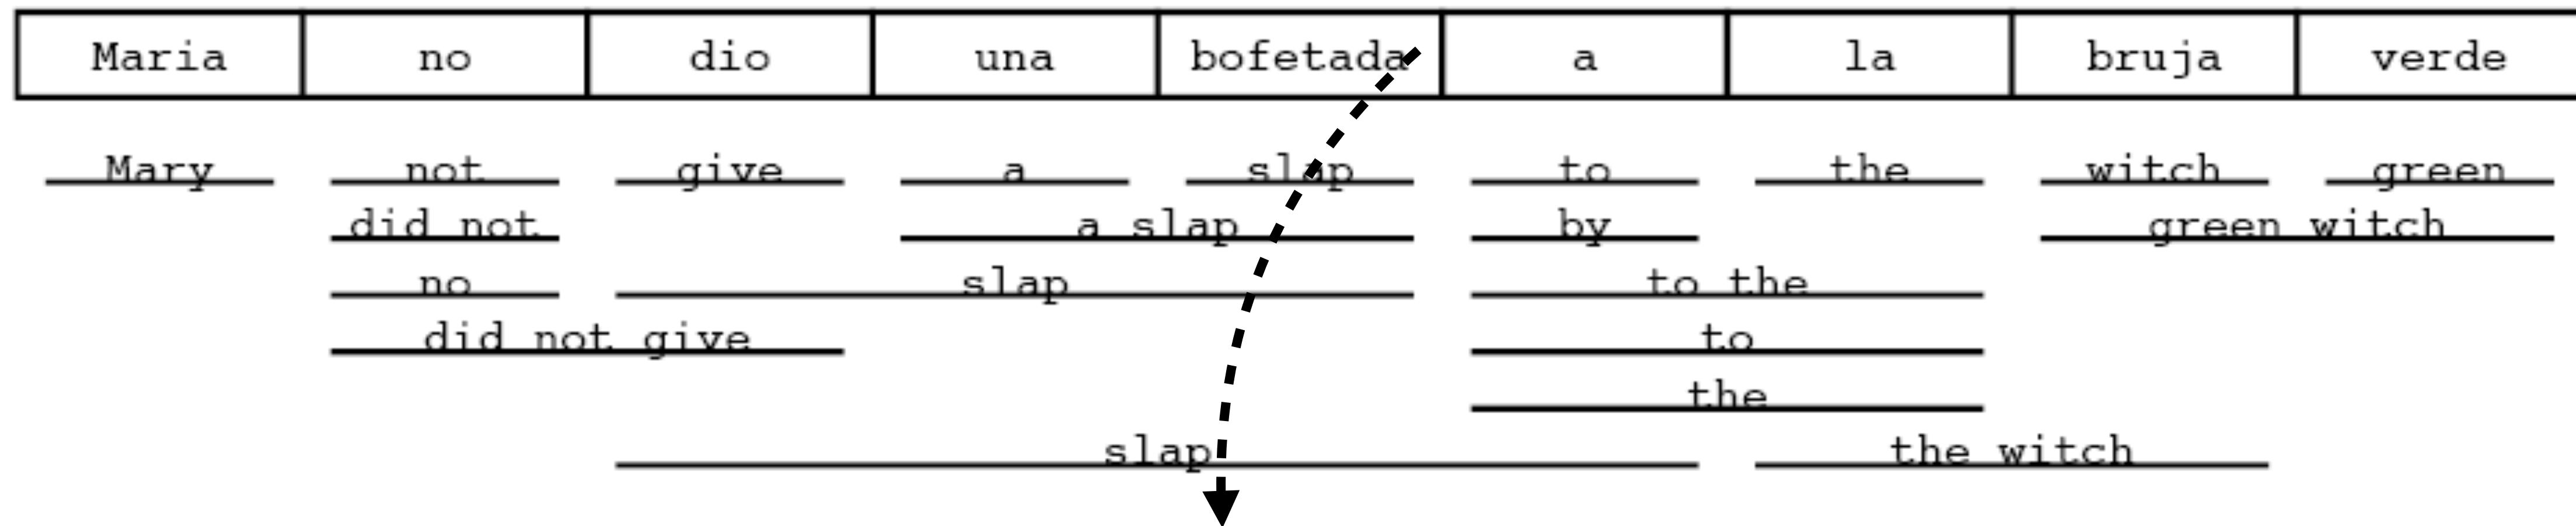
Monotonic Translation



Monotonic Translation

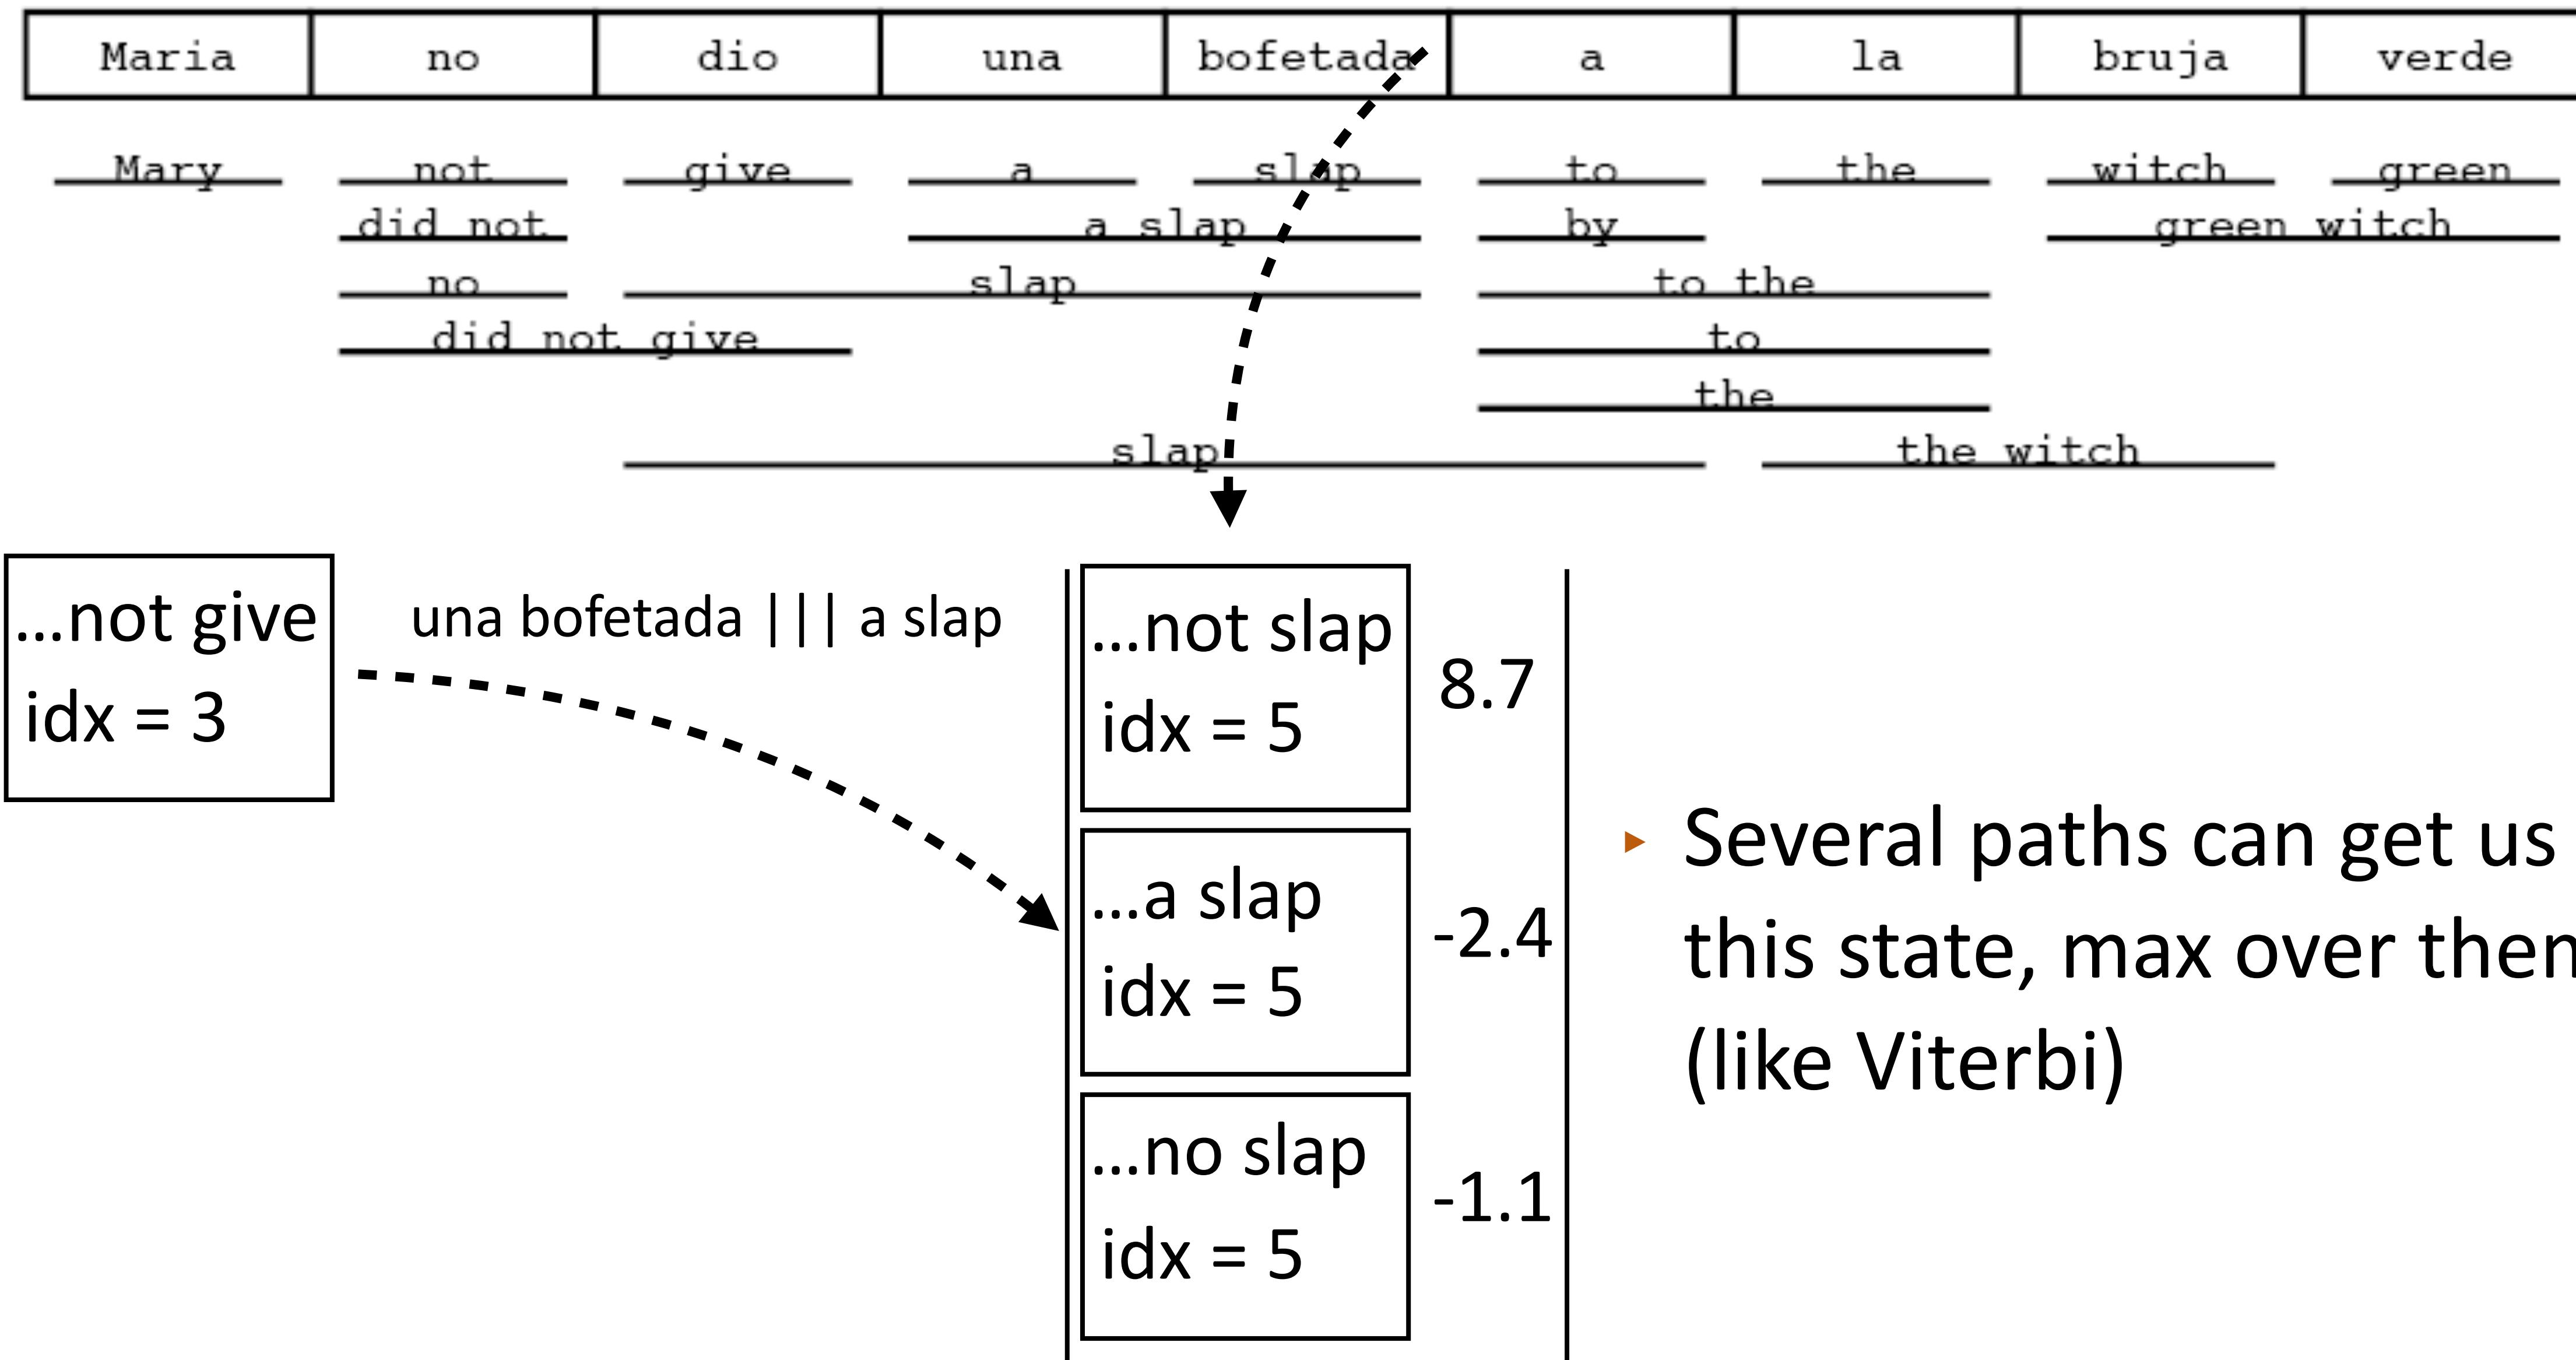


Monotonic Translation

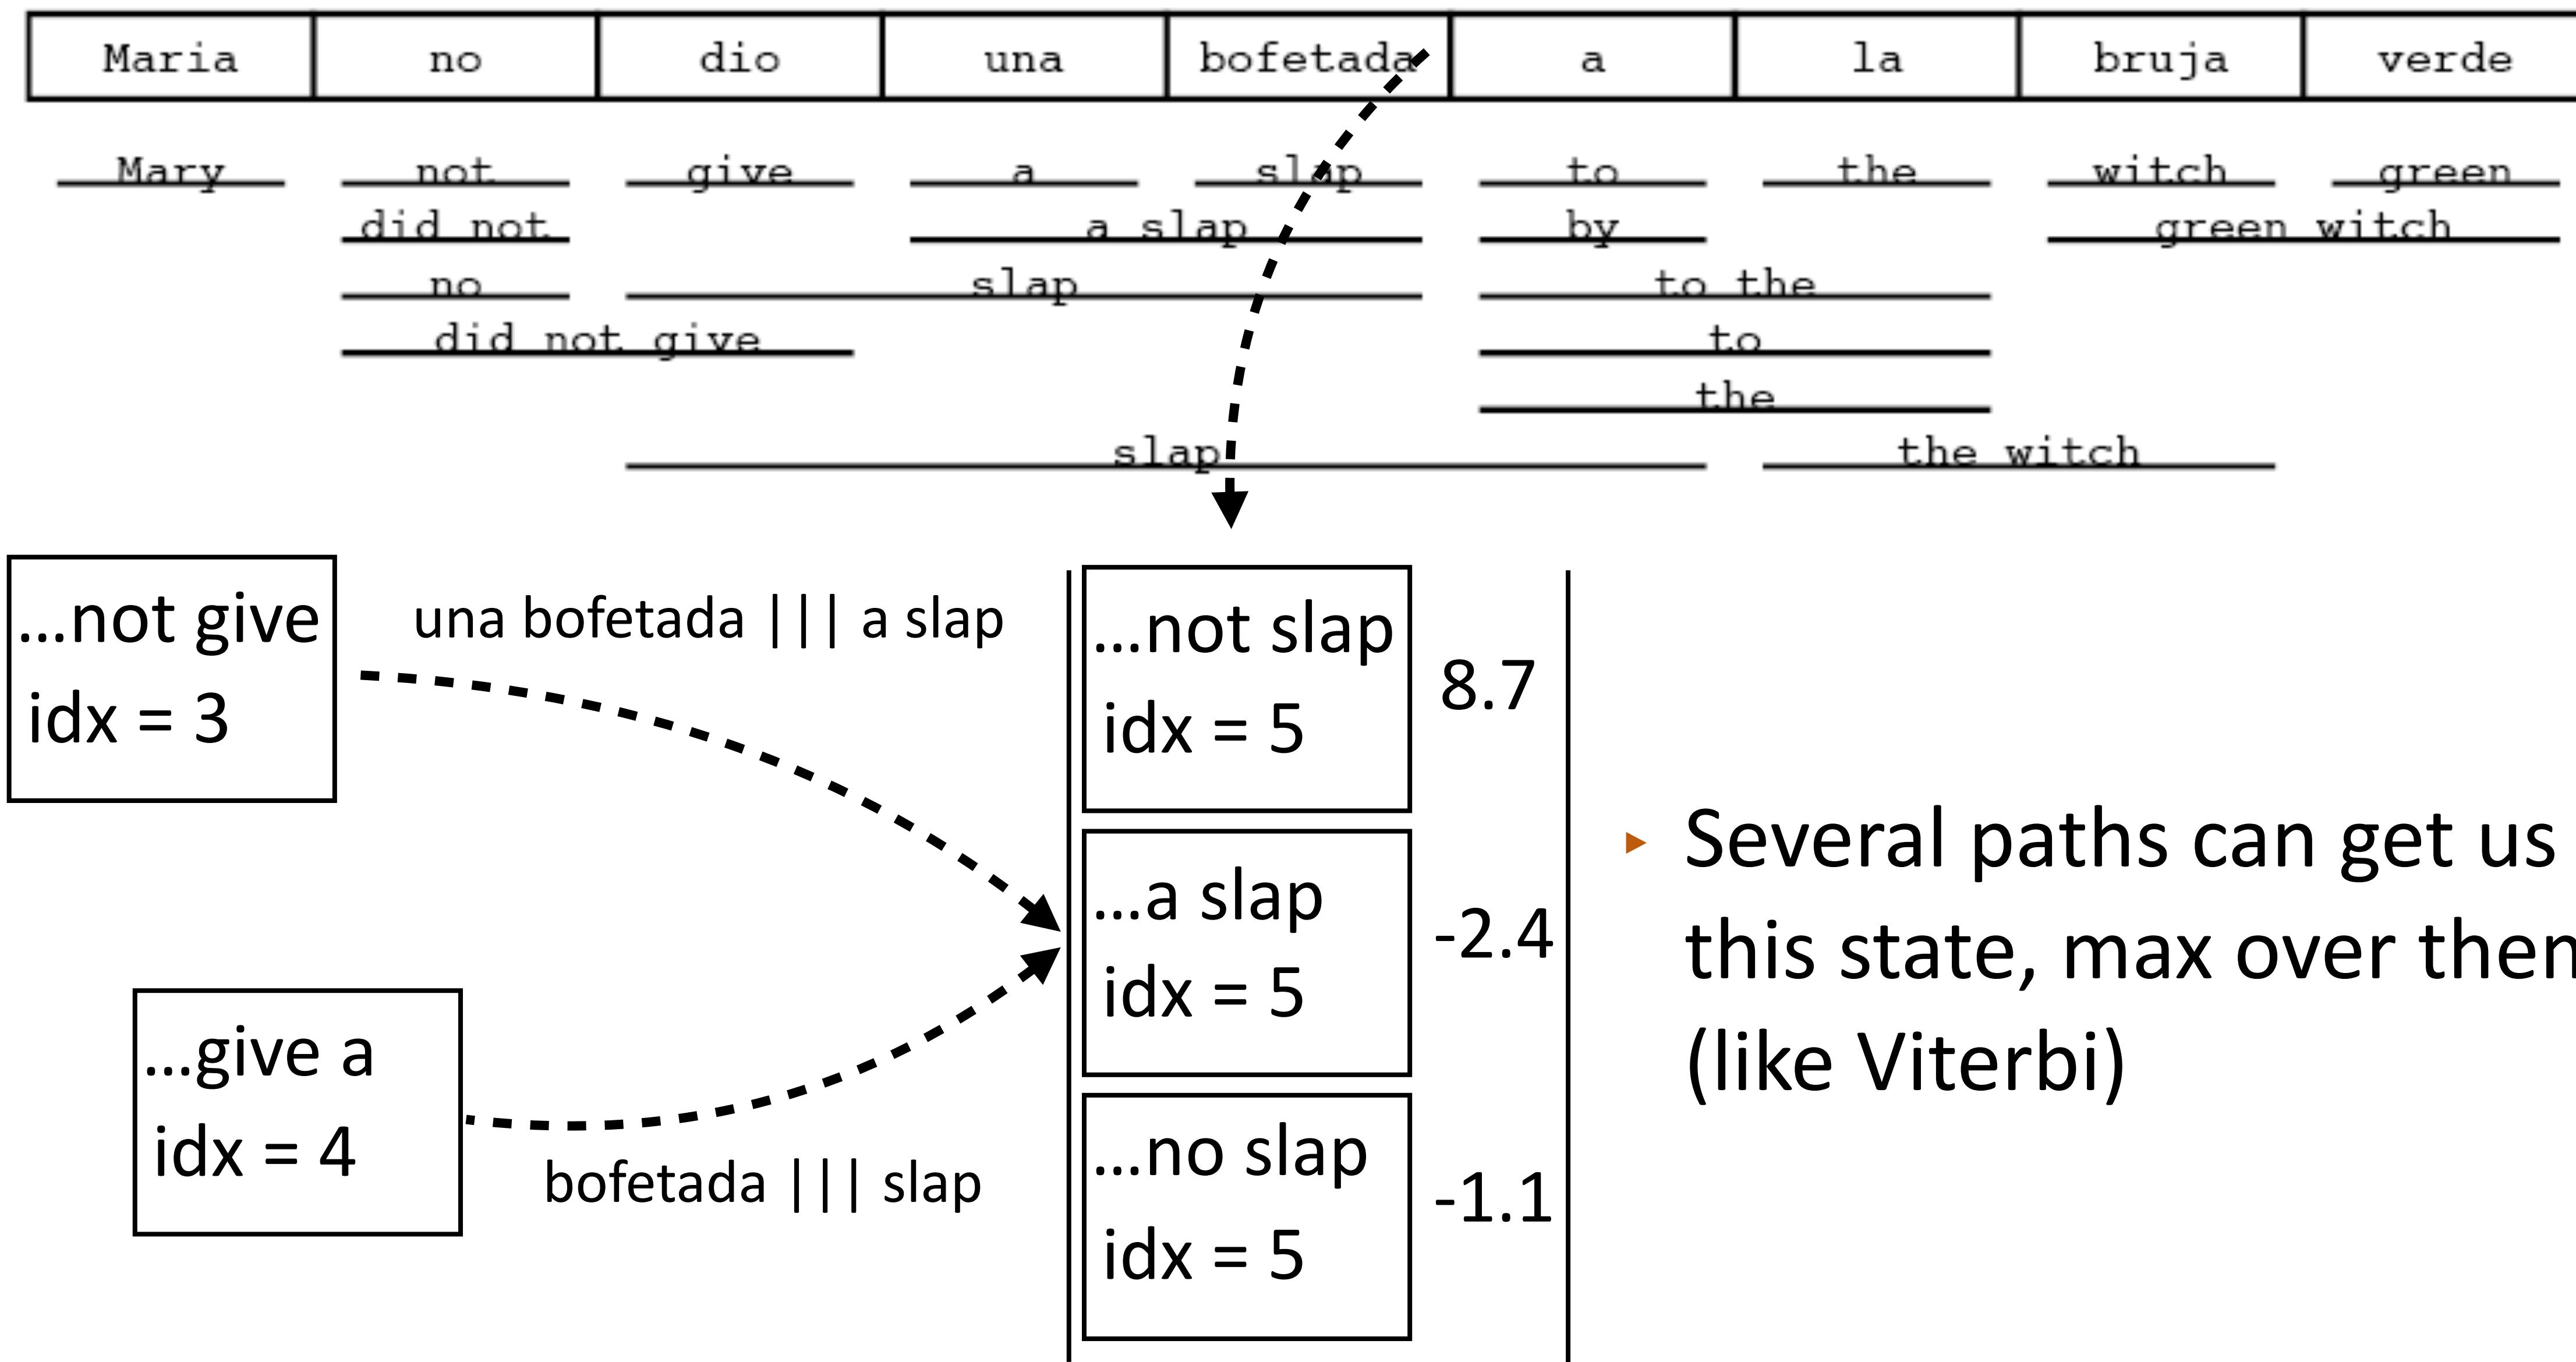


- ▶ Several paths can get us to this state, max over them (like Viterbi)

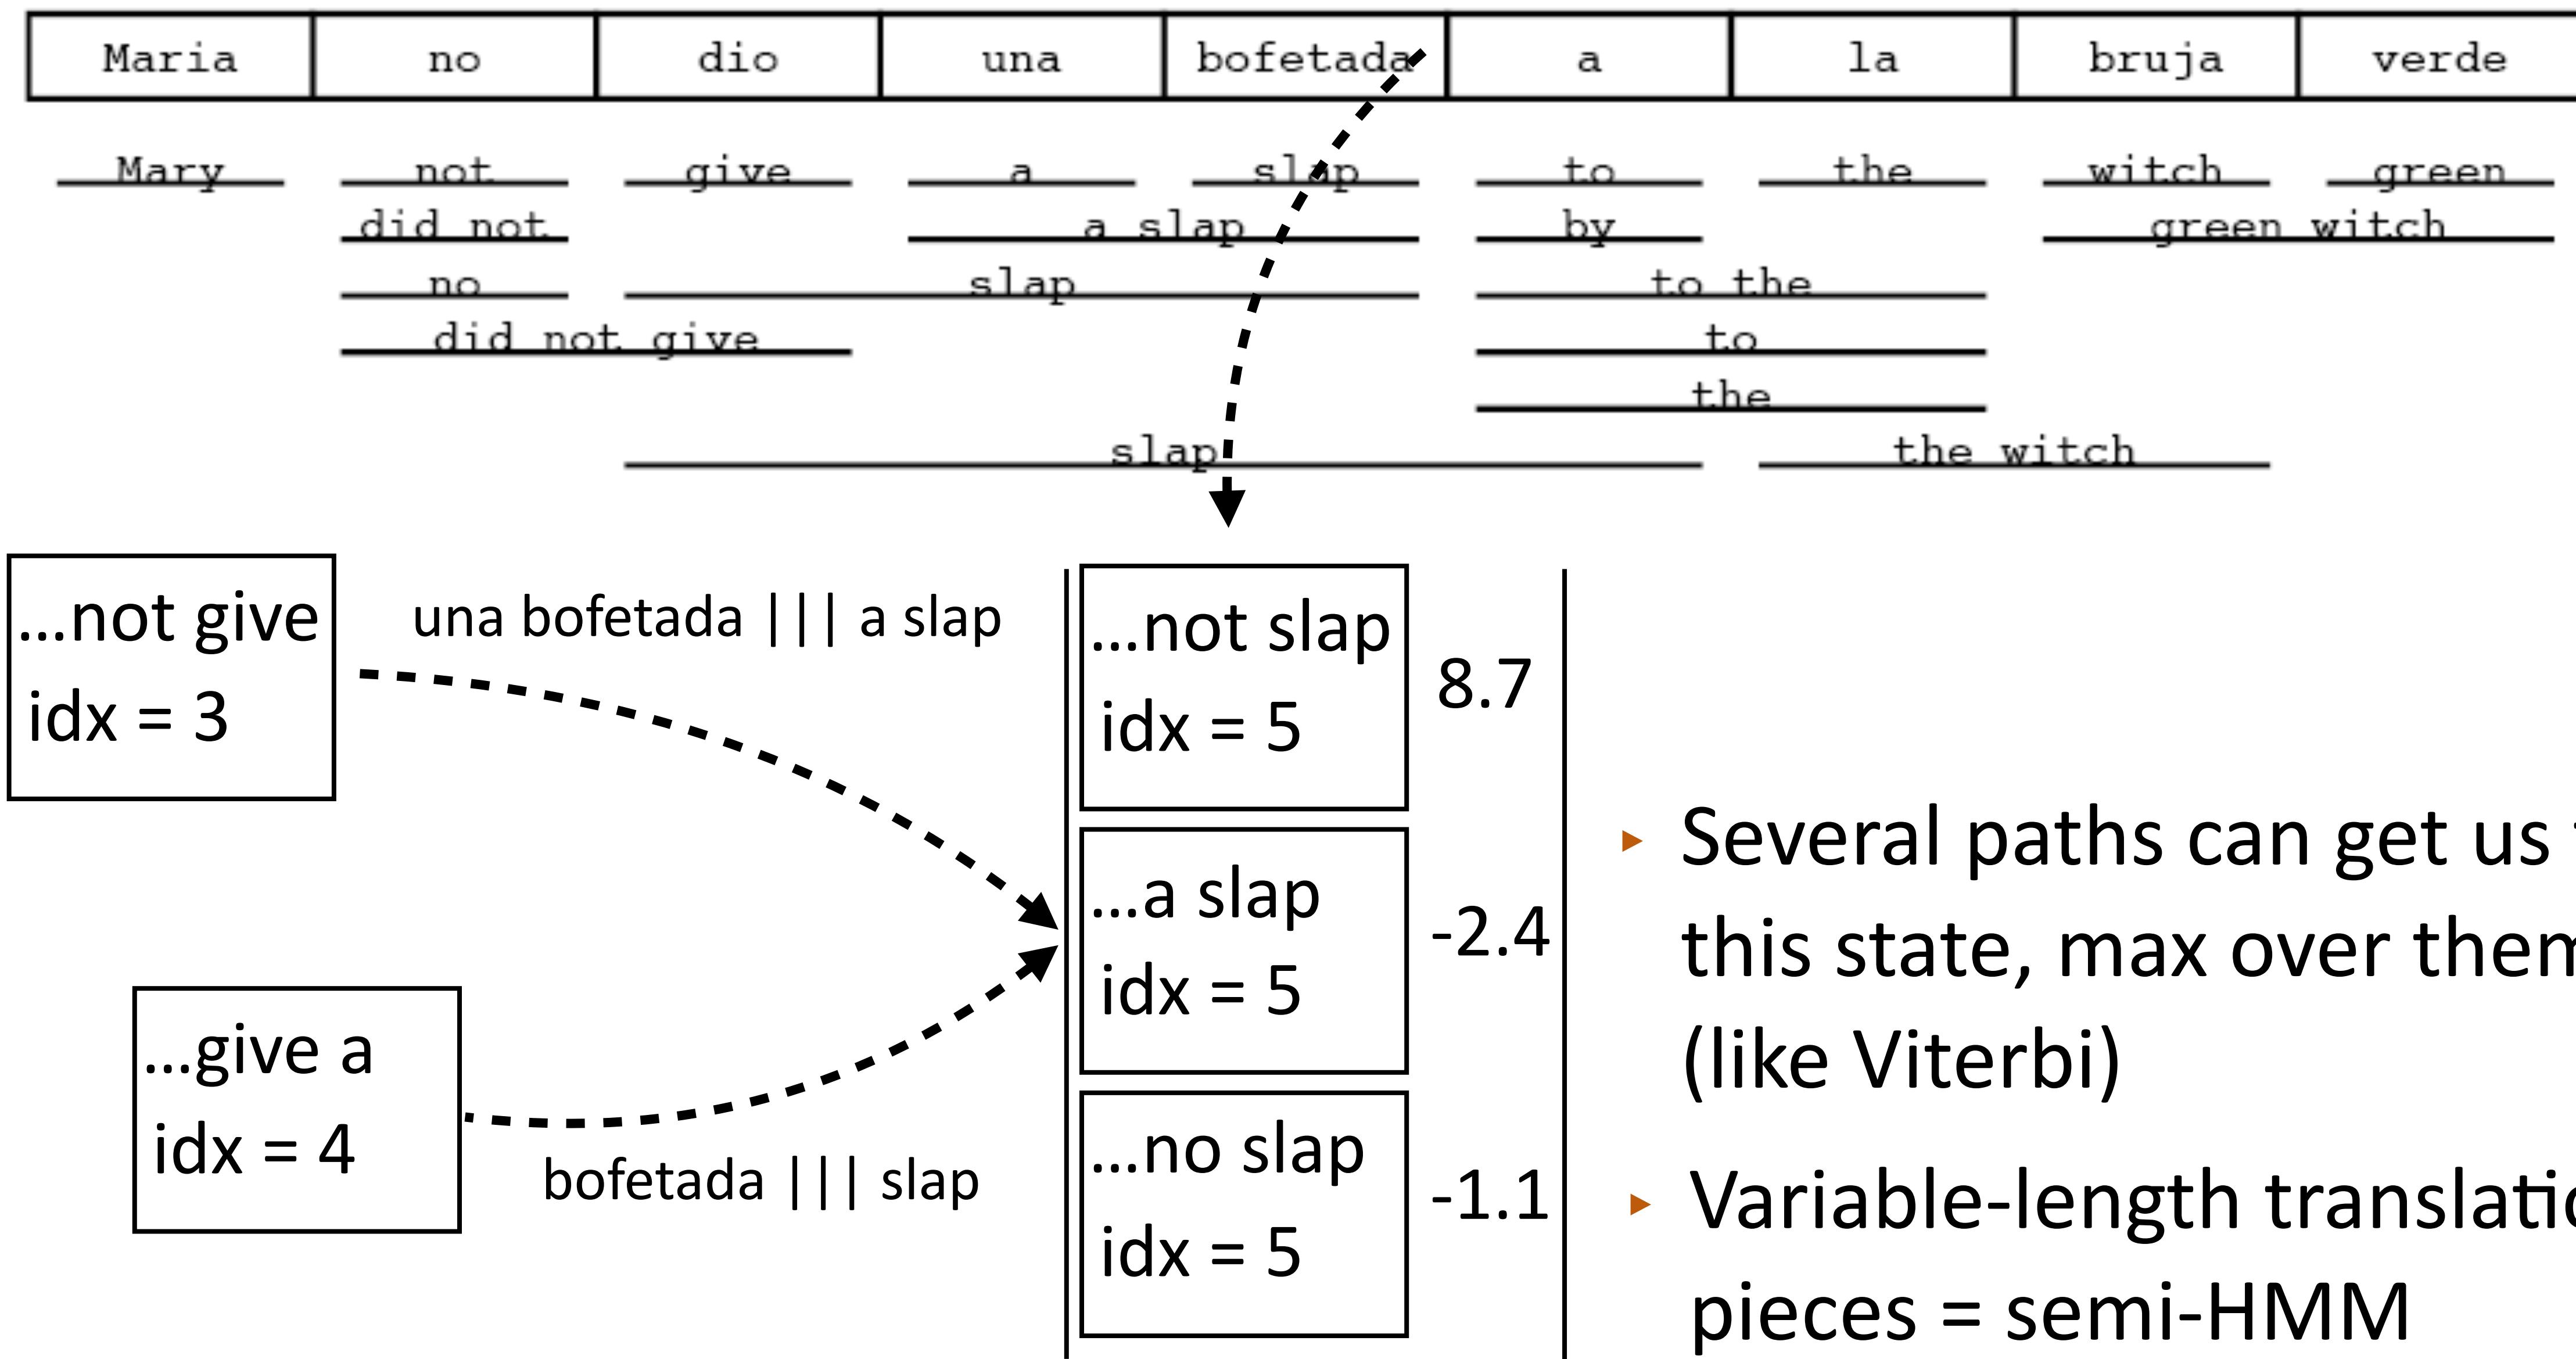
Monotonic Translation



Monotonic Translation



Monotonic Translation



Non-Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
				<u>slap</u>		<u>the</u>		
						<u>the</u>		
							<u>the witch</u>	

- ▶ Non-monotonic translation: can visit source sentence “out of order”

Non-Monotonic Translation

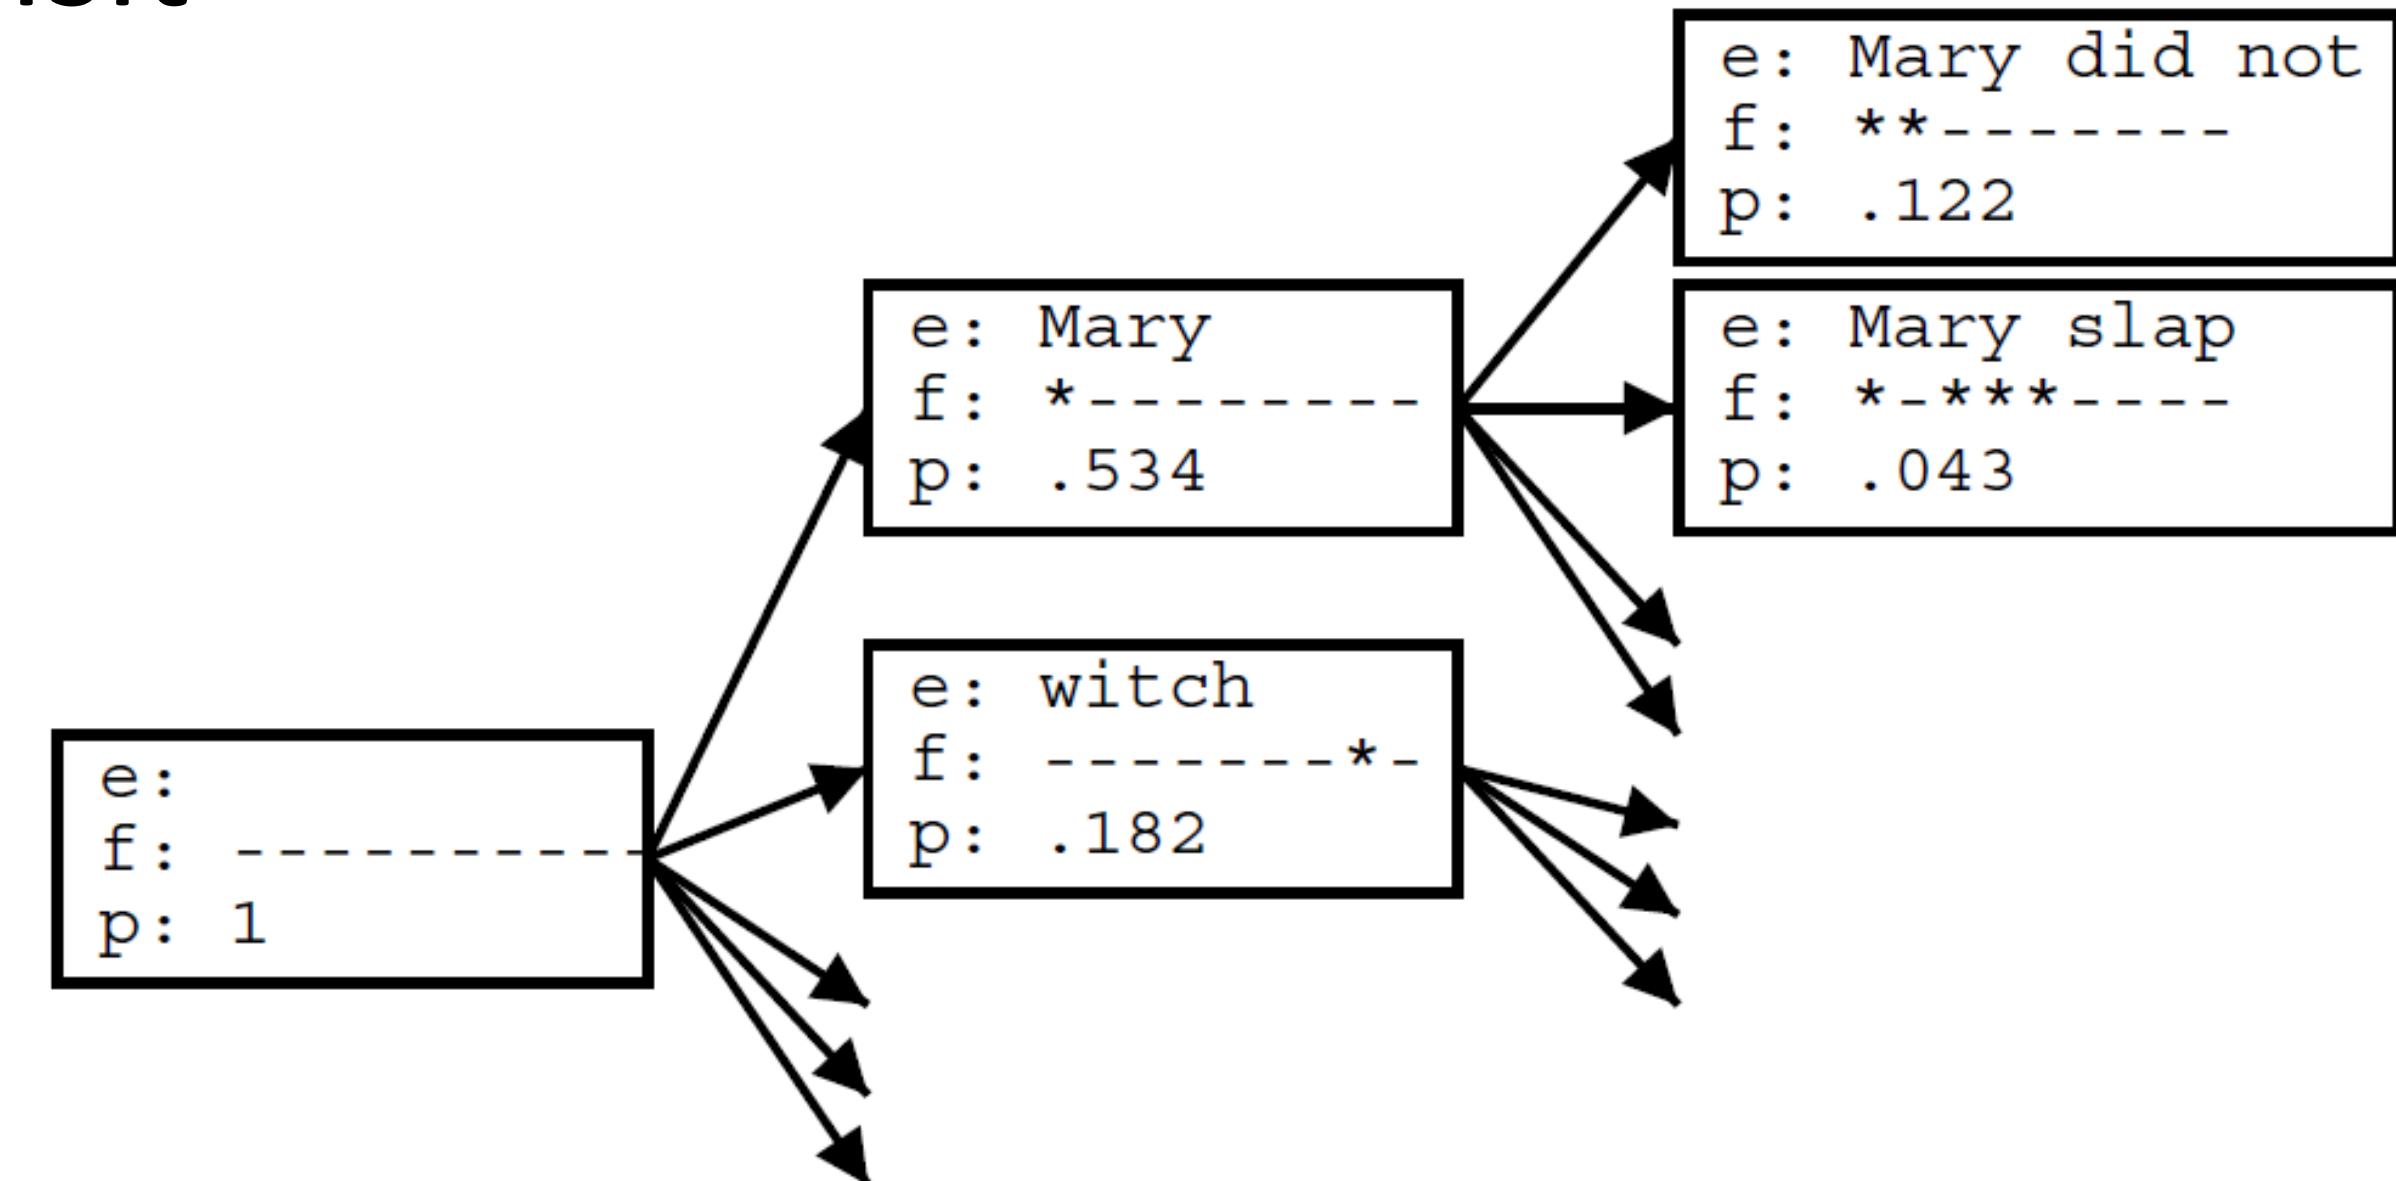
Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>		<u>slap</u>			<u>to the</u>		
		<u>did not give</u>				<u>to</u>		
				<u>slap</u>		<u>the</u>		
						<u>the</u>		
							<u>the witch</u>	

- ▶ Non-monotonic translation: can visit source sentence “out of order”
- ▶ State needs to describe which words have been translated and which haven’t

Non-Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>			<u>slap</u>		<u>to the</u>		
						<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

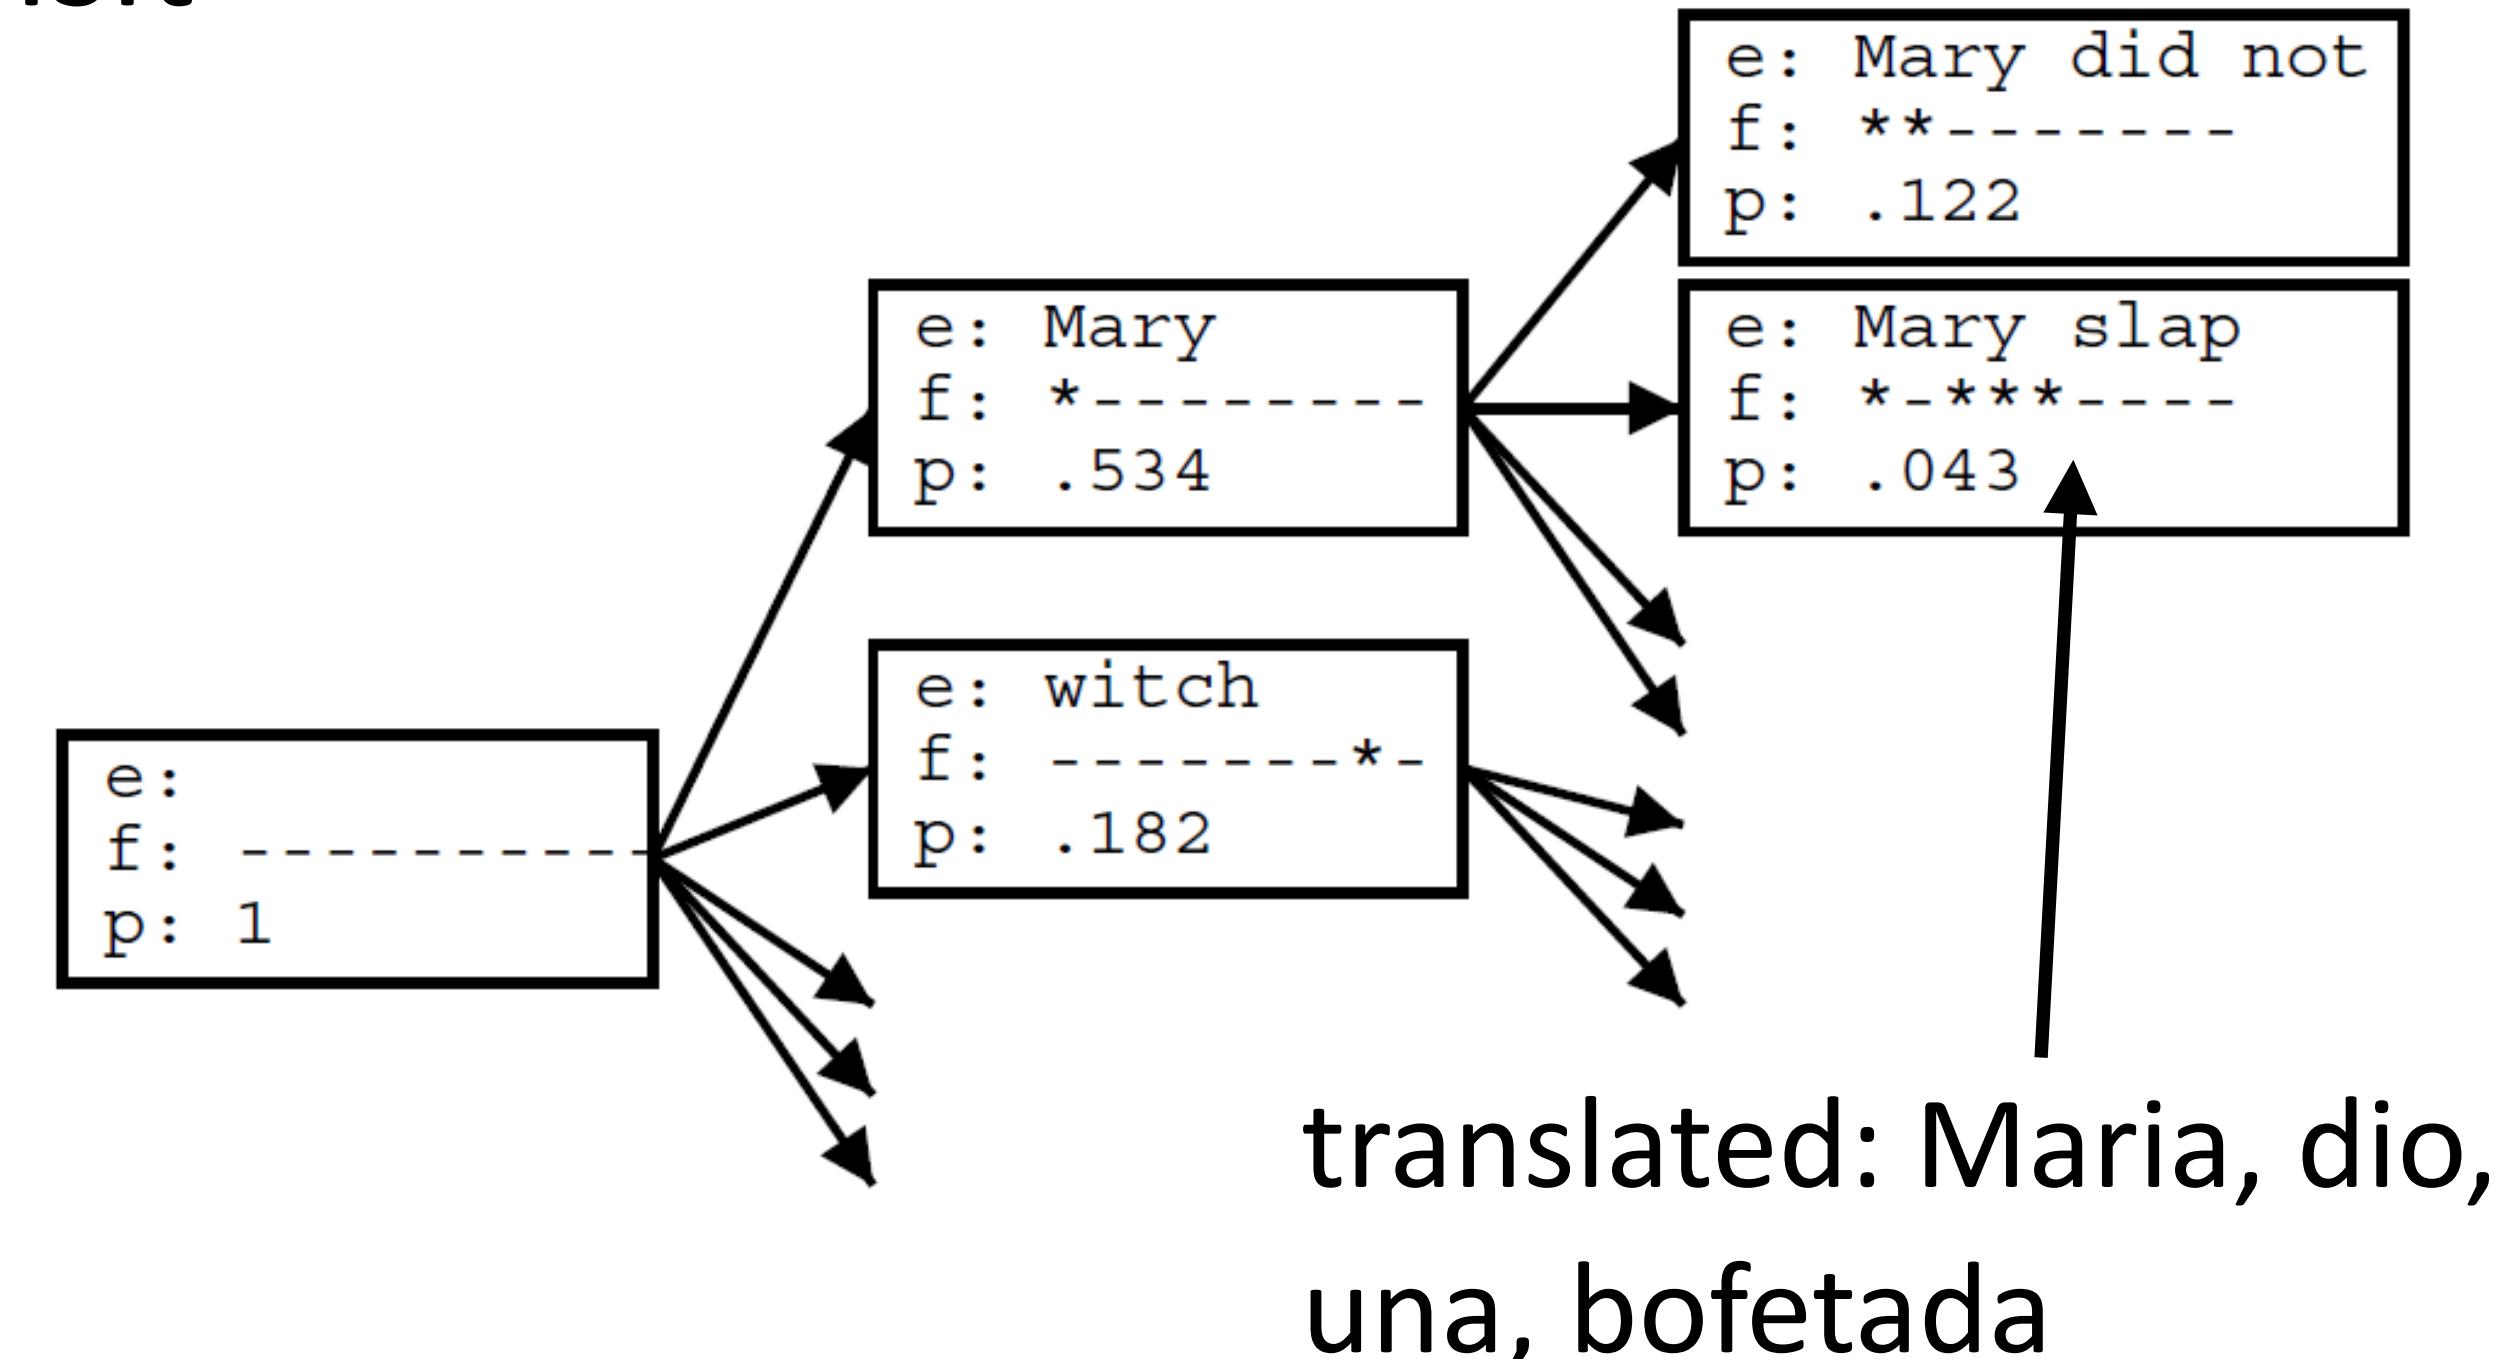
- ▶ Non-monotonic translation: can visit source sentence “out of order”
- ▶ State needs to describe which words have been translated and which haven’t



Non-Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>			<u>slap</u>		<u>to the</u>		
						<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

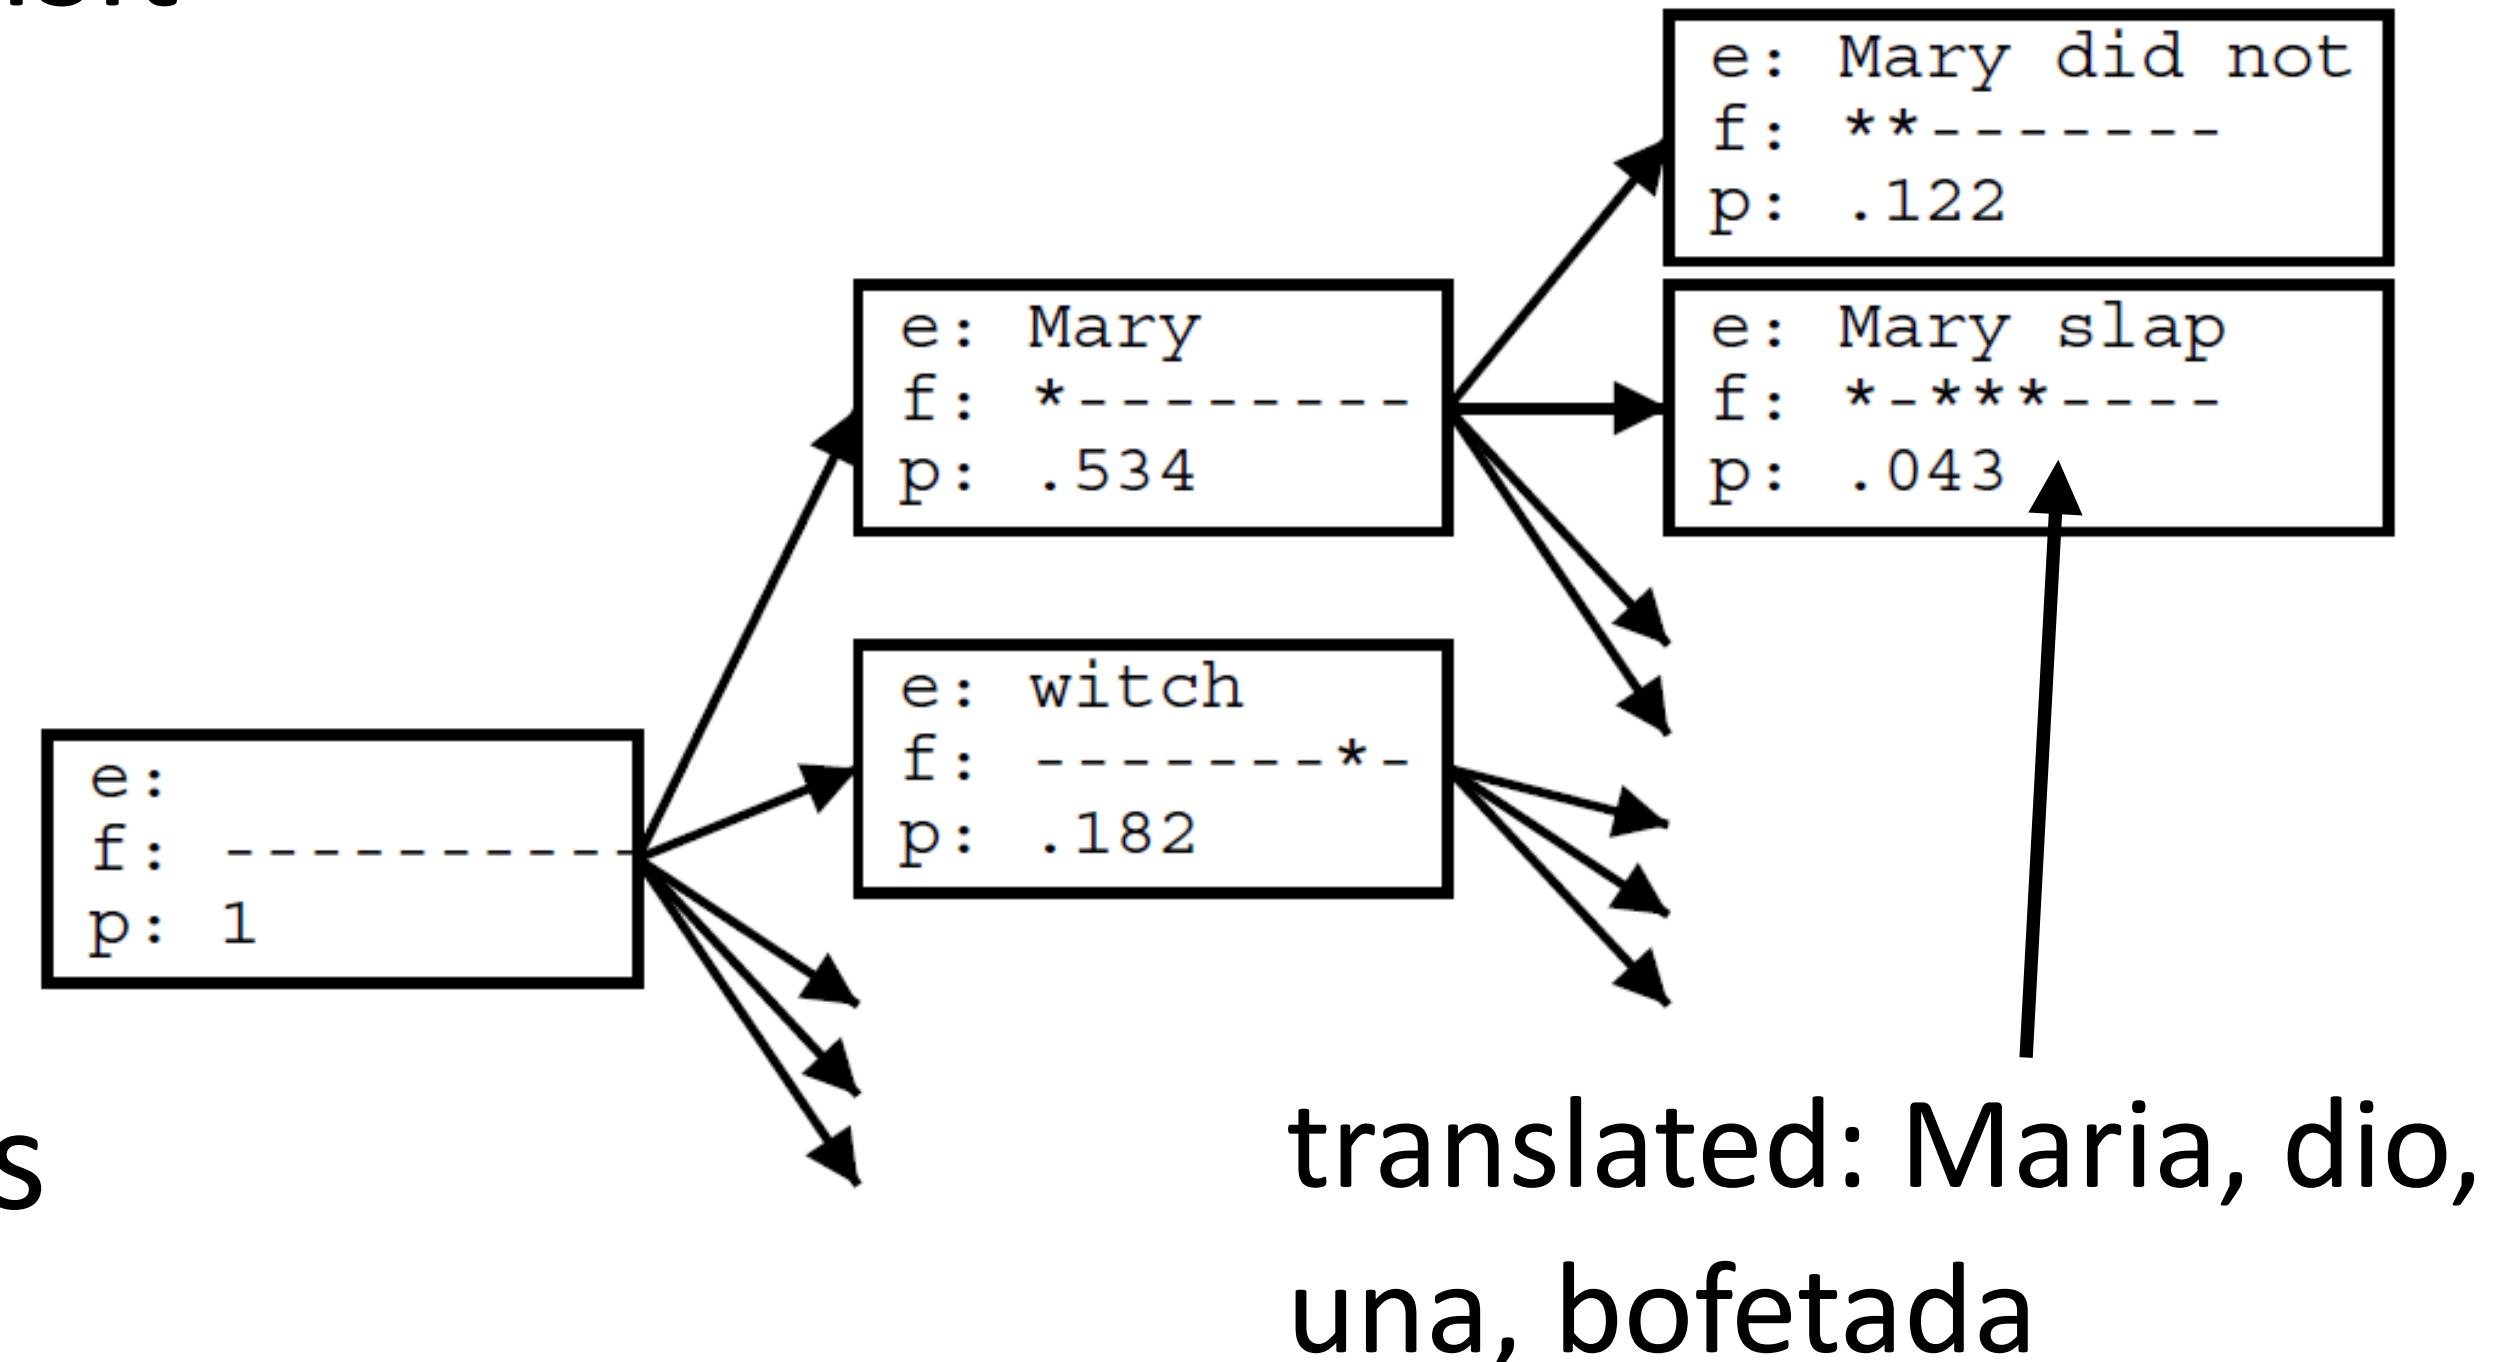
- ▶ Non-monotonic translation: can visit source sentence “out of order”
- ▶ State needs to describe which words have been translated and which haven’t



Non-Monotonic Translation

Maria	no	dio	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>			<u>a slap</u>		<u>by</u>		<u>green witch</u>
	<u>no</u>			<u>slap</u>		<u>to the</u>		
						<u>to</u>		
						<u>the</u>		
				<u>slap</u>			<u>the witch</u>	

- ▶ Non-monotonic translation: can visit source sentence “out of order”
- ▶ State needs to describe which words have been translated and which haven’t
- ▶ Big enough phrases already capture lots of reorderings, so this isn’t as important as you think



Training Decoders

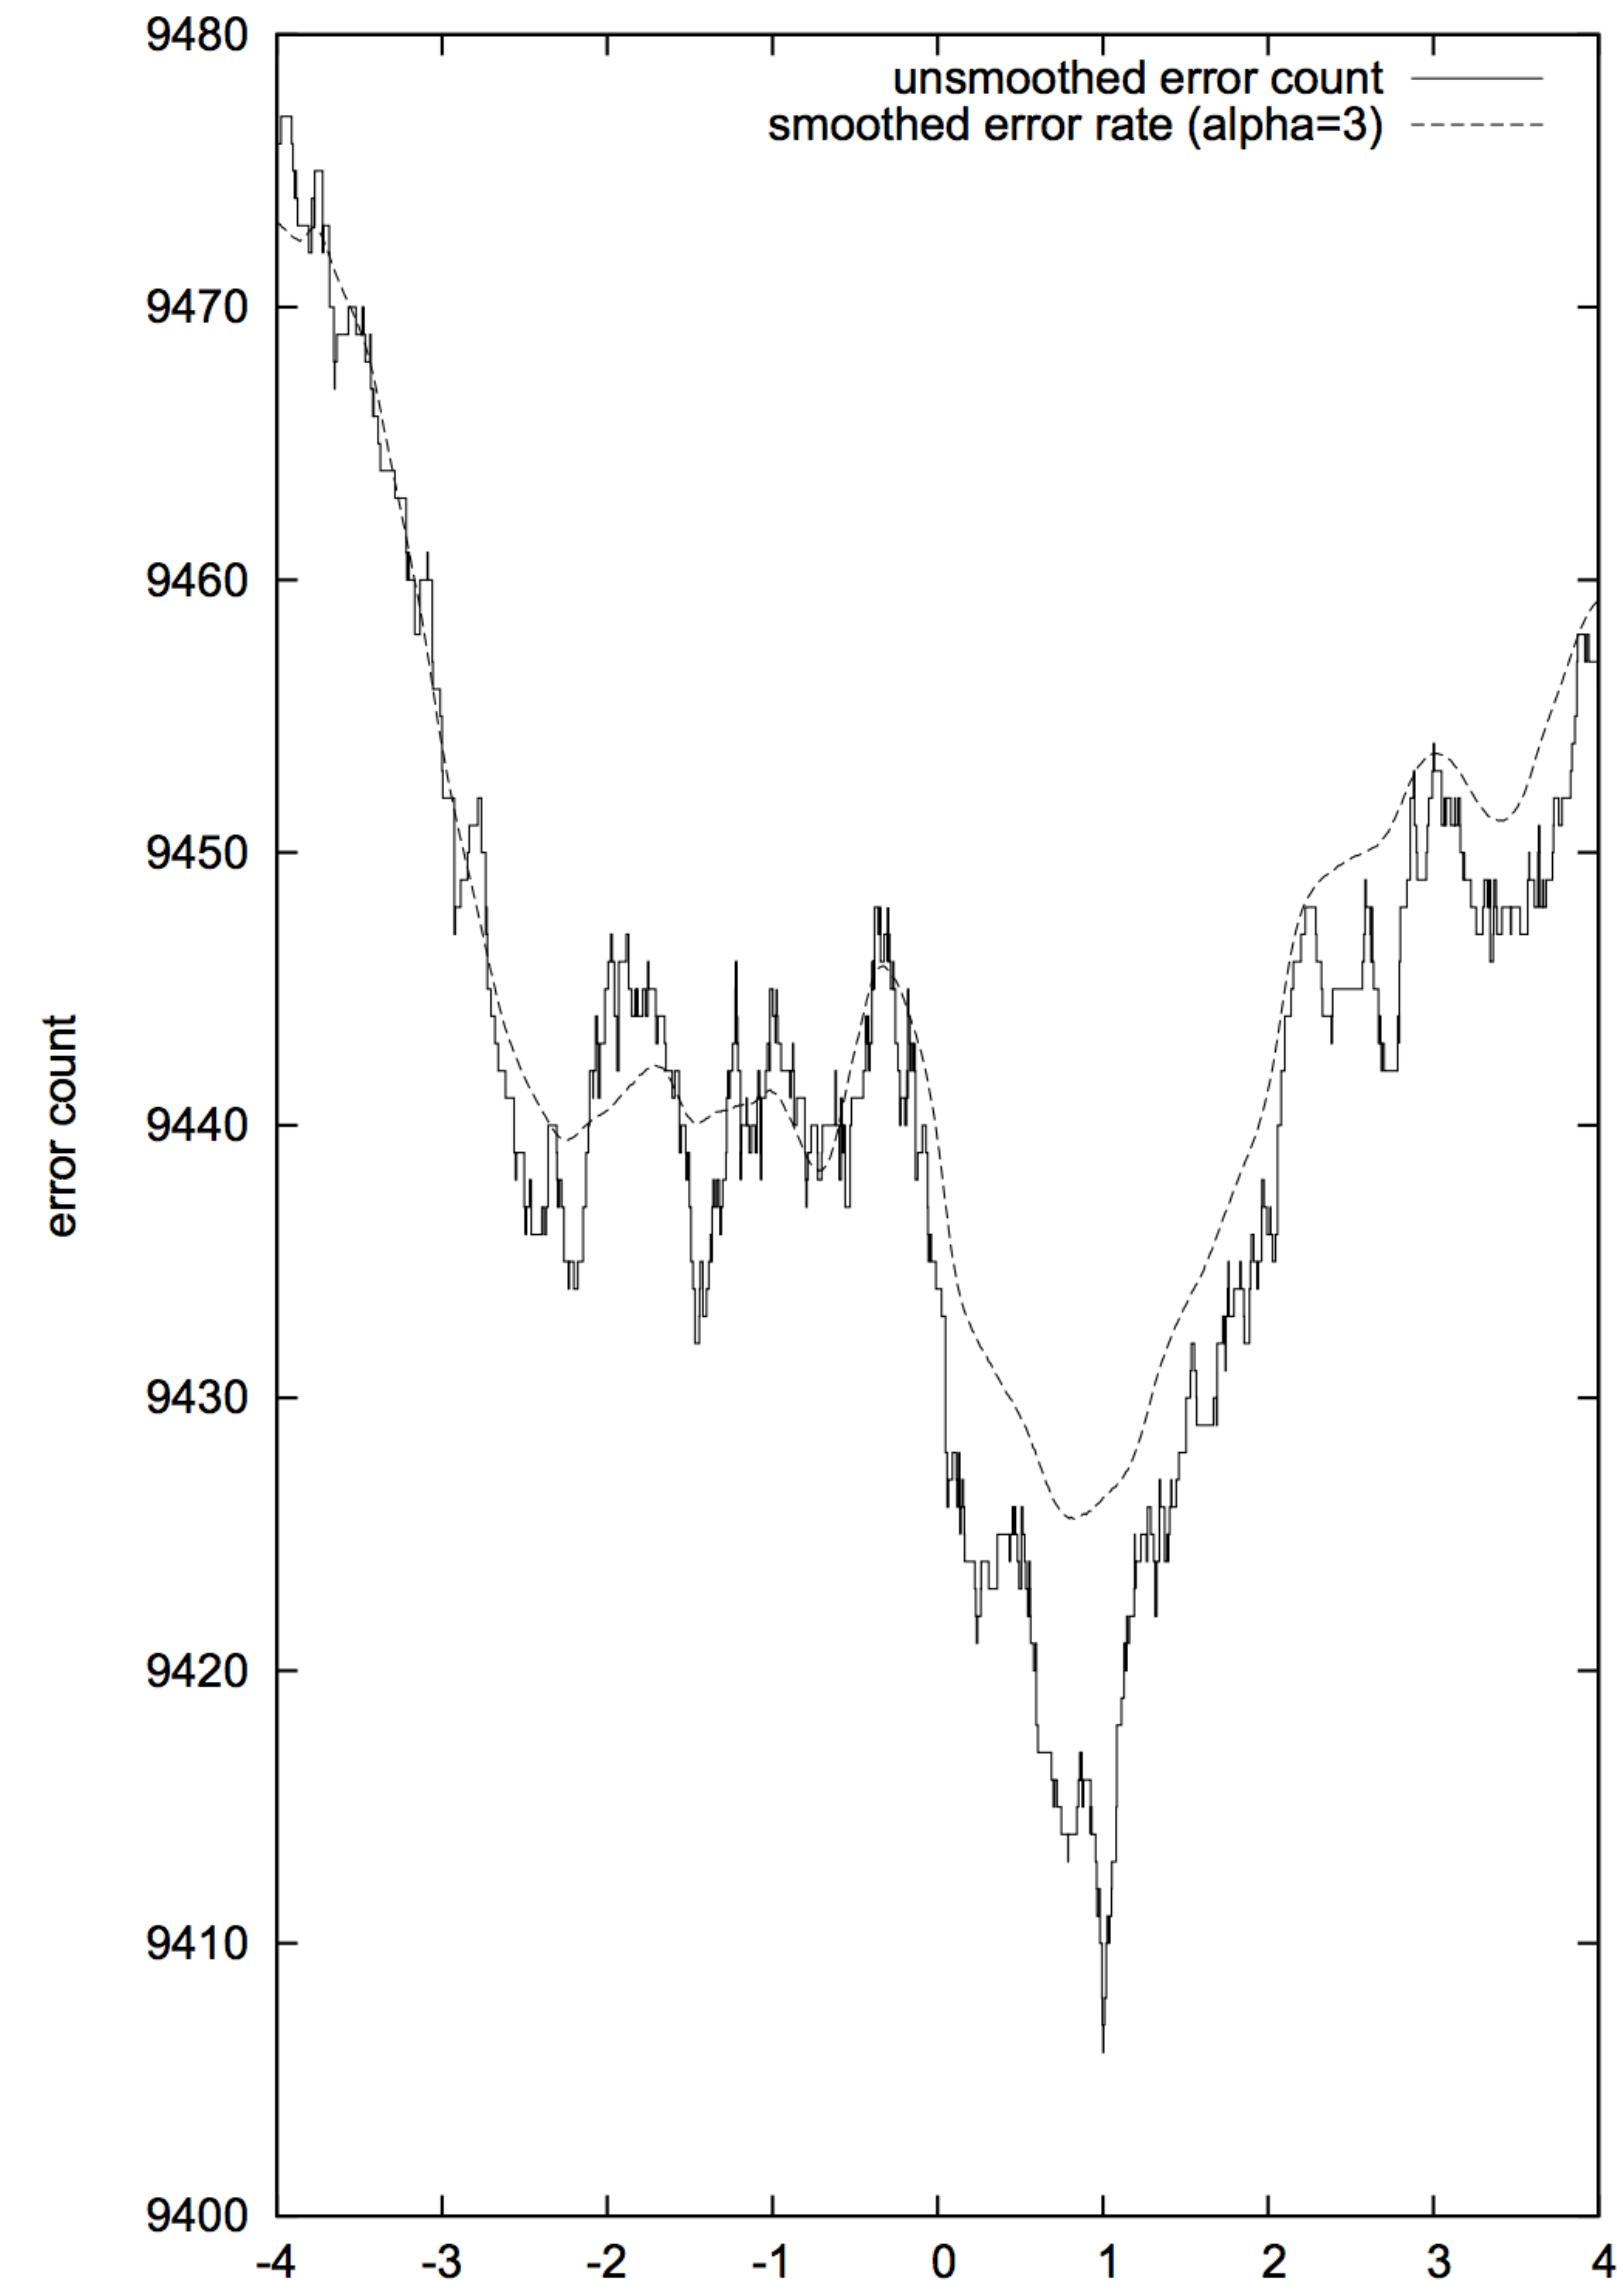
$$\text{score} = \alpha \log P(\text{LM}) + \beta \log P(\text{TM})$$

...and TM is broken down into several features

Training Decoders

$$\text{score} = \alpha \log P(\text{LM}) + \beta \log P(\text{TM})$$

...and TM is broken down into several feature

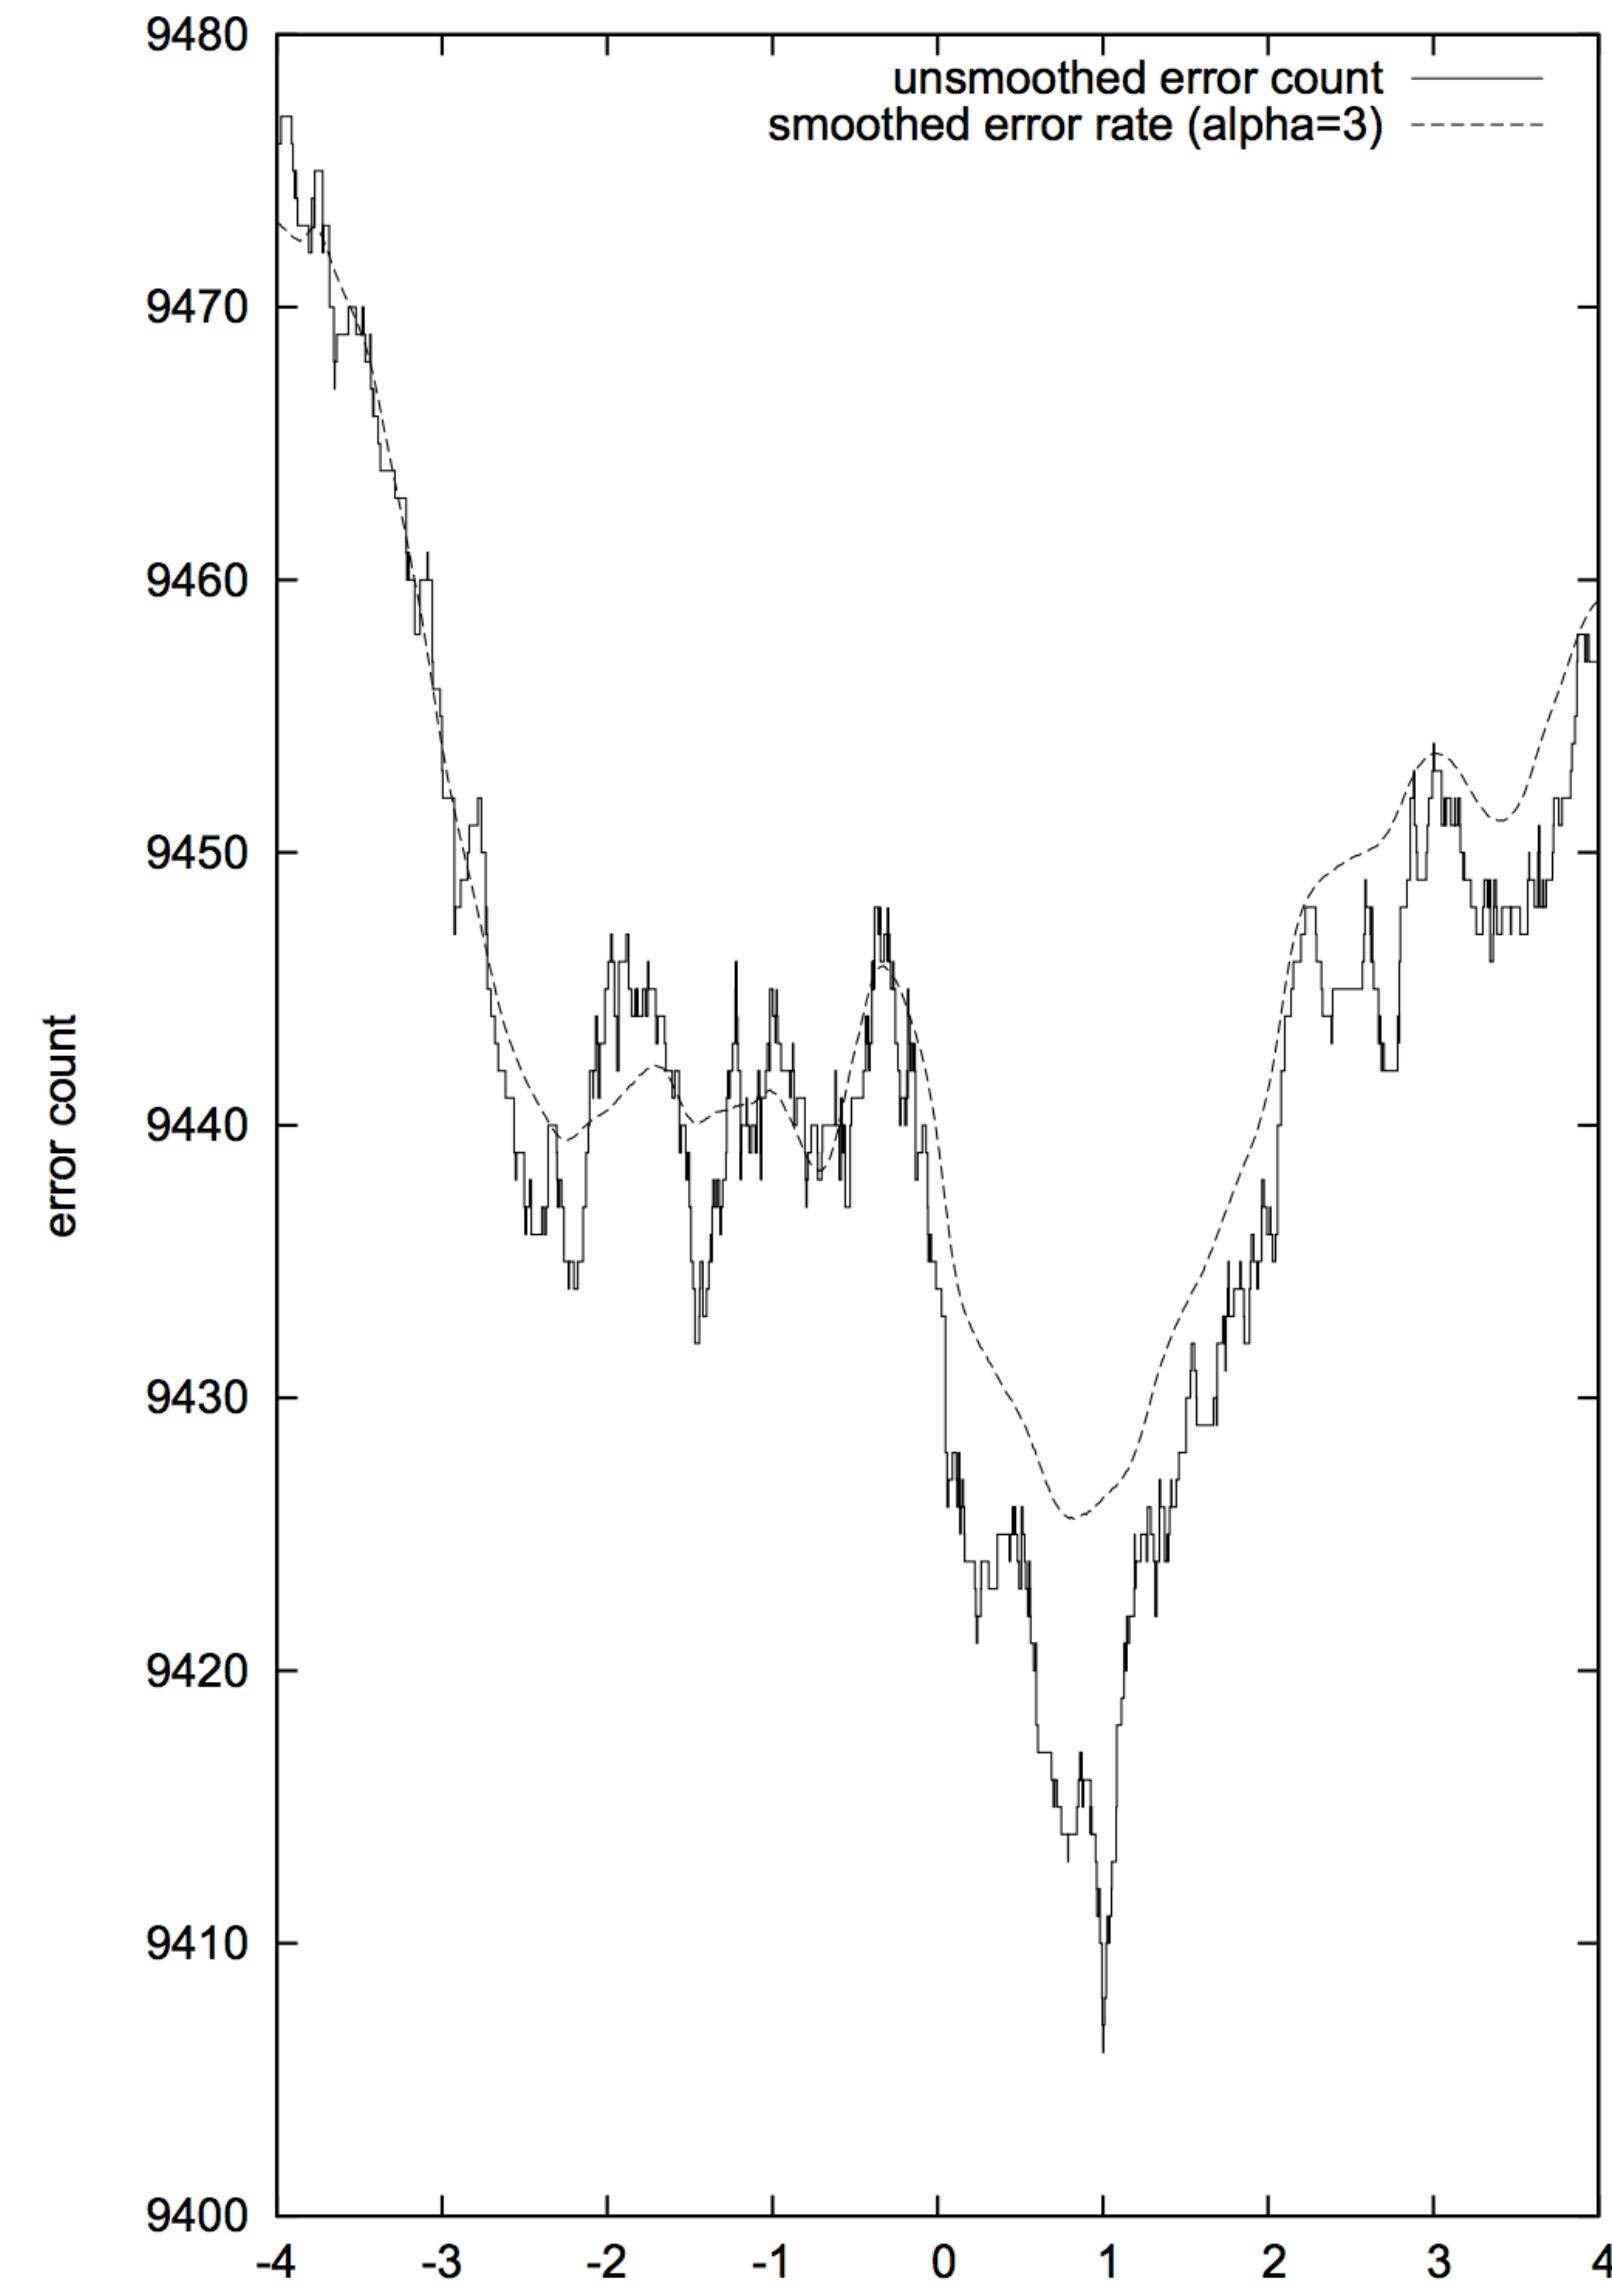


Training Decoders

$$\text{score} = \alpha \log P(\text{LM}) + \beta \log P(\text{TM})$$

...and TM is broken down into several feature

- ▶ Usually 5-20 feature weights to set, want to optimize for BLEU score which is not differentiable
- ▶ MERT (Och 2003): decode to get 1000-best translations for each sentence in a small training set (<1000 sentences), do line search on parameters to directly optimize for BLEU



Moses

- ▶ Toolkit for machine translation due to Philipp Koehn + Hieu Hoang
 - ▶ Pharaoh (Koehn, 2004) is the decoder from Koehn's thesis

Moses

- ▶ Toolkit for machine translation due to Philipp Koehn + Hieu Hoang
 - ▶ Pharaoh (Koehn, 2004) is the decoder from Koehn's thesis
- ▶ Moses implements word alignment, language models, and this decoder, plus *a ton* more stuff
 - ▶ Highly optimized and heavily engineered, could more or less build SOTA translation systems with this from 2007-2013

Moses

- ▶ Toolkit for machine translation due to Philipp Koehn + Hieu Hoang
 - ▶ Pharaoh (Koehn, 2004) is the decoder from Koehn's thesis
- ▶ Moses implements word alignment, language models, and this decoder, plus *a ton* more stuff
 - ▶ Highly optimized and heavily engineered, could more or less build SOTA translation systems with this from 2007-2013
- ▶ Next time: results on these and comparisons to neural methods

Syntax

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

$NN \rightarrow [\text{car}, \text{voiture}]$

Syntactic MT

- ▶ Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the, la}]$

$DT \rightarrow [\text{the, le}]$

$NN \rightarrow [\text{car, voiture}]$

$JJ \rightarrow [\text{yellow, jaune}]$

Syntactic MT

- Rather than use phrases, use a *synchronous context-free grammar*

$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the, la}]$

$DT \rightarrow [\text{the, le}]$

$NN \rightarrow [\text{car, voiture}]$

$JJ \rightarrow [\text{yellow, jaune}]$



Syntactic MT

- Rather than use phrases, use a *synchronous context-free grammar*

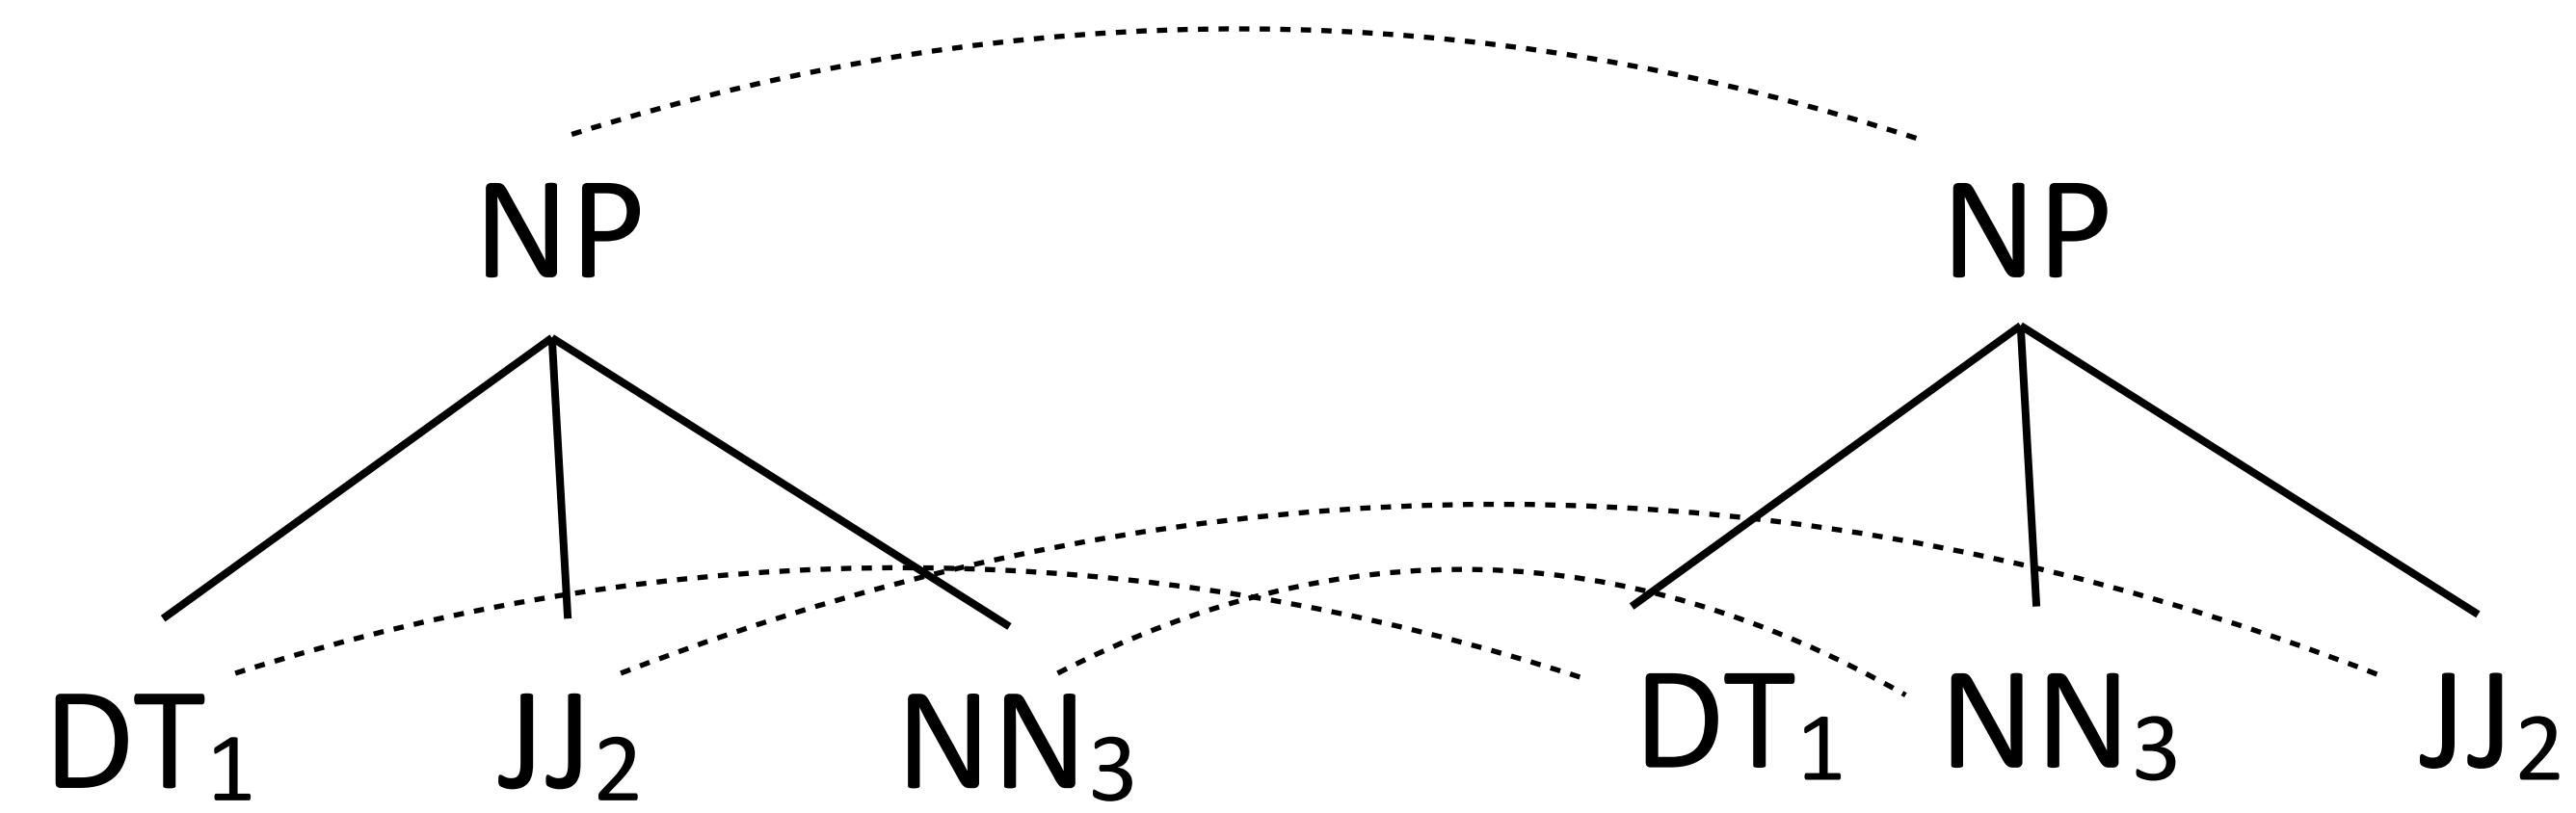
$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

$NN \rightarrow [\text{car}, \text{voiture}]$

$JJ \rightarrow [\text{yellow}, \text{jaune}]$



Syntactic MT

- Rather than use phrases, use a *synchronous context-free grammar*

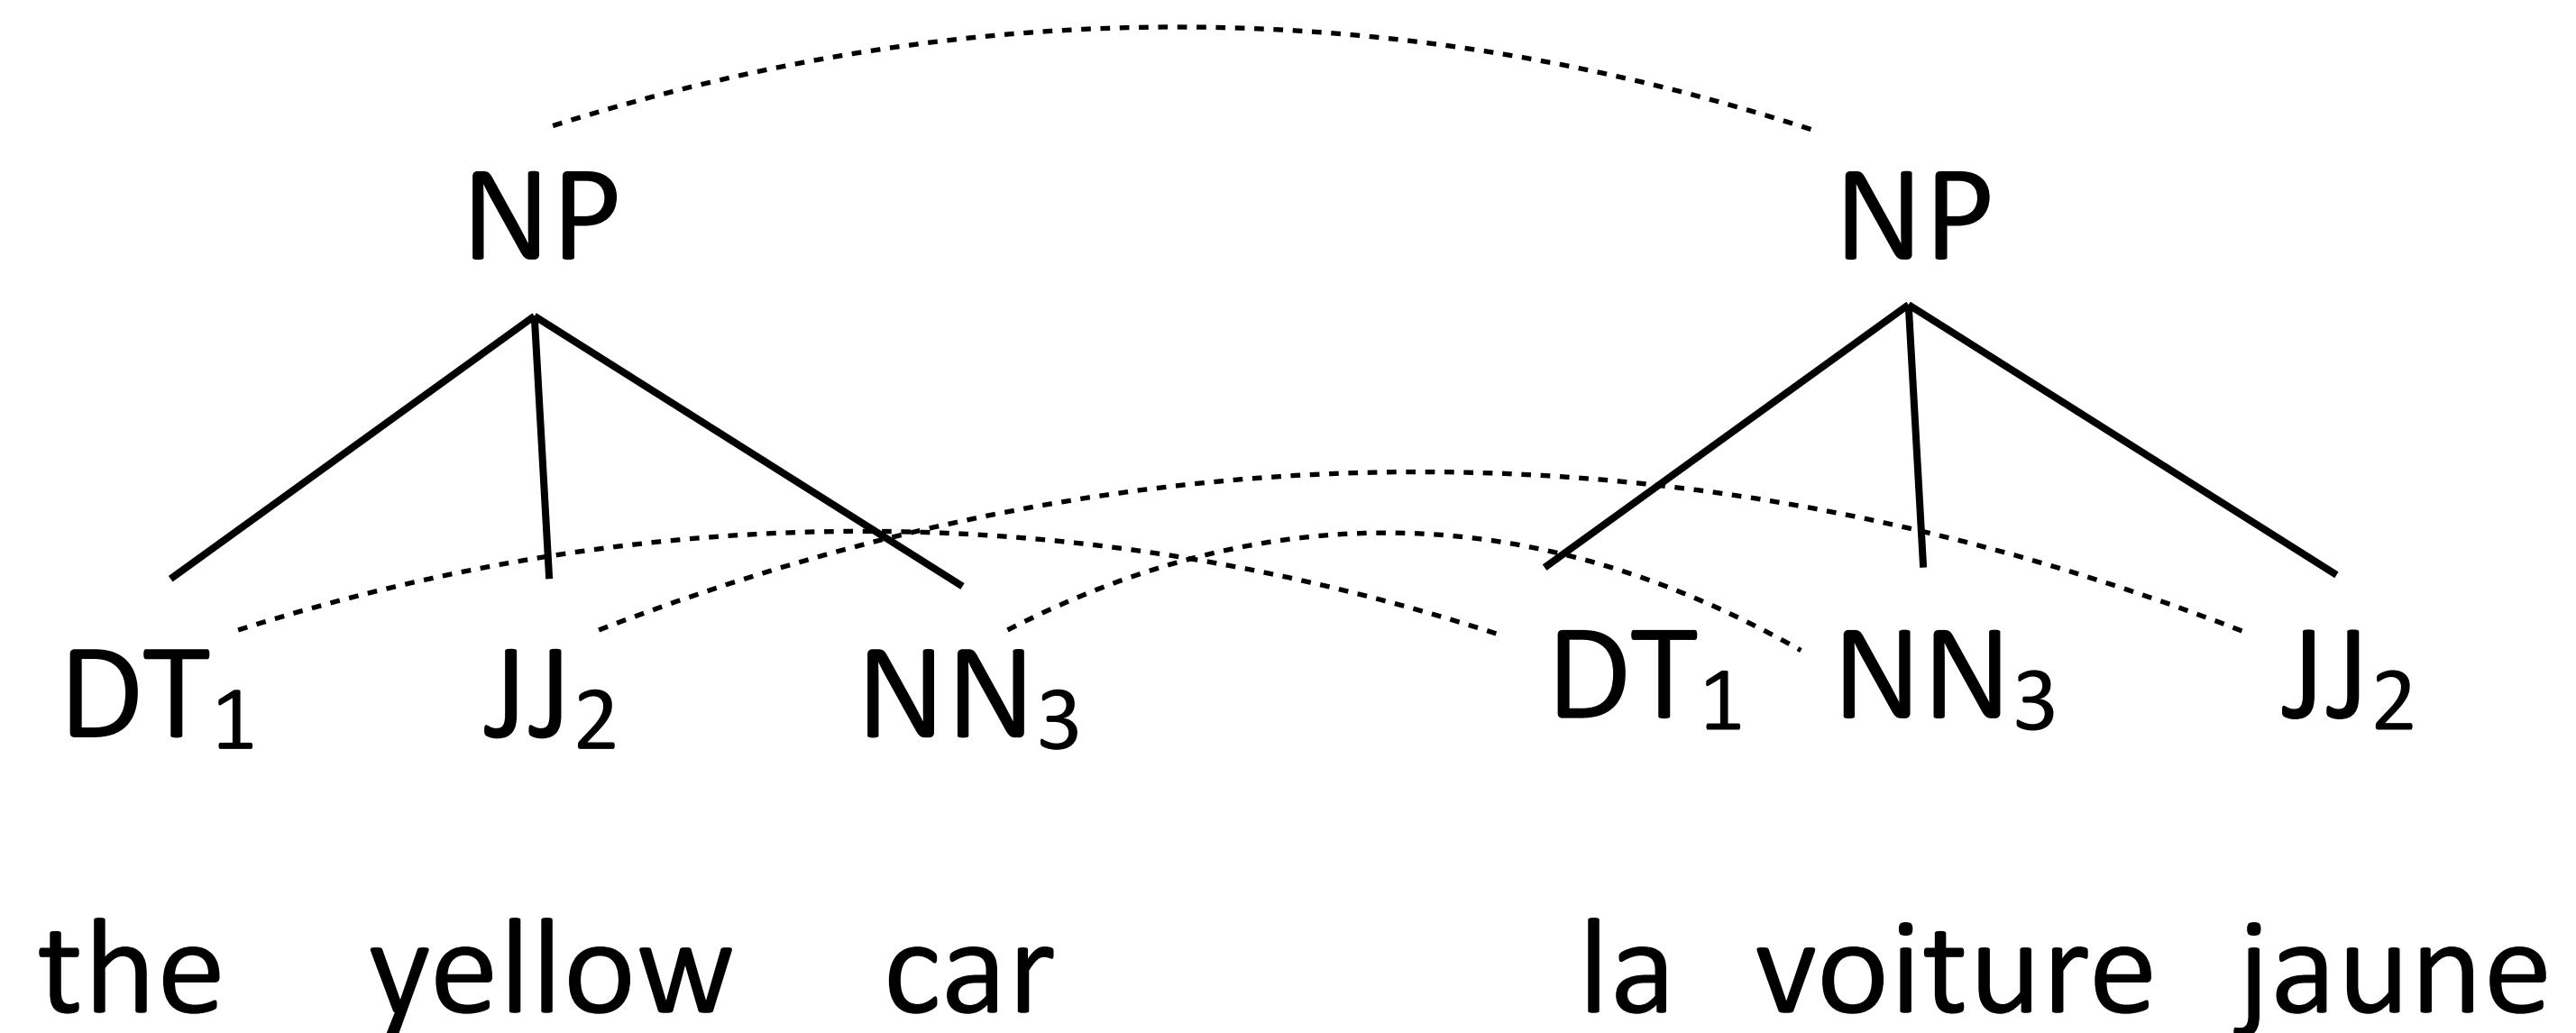
$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

$NN \rightarrow [\text{car}, \text{voiture}]$

$JJ \rightarrow [\text{yellow}, \text{jaune}]$



Syntactic MT

- Rather than use phrases, use a *synchronous context-free grammar*

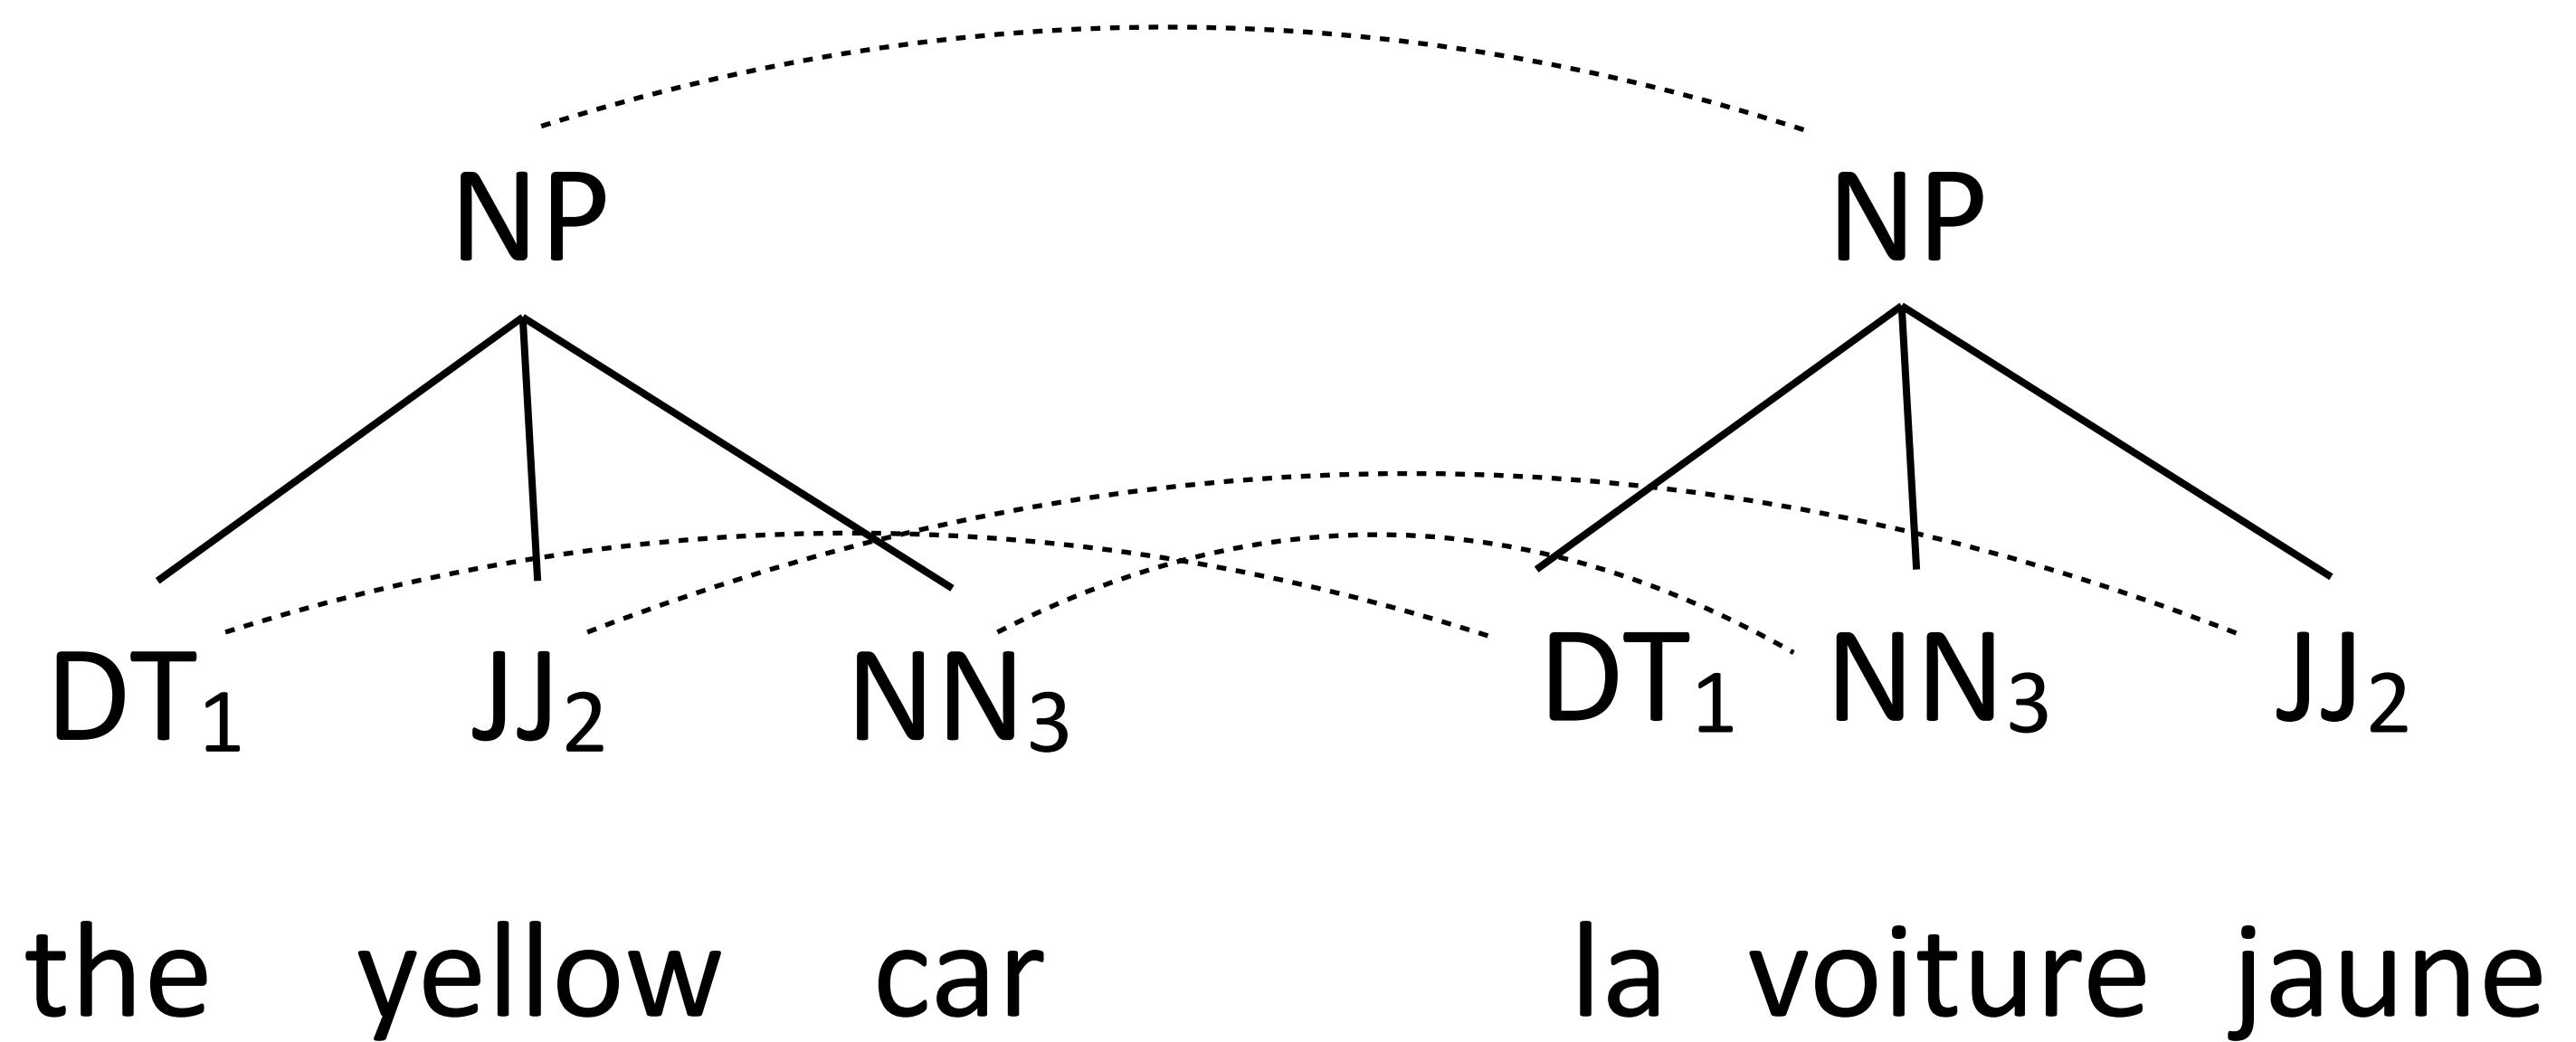
$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

$NN \rightarrow [\text{car}, \text{voiture}]$

$JJ \rightarrow [\text{yellow}, \text{jaune}]$



- Translation = parse the input with “half” of the grammar, read off the other half

Syntactic MT

- Rather than use phrases, use a *synchronous context-free grammar*

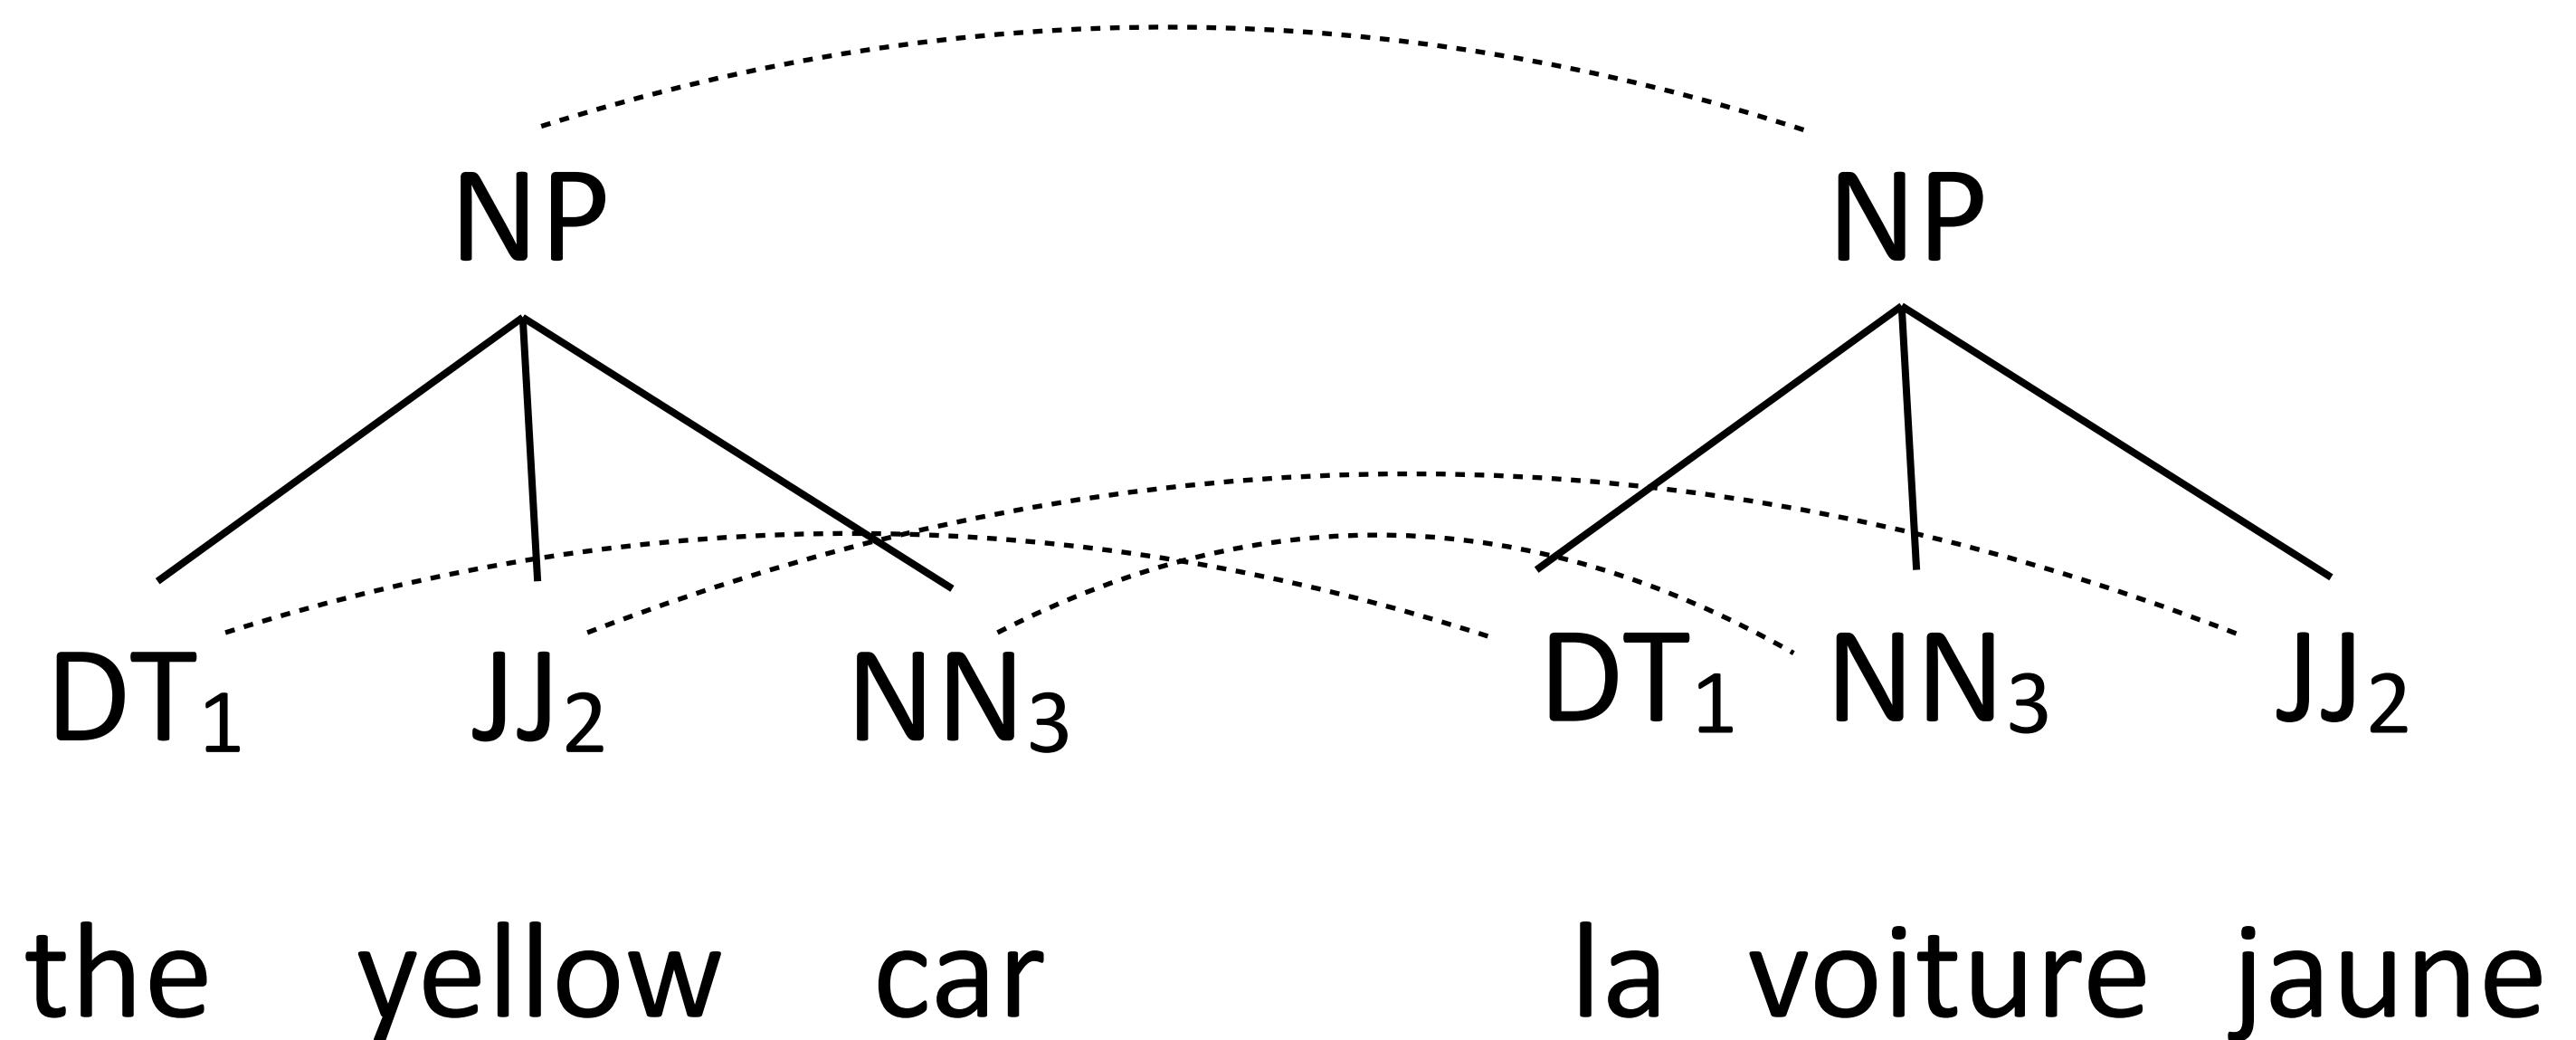
$NP \rightarrow [DT_1\ JJ_2\ NN_3; DT_1\ NN_3\ JJ_2]$

$DT \rightarrow [\text{the}, \text{la}]$

$DT \rightarrow [\text{the}, \text{le}]$

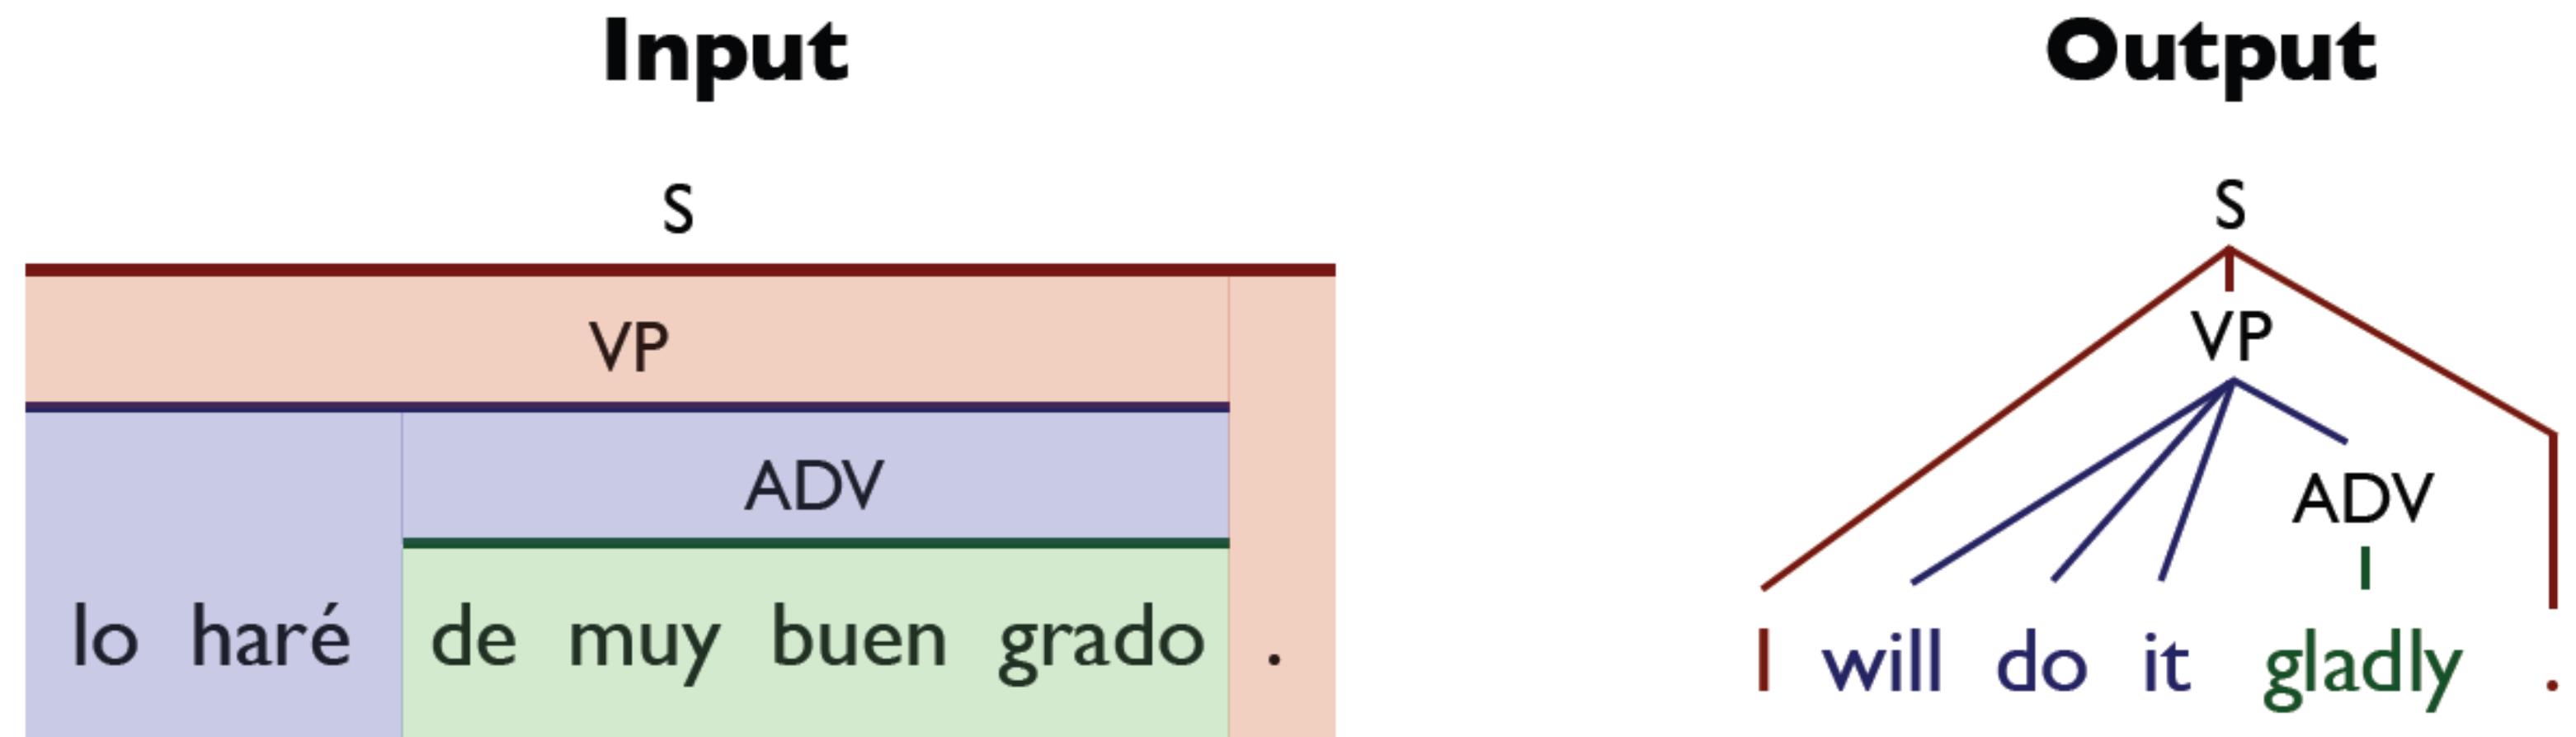
$NN \rightarrow [\text{car}, \text{voiture}]$

$JJ \rightarrow [\text{yellow}, \text{jaune}]$



- Translation = parse the input with “half” of the grammar, read off the other half
- Assumes parallel syntax up to reordering

Syntactic MT



- ▶ Use lexicalized rules, look like “syntactic phrases”
- ▶ Leads to HUGE grammars, parsing is slow

Grammar

$$S \rightarrow \langle VP . ; I VP . \rangle \quad \text{OR} \quad S \rightarrow \langle VP . ; you VP . \rangle$$
$$VP \rightarrow \langle lo\haré\ ADV ; will\do\it\ ADV \rangle$$
$$S \rightarrow \langle lo\haré\ ADV . ; I will\do\it\ ADV . \rangle$$
$$ADV \rightarrow \langle de\muy\buen\grado ; gladly \rangle$$

Slide credit: Dan Klein

Takeaways

- ▶ Phrase-based systems consist of 3 pieces: aligner, language model, decoder
 - ▶ HMMs work well for alignment
 - ▶ N-gram language models are scalable and historically worked well
 - ▶ Decoder requires searching through a complex state space
- ▶ Lots of system variants incorporating syntax
- ▶ Next time: neural MT