

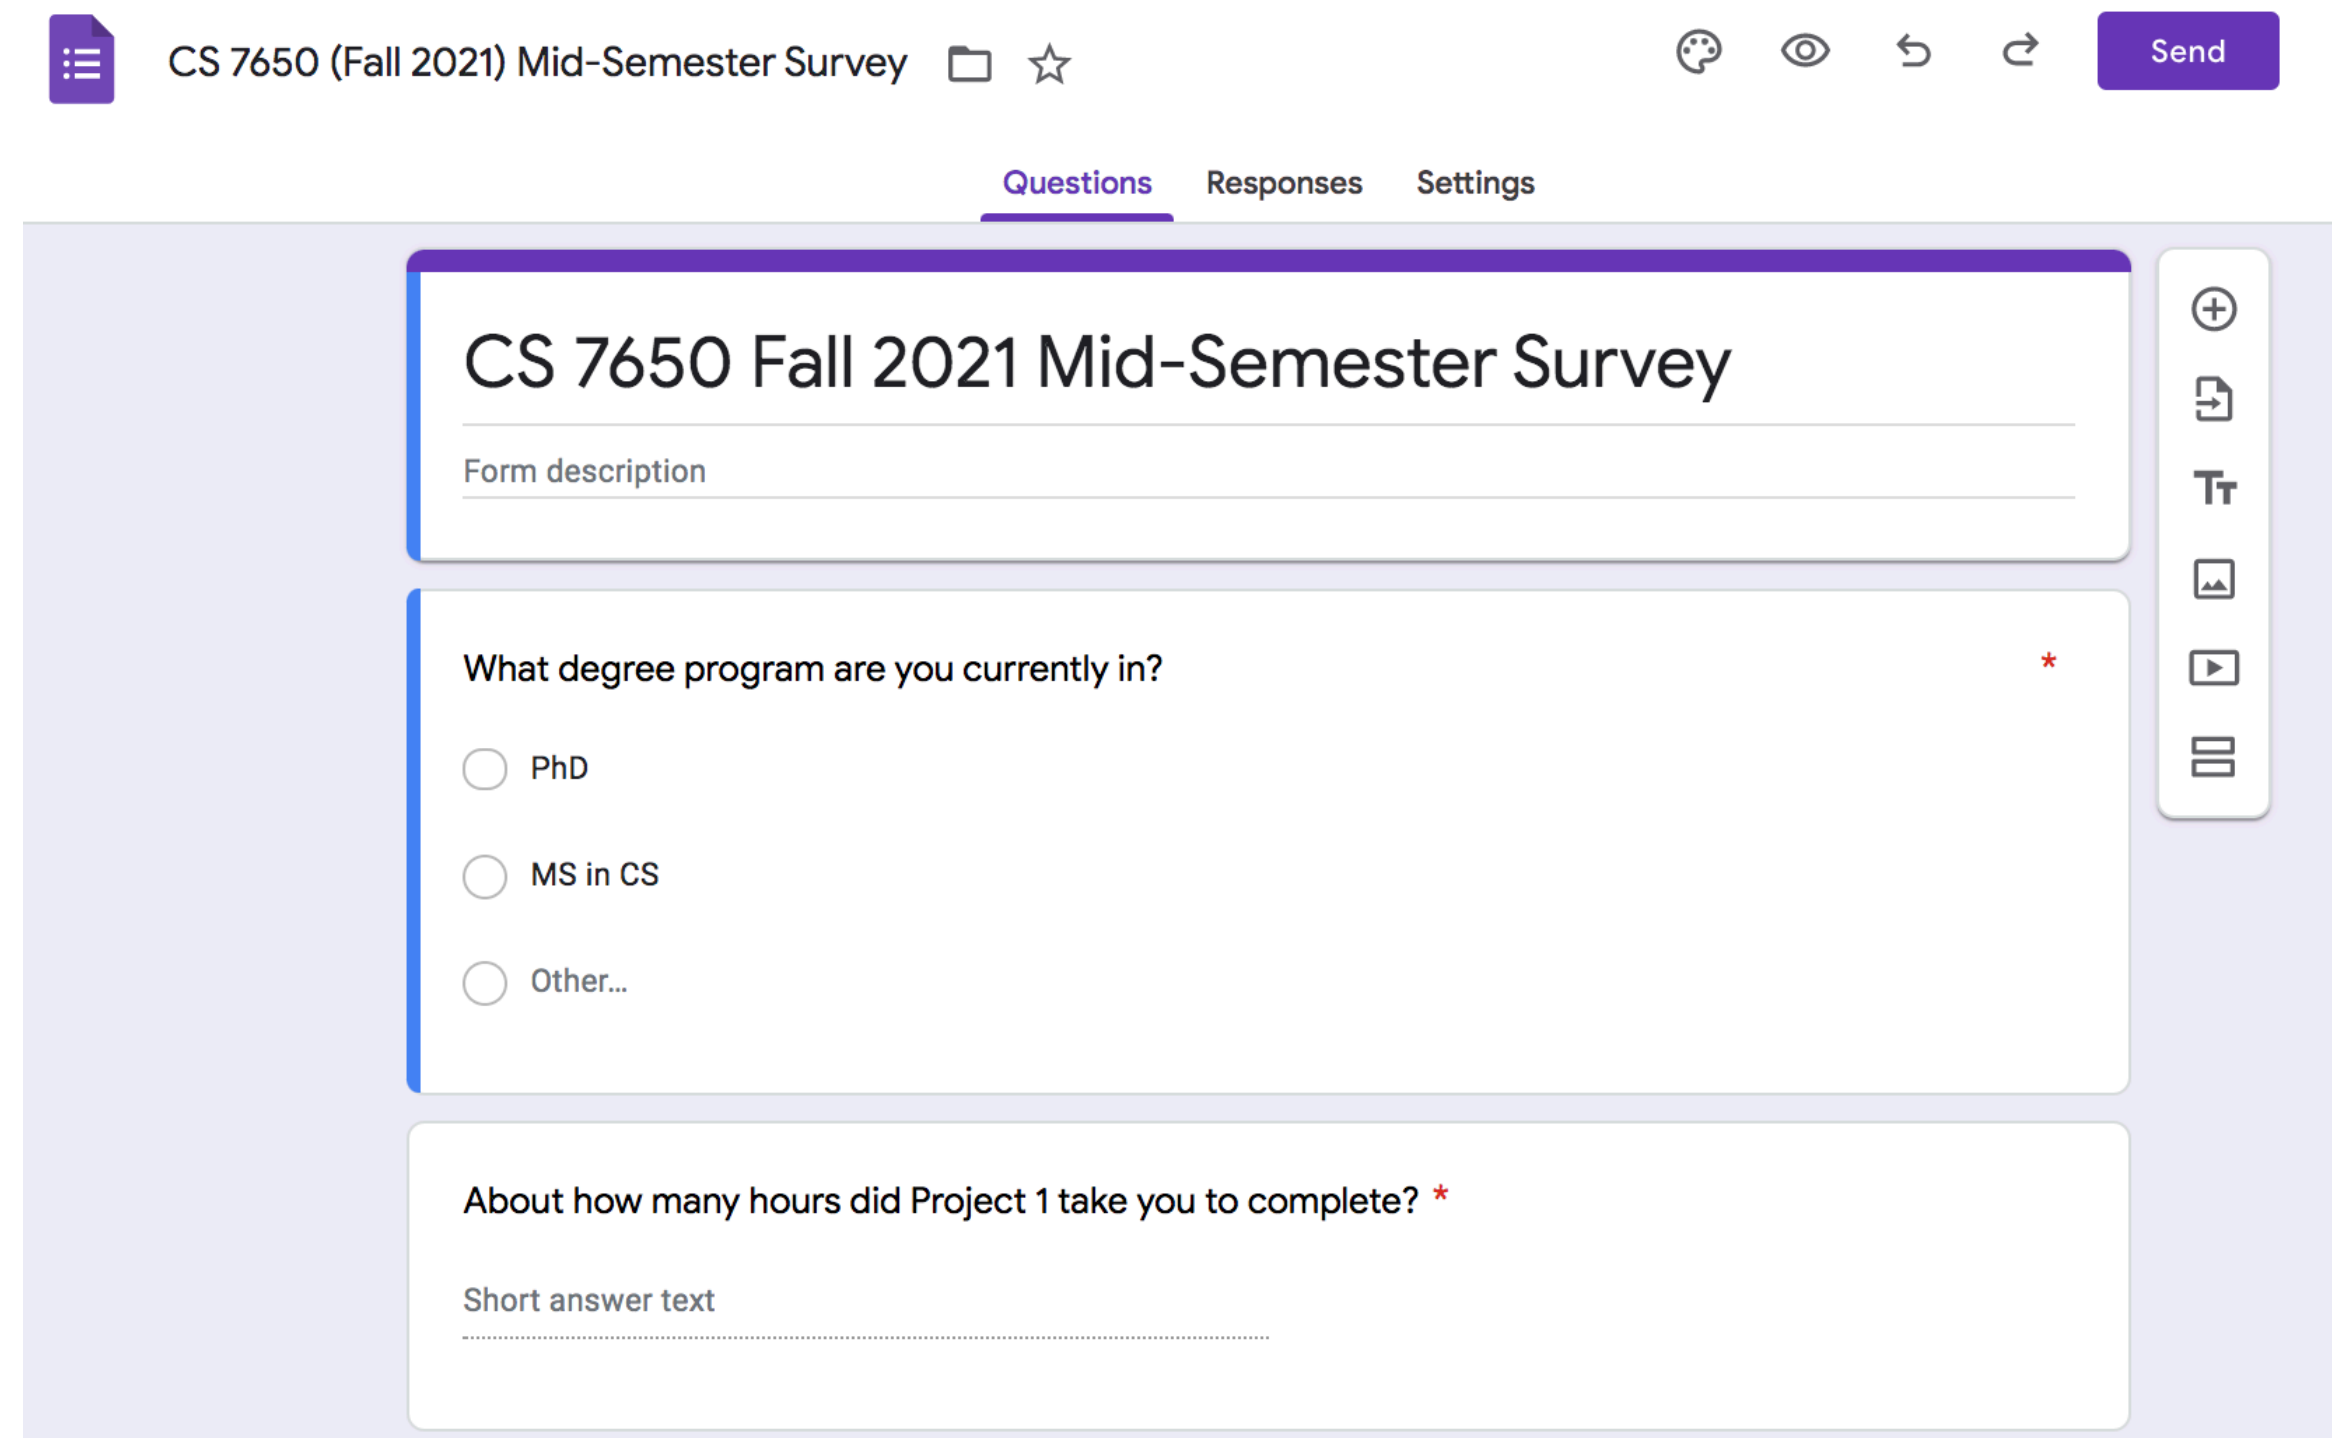
Neural MT + Copy/Pointer

Wei Xu

(many slides from Greg Durrett)

Administrivia

- ▶ Midterm will be released soon (take-home)
- ▶ Mid-semester Survey (by Monday 10/25)



The screenshot shows a Google Forms interface for a survey titled "CS 7650 (Fall 2021) Mid-Semester Survey". The form is titled "CS 7650 Fall 2021 Mid-Semester Survey" and has a description field. The first question is "What degree program are you currently in?" with three radio button options: "PhD", "MS in CS", and "Other...". The second question is "About how many hours did Project 1 take you to complete?" with a short answer text field. The form is set to "Send" and has tabs for "Questions", "Responses", and "Settings".

CS 7650 (Fall 2021) Mid-Semester Survey

Questions Responses Settings

CS 7650 Fall 2021 Mid-Semester Survey

Form description

What degree program are you currently in? *

☐ PhD

☐ MS in CS

☐ Other...


About how many hours did Project 1 take you to complete? *

Short answer text

This Lecture


- ▶ Sequence-to-Sequence Model
- ▶ Attention Mechanism
- ▶ **Neural MT and other Applications**
- ▶ **Copy/Pointer Network**
- ▶ **Transformer Architecture (if time)**

Recap: Seq2Seq Model




- ▶ Encoder: consumes sequence of tokens, produces a vector. Analogous to encoders for classification/tagging tasks $P(y_i | \mathbf{x}, y_1, \dots, y_{i-1}) = \text{softmax}(W \bar{h}_i)$
- ▶ Decoder: separate module, single cell. Takes two inputs: hidden state (vector h or tuple (h, c)) and previous token. Outputs token + new state

Recall: Training Seq2Seq Model



- ▶ Objective: maximize $\sum_{(\mathbf{x}, \mathbf{y})} \sum_{i=1}^n \log P(y_i^* | \mathbf{x}, y_1^*, \dots, y_{i-1}^*)$
- ▶ One loss term for each target-sentence word, feed the correct word regardless of model's prediction

Recall: Attention

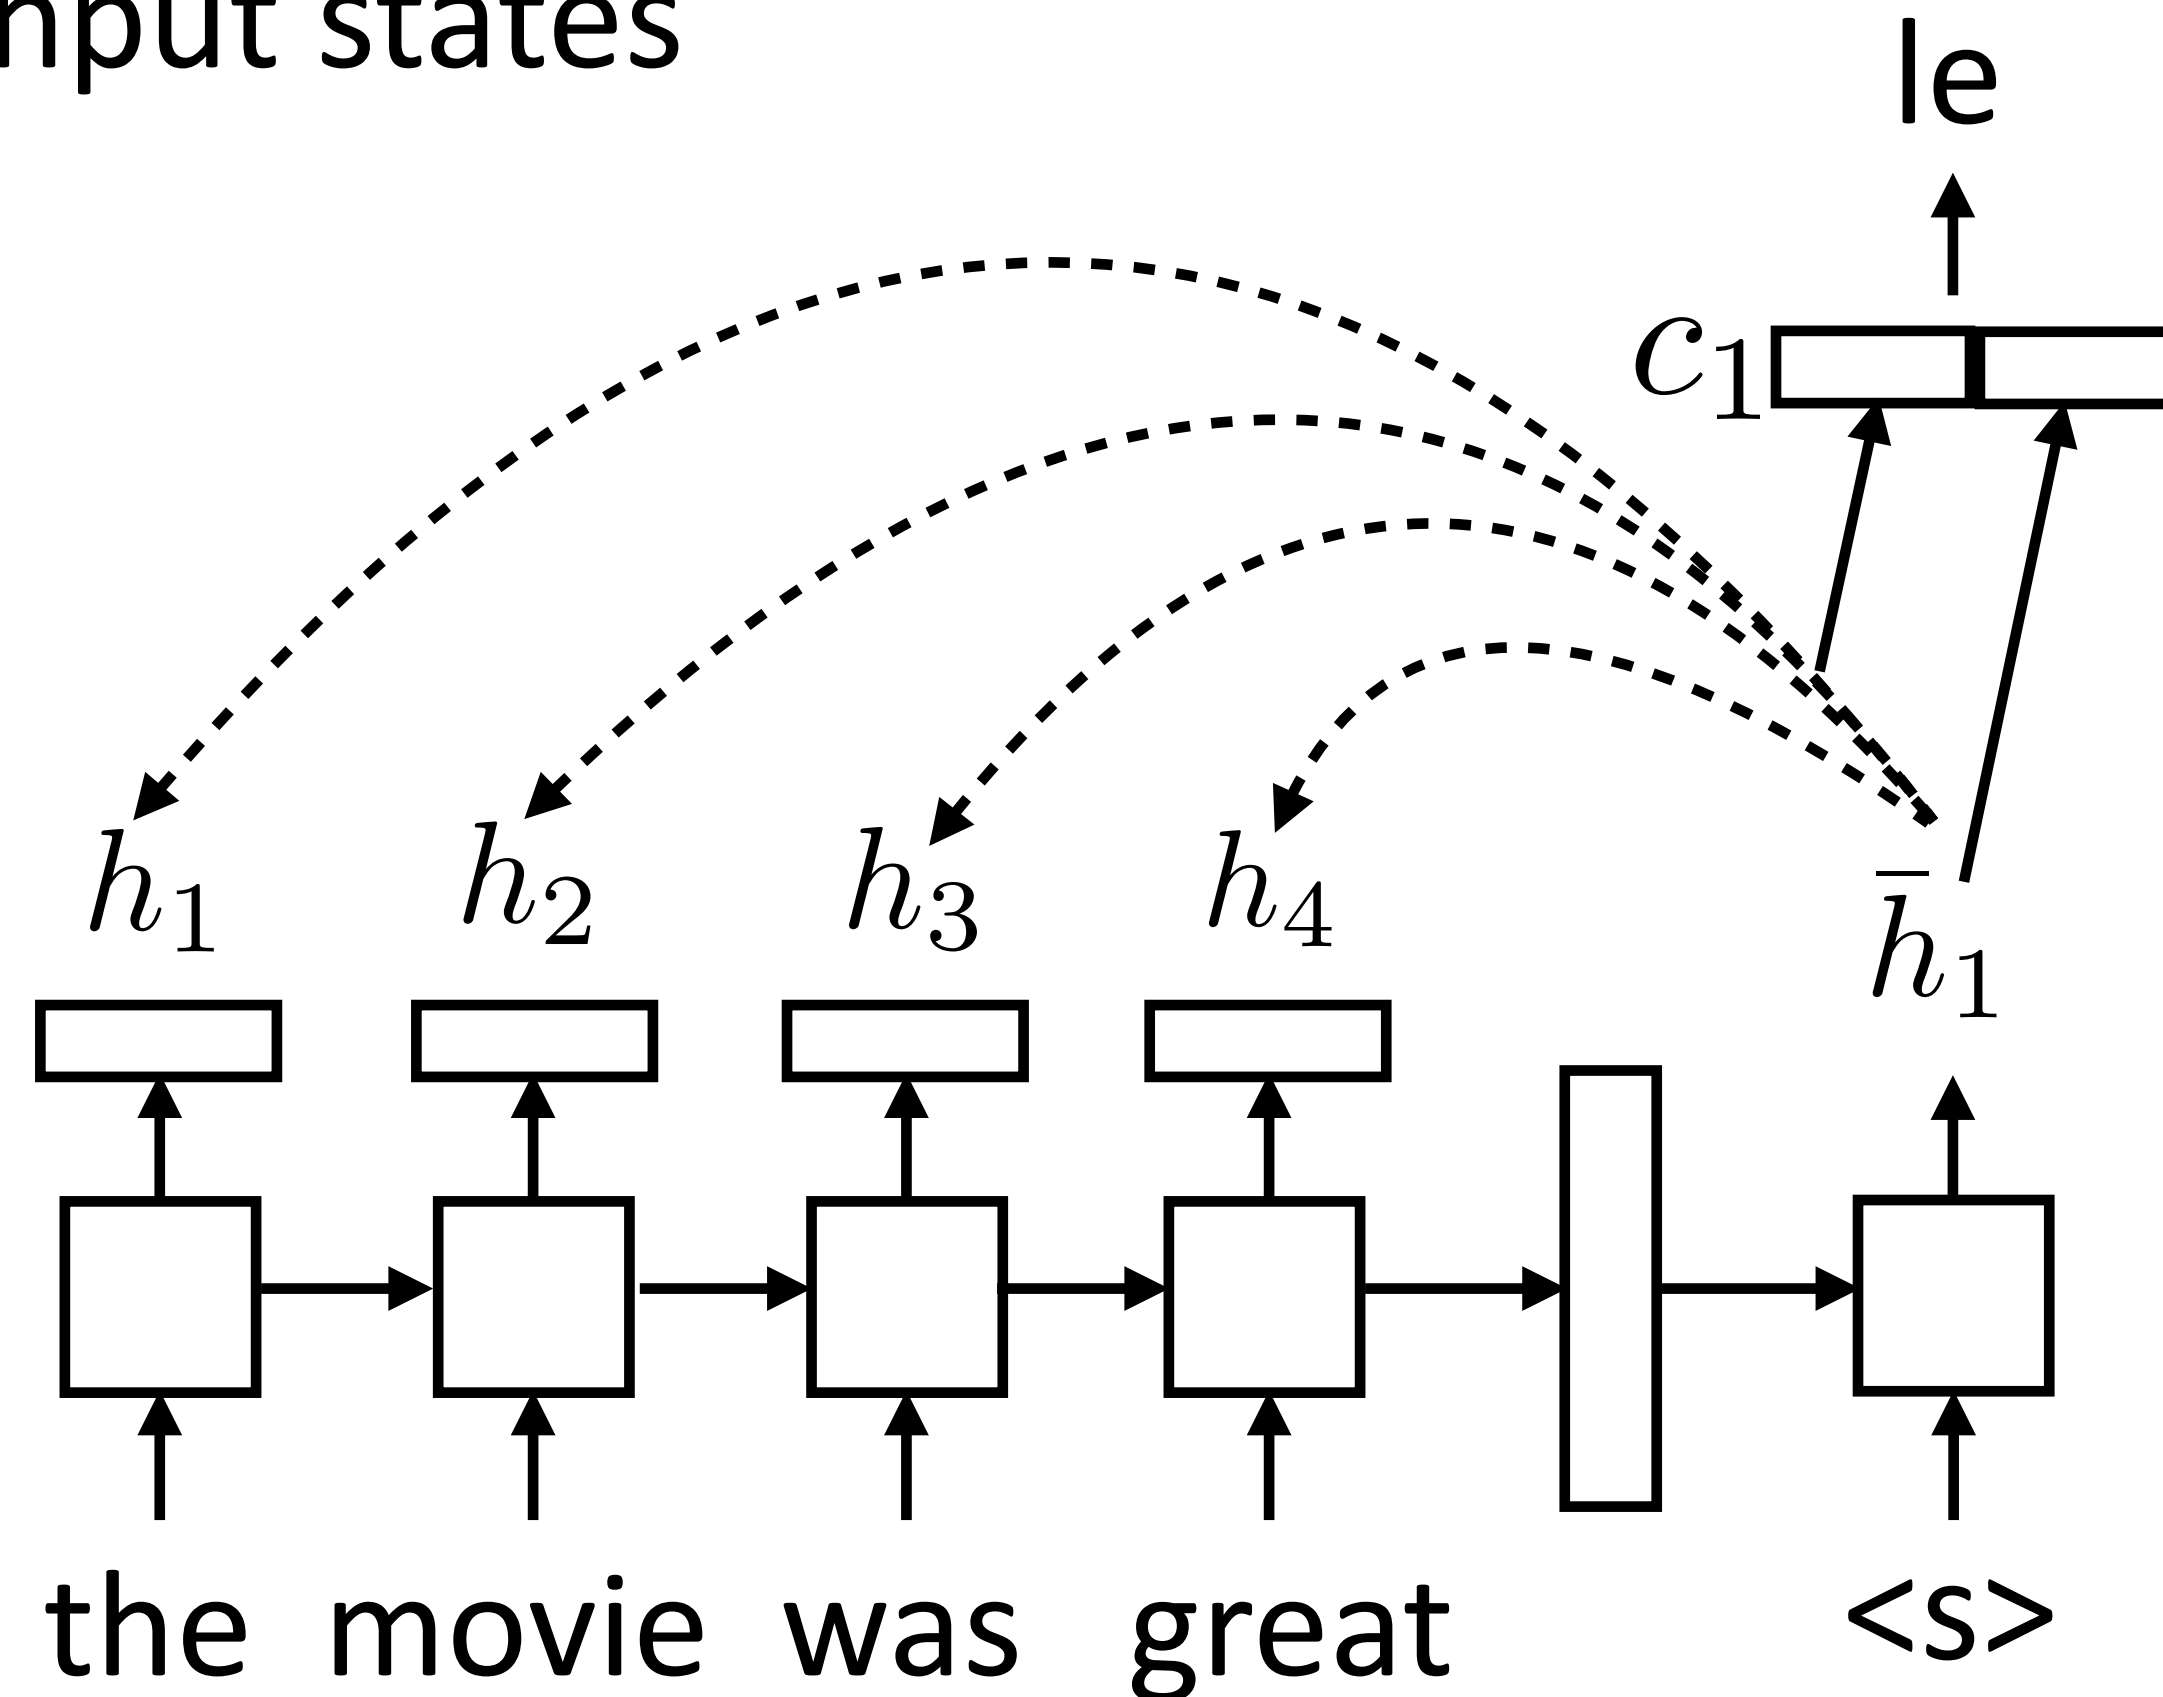


- ▶ At each decoder state, compute a distribution over source inputs based on current decoder state
- ▶ Use that in output layer

Attention

- For each decoder state, compute weighted sum of input states

- No attn: $P(y_i | \mathbf{x}, y_1, \dots, y_{i-1}) = \text{softmax}(W \bar{h}_i)$




$$P(y_i | \mathbf{x}, y_1, \dots, y_{i-1}) = \text{softmax}(W[c_i; \bar{h}_i])$$

$$c_i = \sum_j \alpha_{ij} h_j$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{j'} \exp(e_{ij'})}$$


$$e_{ij} = f(\bar{h}_i, h_j)$$

- Weighted sum of input hidden states (vector)



- Some function f (next slide)

Attention



$$c_i = \sum_j \alpha_{ij} h_j$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{j'} \exp(e_{ij'})}$$

$$e_{ij} = f(\bar{h}_i, h_j)$$

$$f(\bar{h}_i, h_j) = \tanh(W[\bar{h}_i, h_j])$$

► Bahdanau+ (2014): additive

$$f(\bar{h}_i, h_j) = \bar{h}_i \cdot h_j$$

► Luong+ (2015): dot product

$$f(\bar{h}_i, h_j) = \bar{h}_i^\top W h_j$$


► Luong+ (2015): bilinear

► Note that this all uses outputs of hidden layers

Neural MT


Encoder-Decoder MT

- ▶ encoder-decoder with attention and copying for rare words



Rare Words: Character Models

- ▶ If we predict an unk token, generate the results from a character LSTM
- ▶ Can potentially transliterate new concepts, but architecture is more complicated and slower to train
- ▶ Models like this in part contributed to dynamic computation graph frameworks becoming popular



Handling Rare Words

- ▶ Words are a difficult unit to work with: copying can be cumbersome, word vocabularies get very large
- ▶ Character-level models don't work well
- ▶ Solution: “word pieces” (which may be full words but may be subwords)

Input: *_the_ **eco tax** _port i co _in_ _Po nt - de - Bu is ...*

Output: *_le _port ique_ **éco taxe** _de_ _Pont - de - Bui s*

- ▶ Can help with transliteration; capture shared linguistic characteristics between languages (e.g., transliteration, shared word root, etc.)

Wu et al. (2016)

Byte Pair Encoding (BPE)

- ▶ Start with every individual byte (basically character) as its own symbol

```
for i in range(num_merges):
```

```
    pairs = get_stats(vocab)
```

```
    best = max(pairs, key=pairs.get)
```

```
    vocab = merge_vocab(best, vocab)
```

- ▶ Count bigram character cooccurrences

- ▶ Merge the most frequent pair of adjacent characters

- ▶ Do this either over your vocabulary (original version) or over a large corpus (more common version)
- ▶ Final vocabulary size is often in 10k ~ 30k range for each language
- ▶ Most SOTA NMT systems use this on both source + target

Word Pieces

while voc size < target voc size:

Build a language model over your corpus

Merge pieces that lead to highest improvement in language model perplexity

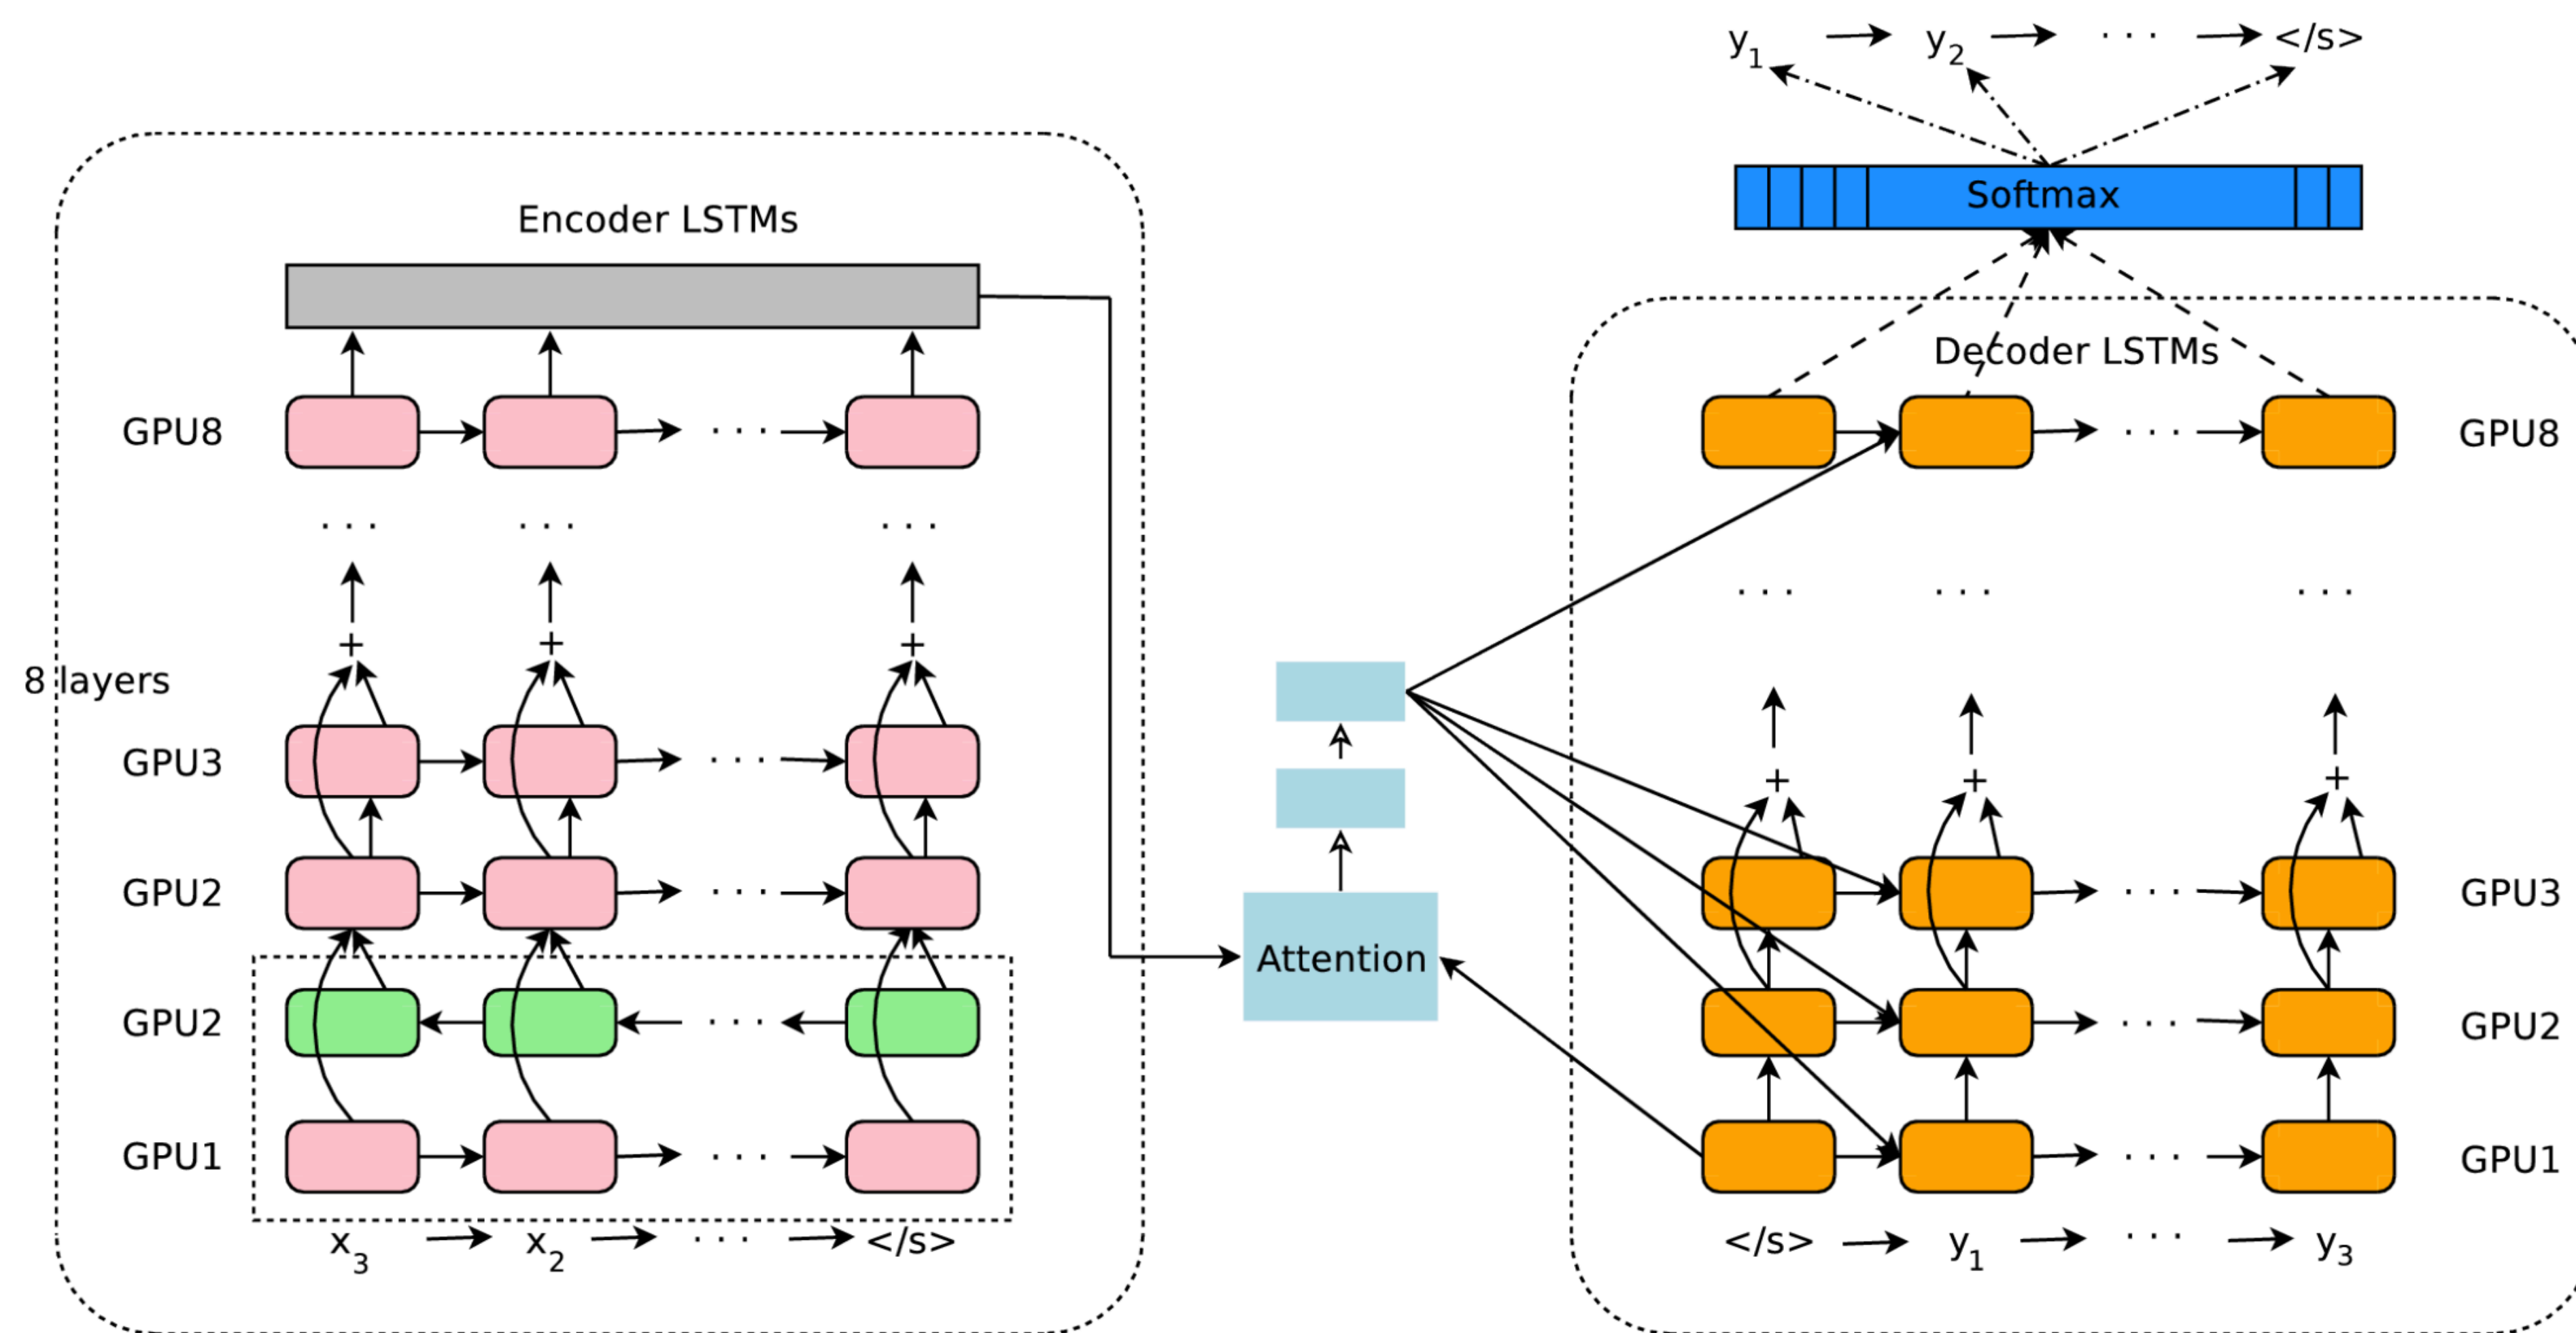
- ▶ SentencePiece library from Google: unigram LM
- ▶ Result: way of segmenting input appropriate for translation

Comparison

(a)	Original:	furiously		(b)	Original:	tricycles			
	BPE:	_fur	iously		BPE:	_t	ric	y	cles
	Unigram LM:	_fur	ious ly		Unigram LM:	_tri	cycle	s	
(c)	Original:	Completely preposterous suggestions							
	BPE:	_Comple	t	ely	_prep	ost	erous	_suggest	ions
	Unigram LM:	_Complete	ly	_pre	post	er	ous	_suggestion	s

- ▶ BPE produces less linguistically plausible units than word pieces (unigram LM)
- ▶ Some evidence that unigram LM works better in pre-trained transformer models

Google's NMT System



- ▶ 8-layer LSTM encoder-decoder with attention, word piece vocabulary of 8k-32k

Google's NMT System

English-French:

Google's phrase-based system: 37.0 BLEU

Luong+ (2015) seq2seq ensemble with rare word handling: 37.5 BLEU

Google's 32k word pieces: 38.95 BLEU

English-German:

Google's phrase-based system: 20.7 BLEU

Luong+ (2015) seq2seq ensemble with rare word handling: 23.0 BLEU

Google's 32k word pieces: 24.2 BLEU

Google's NMT System


Source	She was spotted three days later by a dog walker trapped in the quarry	
PBMT	Elle a été repéré trois jours plus tard par un promeneur de chien piégé dans la carrière	6.0
GNMT	Elle a été repérée trois jours plus tard par un traîneau à chiens piégé dans la carrière.	2.0
Human	Elle a été repérée trois jours plus tard par une personne qui promenait son chien coincée dans la carrière	5.0

Gender is correct in GNMT
but not in PBMT

“sled” “walker”

The right-most column shows the human ratings on a scale of 0 (complete nonsense) to 6 (perfect translation)

Human Evaluation (En-Es)




- ▶ Similar to human-level performance *on English-Spanish*

Figure 6: Histogram of side-by-side scores on 500 sampled sentences from Wikipedia and news websites for a typical language pair, here English → Spanish (PBMT blue, GNMT red, Human orange). It can be seen that there is a wide distribution in scores, even for the human translation when rated by other humans, which shows how ambiguous the task is. It is clear that GNMT is much more accurate than PBMT.

Frontiers in MT

Non-Autoregressive NMT



- Q: why non-autoregressive? Pros and cons?

Gu et al. (2018), Ghazvininejad et al. (2019), Kasai et al. (2020)

Low-Resource MT

- ▶ Particular interest in deploying MT systems for languages with little or no parallel data

- ▶ BPE allows us to transfer models even without training on a specific language

- ▶ Pre-trained models can help further

Burmese, Indonesian, Turkish
BLEU

Transfer	My→En	Id→En	Tr→En
baseline (no transfer)	4.0	20.6	19.0
transfer, train	17.8	27.4	20.3
transfer, train, reset emb, train	13.3	25.0	20.0
transfer, train, reset inner, train	3.6	18.0	19.1

Table 3: Investigating the model’s capability to restore its quality if we reset the parameters. We use En→De as the parent.

Unsupervised MT

Approach	Train/Val	Test	Loss
Supervised MT	L1-L2	L1-L2	$\mathcal{L}_{x \rightarrow y}^{MT} = \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim (\mathcal{X}, \mathcal{Y})} [-\log p_{x \rightarrow y}(\mathbf{y} \mathbf{x})]$
Unsupervised MT	L1, L2	L1-L2	$\mathcal{L}_{x \leftrightarrow y}^{BT} = \mathbb{E}_{\mathbf{x} \sim \mathcal{X}} [-\log p_{y \rightarrow x}(\mathbf{x} g^*(\mathbf{x}))]$ $+ \mathbb{E}_{\mathbf{y} \sim \mathcal{Y}} [-\log p_{x \rightarrow y}(\mathbf{y} h^*(\mathbf{y}))]$ $g^*, h^*: \text{sentence predictors}$

- ▶ Common principles of unsupervised MT
 - ▶ Language models
 - ▶ (Iterative) Back-translation!

Takeaways

- ▶ Can build MT systems with LSTM encoder-decoders, CNNs, or transformers
- ▶ Word piece / byte pair models are really effective and easy to use
- ▶ State of the art systems are getting pretty good, but lots of challenges remain, especially for low-resource settings